

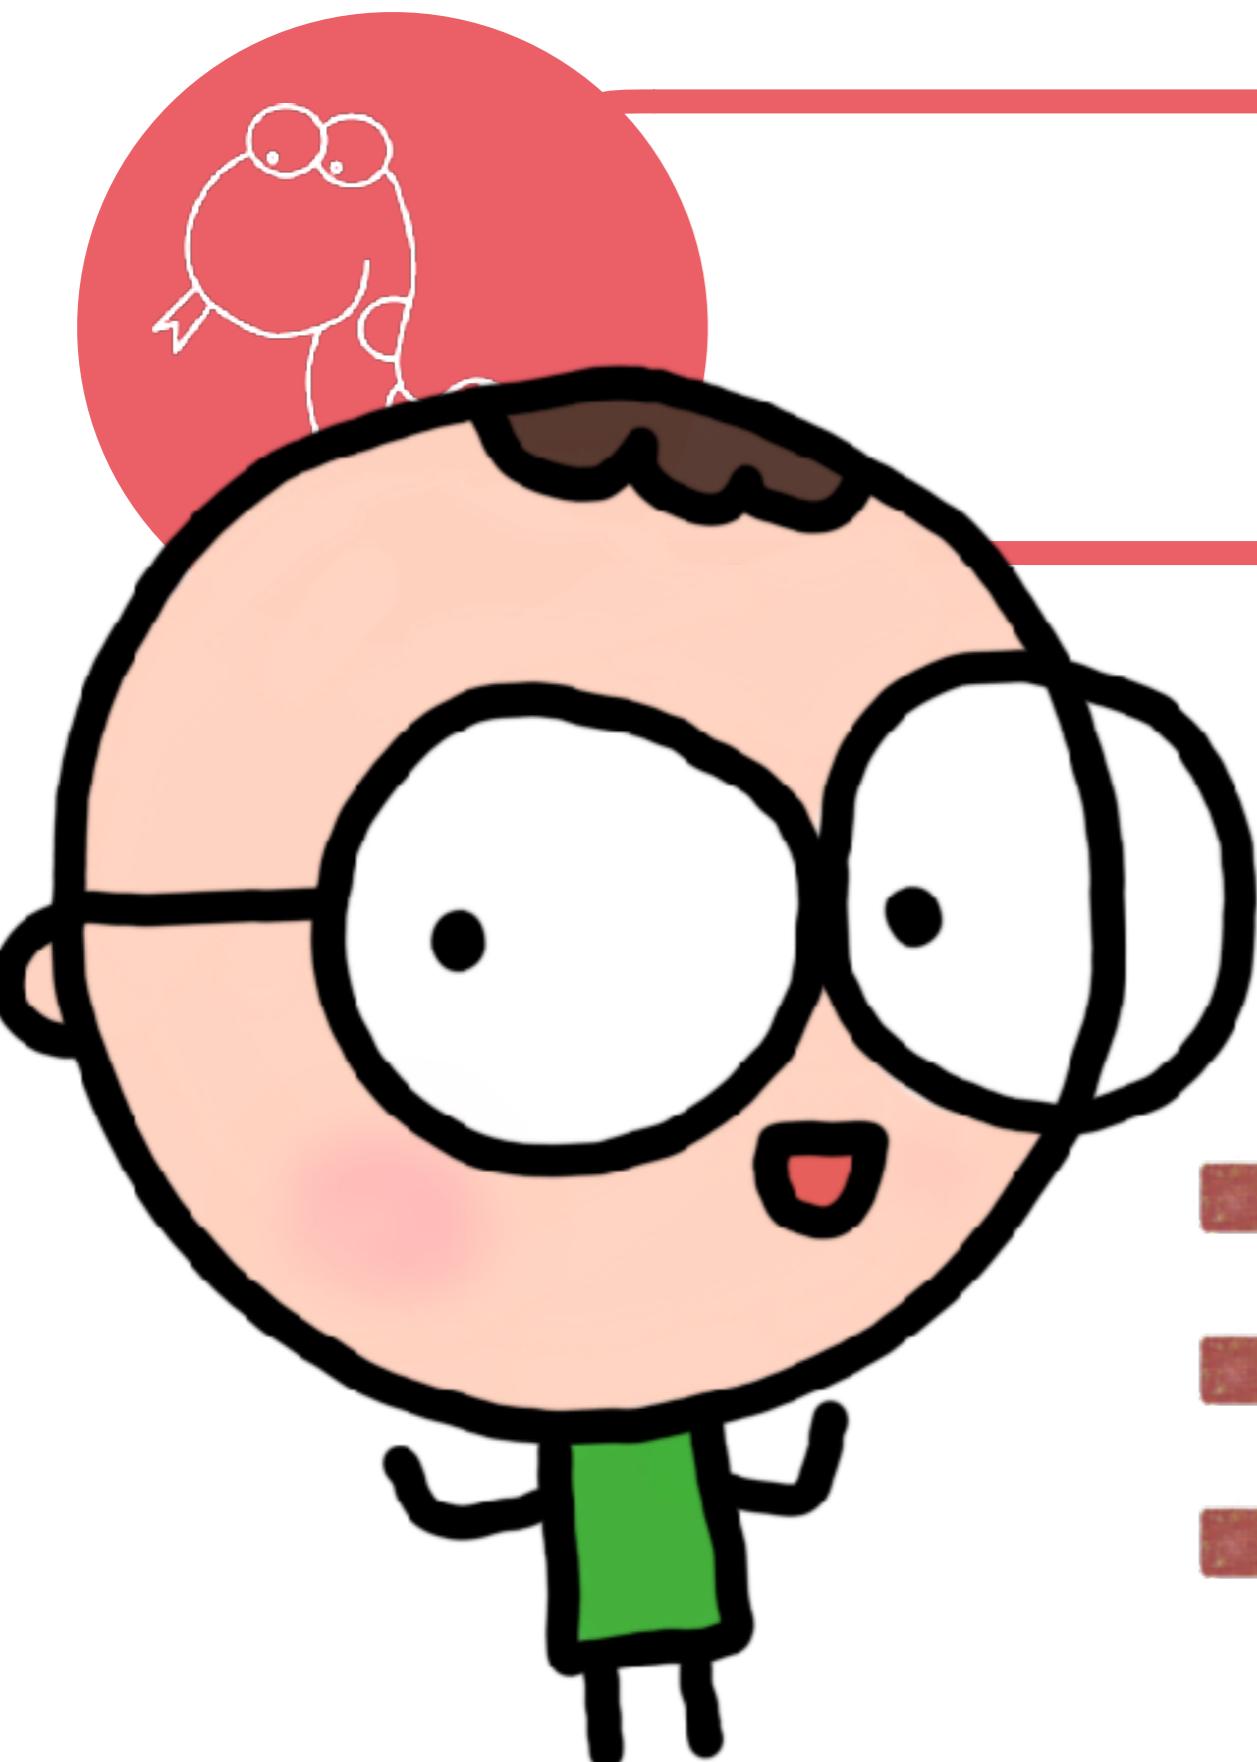
Seminar in Information Systems and Applications

Deep Learning and Design Thinking



Yen-lung Tsai

Department of Mathematical Sciences
National Chengchi University



Outline

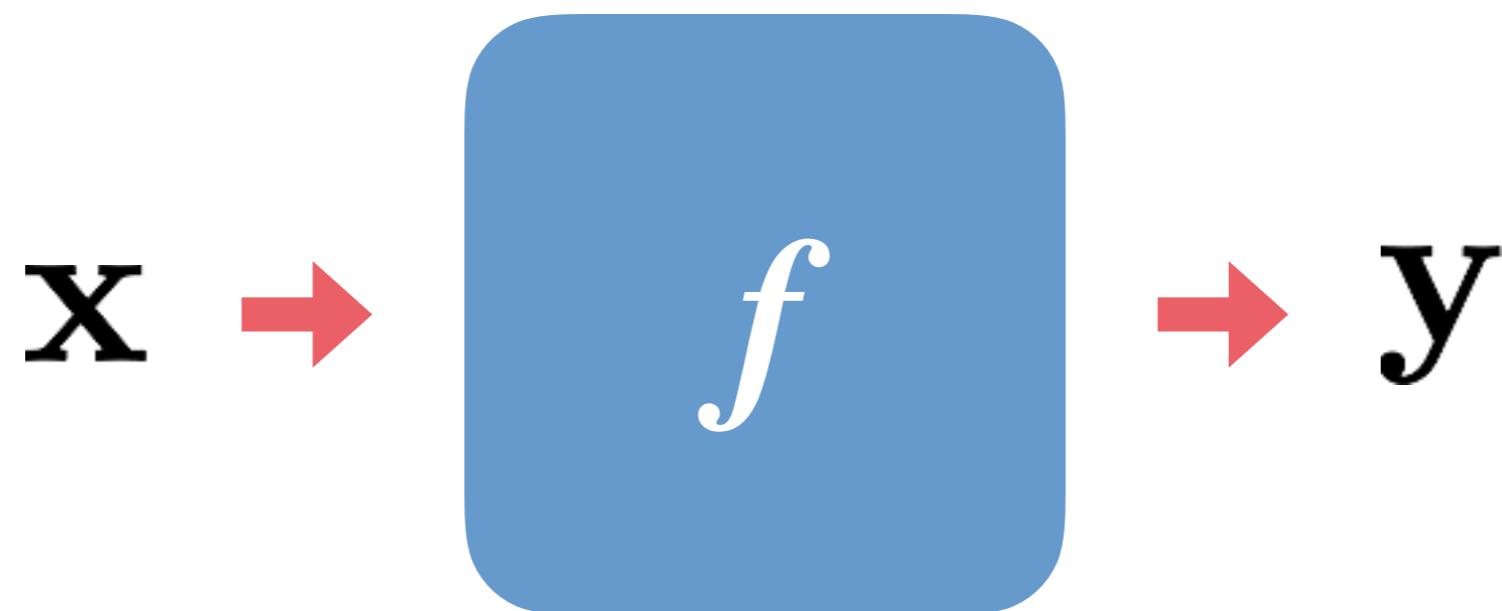
- Asking a good question
- Learning Neural Networks
- Applying Design Thinking

1

Asking a Good Question



What we want to do is just transforming the real world problem to a function. Then we **use deep learning techniques** to find the function.



**Functions are solution
manuals**

I wan to know what
the animal in the
picture is.



Question



Answer

**Formosan
black bear**



Python



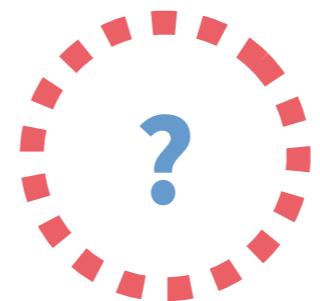
We have partial answers

Quesiton



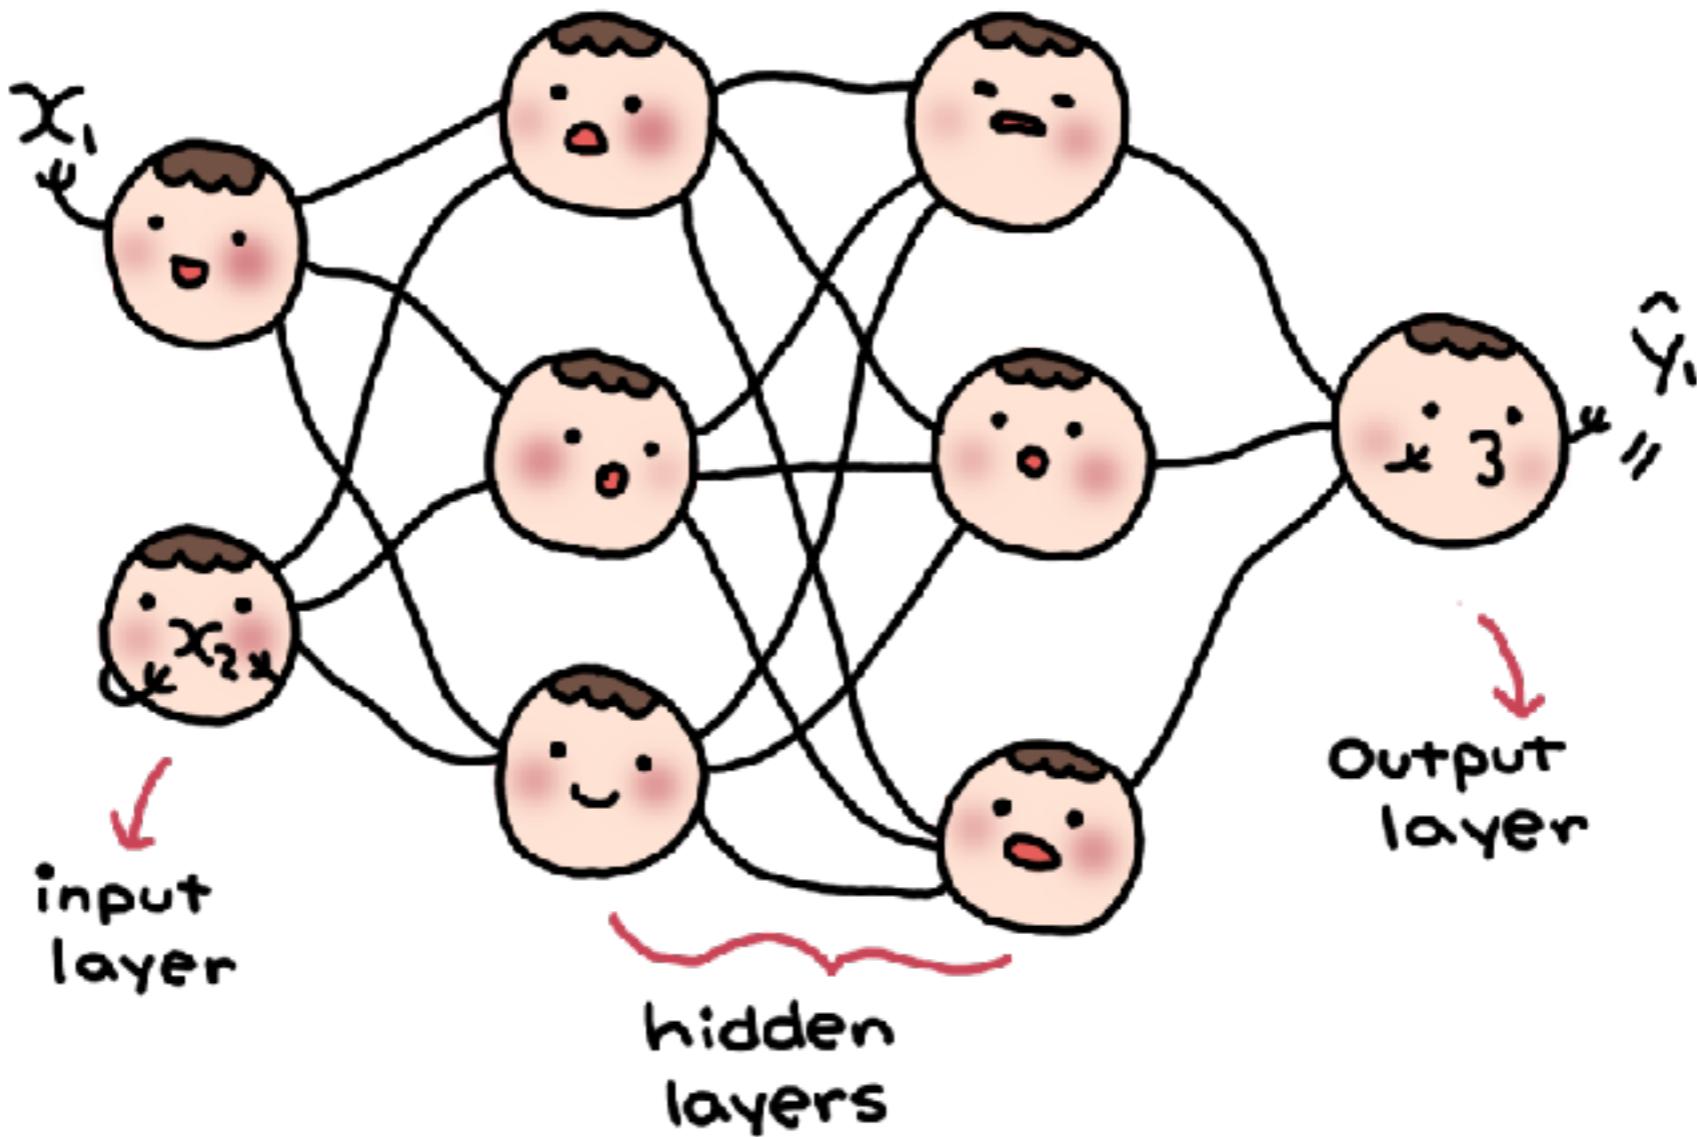
Answer

**Formosan
black bear**



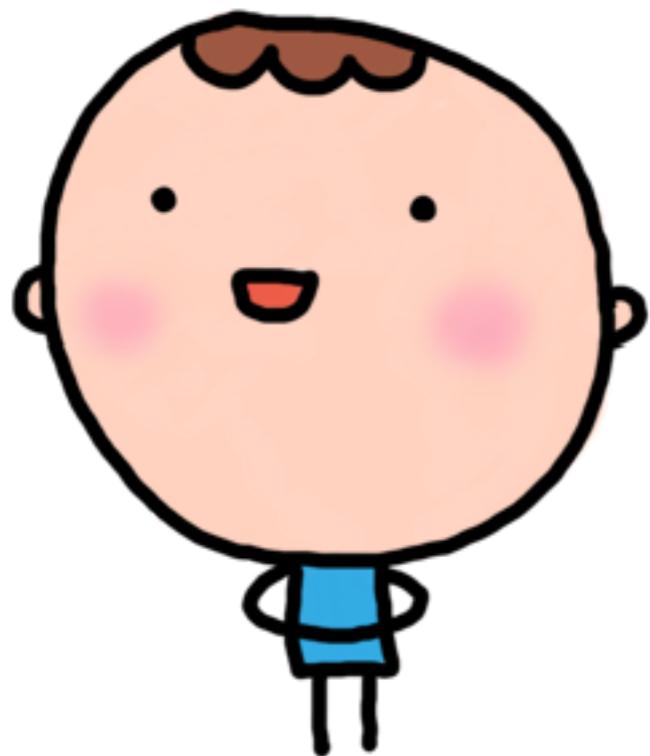
Python

**There might be infinitely many possible
cases that we haven't seen before**



Finding our functions by Neural Networks!

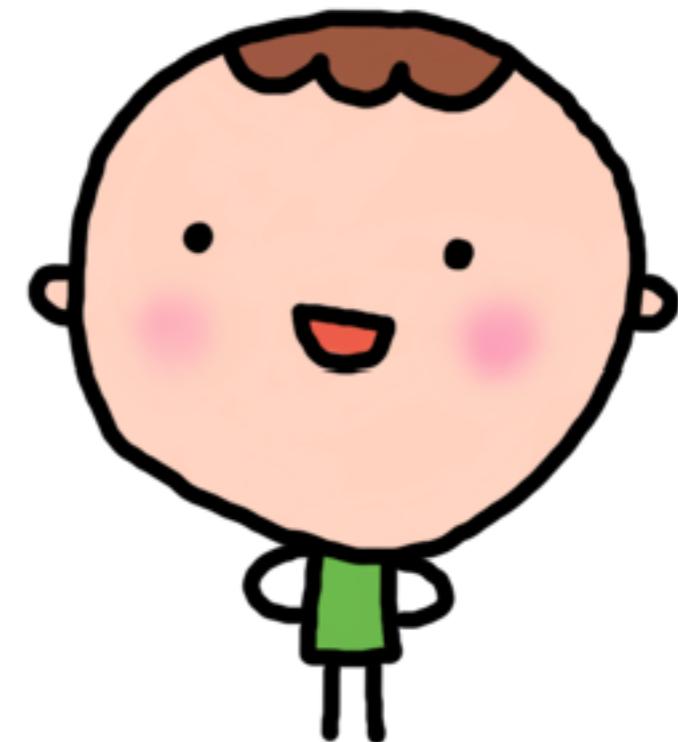
“Big 3” of Deep Learning (Neural Works)



Standard NN



CNN



RNN

1

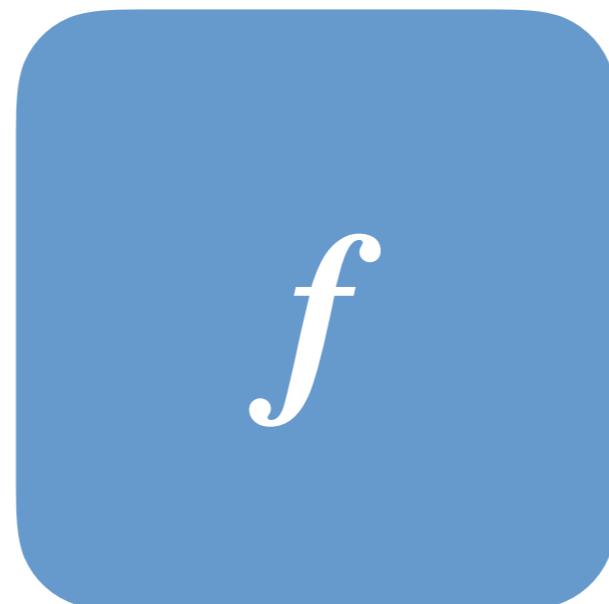
Asking a Question

**Seeing an animal in wild, we want to
know what is is?**



2

Transforming our question into a function



**Formosan
black bear**

3

Preparing Data Sets for Training

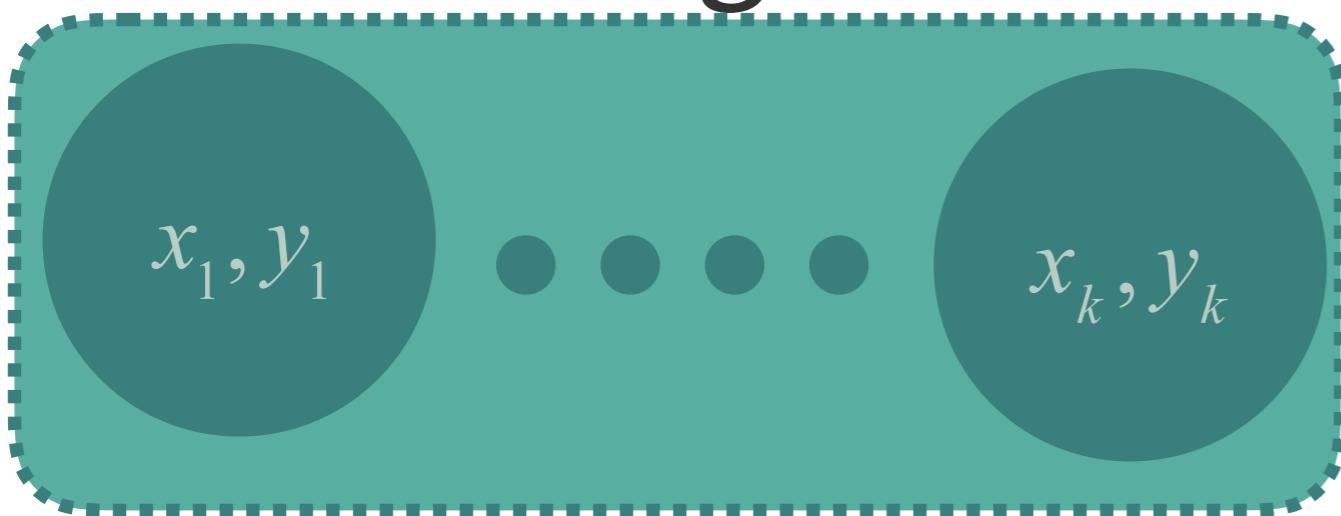


“Formosan
black bear”)

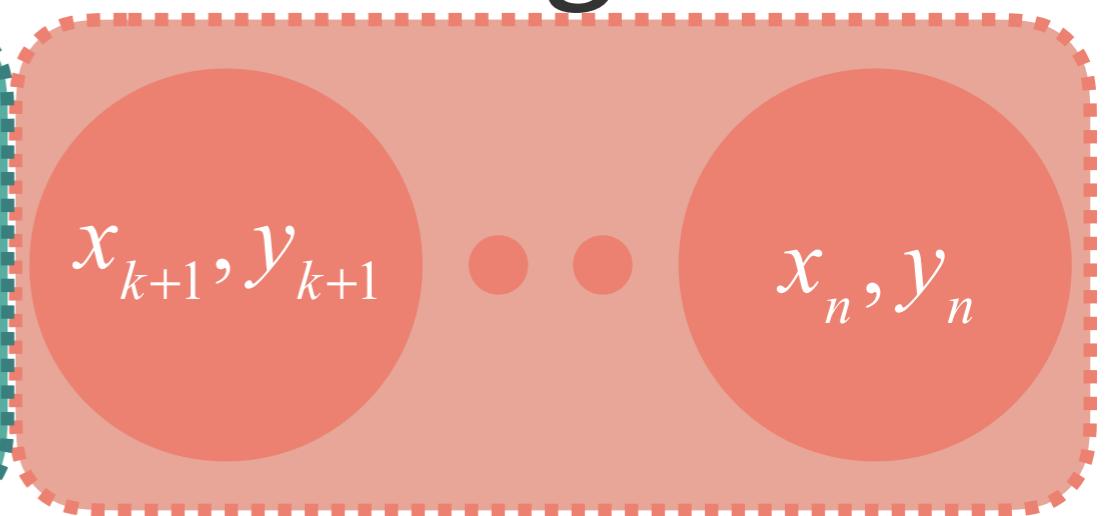


, “Python”), ...

Training Data



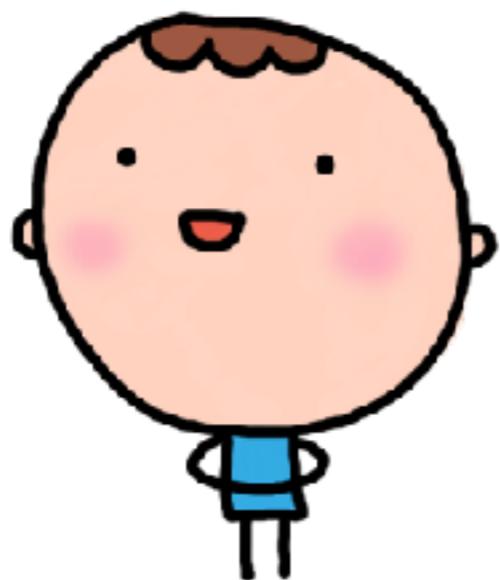
Testing Data



4

Constructing NN

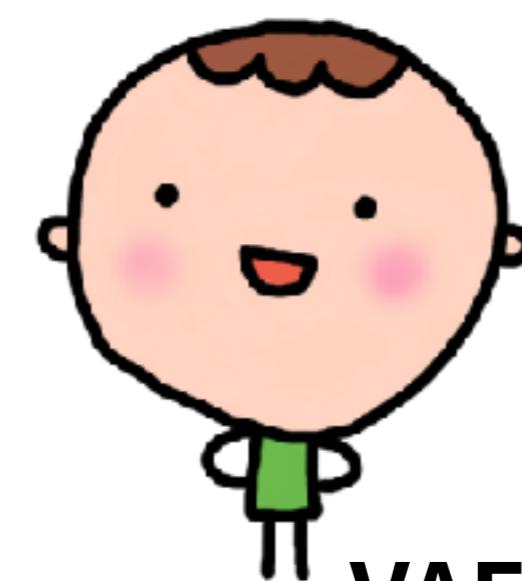
Standard NN



CNN



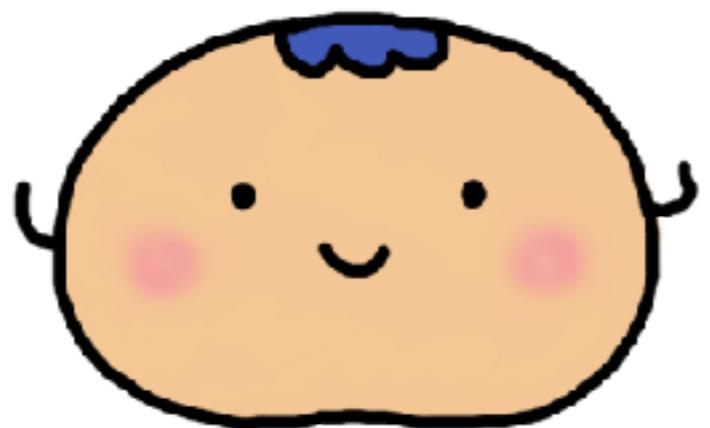
RNN



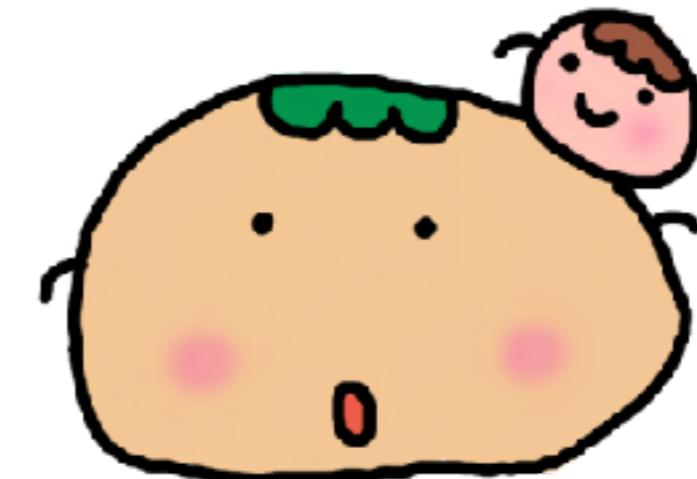
VAE



Capsule



Reinforcement
Learning (RL)



Generative Adversarial
Network (GAN)

Once we decide our structure of neural network, we have a set of parameters need to adjust.

$$\theta = \{w_i, b_j\}$$

Once we determine the parameters, we have a function:

$$f_{\theta}$$

Learning

The learning is sent to our neural network using our training data, adjusting our parameters, and then using a loss function to see how much difference we have and the observed value.

$$L(\theta)$$

Basically, we use the method called:

gradient descent

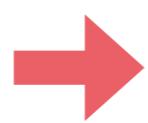
For Neural Networks, the method is also called:

backpropagation



I want to know the
closing price of a
stock tomorrow.

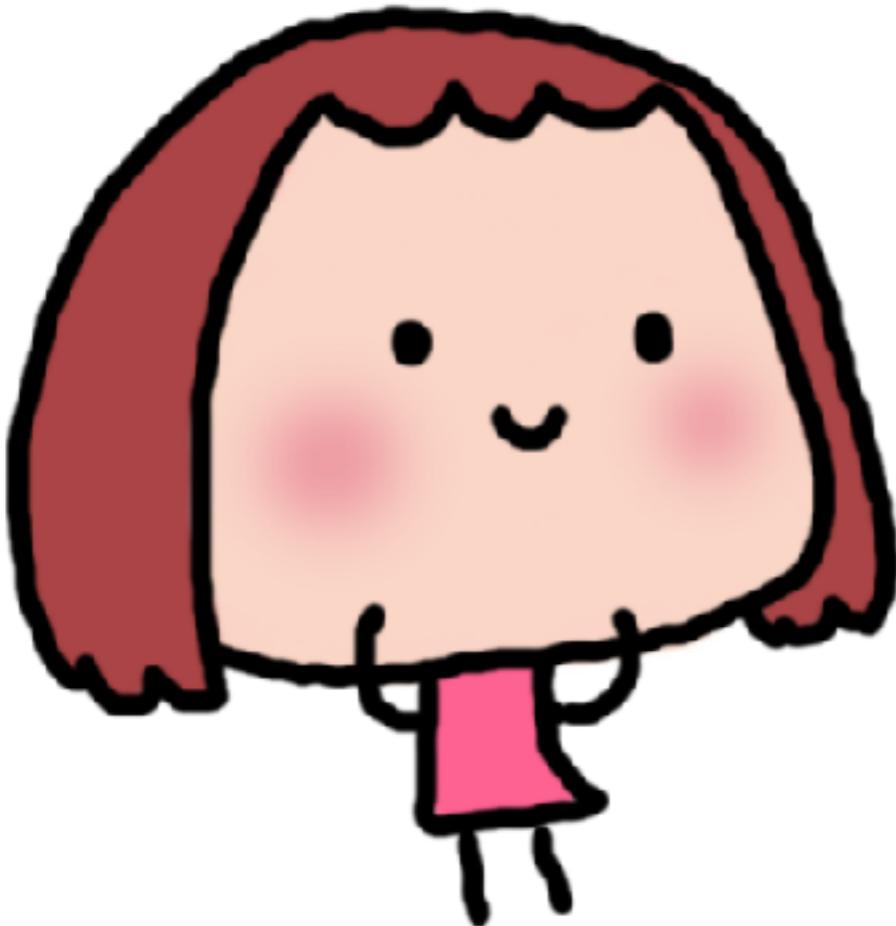
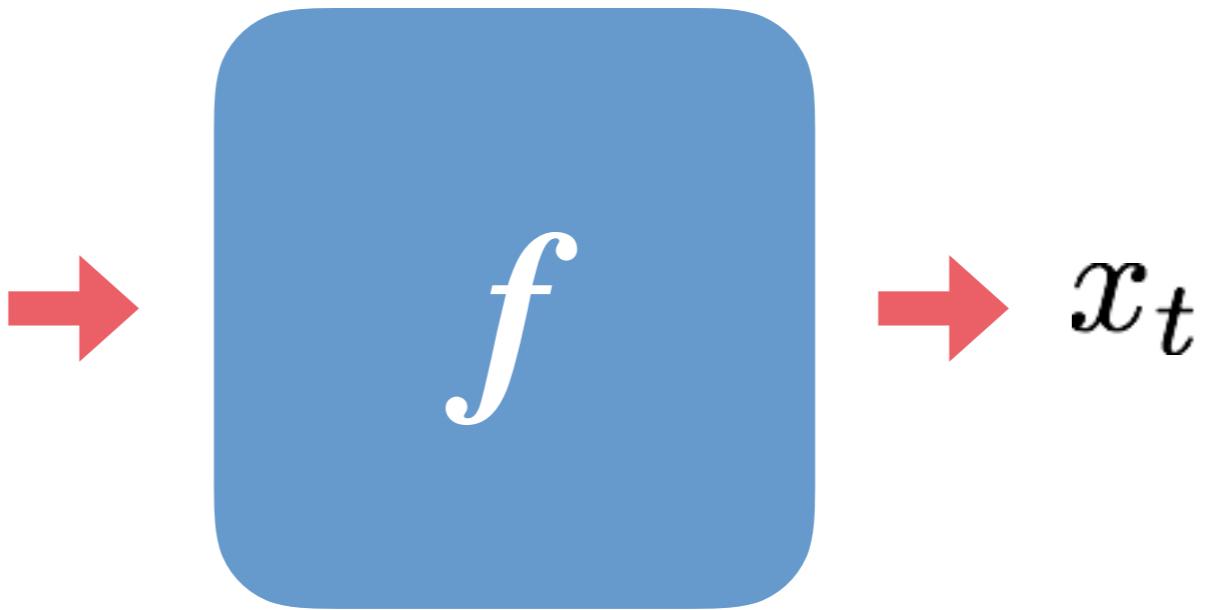
Date x →



**the closing price of
a stock x**

NN, CNN, RNN

$x_{t-1}, x_{t-2}, x_{t-3},$
 x_{t-4}, x_{t-5}



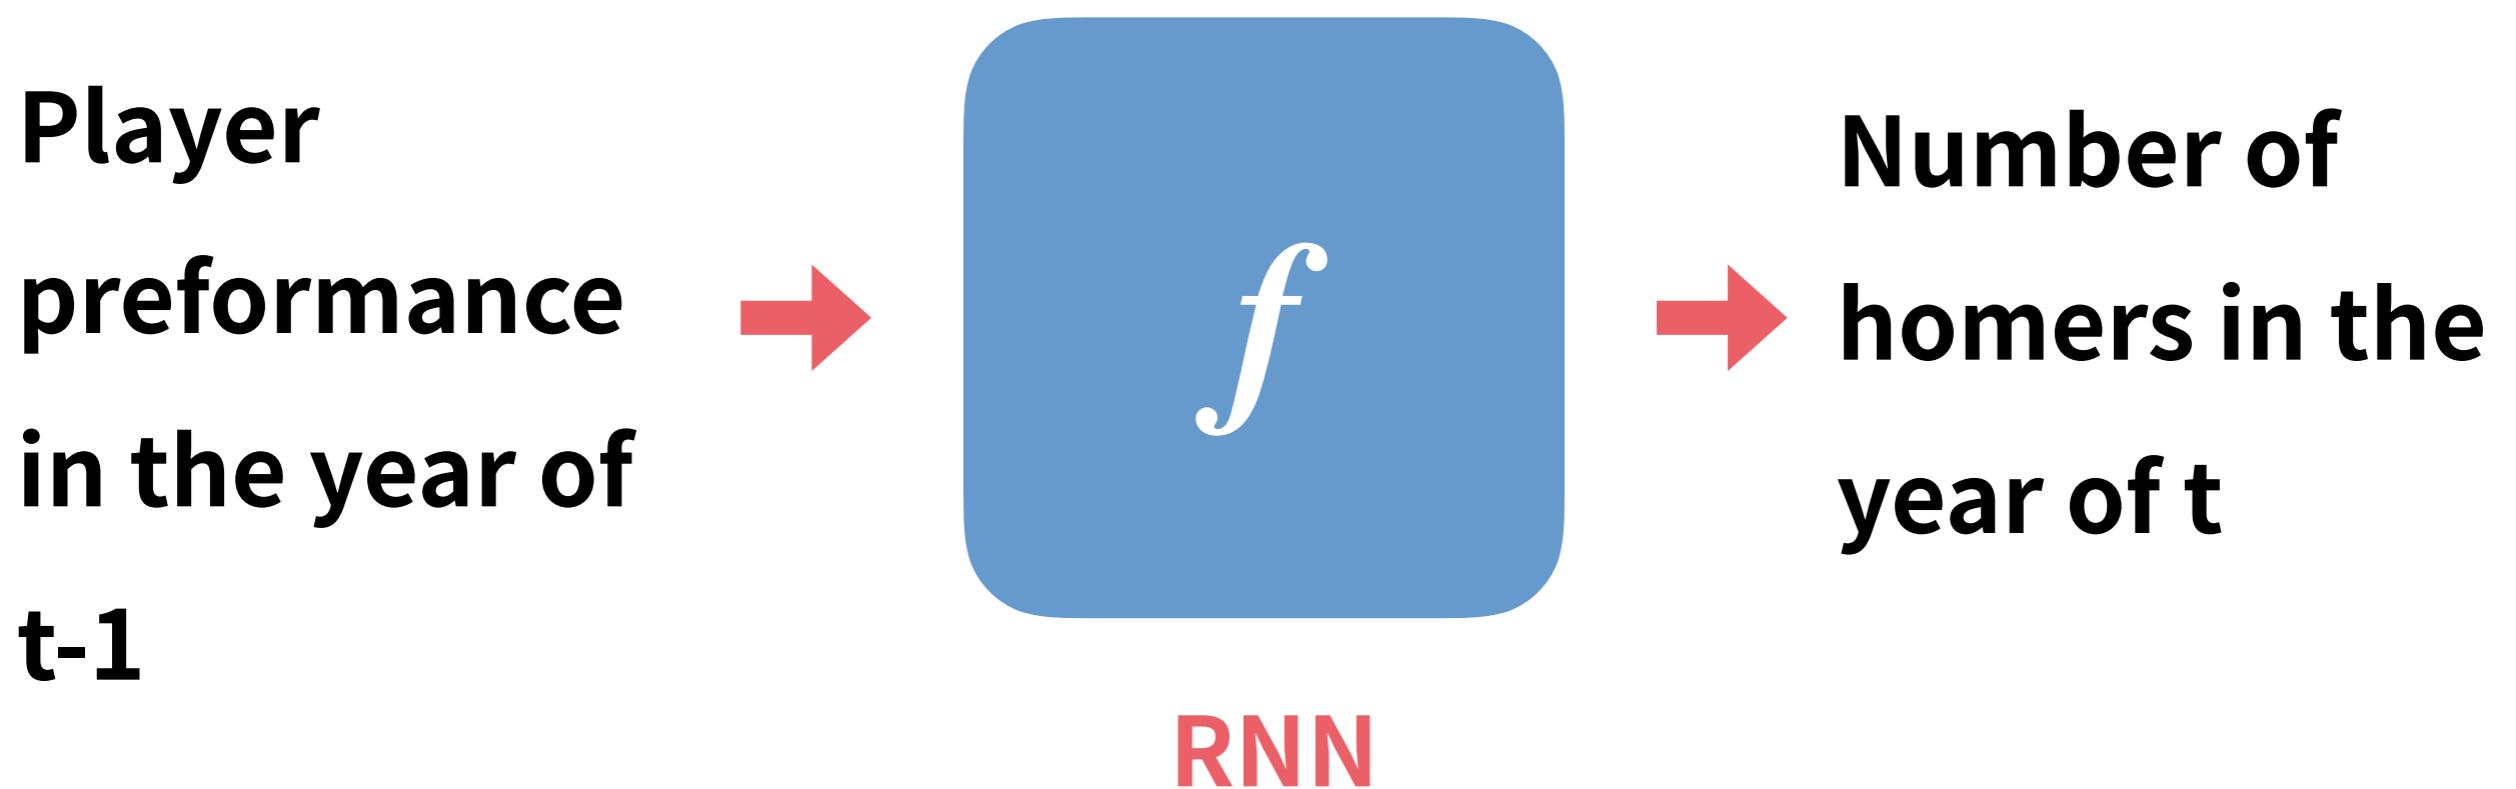
Use the prices of the
previous week to predict
the price for the next day.

I want to know how
many homers a
particular MLB
player can hit in 2018
season.

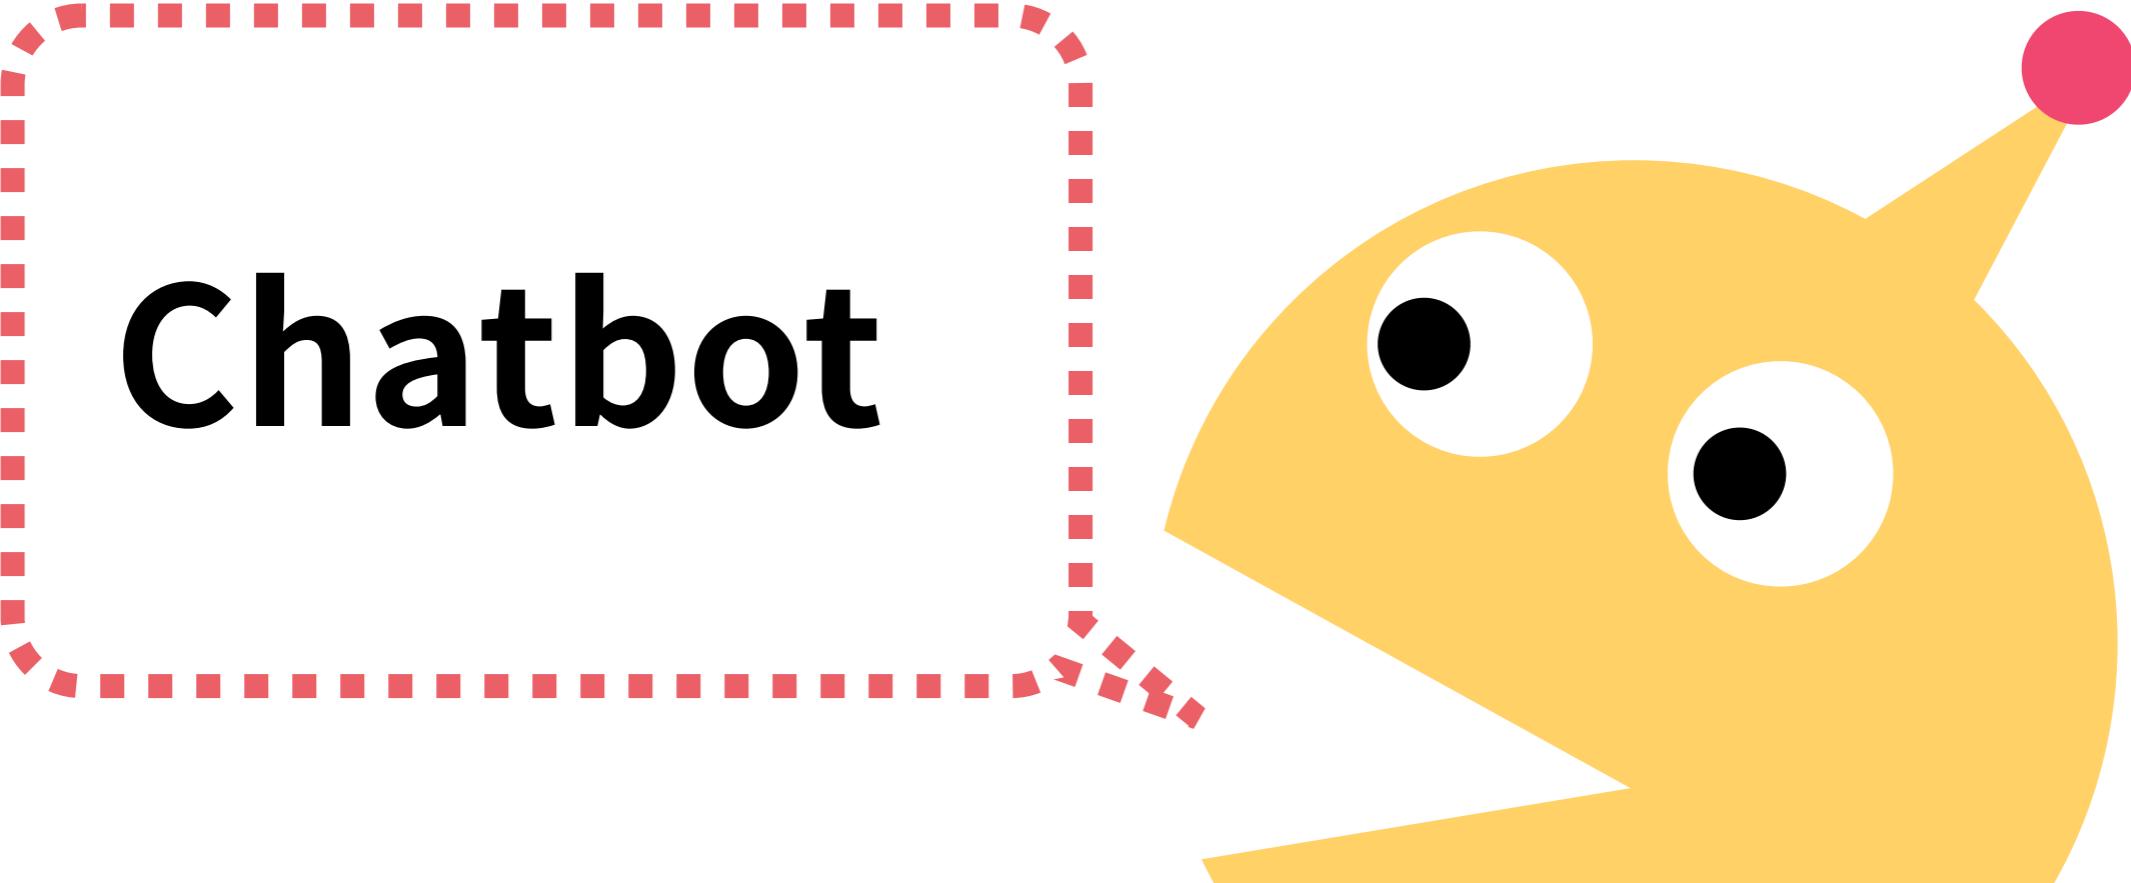


15 features!

[Age, G, PA, AB, R, H, 2B, 3B, HR,
RBI, SB, BB, SO, OPS+, TB]



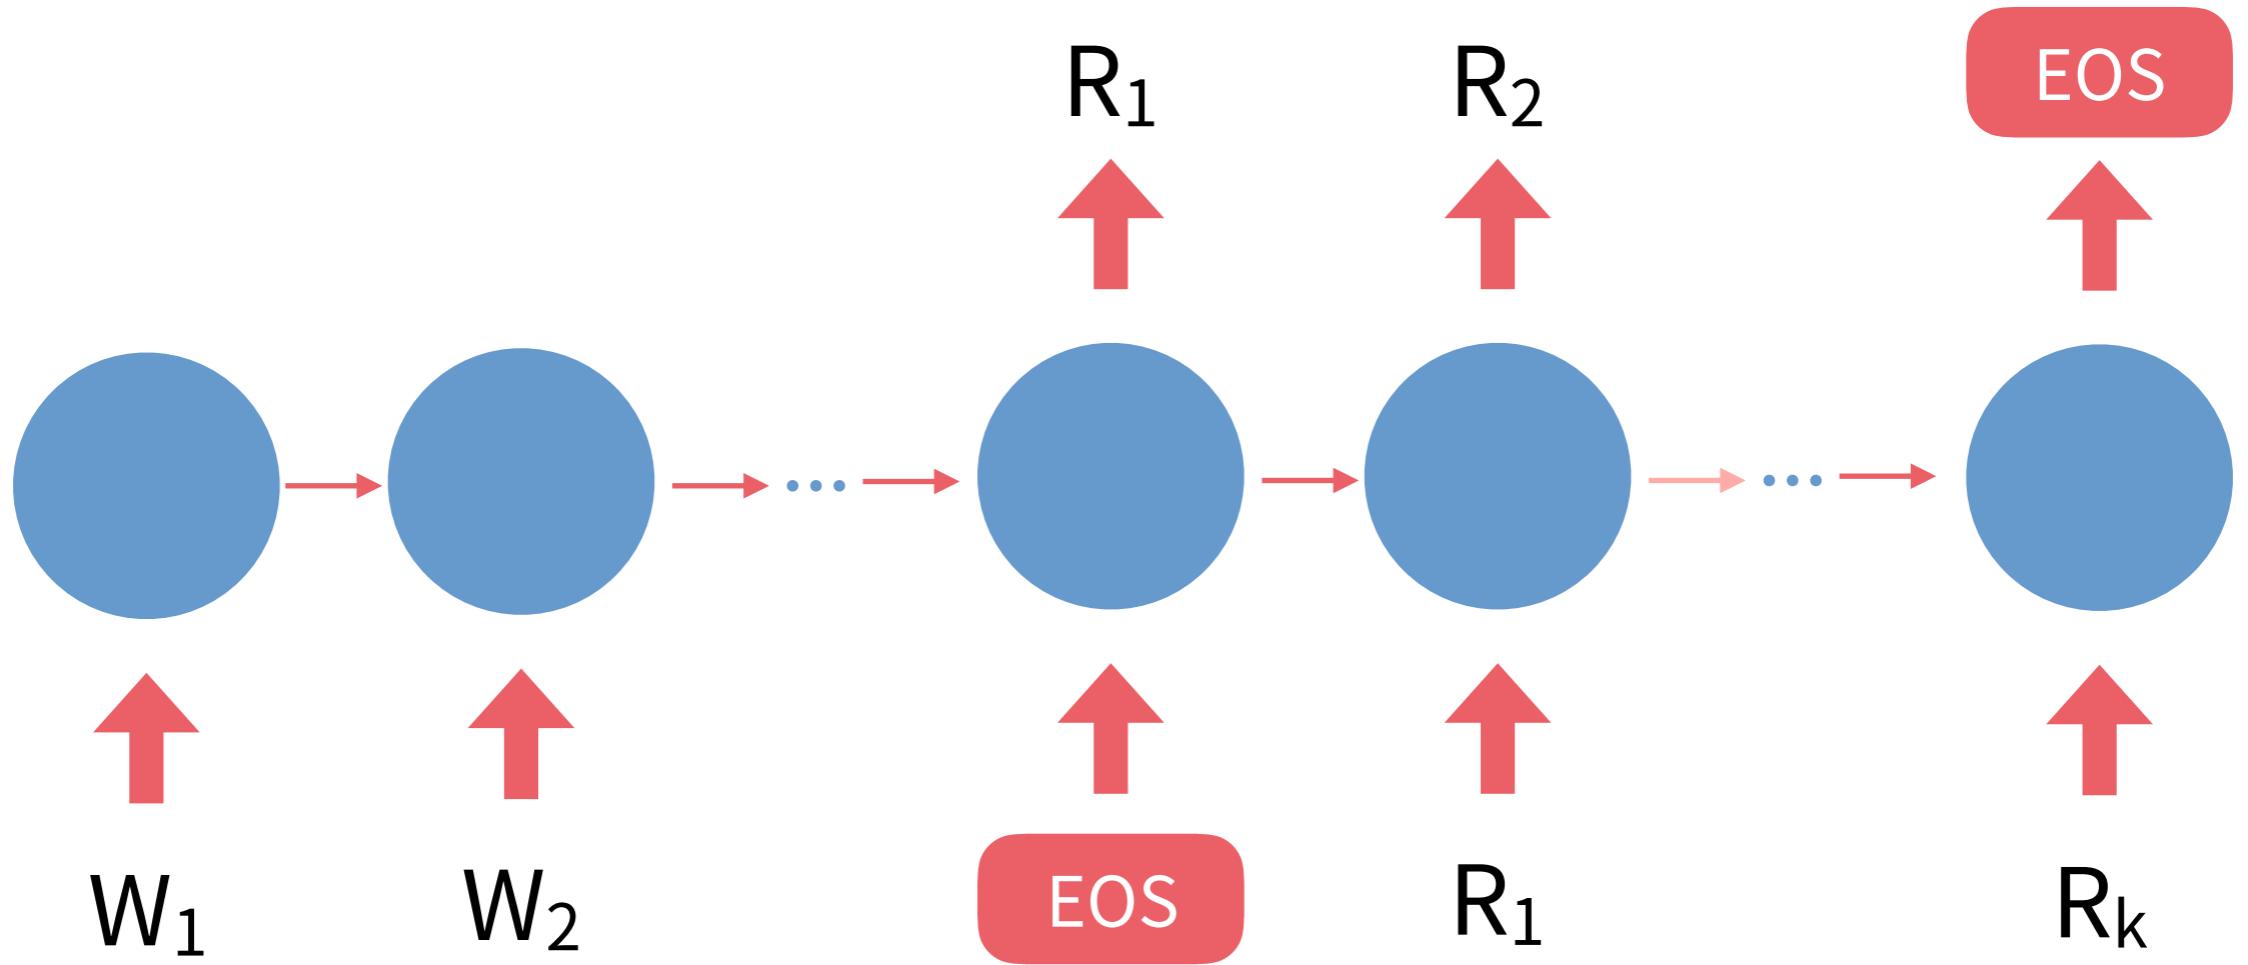
Chatbot



Current word →



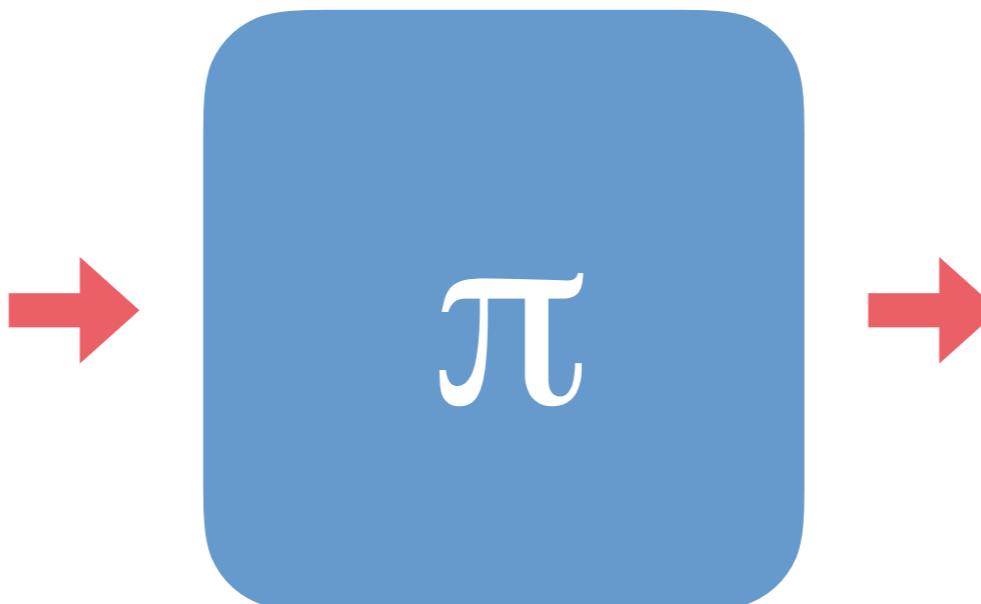
→ **Next word**



**Note that in this mode, each input and output
is not of fixed length!**

I want to use AI to
play games (drive a
car, brew coffee...)

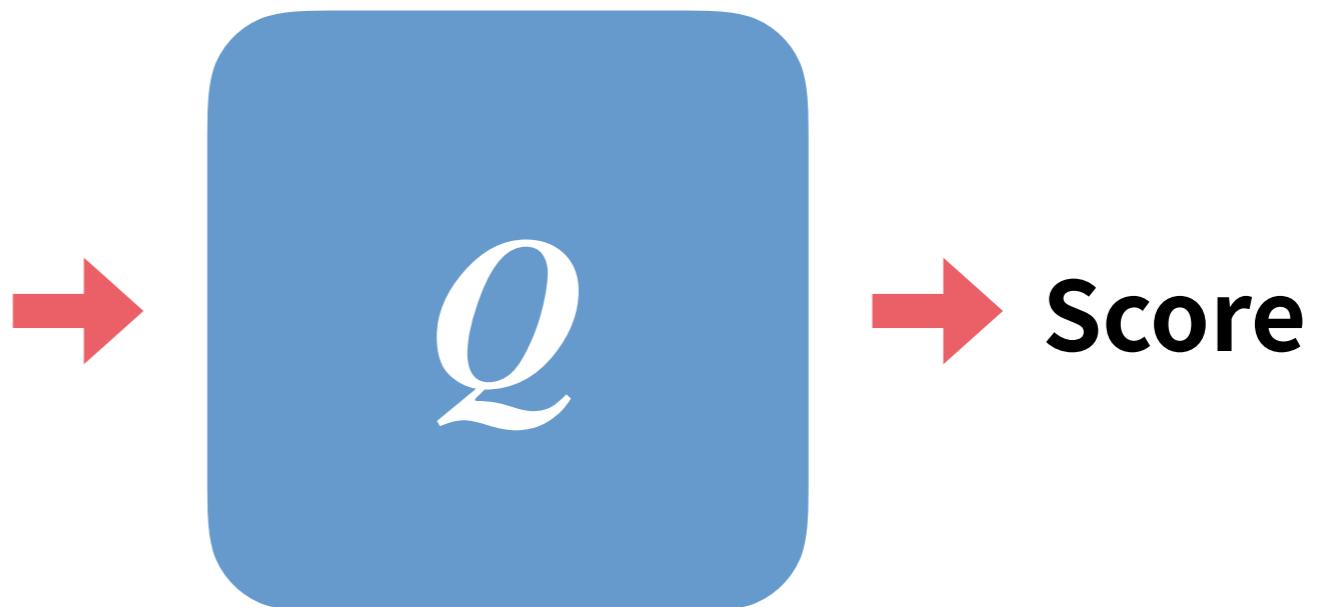




CNN + NN

The best
action

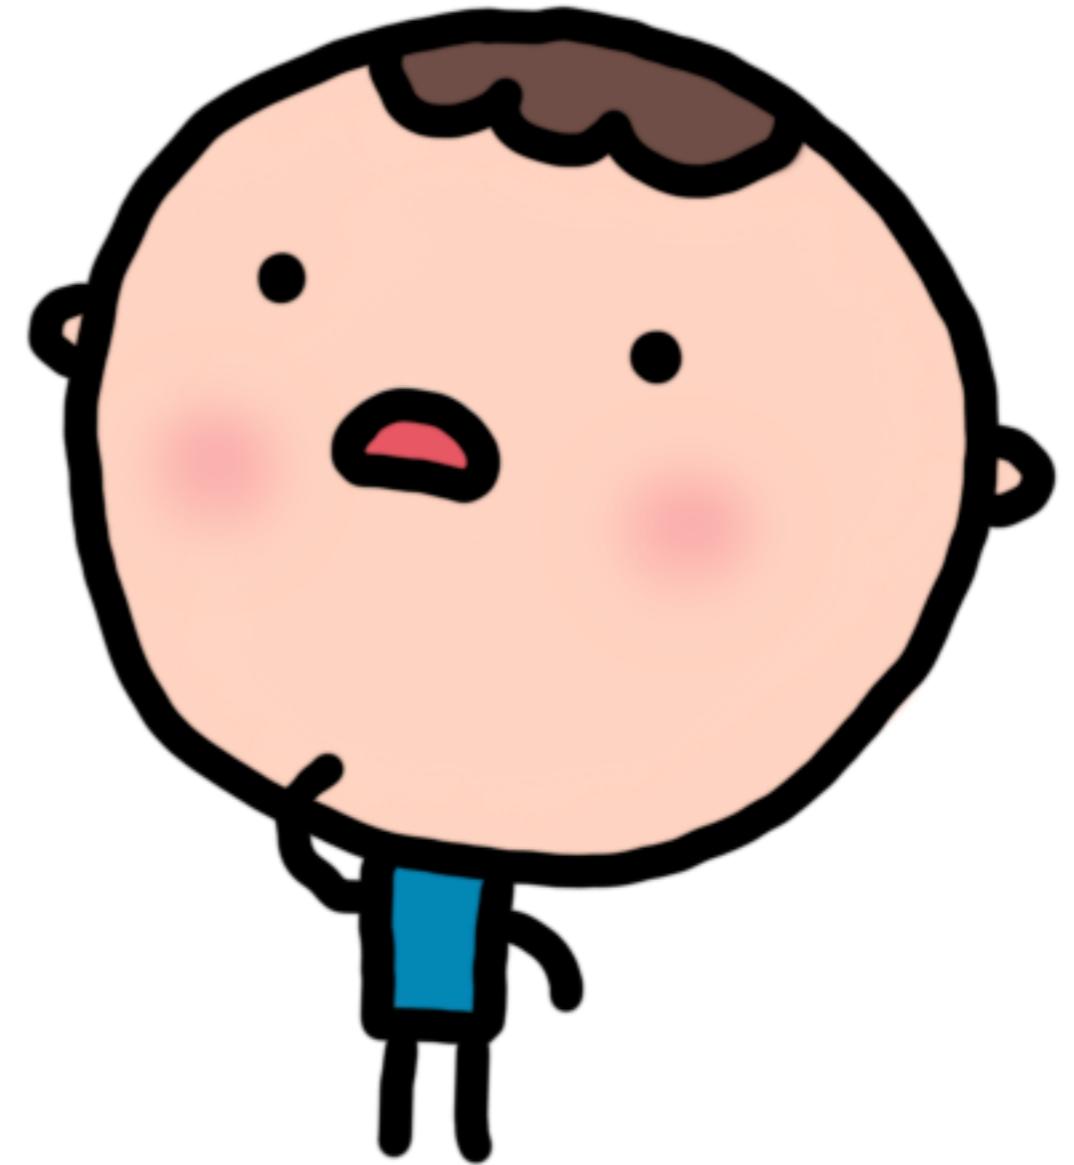
Reinforcement Learning



+

Left

**some characters
are missing in a
font I like...**



Font A



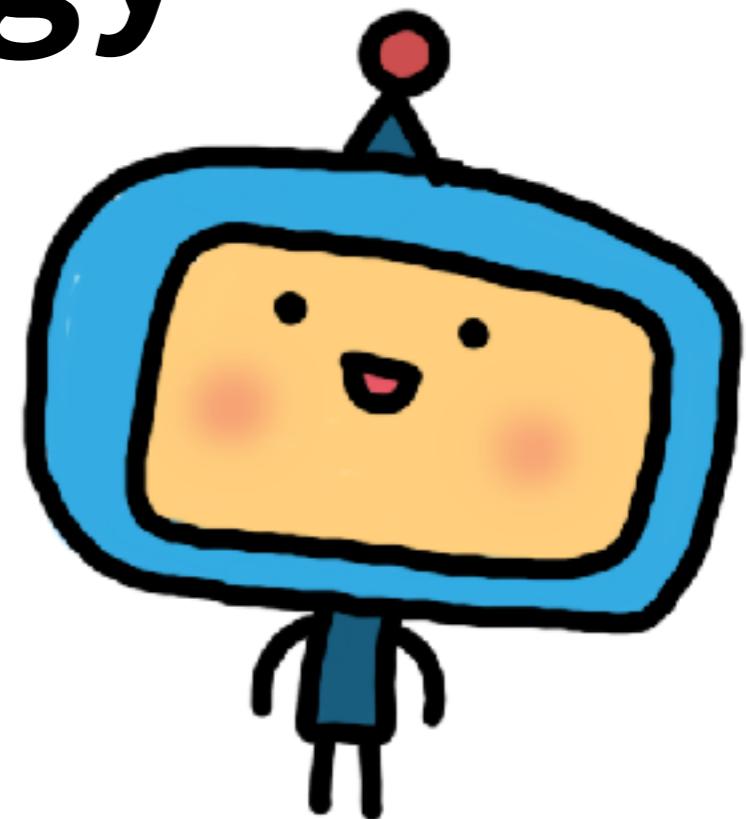
Font B



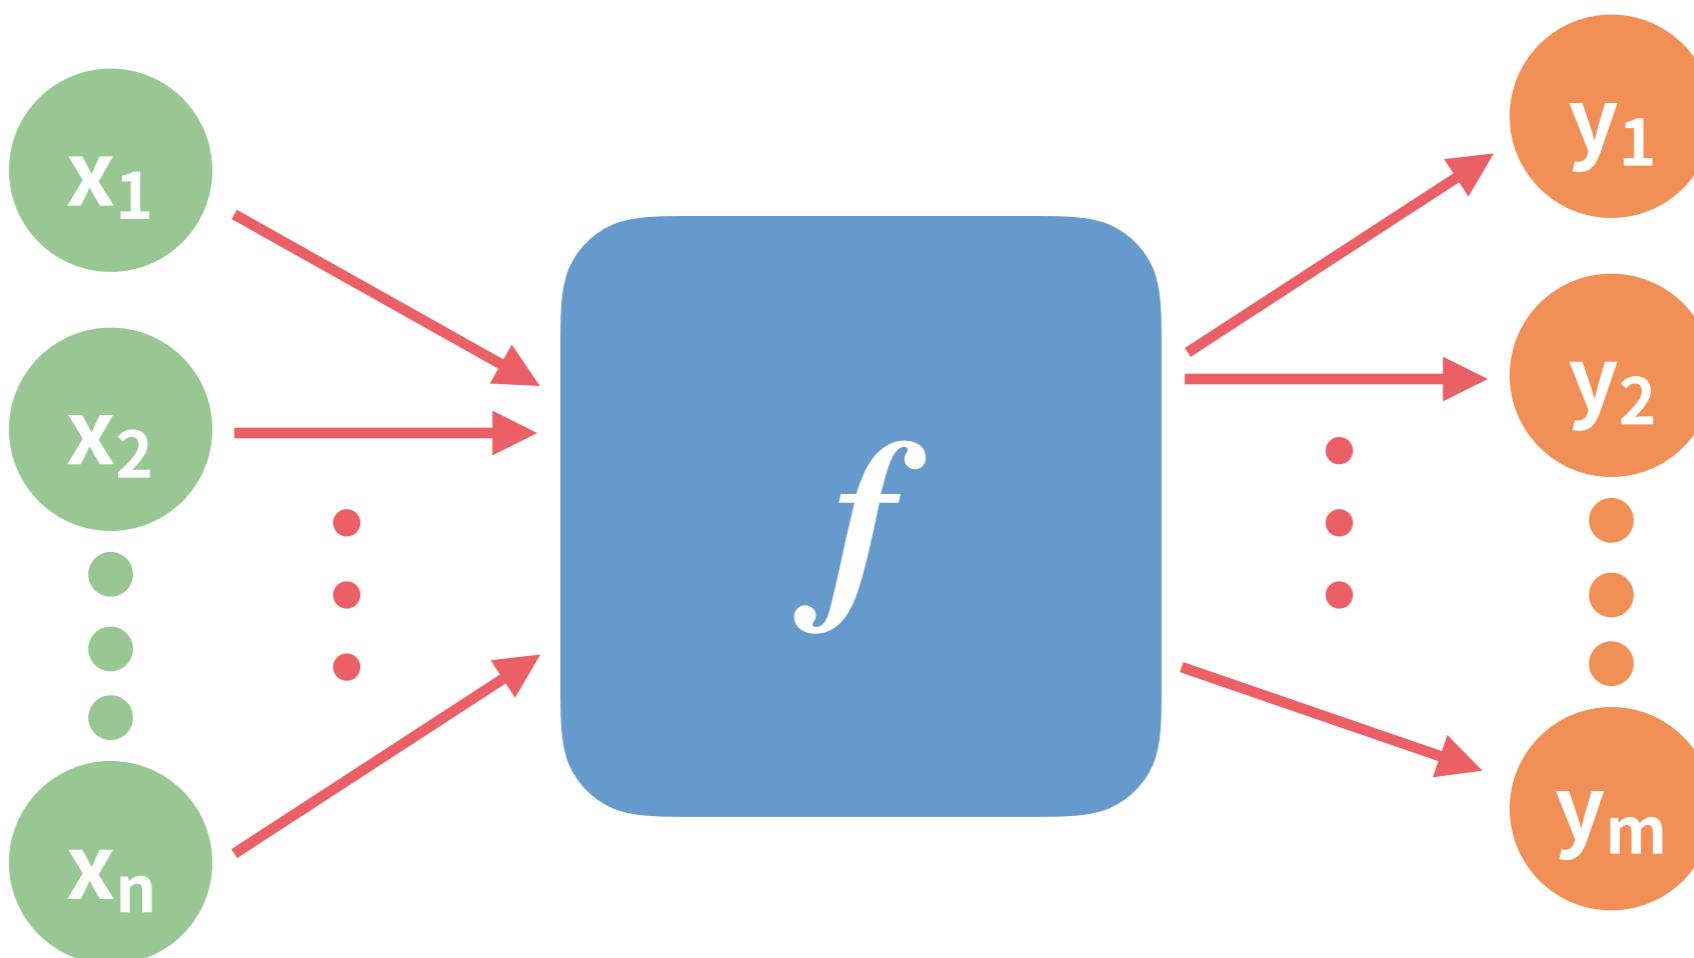
CNN, VAE, GAN

2

AI Core Technology - Principles of Neural Networks



Remember we just have to learn a function



Three Steps of Learning Functions

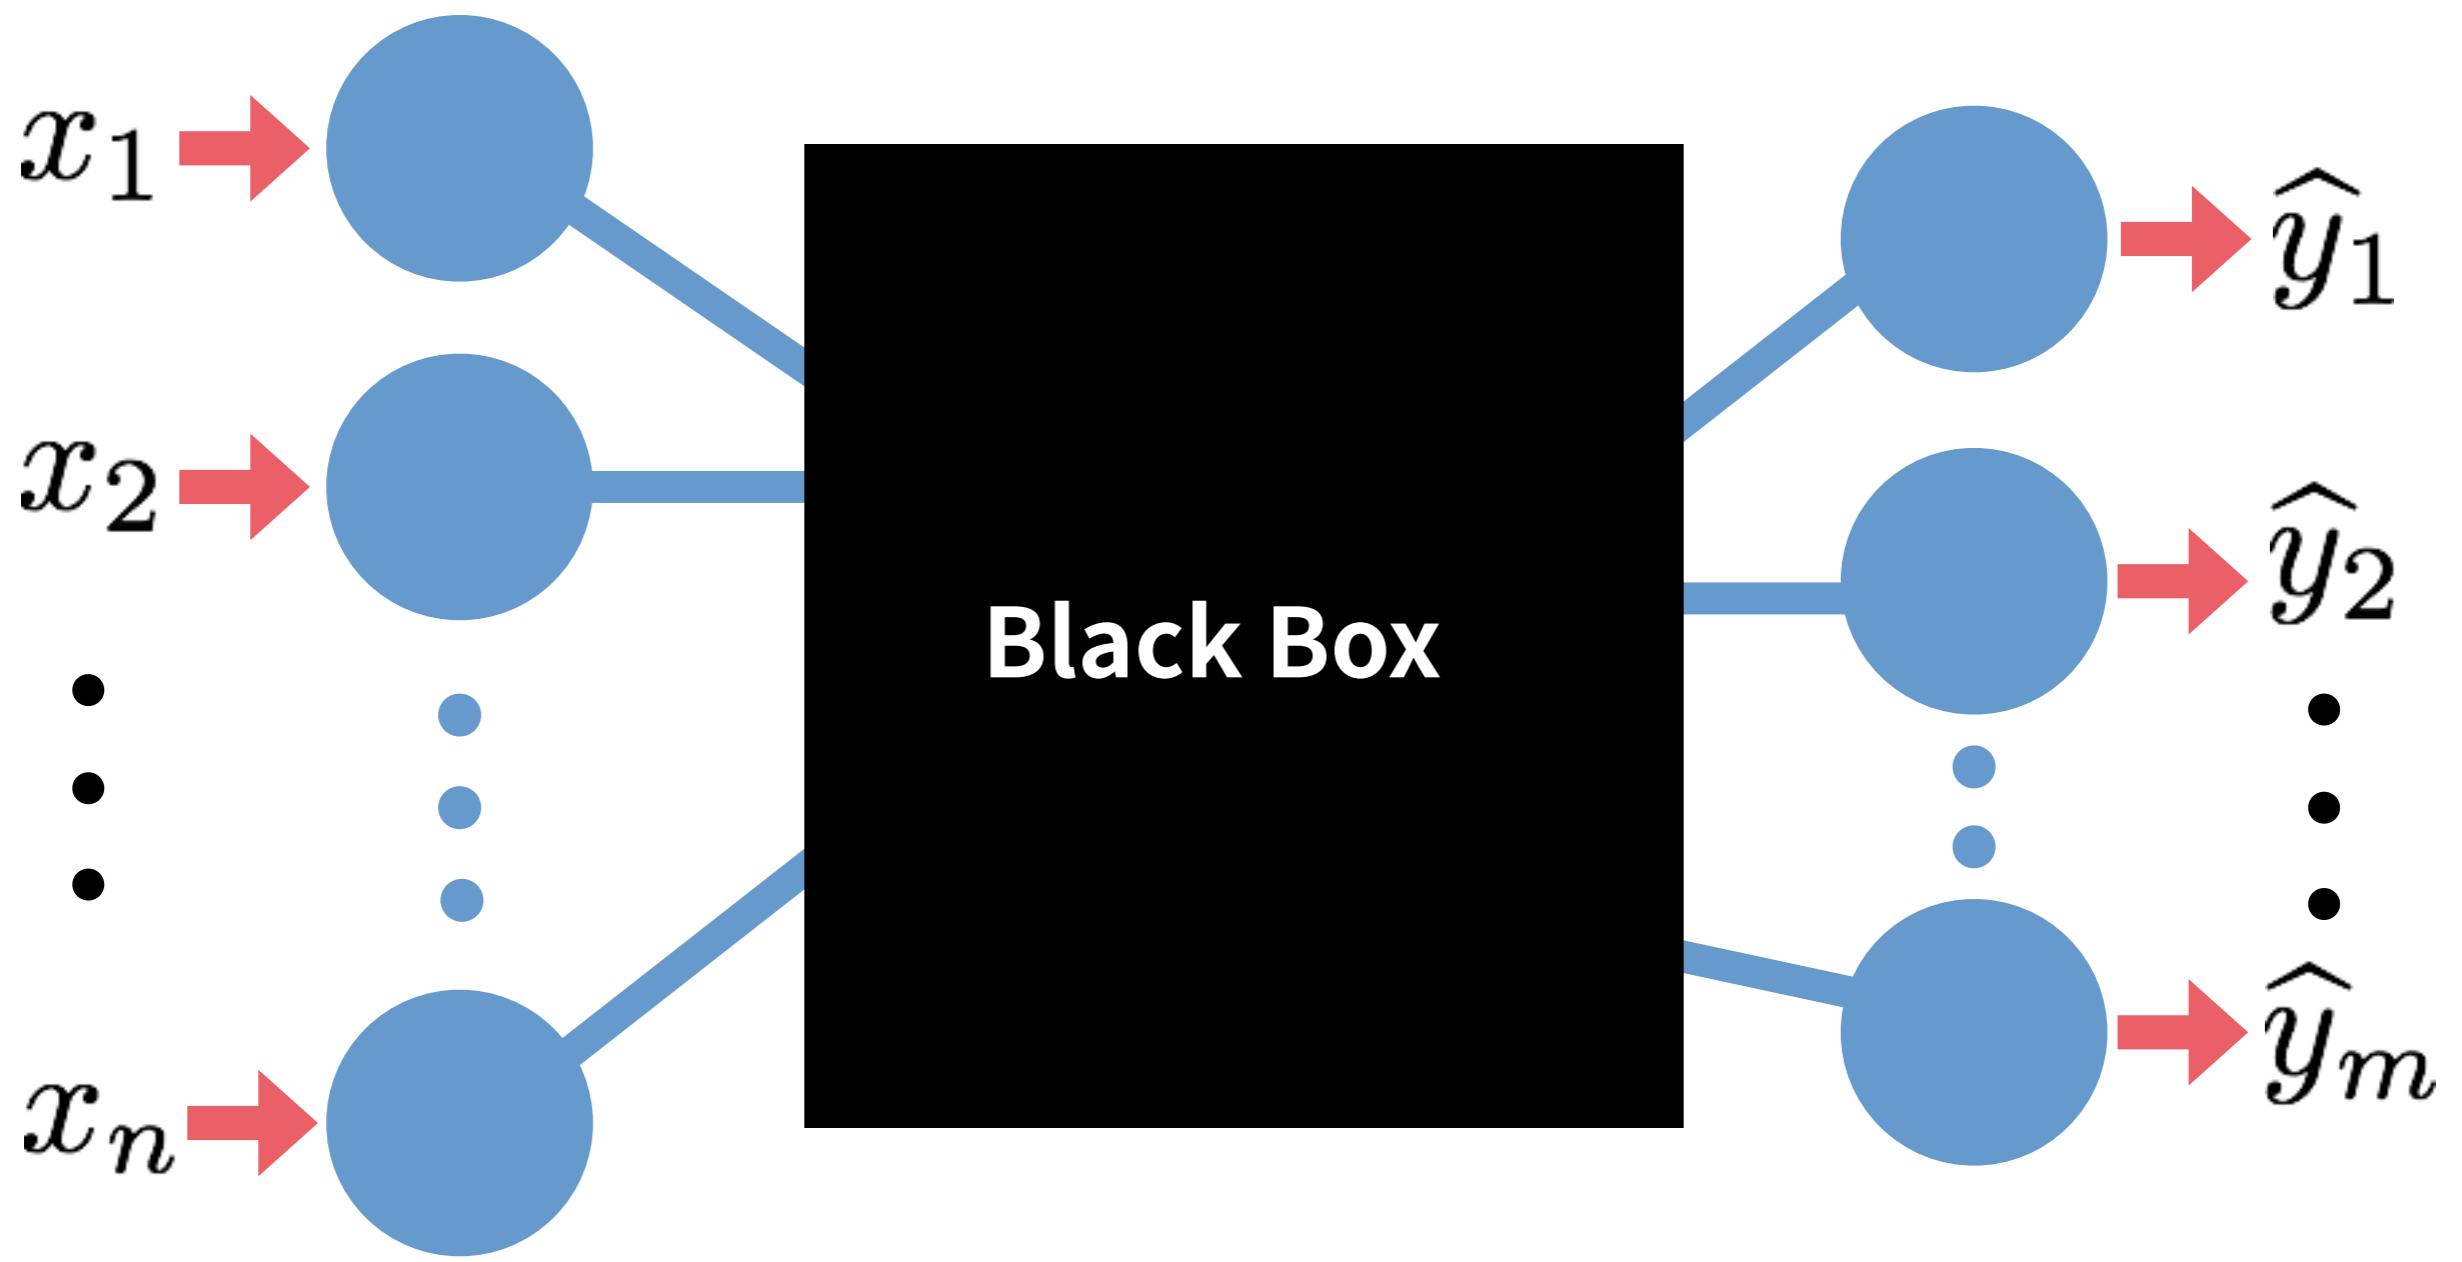
- Transfer a real world problem into a function.
- Collect training materials that we know "correct answers."
- Find the function!

Black Box Learning

There are really techniques for learning
arbitrary functions

neural networks

**In the 1980-1990 or so, it
is a pretty fancy stuff.**



Input
Layer

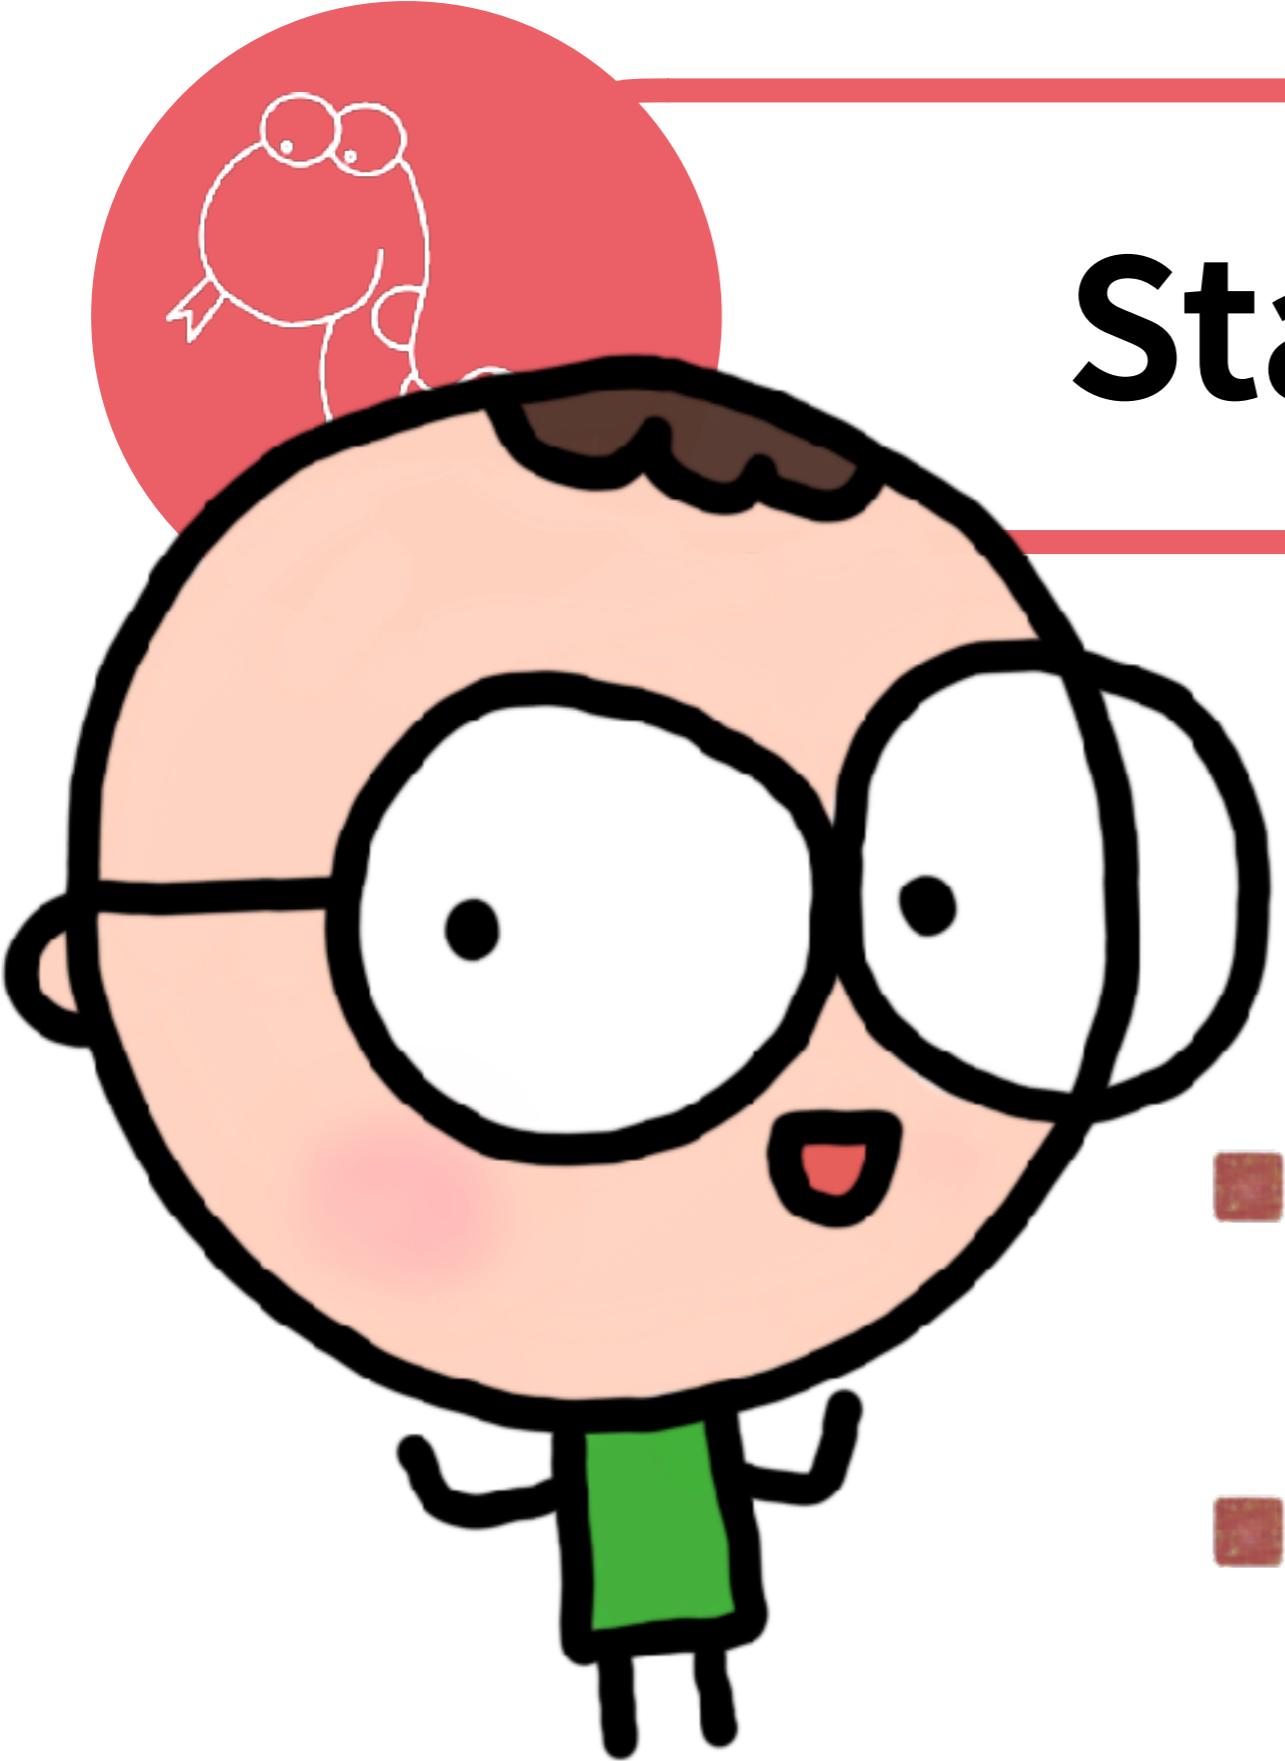
Hidden
Layer

Output
Layer

What is powerful is that neural networks will learn everything!

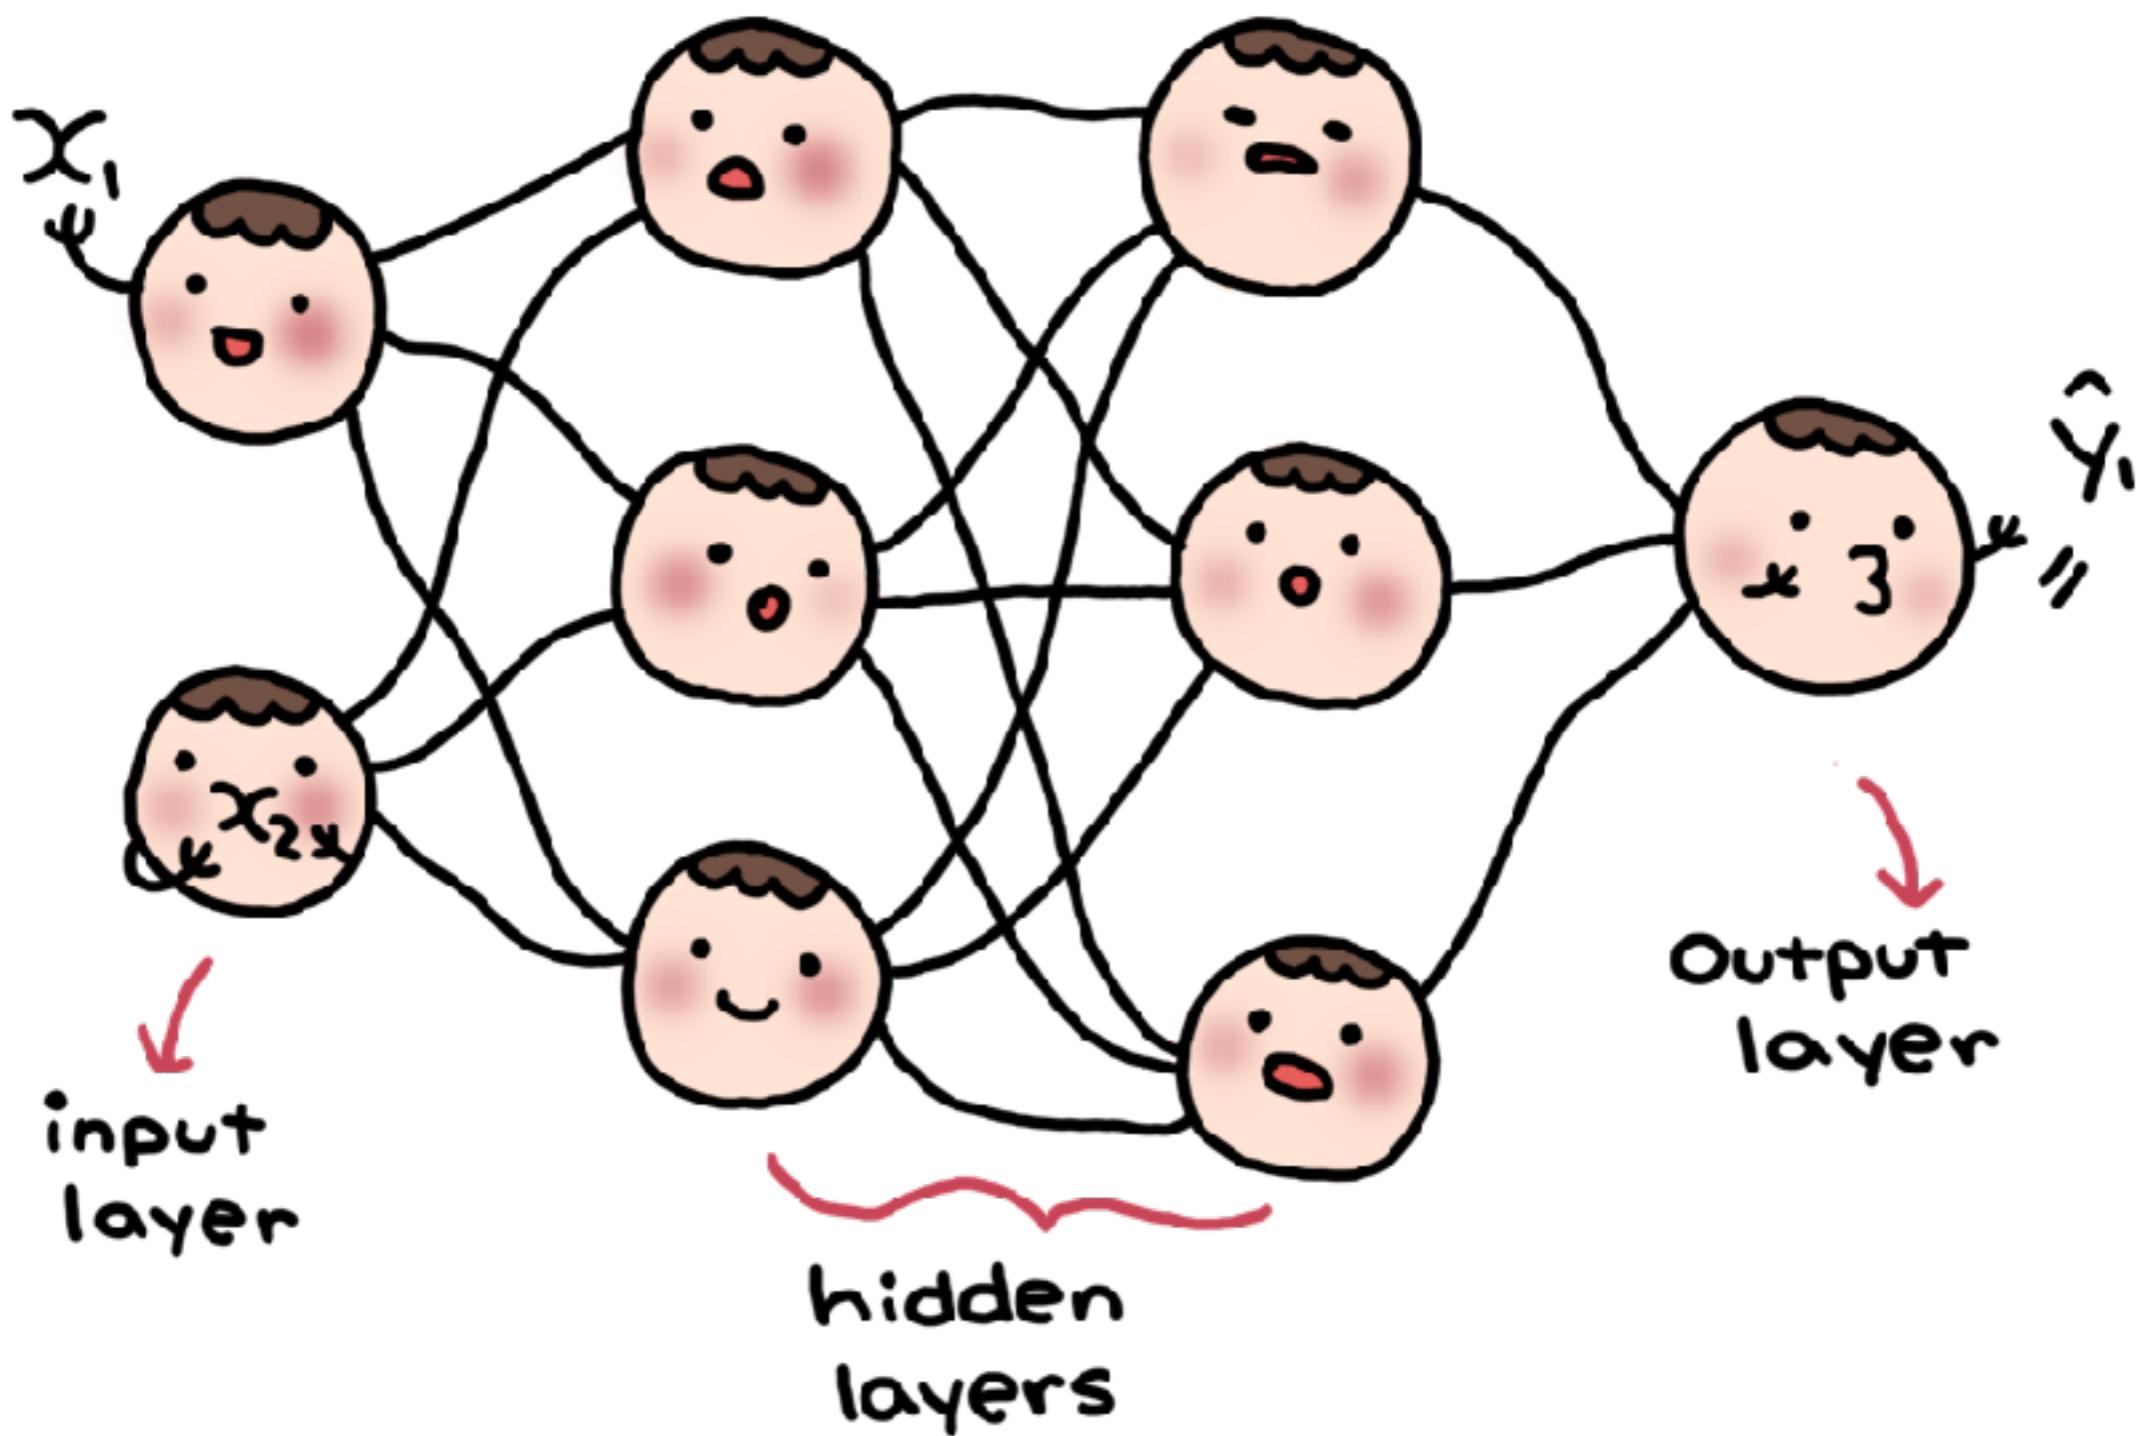
And you don't have to tell it what the function should look like: linear, quadratic polynomial, and so on.

**Open the
Black Box!**

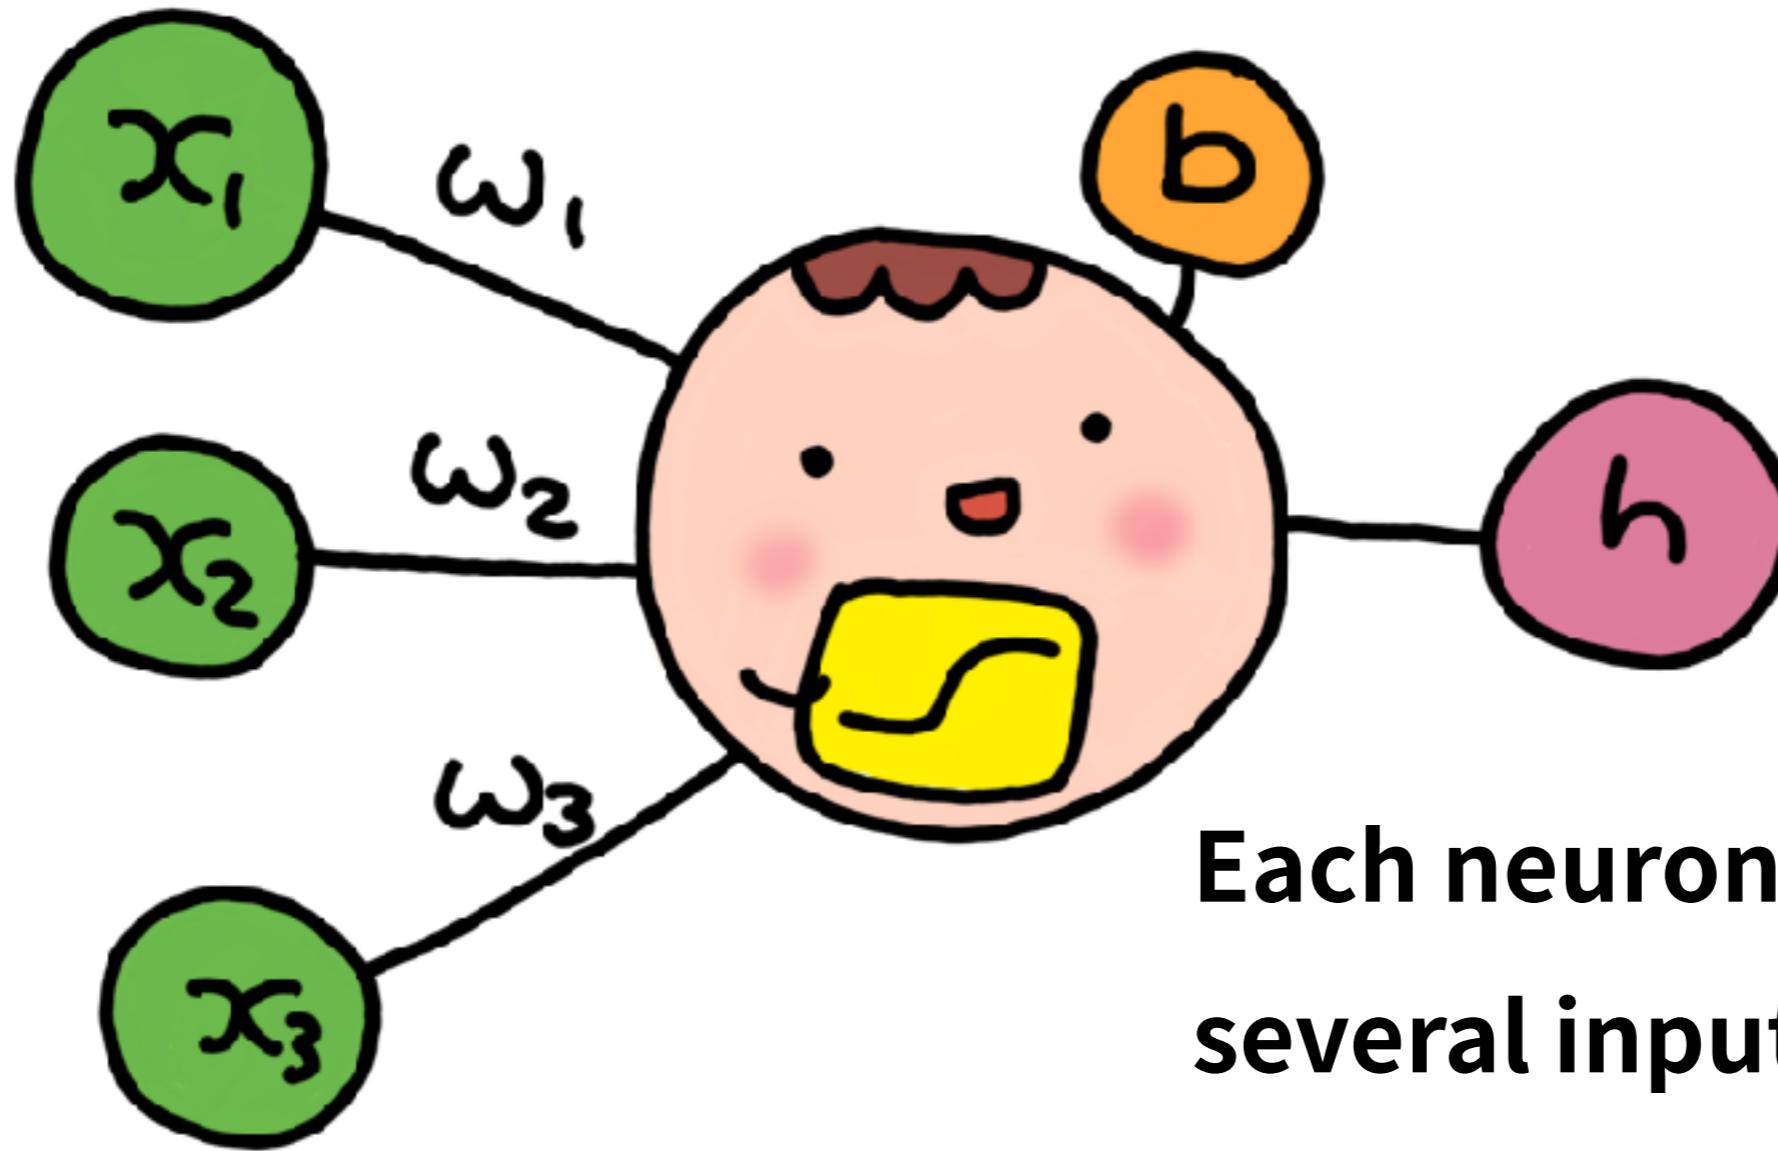


Standard NN

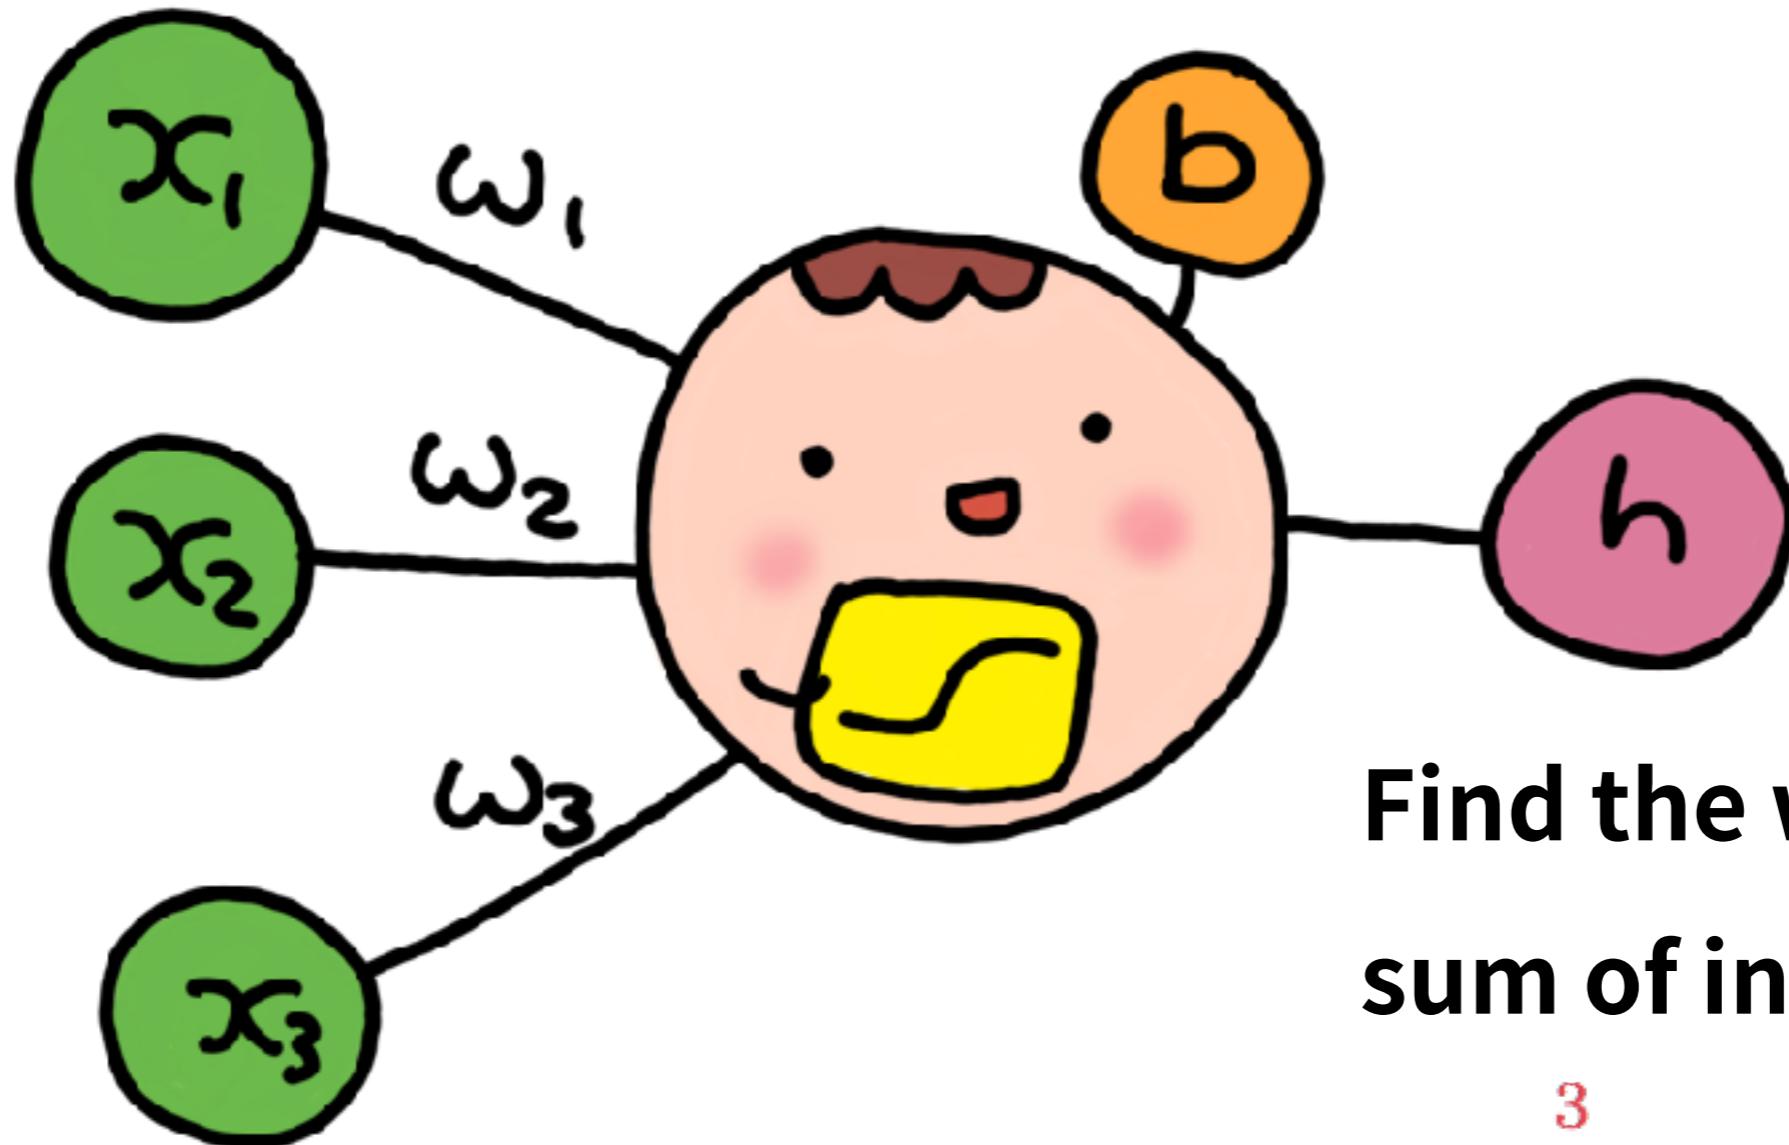
- Fully Connected Neural Networks
- Very popular since 1980s



**The action of every neuron
is basically the same!**

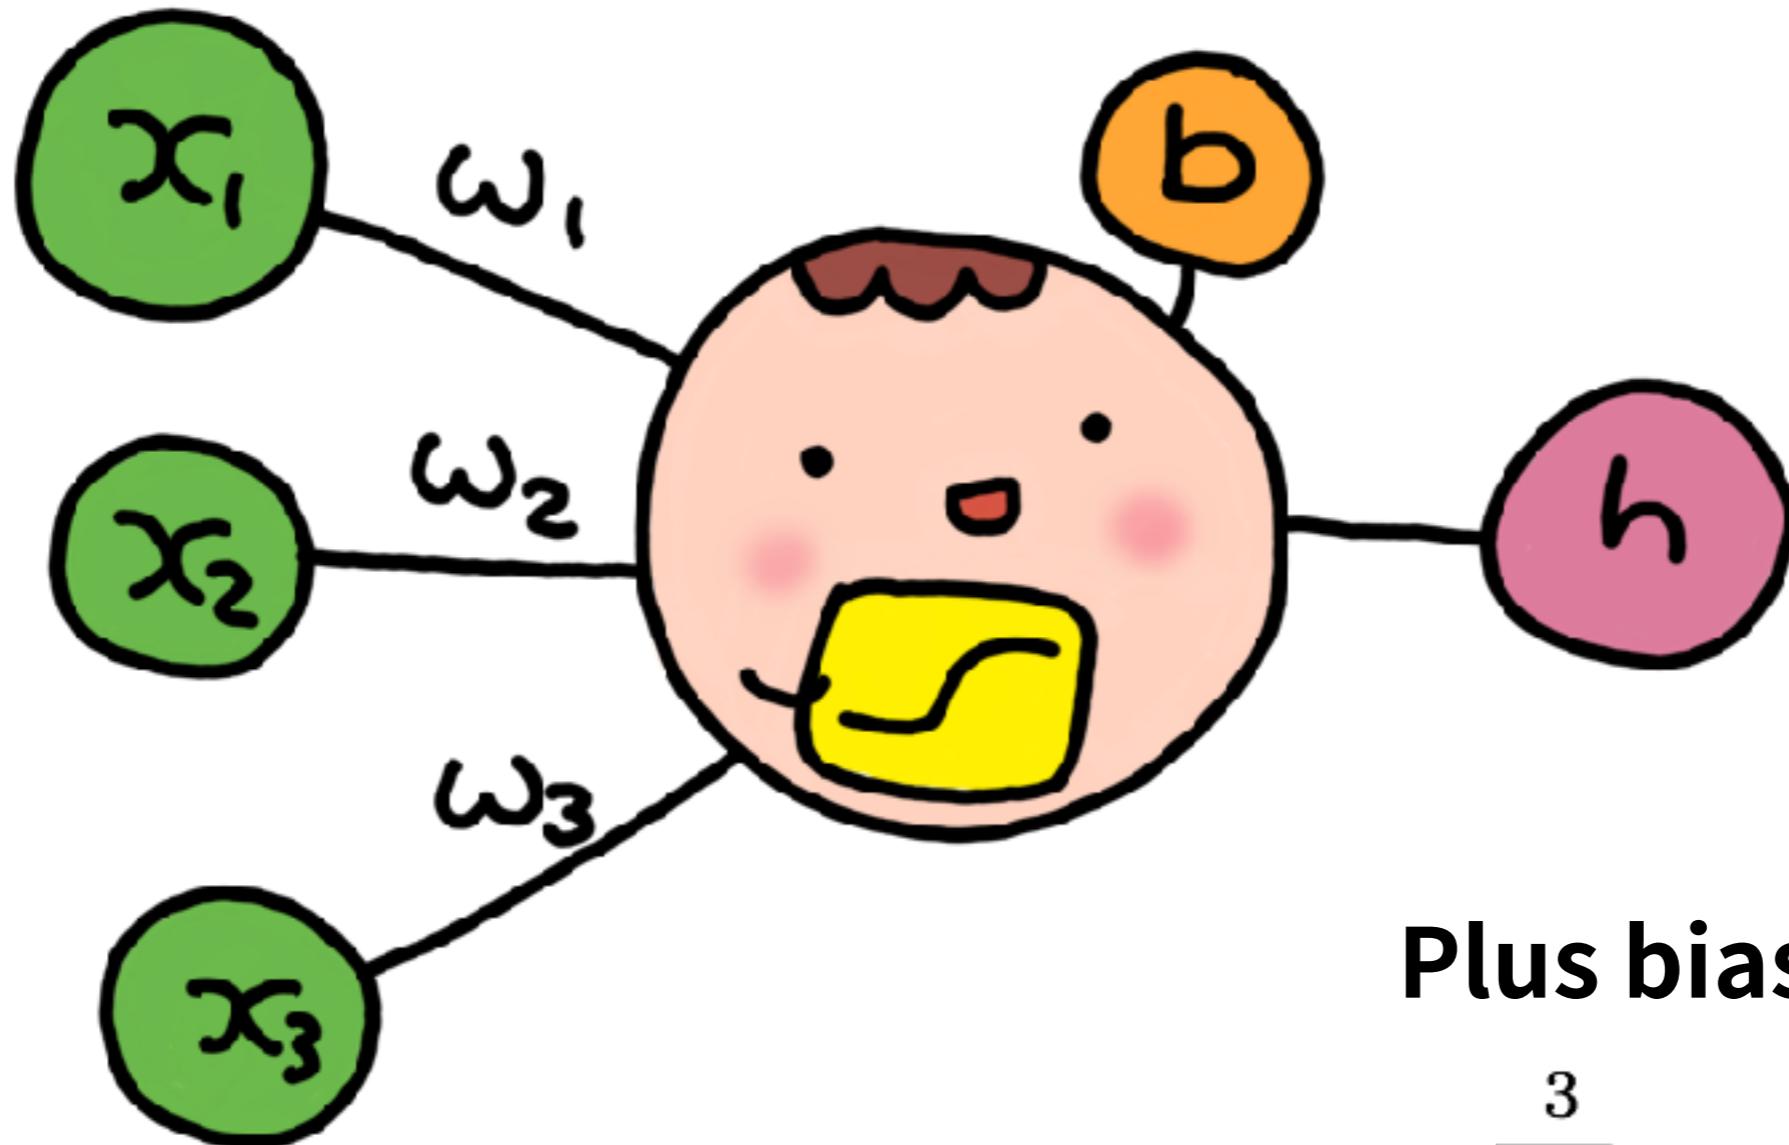


**Each neuron accepts
several inputs and then
sends one output.**



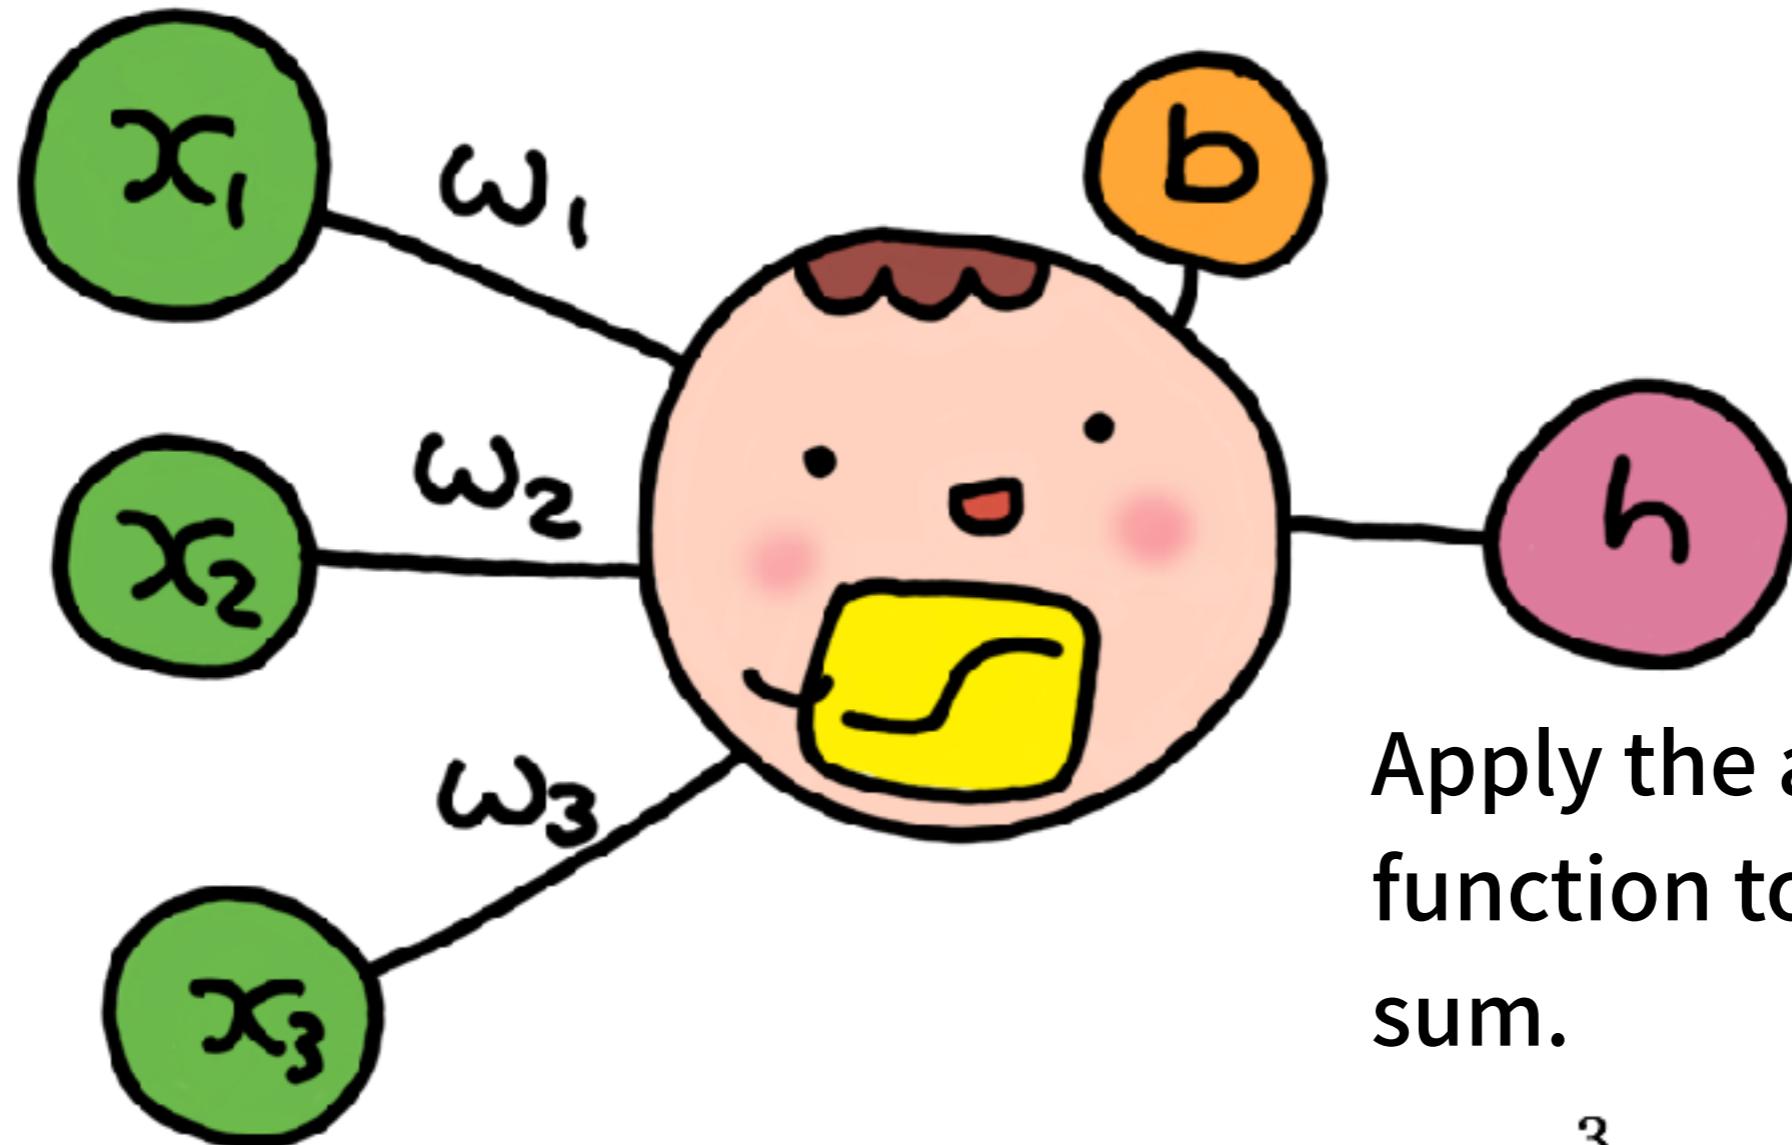
Find the weighted
sum of inputs.

$$\sum_{i=1}^3 w_i x_i$$



Plus bias.

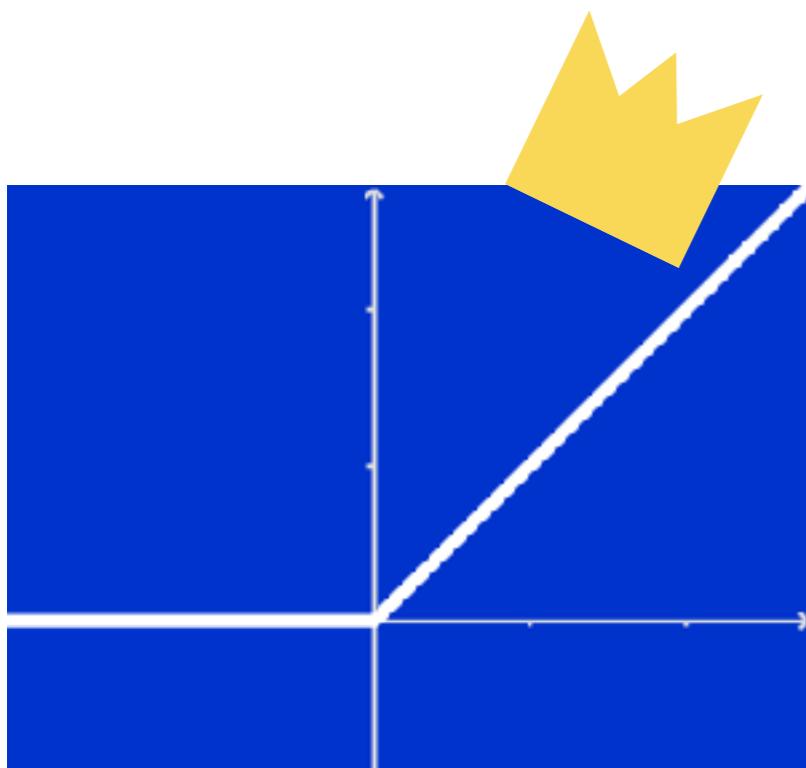
$$\sum_{i=1}^3 w_i x_i + b$$



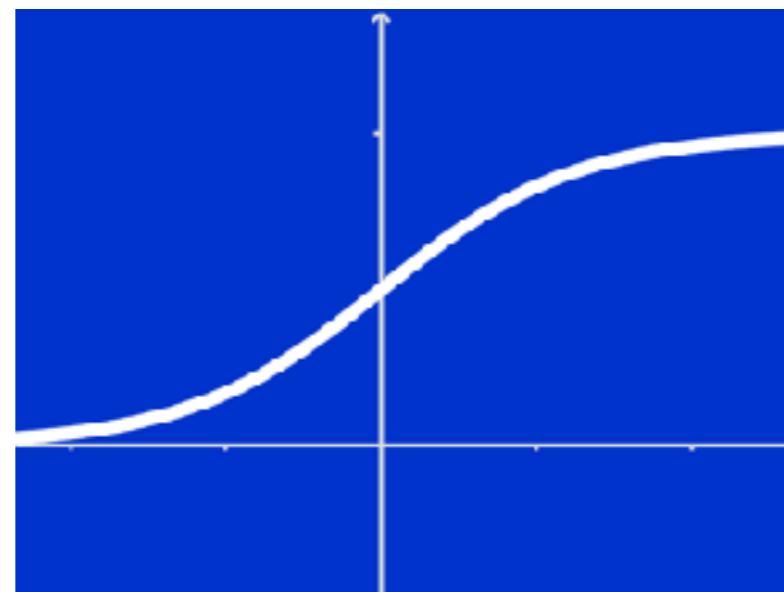
Apply the activation
function to the weighted
sum.

$$\varphi\left(\sum_{i=1}^3 w_i x_i + b\right) = h$$

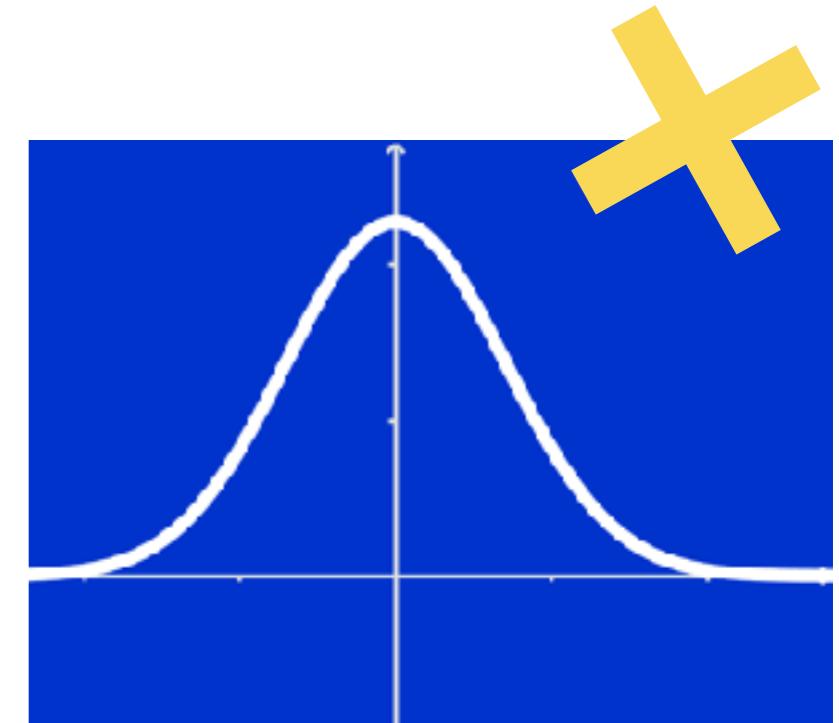
Popular activation functions



ReLU

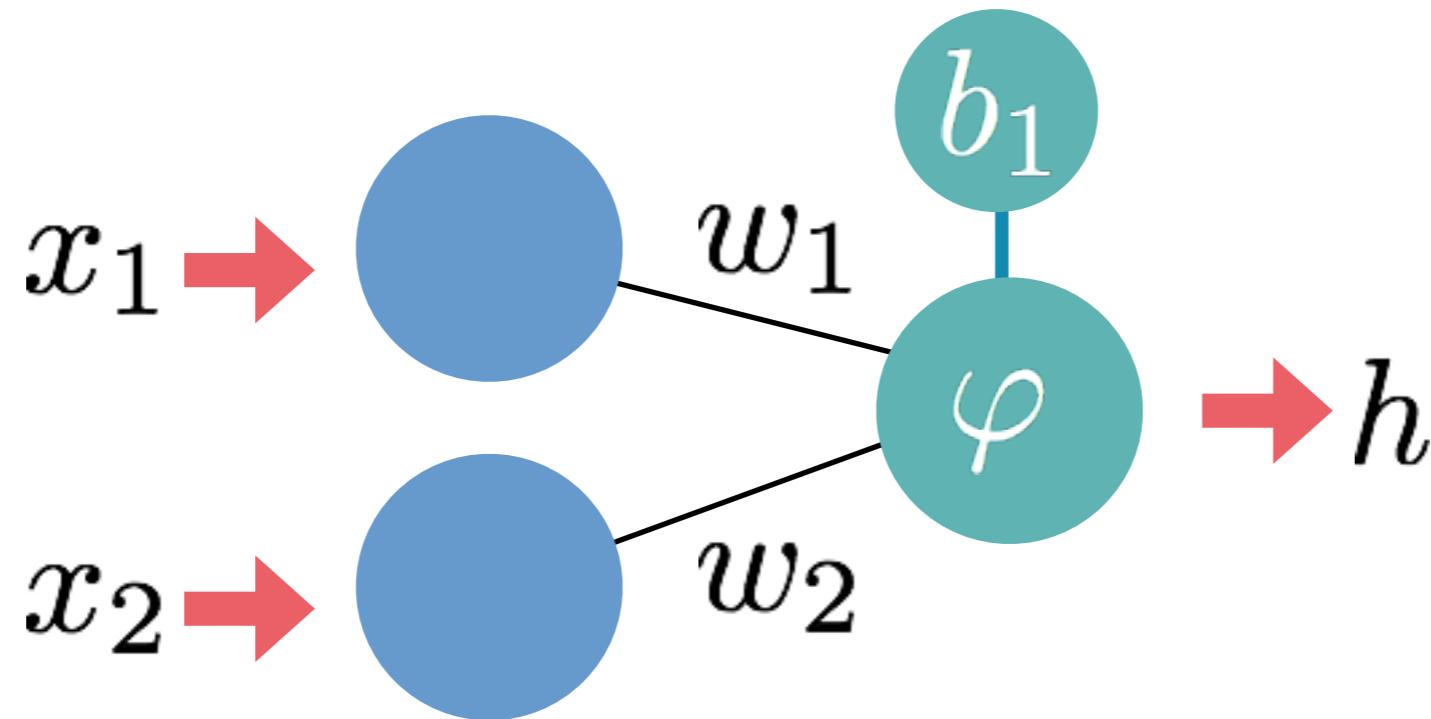


Sigmoid



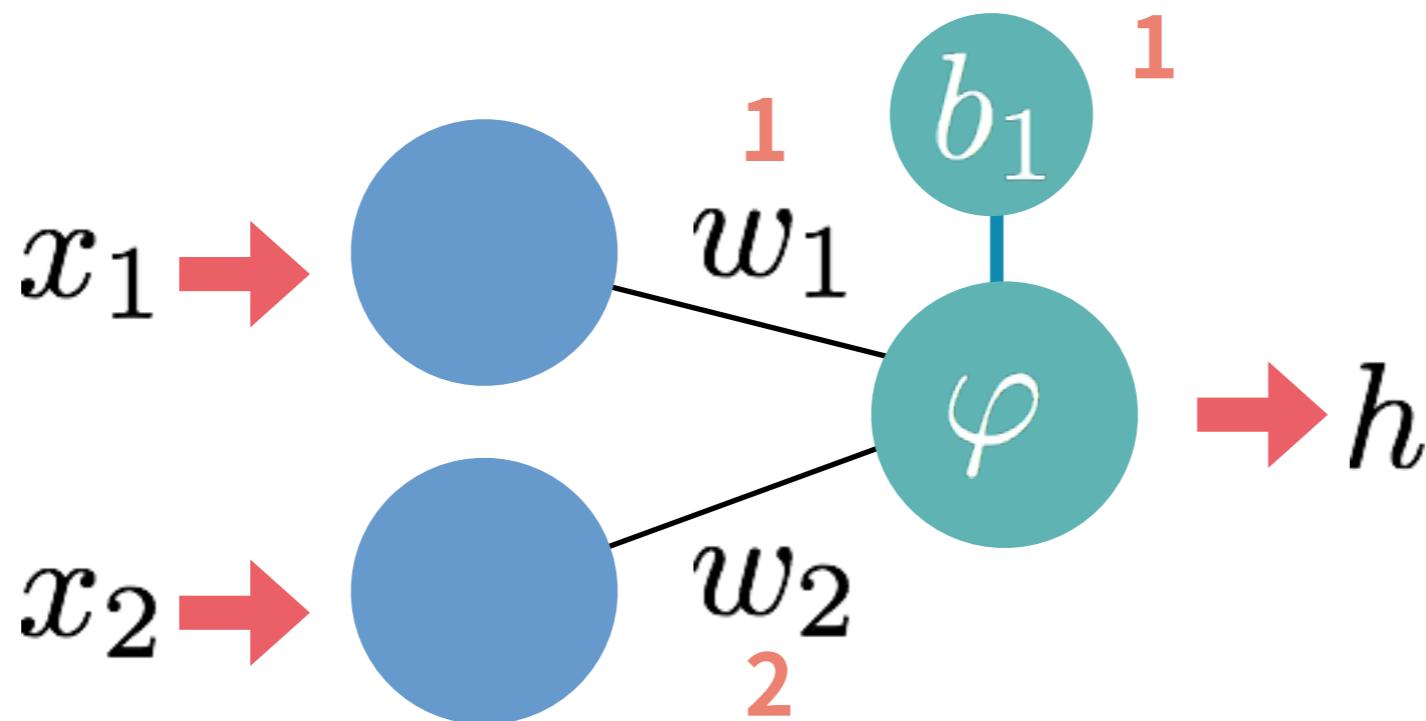
Gaussian

Parameters are weights, biases



$$\varphi(w_1x_1 + w_2x_2 + b_1) = h$$

“Learned” neural network

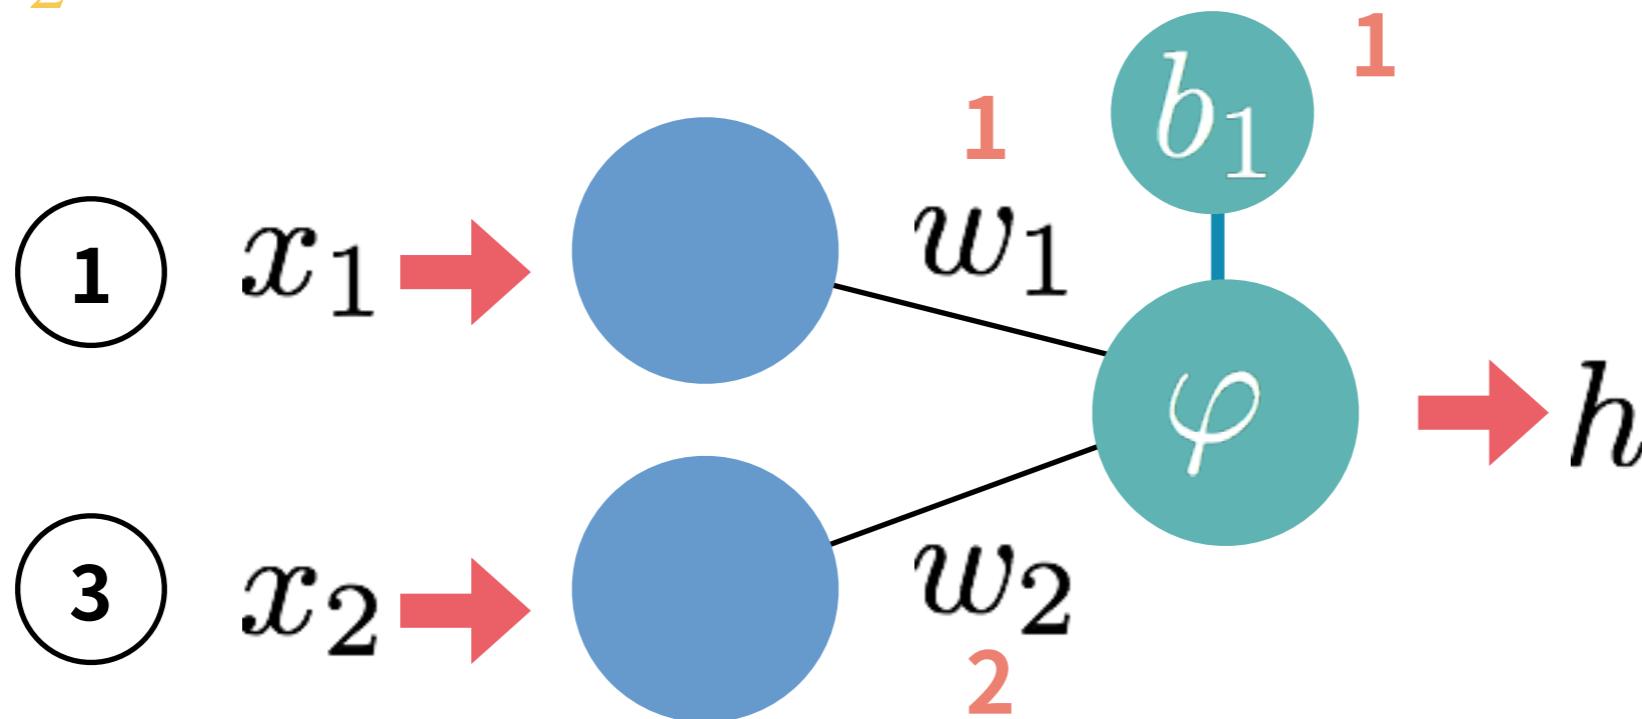


$$\varphi(w_1 x_1 + w_2 x_2 + b_1) = h$$

1 2 1

Suppose we have the input

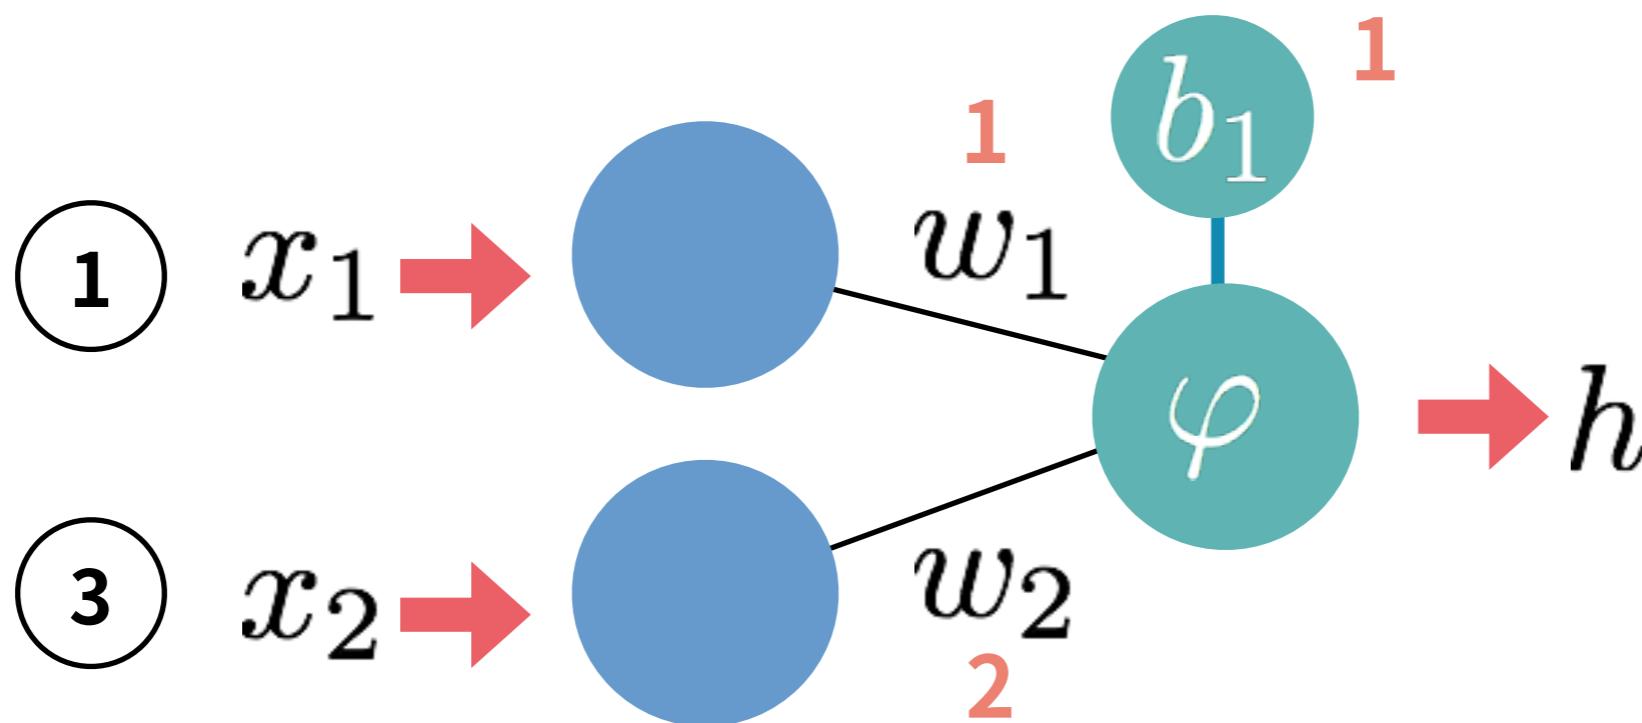
$$(x_1, x_2) = (1, 3)$$



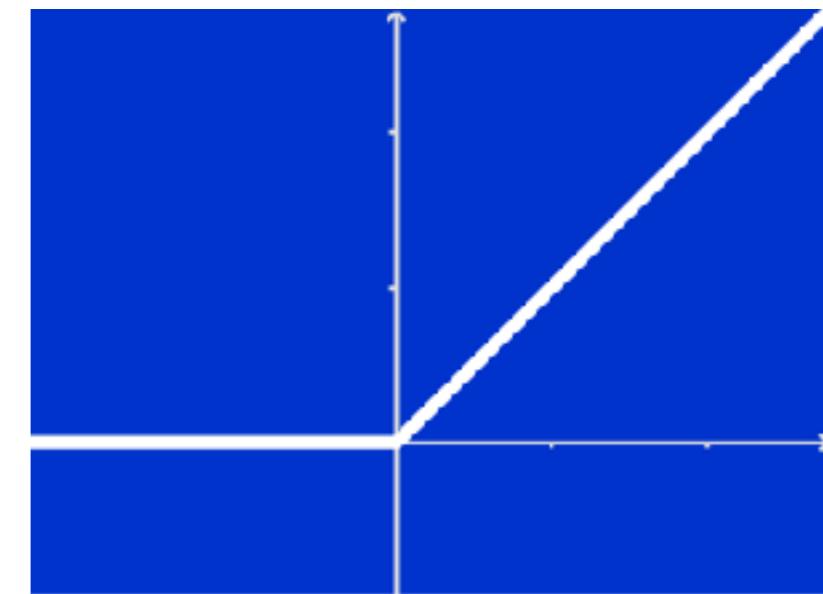
$$\varphi(w_1 x_1 + w_2 x_2 + b_1) = h$$

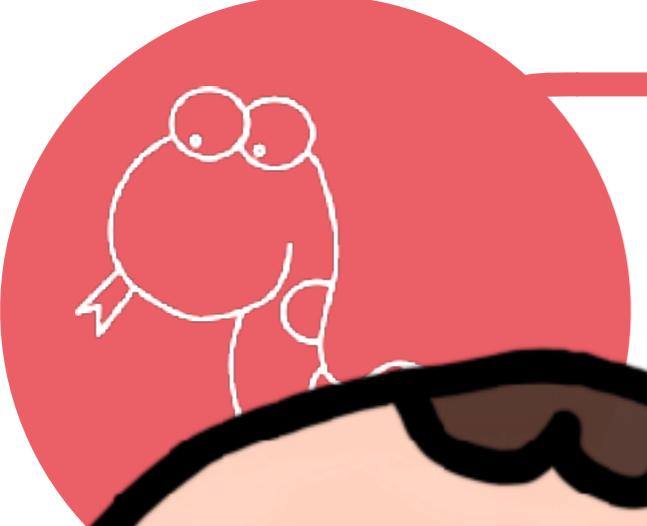
1 1 2 3 1

Using ReLU as activation functions

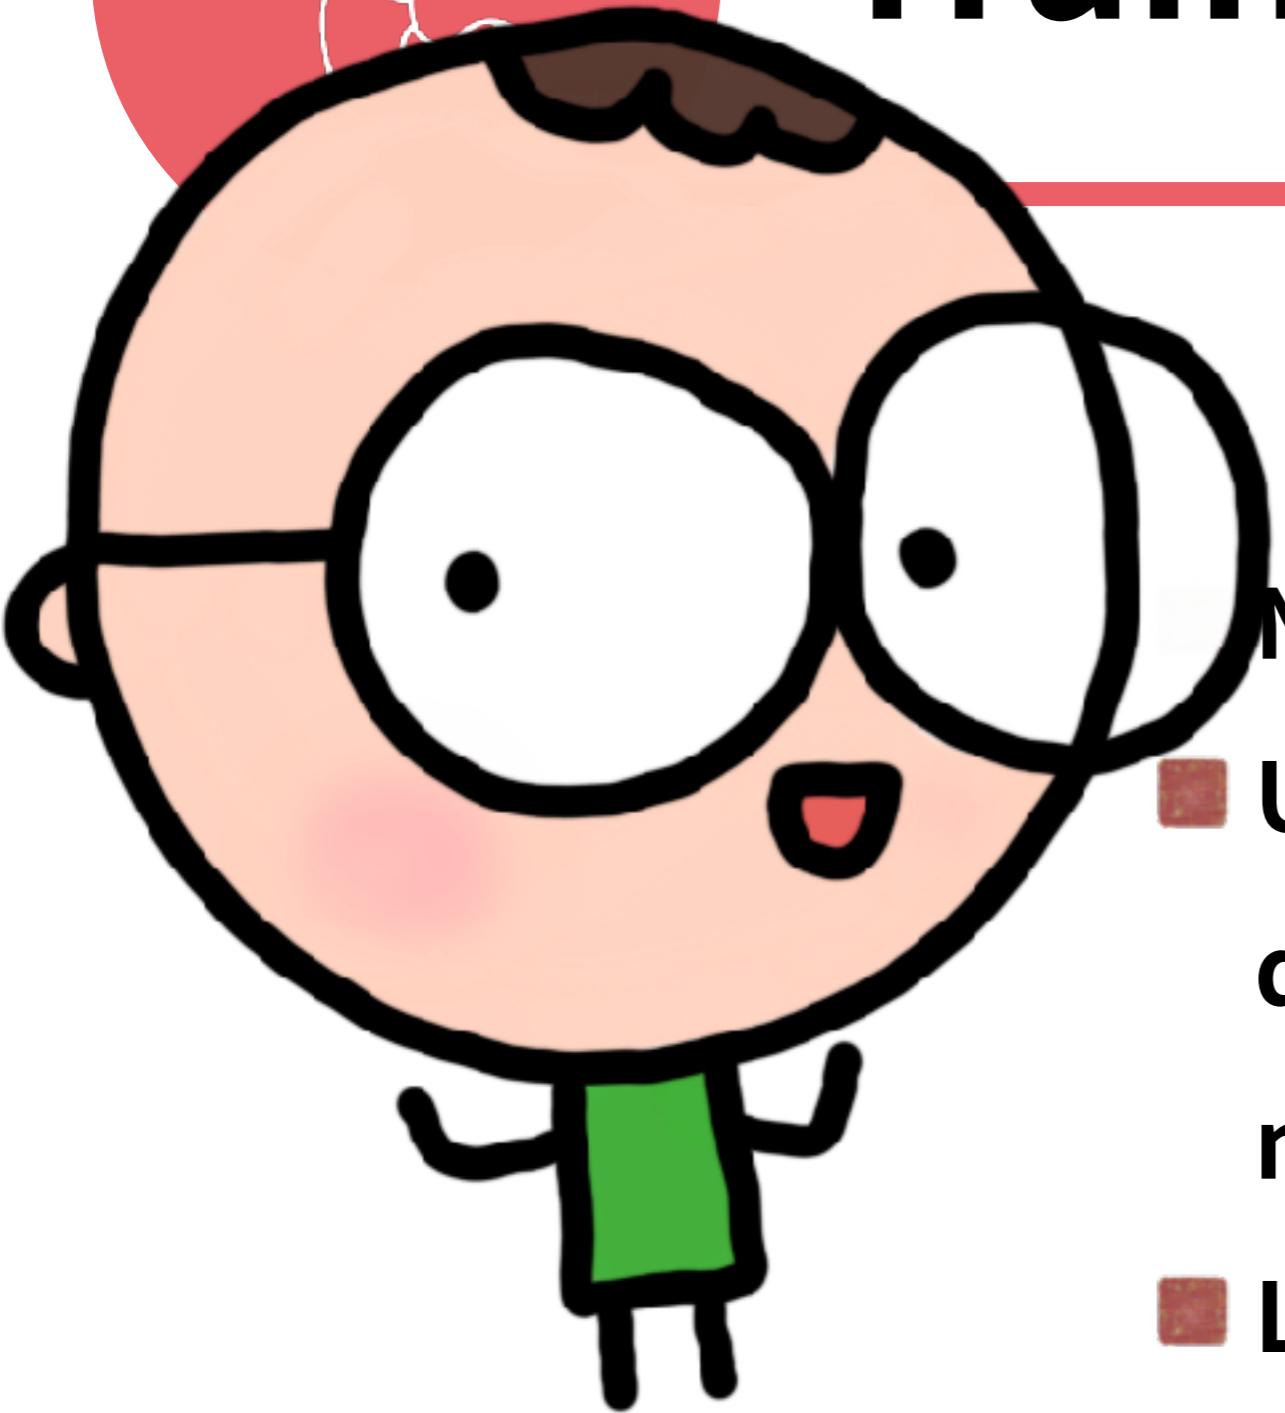


$$\varphi(8) = 8$$





Training Methods



- NN needs to be trained
 - Using “previous exam questions” to train our neural network
 - Learning method is called backpropagation

The Function Space of our NN

When a neural network structure is determined and activation functions are also determined, it can be adjusted by weights, biases. We call the set of these parameters θ , each of which defines a function, and we treat it as a set.

$$\{F_\theta\}$$

We are looking for θ^*

Makes F_{θ^*} the closest to the objective function

What does “the closest” mean

That means the value of “loss function” is minimum

**Suppose we have the
training data**

$$\{(x_1, y_1), (x_2, y_2), \dots, (x_k, y_k)\}$$

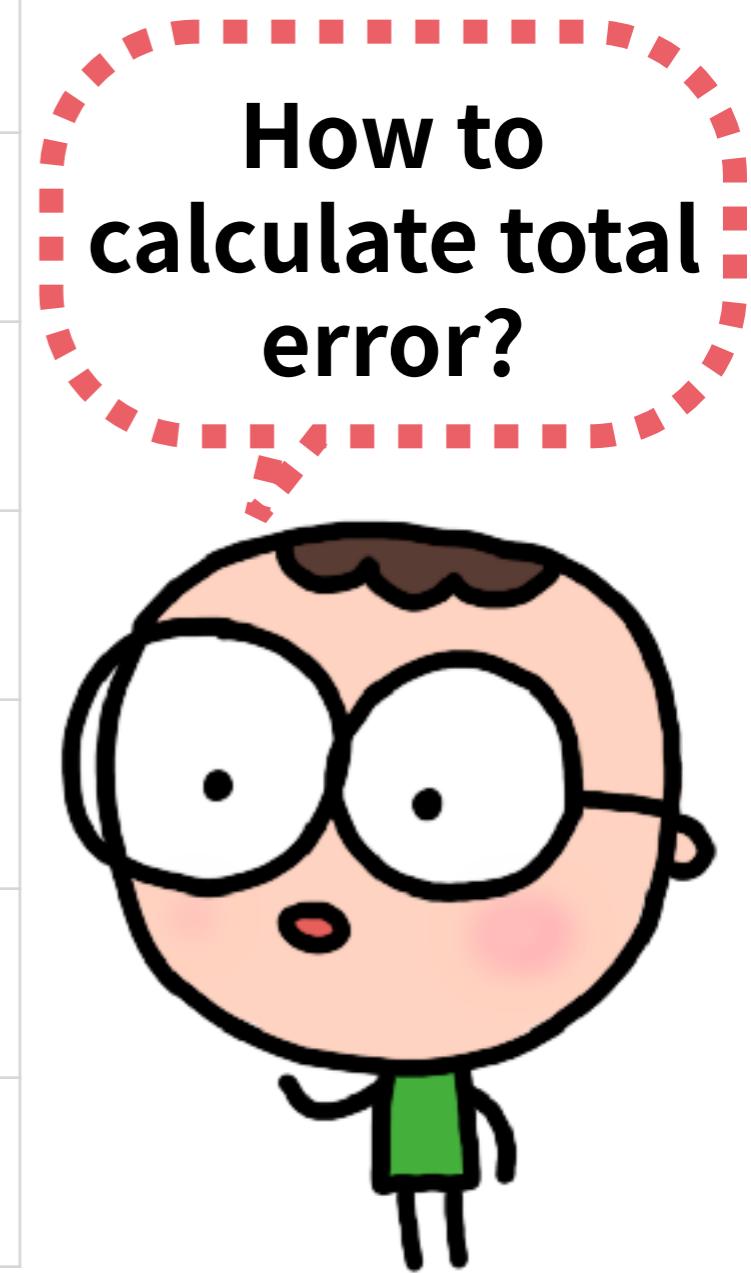
The most popular loss function

We the value is as small as possible

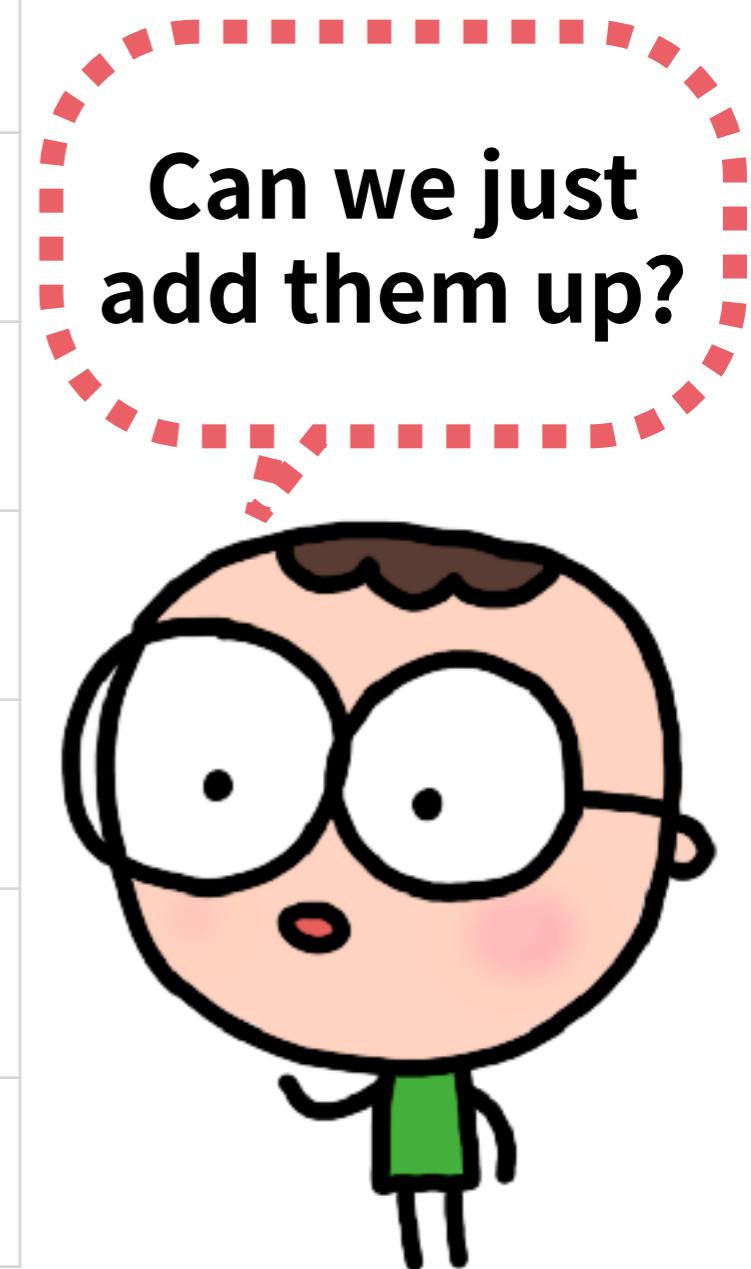
$$L(\theta) = \frac{1}{2} \sum_{i=1}^k \|y_i - F_\theta(x_i)\|^2$$



x	網路輸出	真正的	誤差
1	2	3	1
2	5	2	-3
3	6	8	2
4	7	2	-5
5	6	6	0
6	4	8	4
7	8	9	1

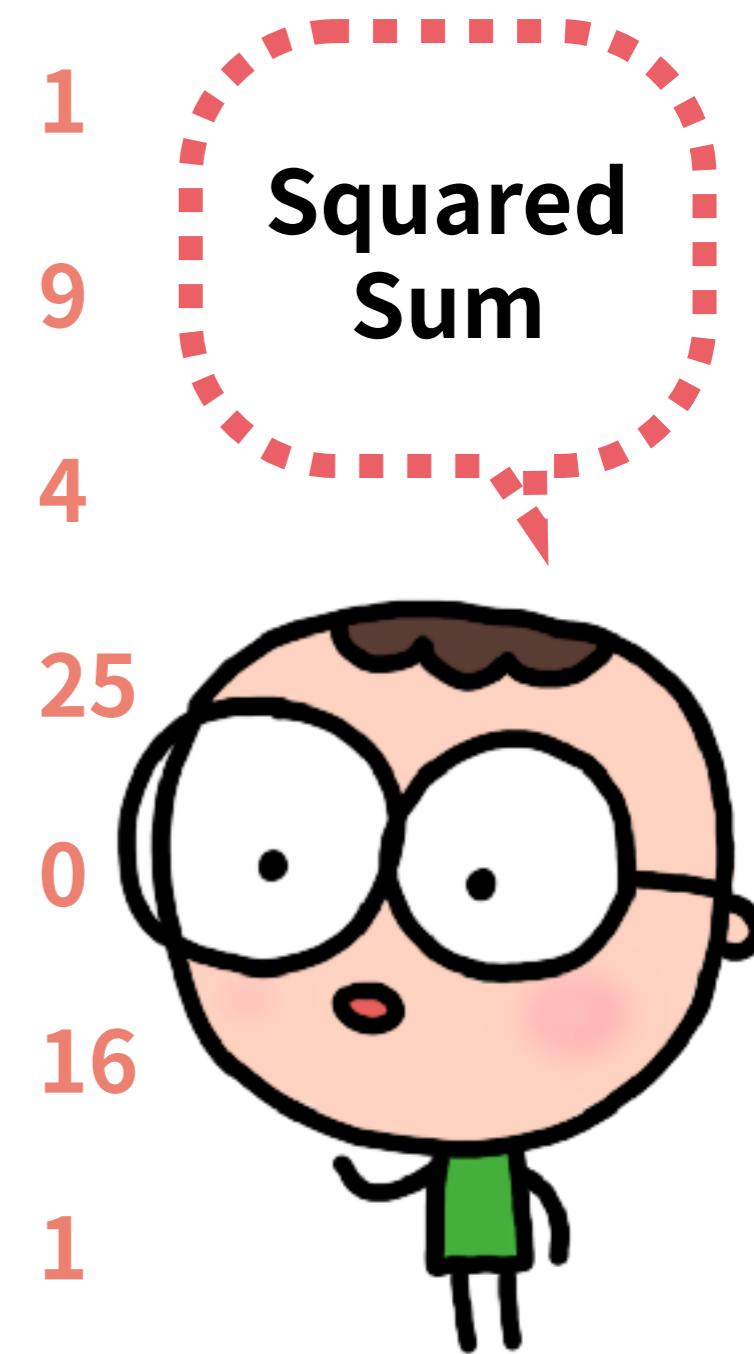


x	網路輸出	真正的	誤差
1	2	3	1
2	5	2	-3
3	6	8	2
4	7	2	-5
5	6	6	0
6	4	8	4
7	8	9	1



Total = 0!! 59

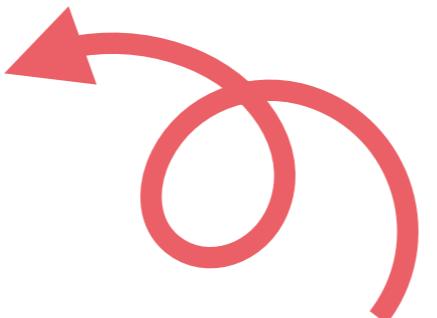
x	網路輸出	真正的	誤差
1	2	3	1
2	5	2	-3
3	6	8	2
4	7	2	-5
5	6	6	0
6	4	8	4
7	8	9	1



56 in total

60

Basically, do such
adjustment

learning rate 

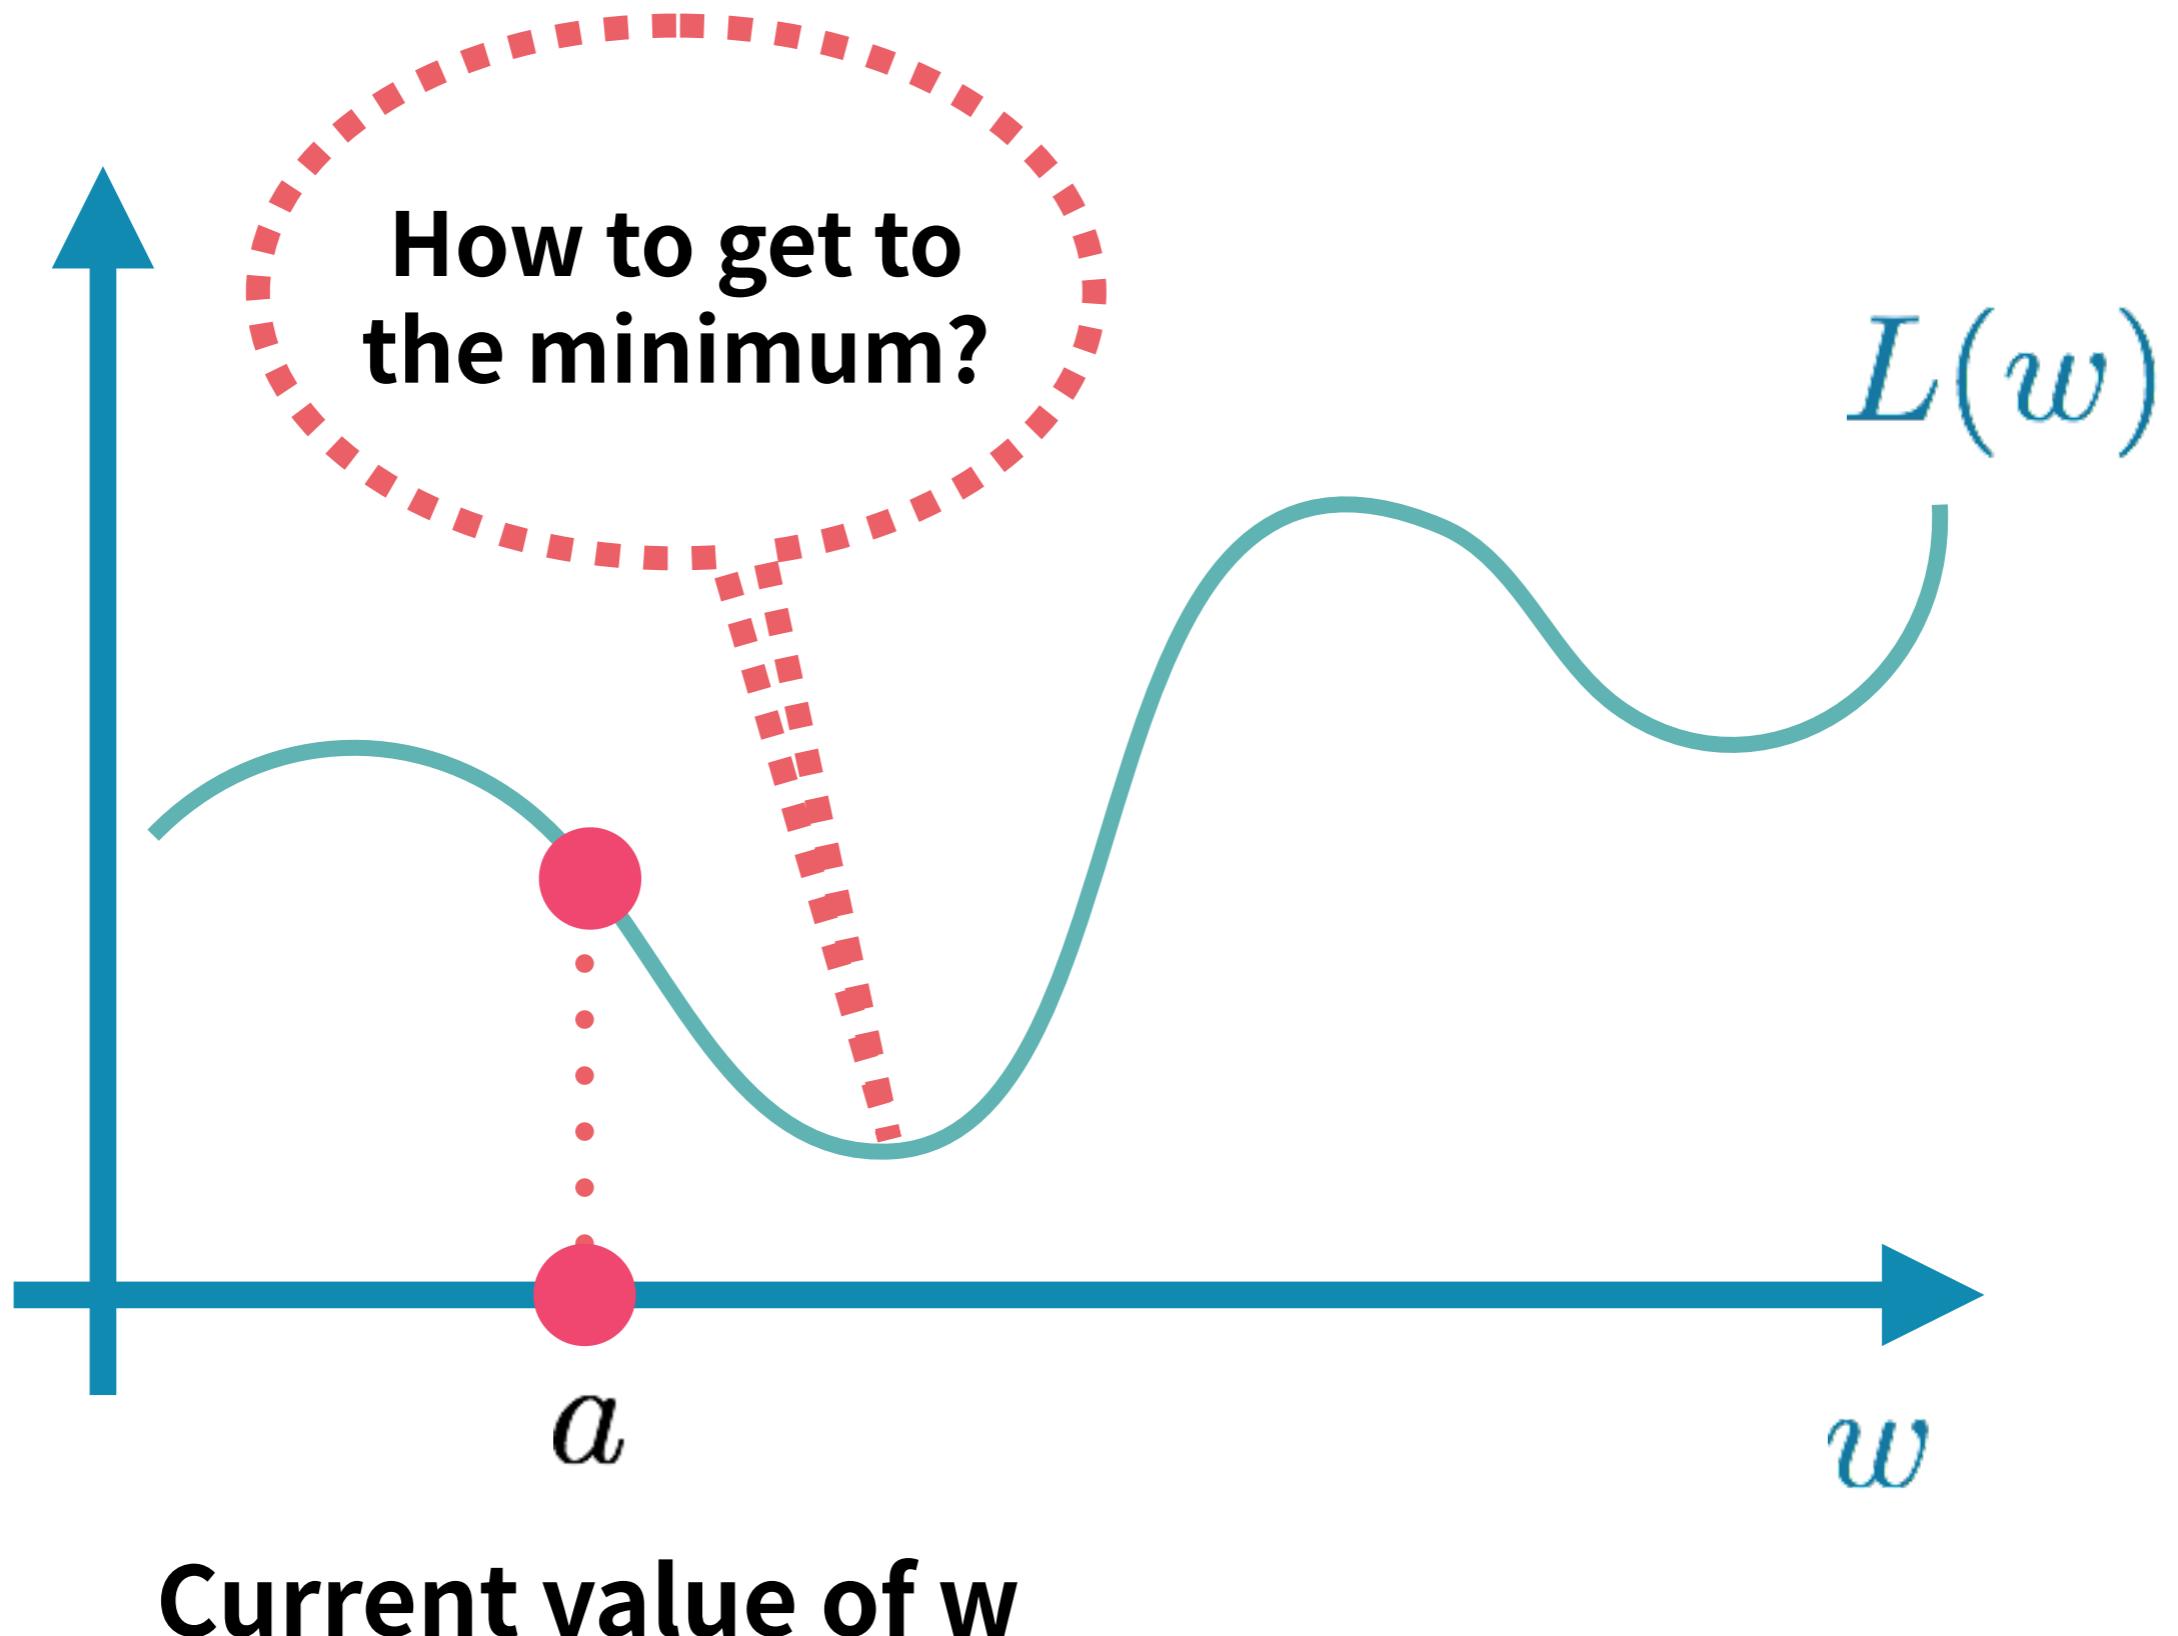
$$-\eta \frac{\partial L}{\partial w_{ij}}$$

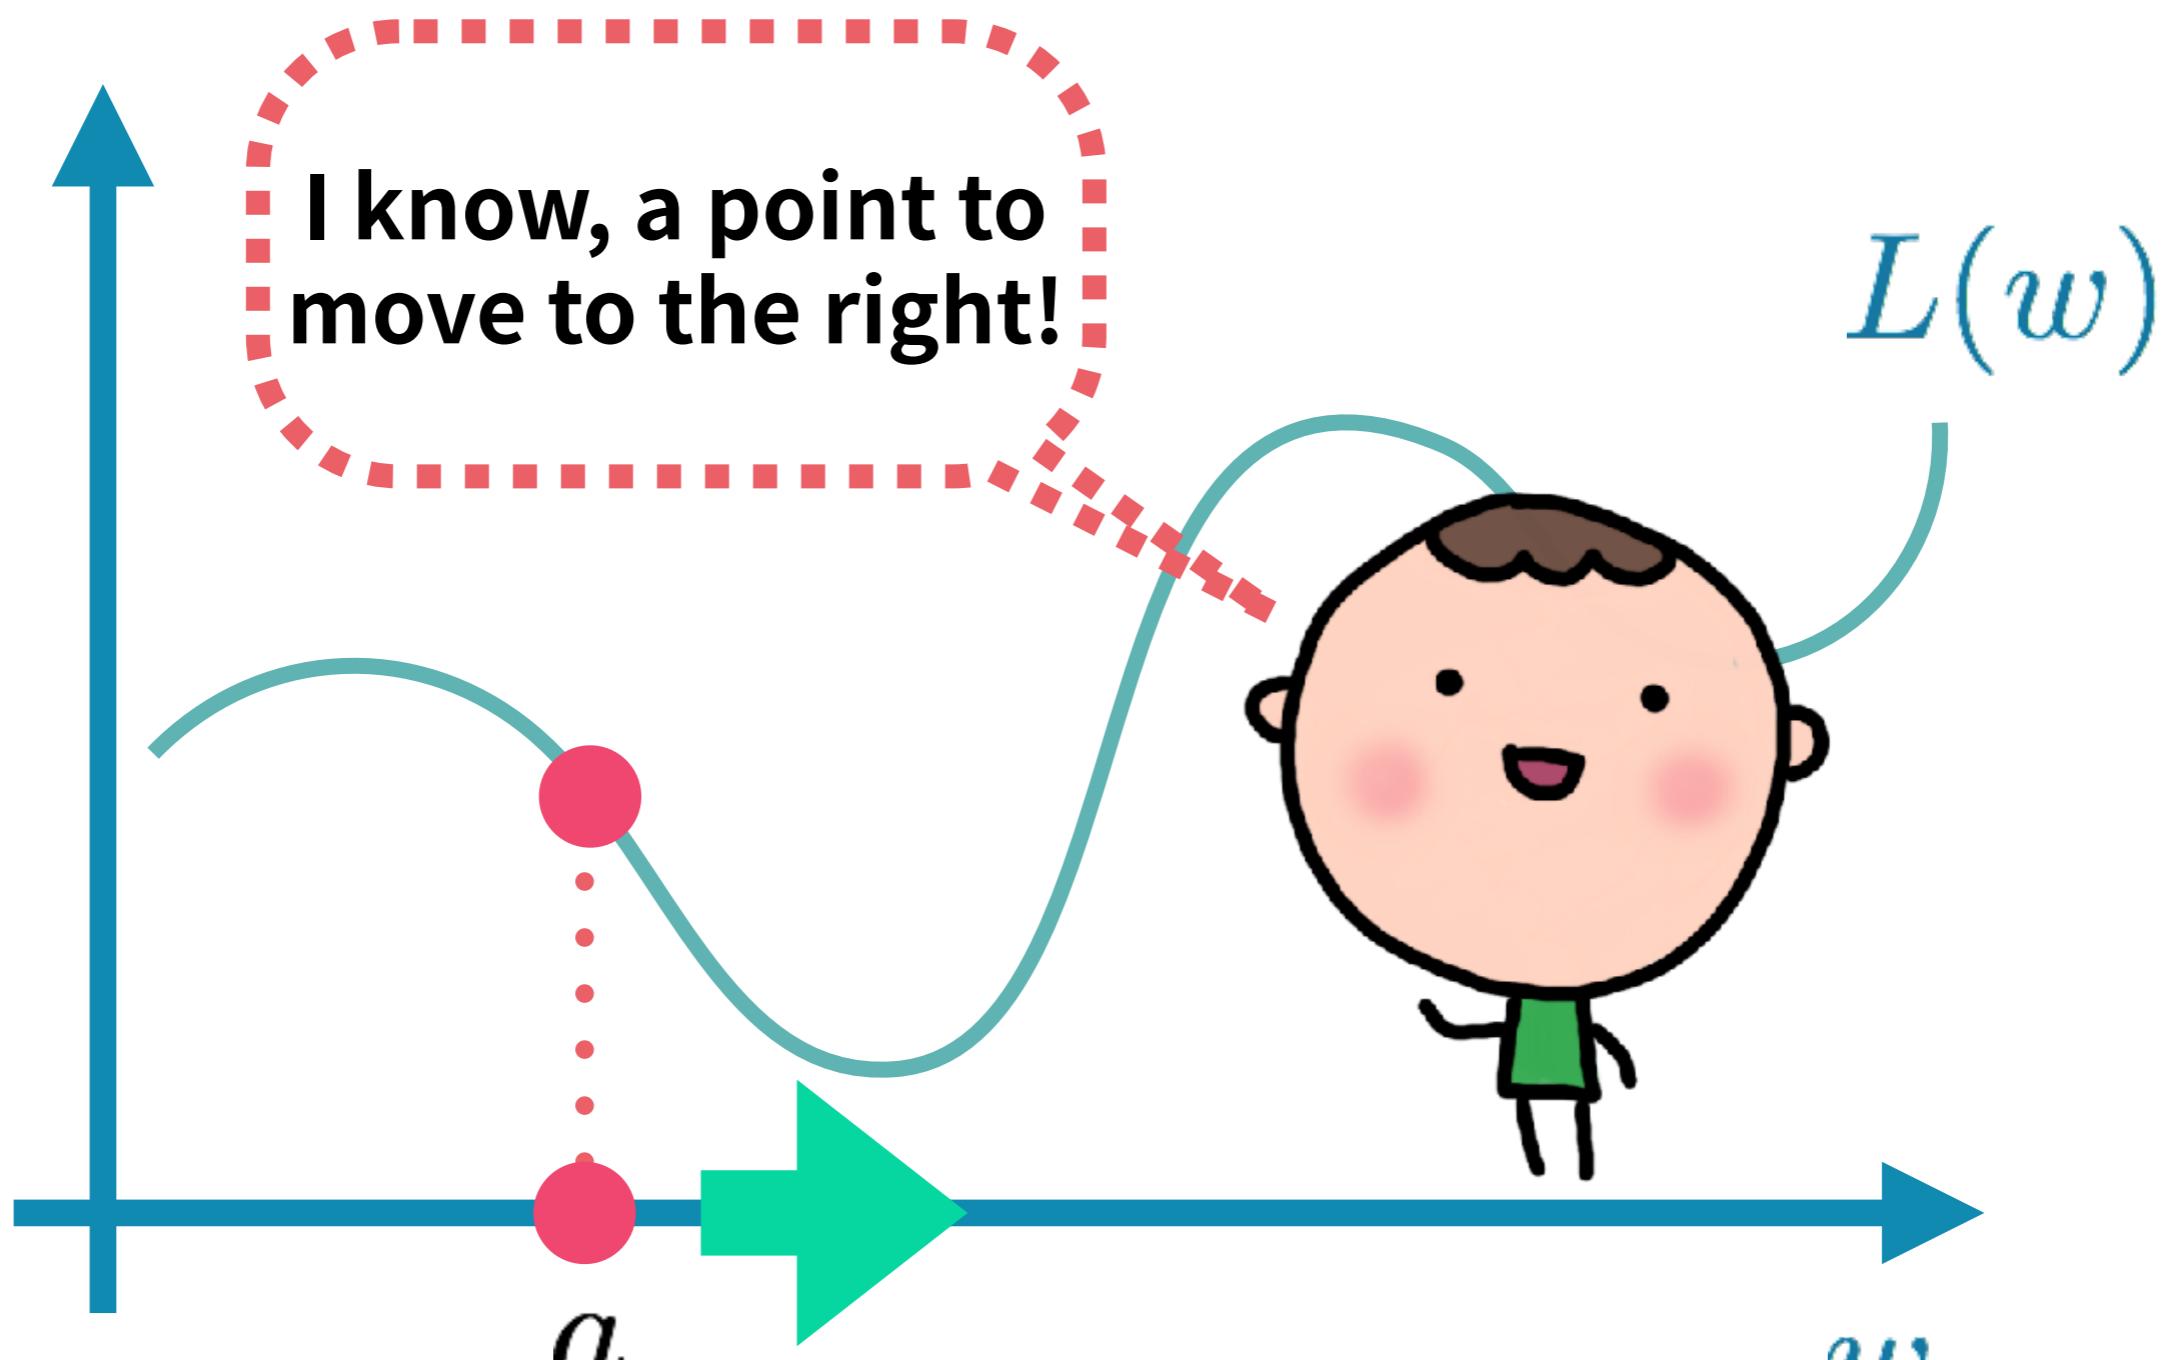
**what does this
terrible thing mean?**

**Remember that L is a
function of w_1, w_2, b_1, \dots**

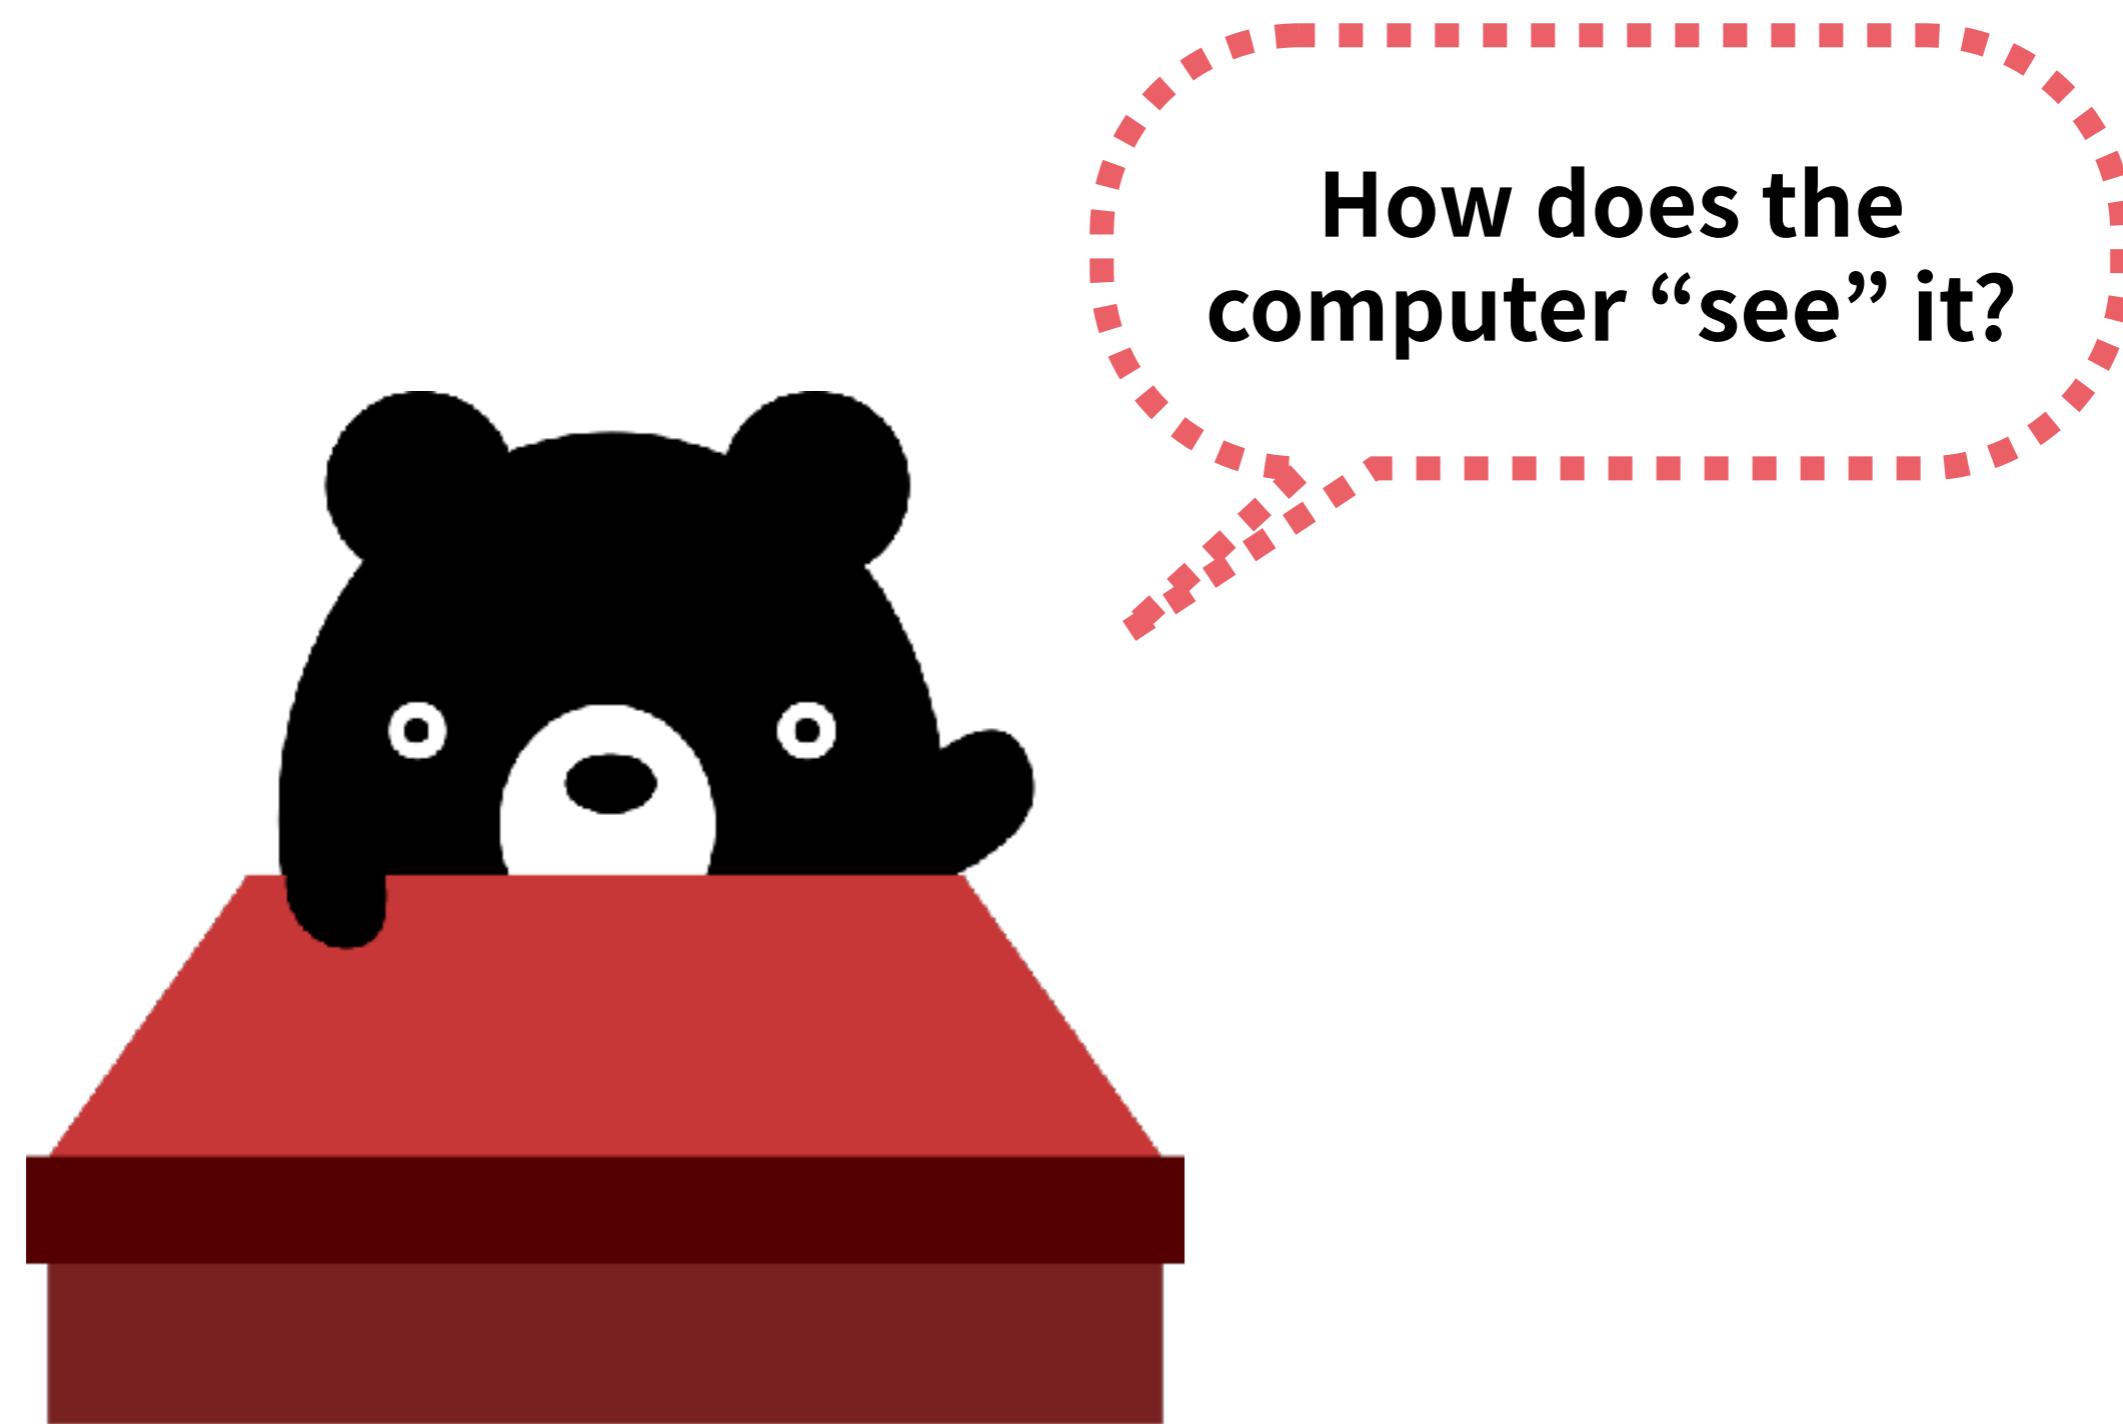


**For simplification, we first
think of L has only one
variable w**

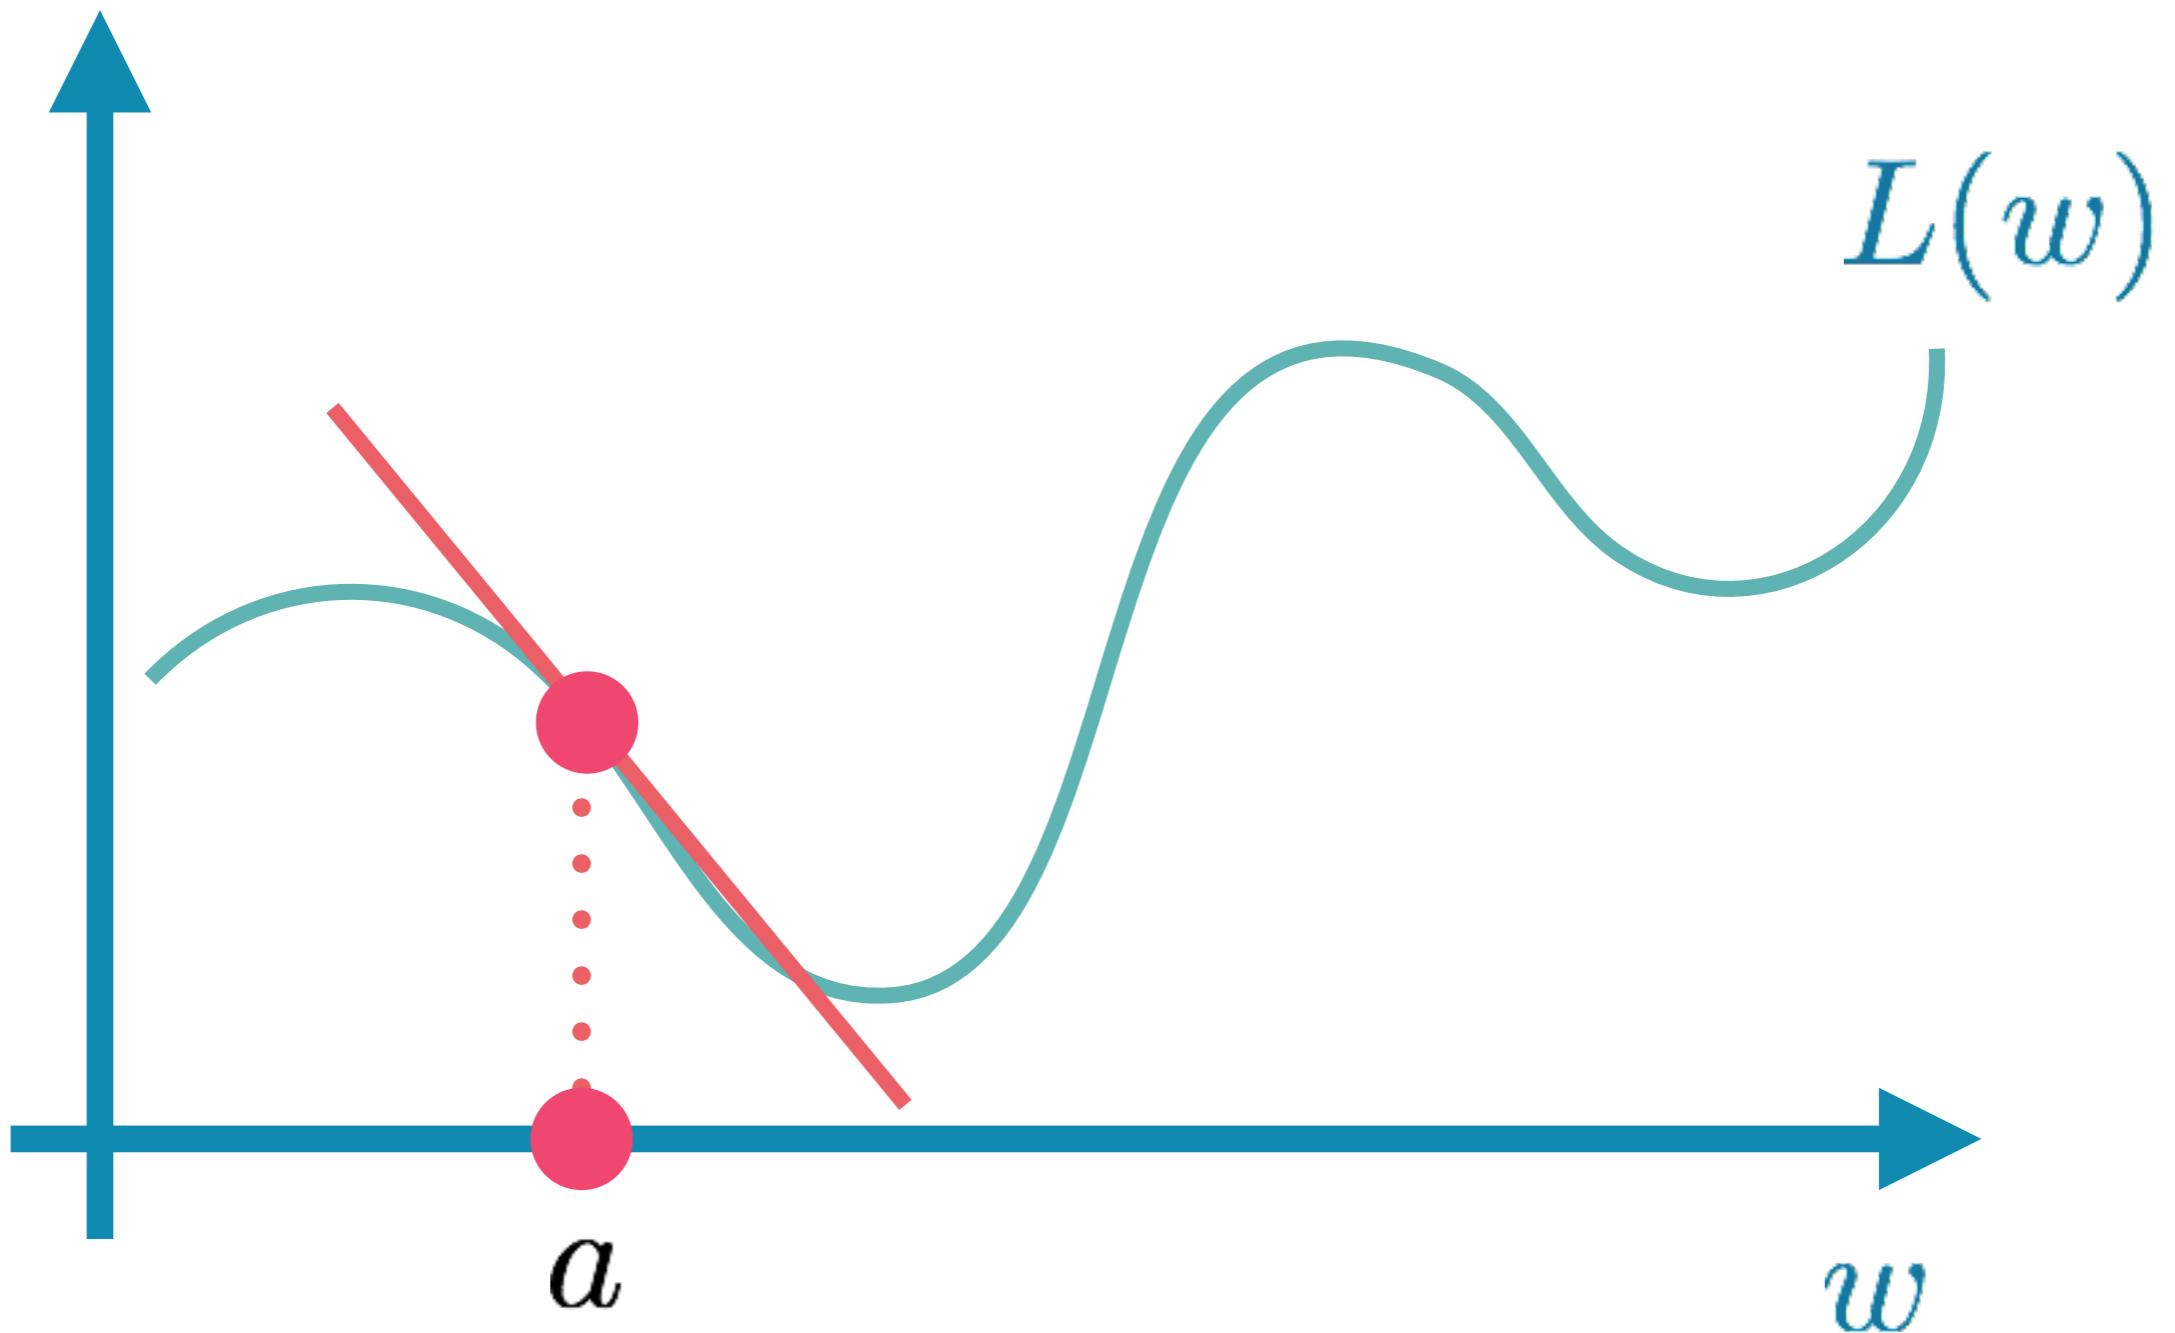




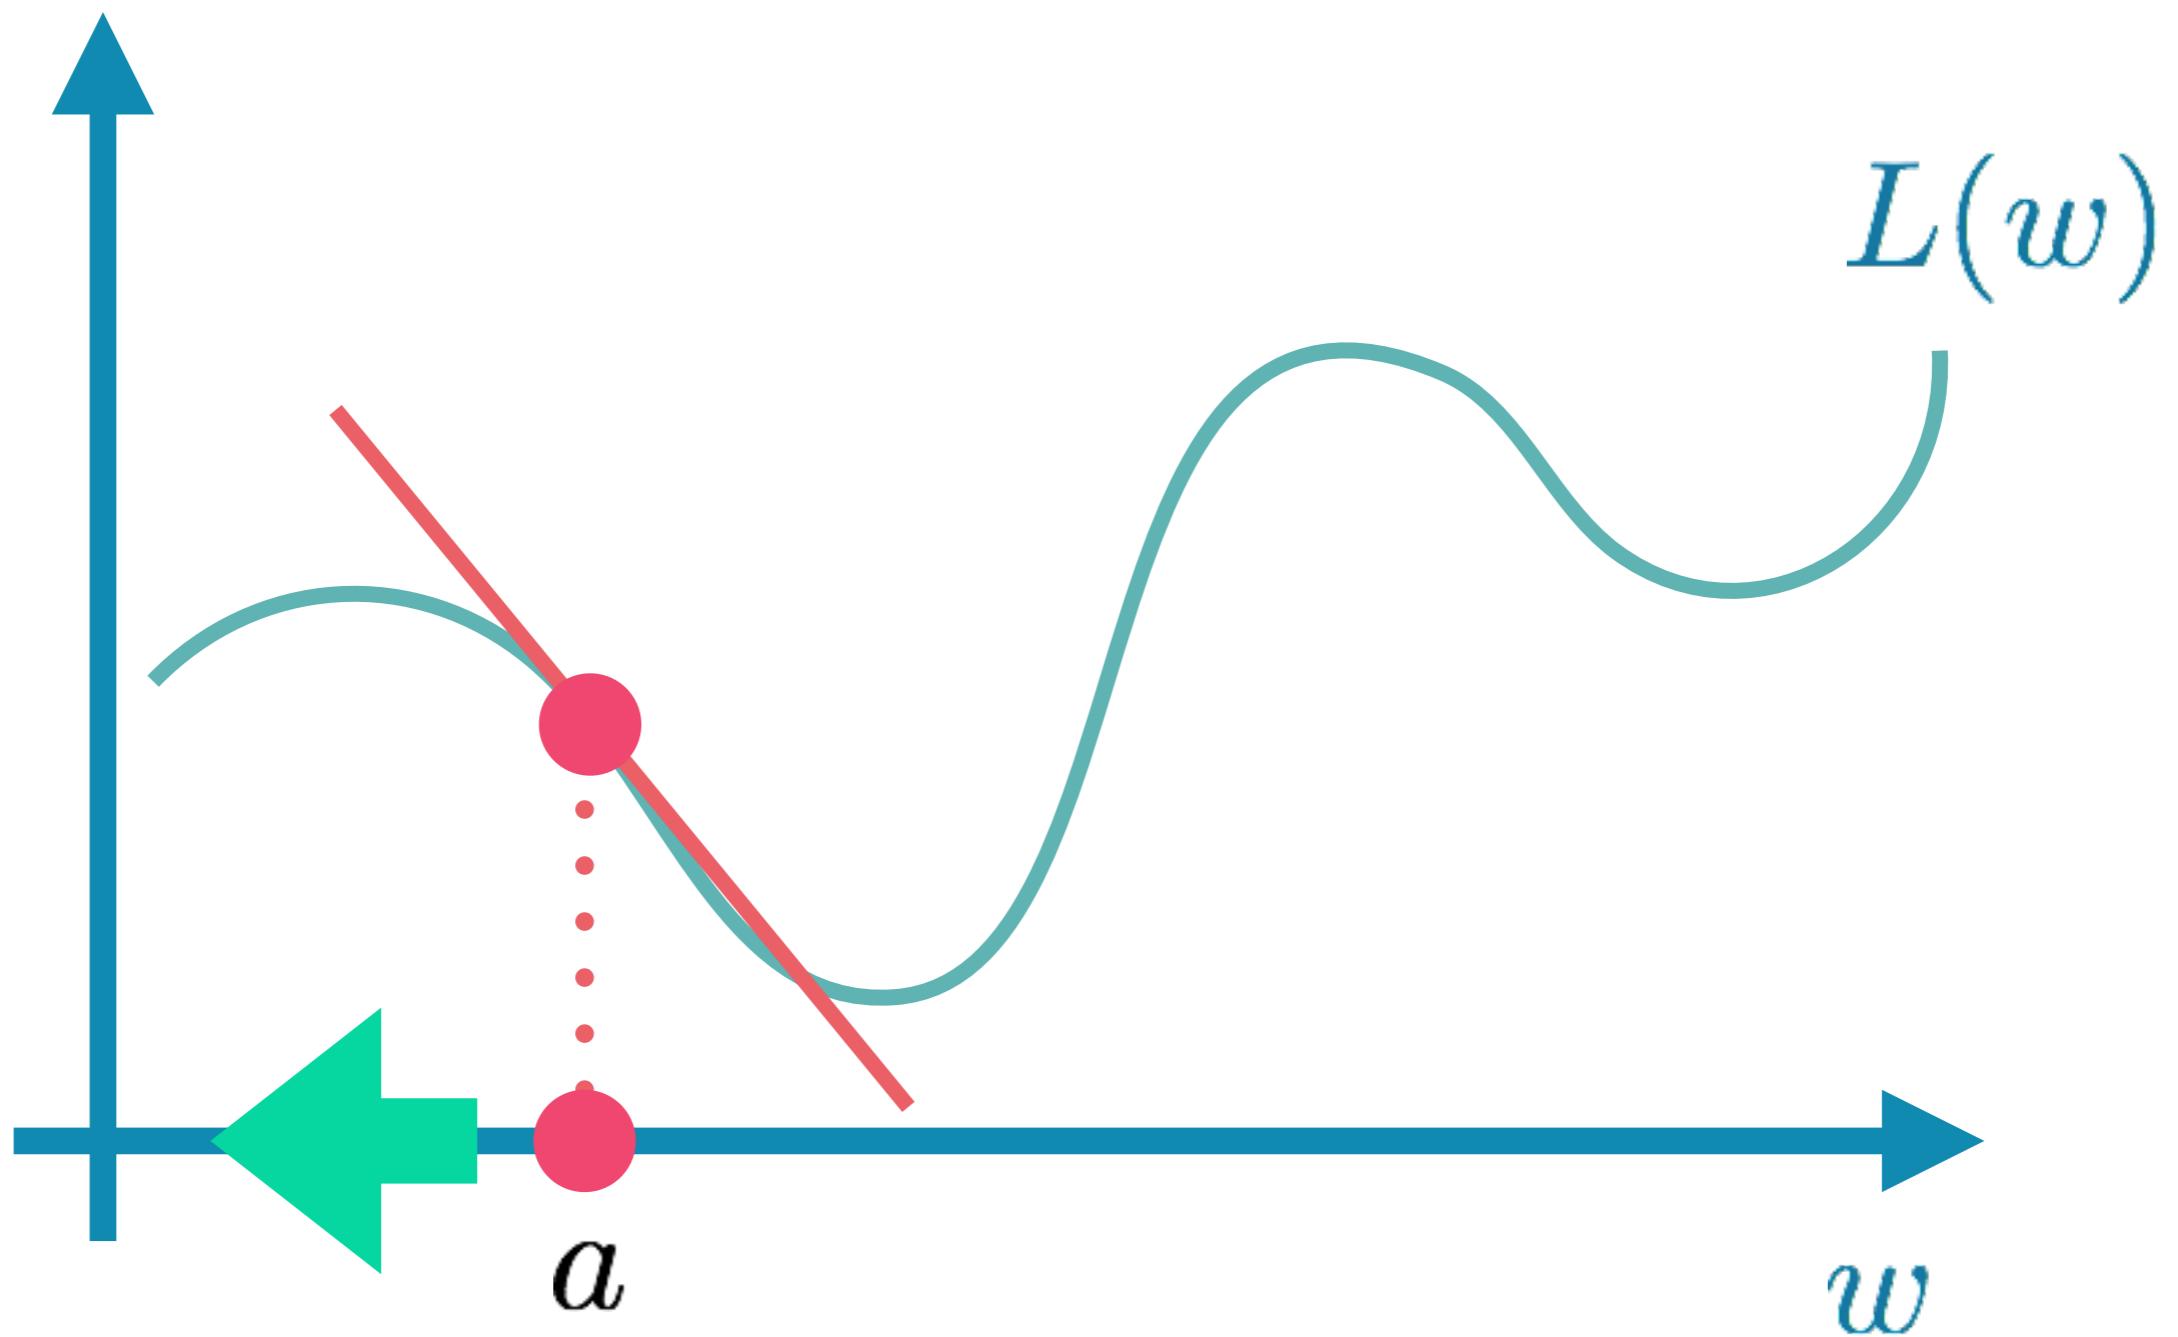
Current value of w



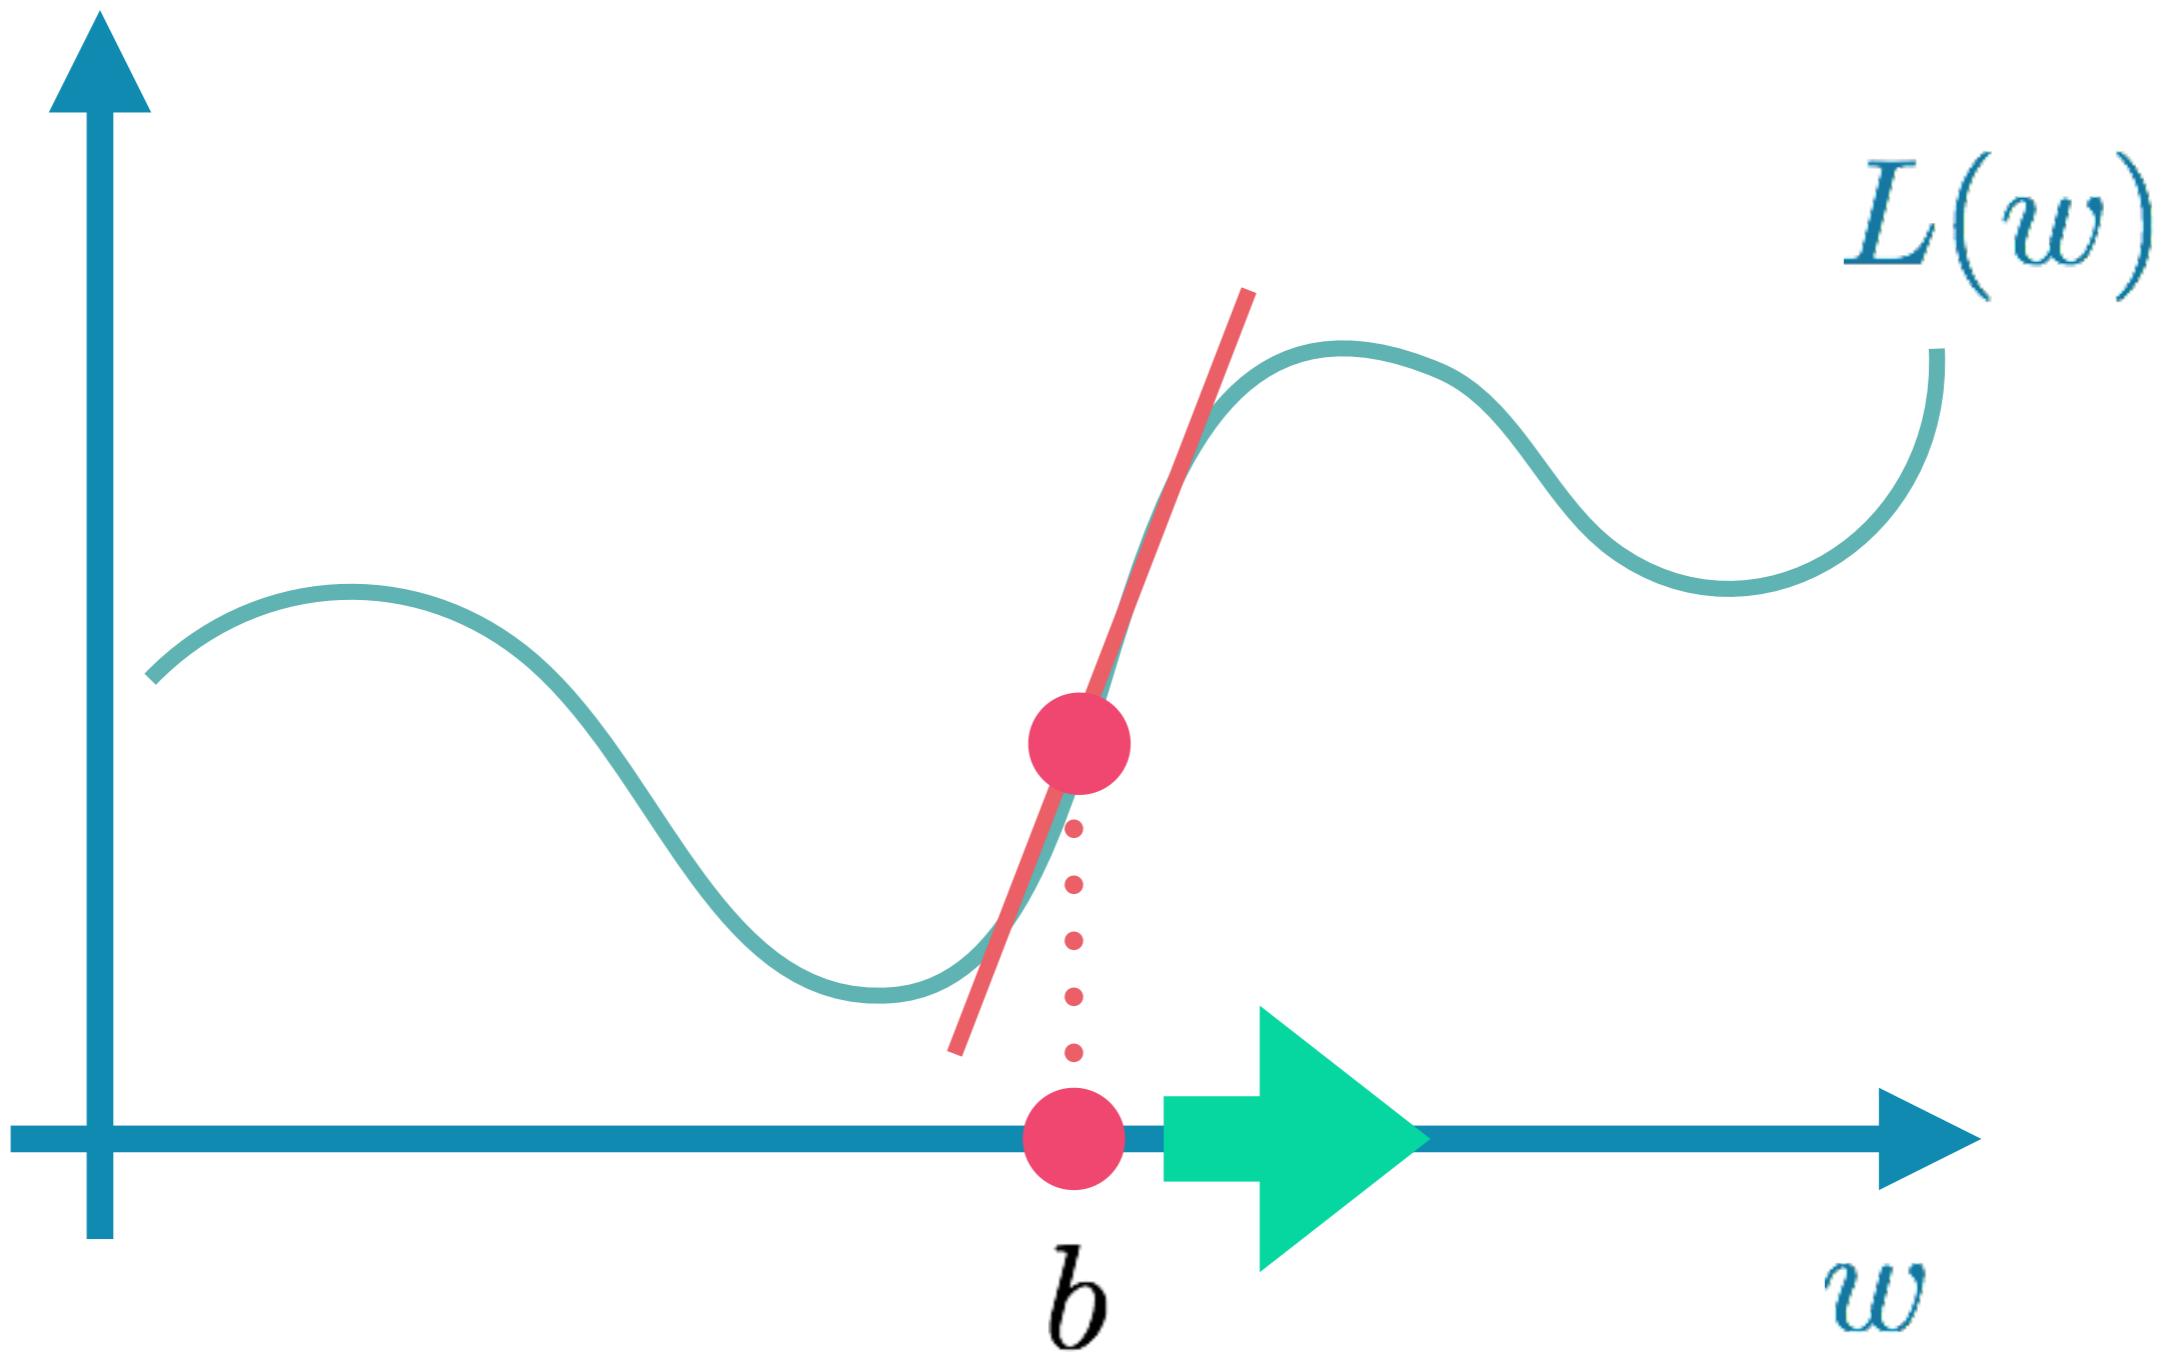
Tangent is the key!



Tangent slope < 0

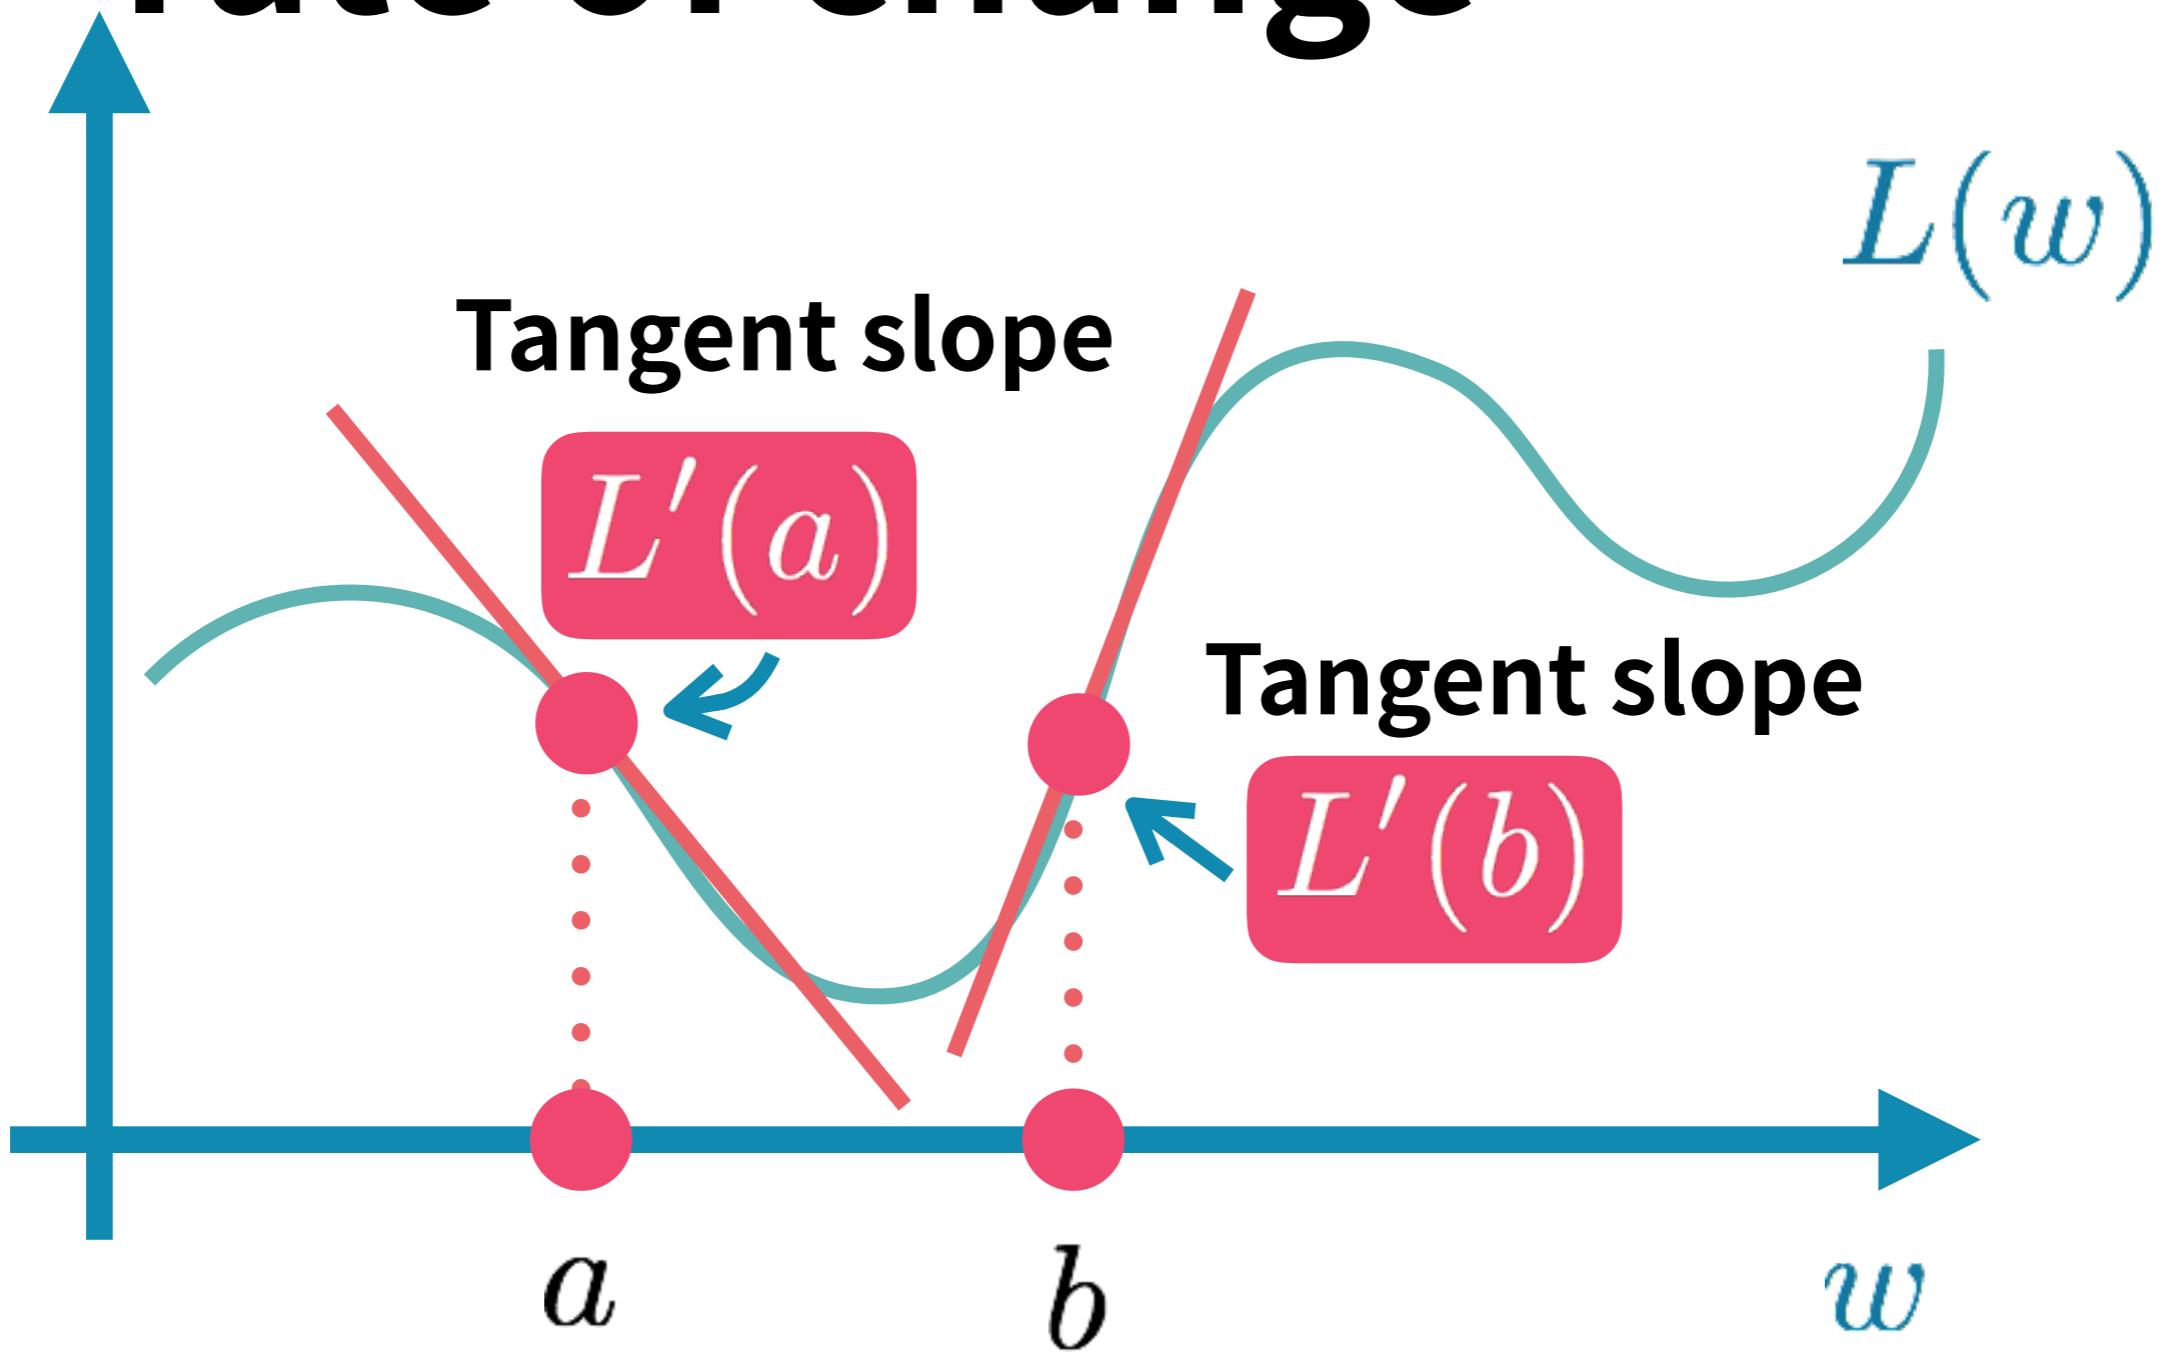


Tangent slope > 0

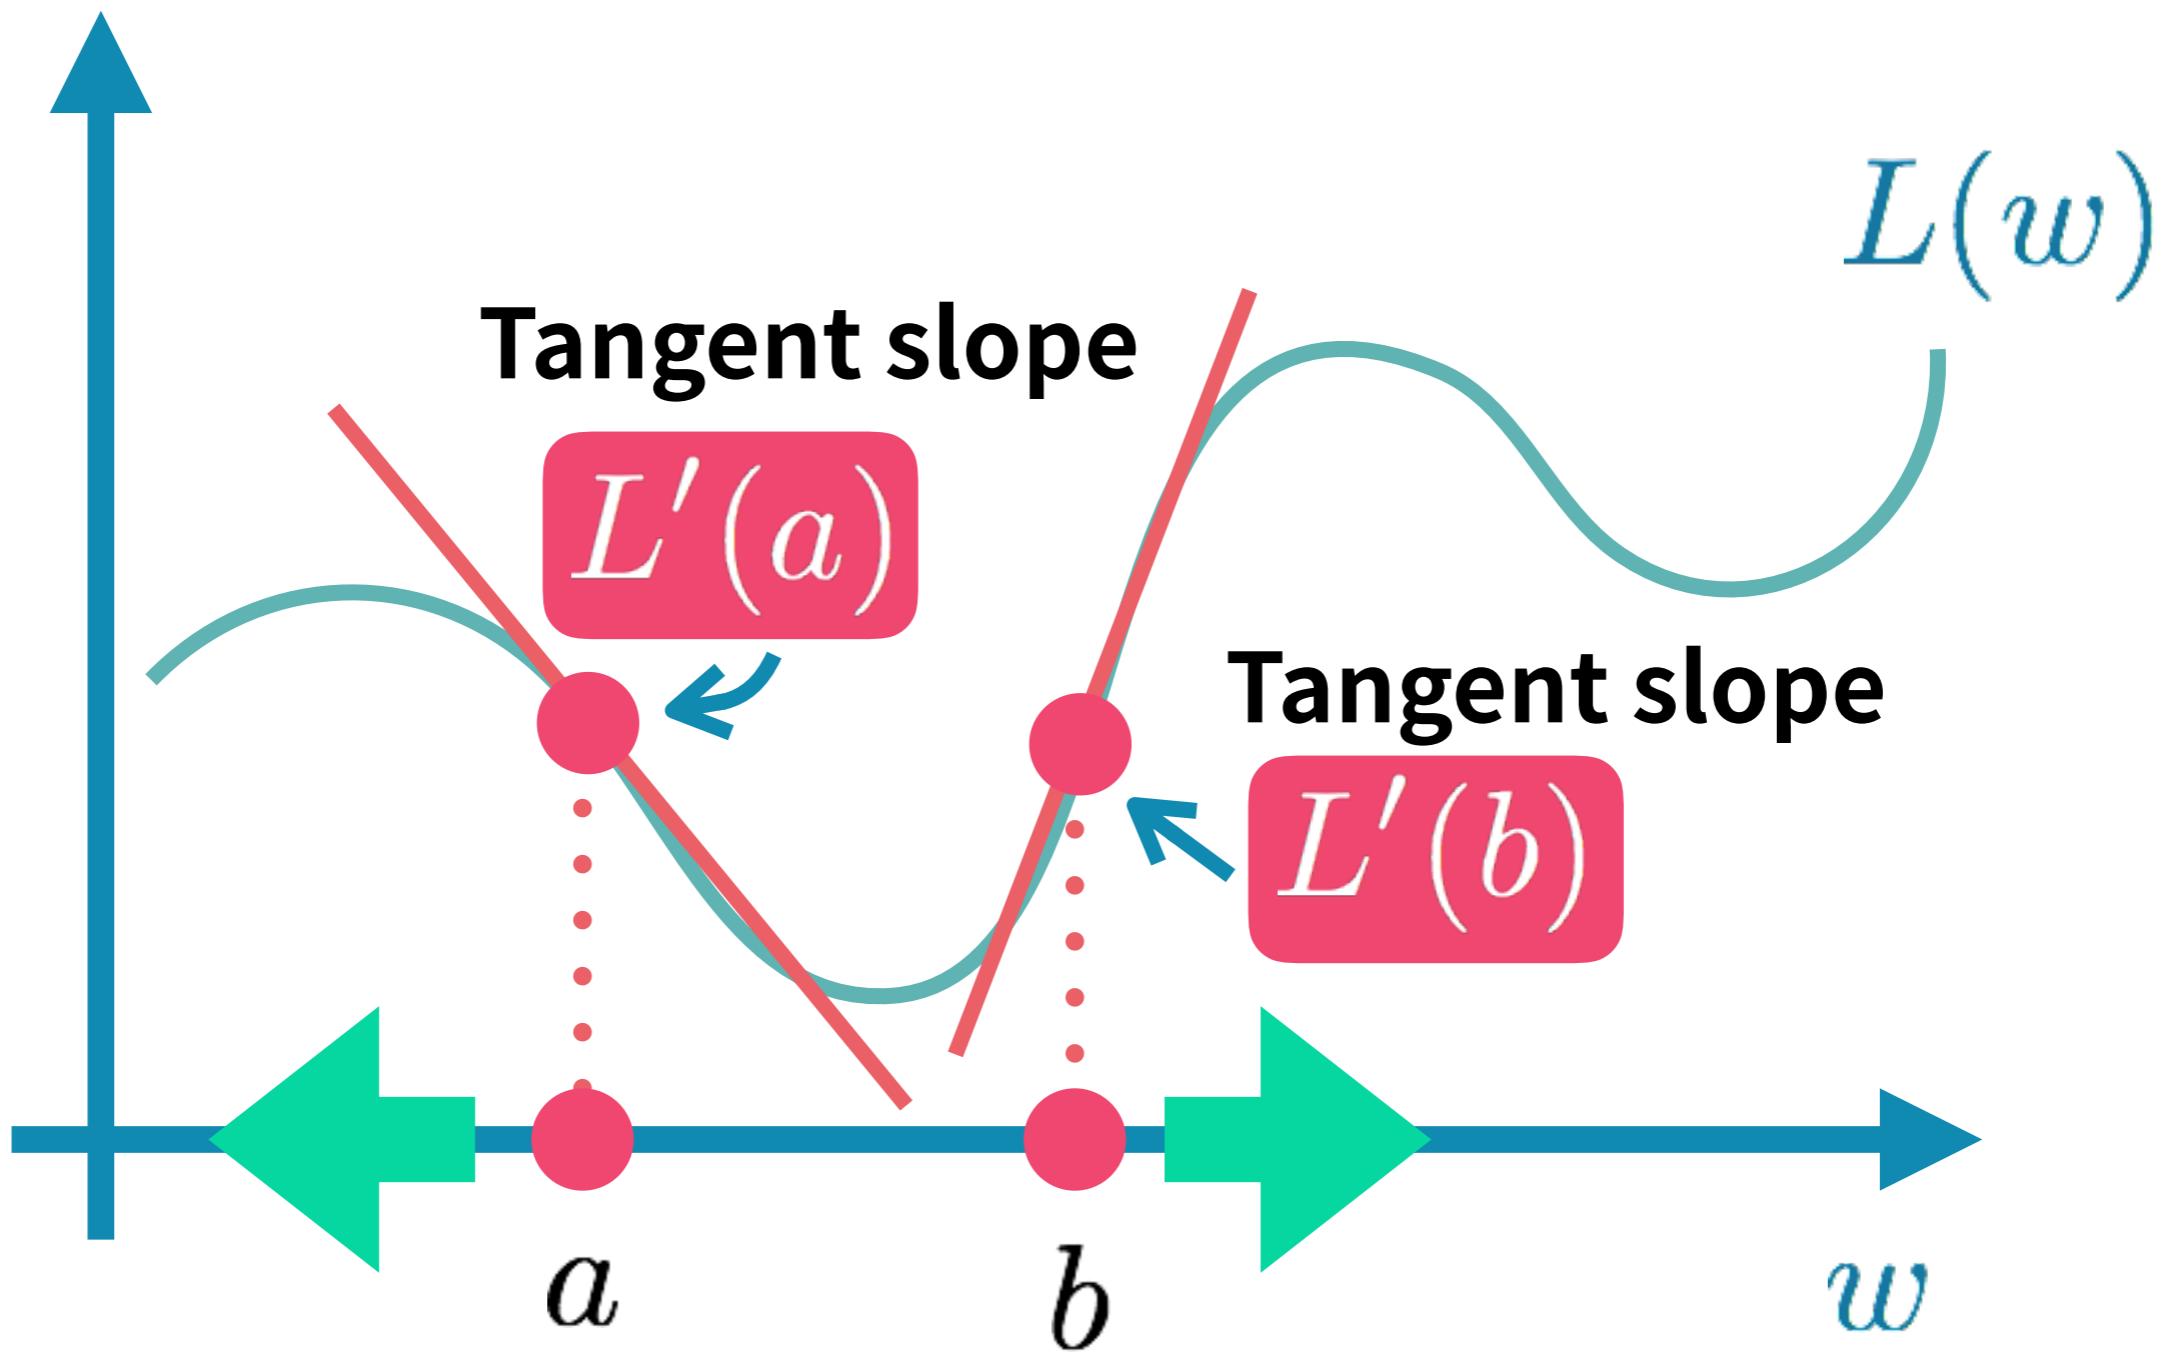


**The slope of the tangent
points to the (local)
maximum direction!**

Tangent slope is the rate of change



Pointing (local) maximum



Key

Moving toward (local) minima

**In order to make the value of the loss
function L smaller, we should adjust the
weight w:**

$$w - \frac{dL}{dw}$$

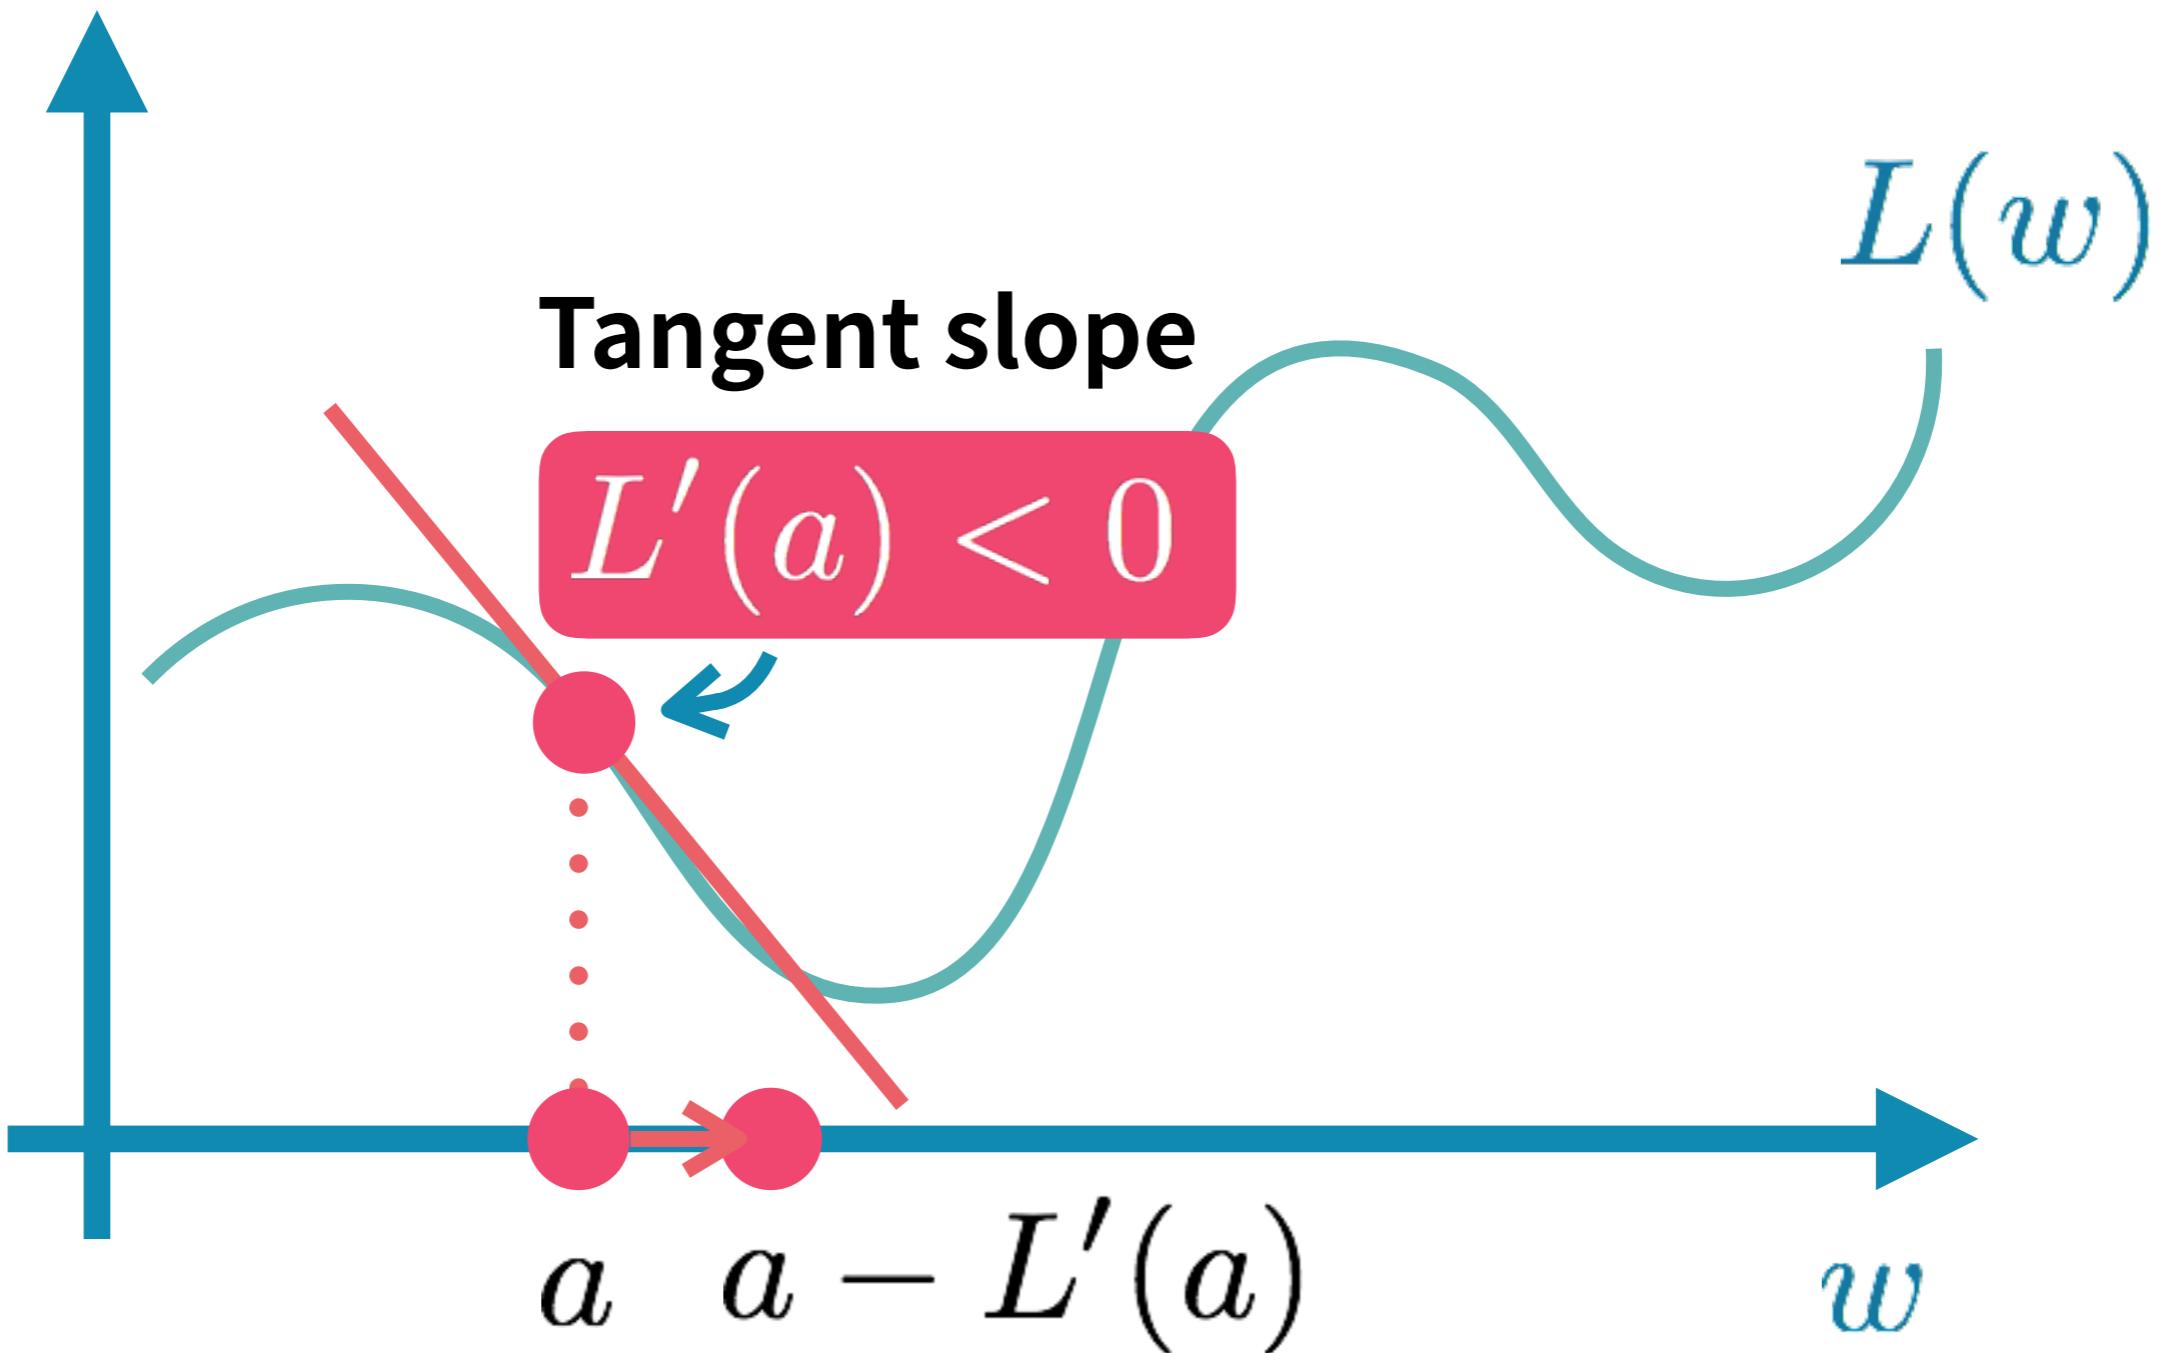
Key

Moving toward (local) minima

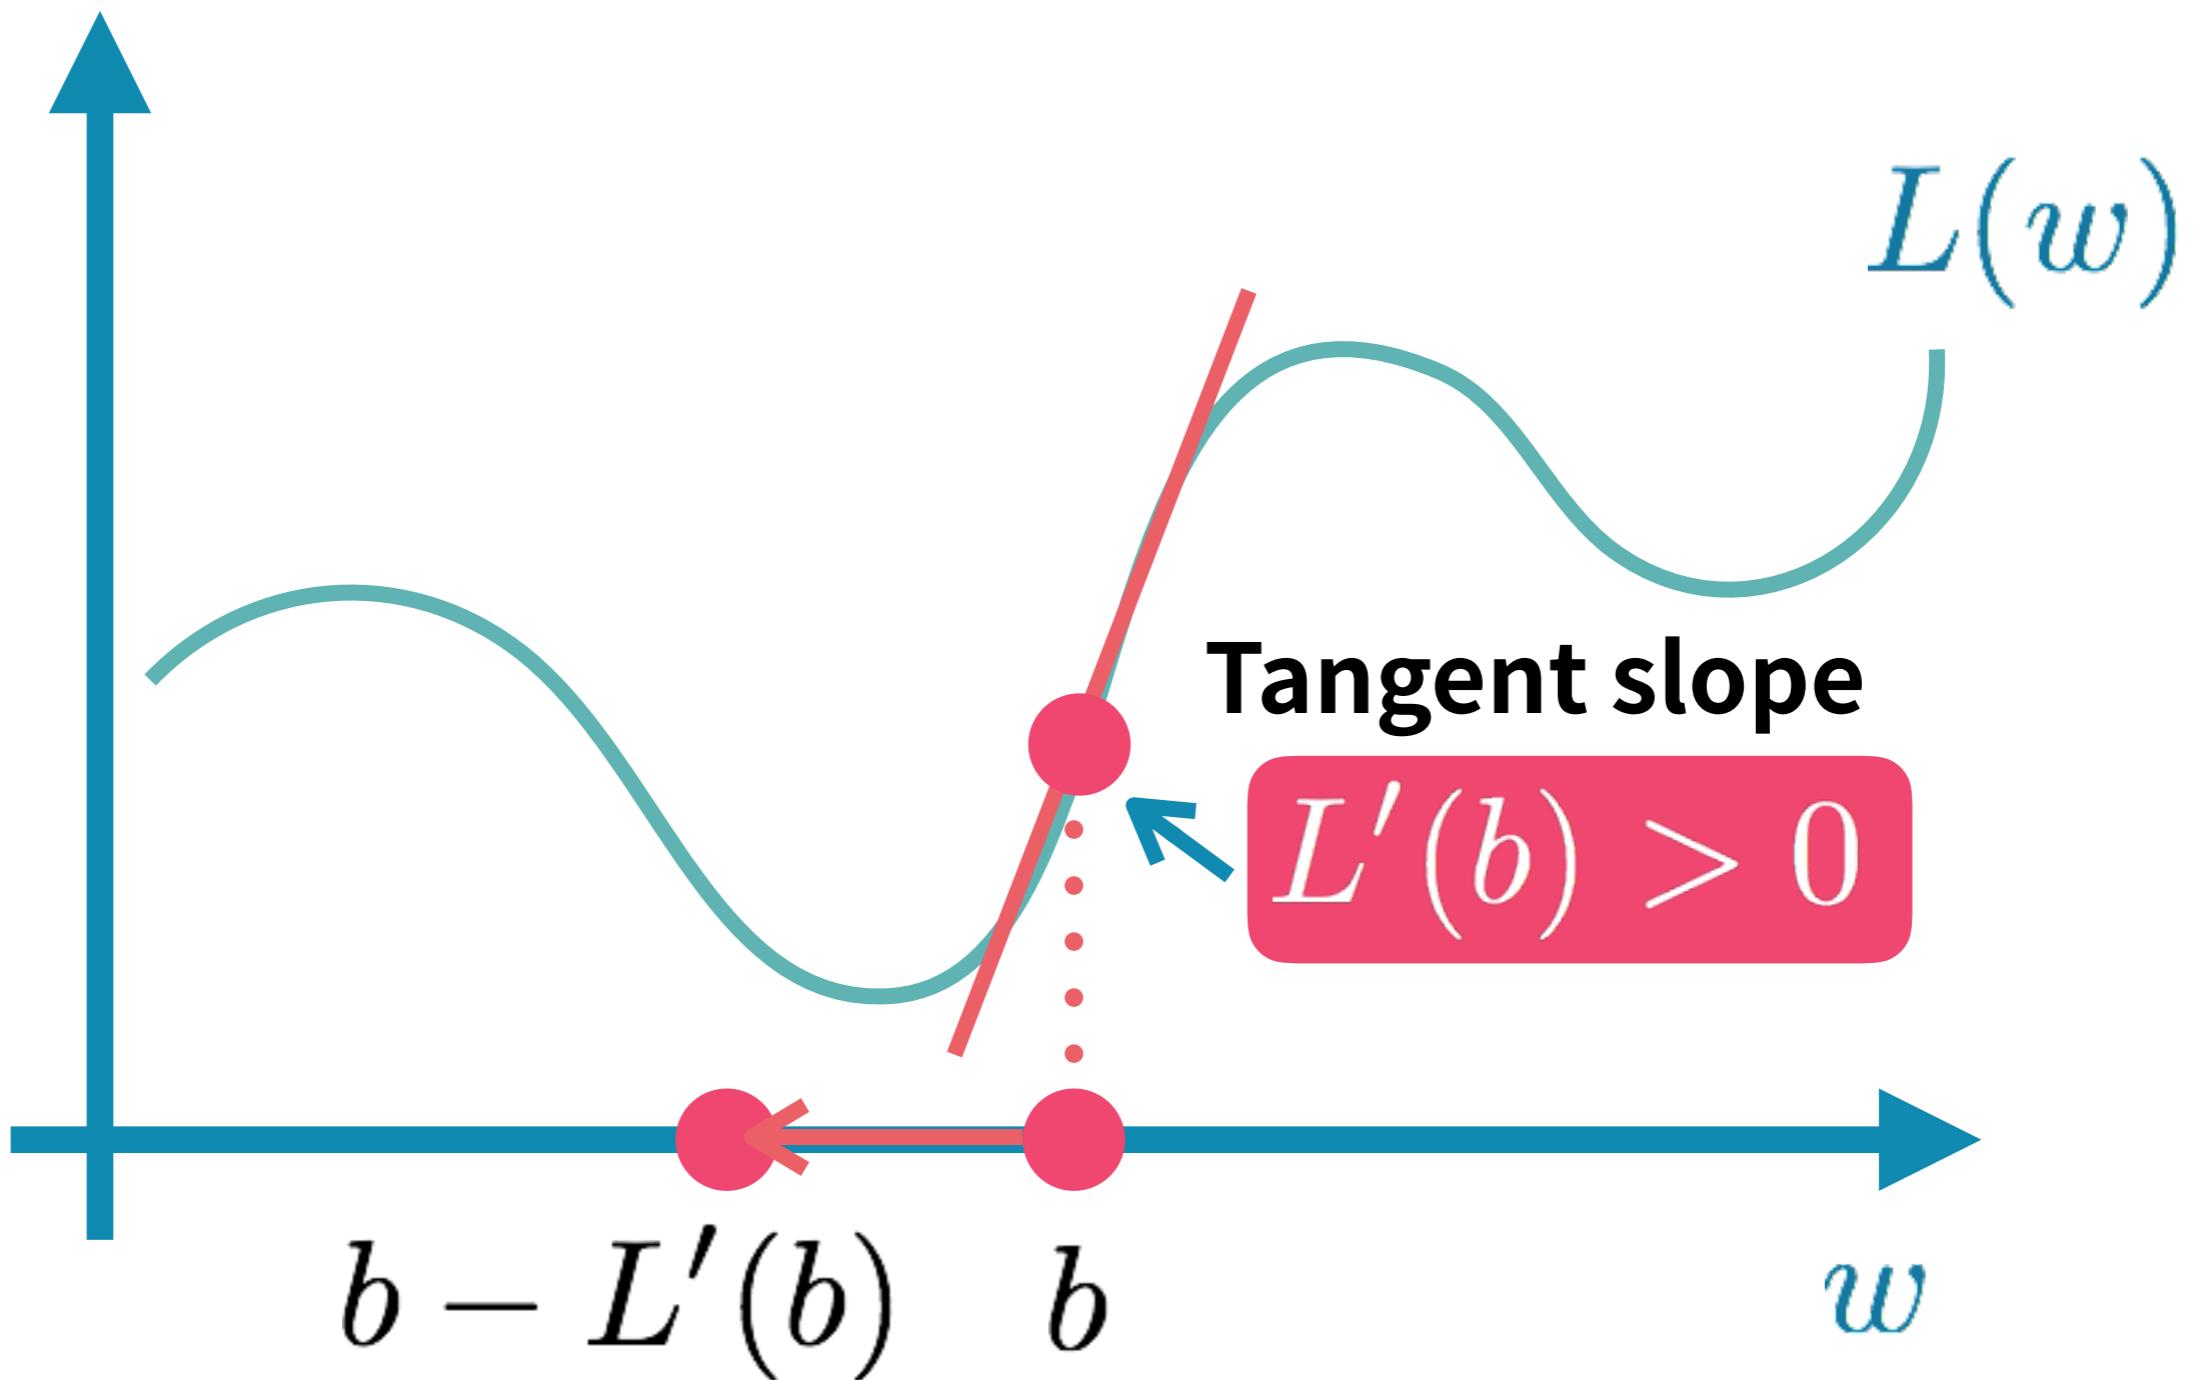
Suppose current $w=a$, we should update our w as:

$$a - L'(a)$$

Adjusting the Weight



Sometimes it will run over!



Key

Moving toward (local) minima

In order not to make too much adjustment
at one time, we will multiply a small
number called Leraning Rate:

$$w - \eta \frac{dL}{dw}$$

**However, there are
more than one
parameter...**

Example

$$L(w_1, w_2, b_1) = (b_1 + 2w_1 - w_2 - 3)^2$$

Suppose that

$$w_1 = 1, w_2 = -1, b_1 = 2$$

And we are moving toward to local minimum.

**Pretend to have only
one parameter!**

say, it is w_1

例子

$$L(w_1, w_2, b_1) = (b_1 + 2w_1 - w_2 - 3)^2$$

$$w_1 = 1, w_2 = -1, b_1 = 2$$

All parameters are fixed, except w_1 ...

$$L_{w_1}(w_1) = L(w_1, -1, 2) = 4w_1^2$$



Example

$$L(w_1, w_2, b_1) = (b_1 + 2w_1 - w_2 - 3)^2$$

$$w_1 = 1, w_2 = -1, b_1 = 2$$

Similarly, we can also define

$$L_{w_2}(w_2) = L(1, w_2, 2) = (-w_2 + 1)^2$$

$$L_{b_1}(b_1) = L(1, -1, b_1) = b_1^2$$

Example

So we can adjust w_1 , w_2 , b_1 , and make L smaller and smaller...

w_1

$$w_1 \leftarrow w_1 - \eta \frac{dL_{w_1}}{dw_1}$$

b_1

$$b_1 \leftarrow b_1 - \eta \frac{dL_{b_1}}{db_1}$$

w_2

$$w_2 \leftarrow w_2 - \eta \frac{dL_{w_2}}{dw_2}$$

Example

Writing together seems more learned!

$$\begin{bmatrix} w_1 \\ w_2 \\ w_3 \end{bmatrix} = \begin{bmatrix} w_1 \\ w_2 \\ w_3 \end{bmatrix} - \eta \begin{bmatrix} \frac{dL_{w_1}}{dw_1} \\ \frac{dL_{w_2}}{dw_2} \\ \frac{dL_{b_1}}{db_1} \end{bmatrix}$$

Updated

We call this the
gradient of L

Definition

Partial Derivative

What we have done is just so called partial derivatives.

$$\frac{\partial L}{\partial w_1} = \frac{dL_{w_1}}{dw_1}$$

Definition

Partial Derivative

Similarly,

$$\frac{\partial L}{\partial w_2} = \frac{dL_{w_2}}{dw_2}$$

$$\frac{\partial L}{\partial b_1} = \frac{dL_{b_1}}{db_1}$$

Notation

gradient

Recall the gradient of L is:

$$\begin{bmatrix} \frac{\partial L}{\partial w_1} \\ \frac{\partial L}{\partial w_2} \\ \frac{\partial L}{\partial b_1} \end{bmatrix}$$


$$\begin{bmatrix} \frac{dL_{w_1}}{dw_1} \\ \frac{dL_{w_2}}{dw_2} \\ \frac{dL_{b_1}}{db_1} \end{bmatrix}$$

Notation

gradient

We should have cool notation for the gradient:

$$\nabla L = \begin{bmatrix} \frac{\partial L}{\partial w_1} \\ \frac{\partial L}{\partial w_2} \\ \frac{\partial L}{\partial b_1} \end{bmatrix}$$

Notation

gradient

So we can update all parameters by the formula:

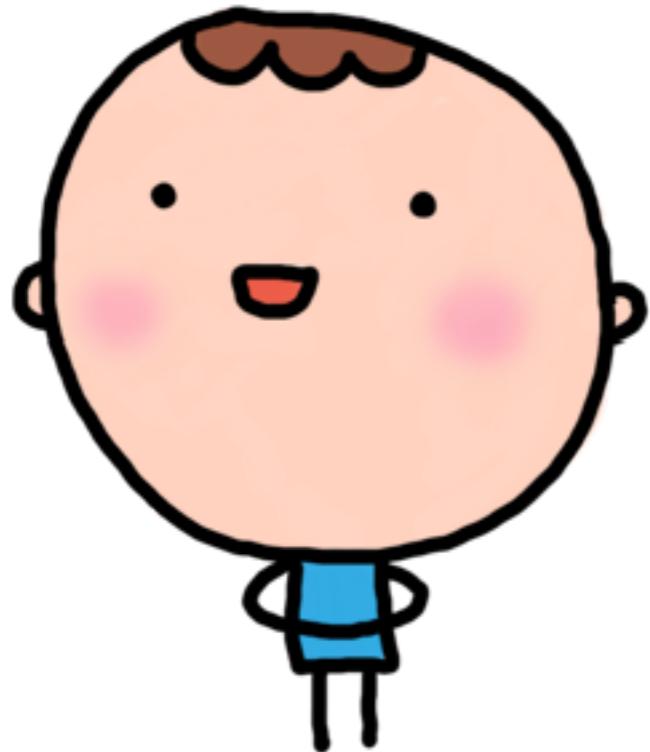
$$\begin{bmatrix} w_1 \\ w_2 \\ b_1 \end{bmatrix} - \eta \nabla L$$

This “learning method” has a
fancy name

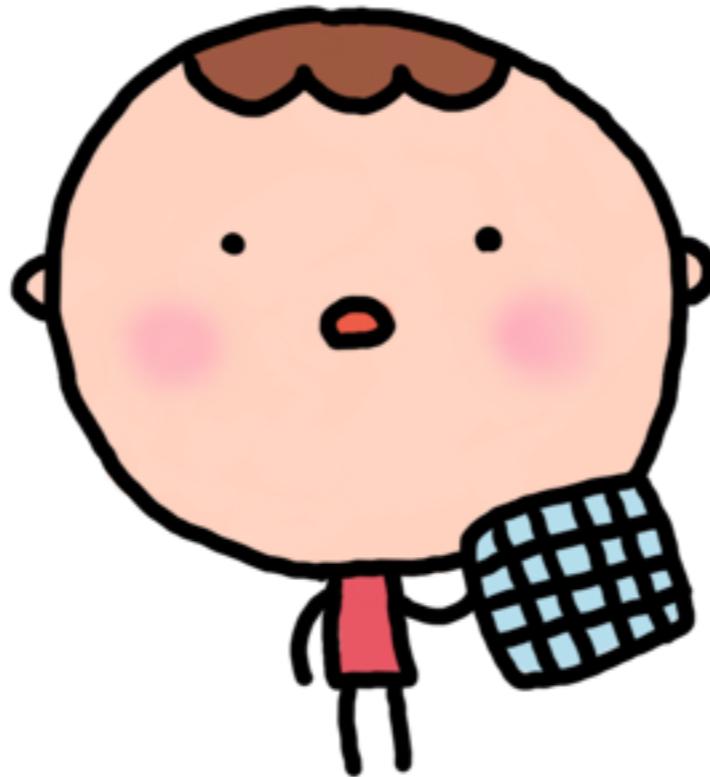
Gradient Descent

Deep Learning

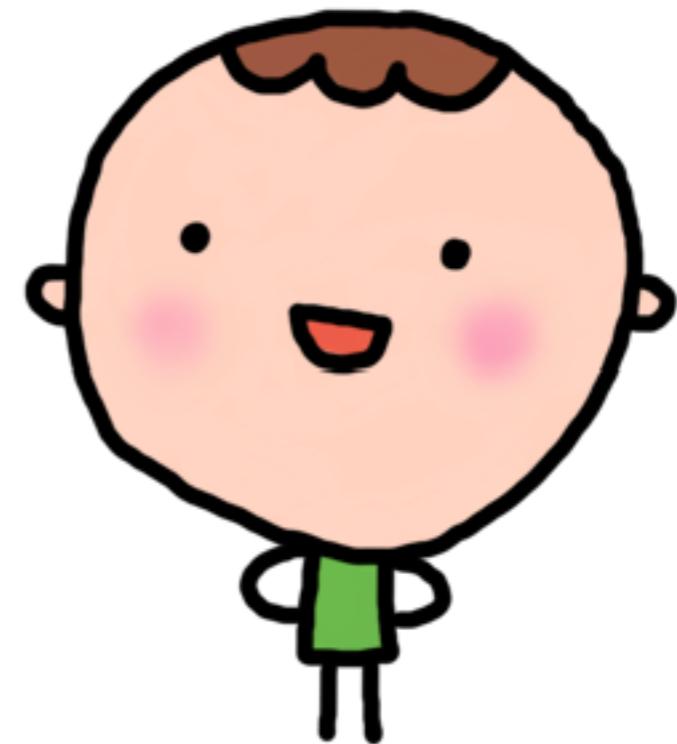
Big 3



Standard NN



CNN



RNN

No Kidding!



eder @edersantana · 1天

When you gotta hustle selling neural nets
in China. Already profitable AI startup!

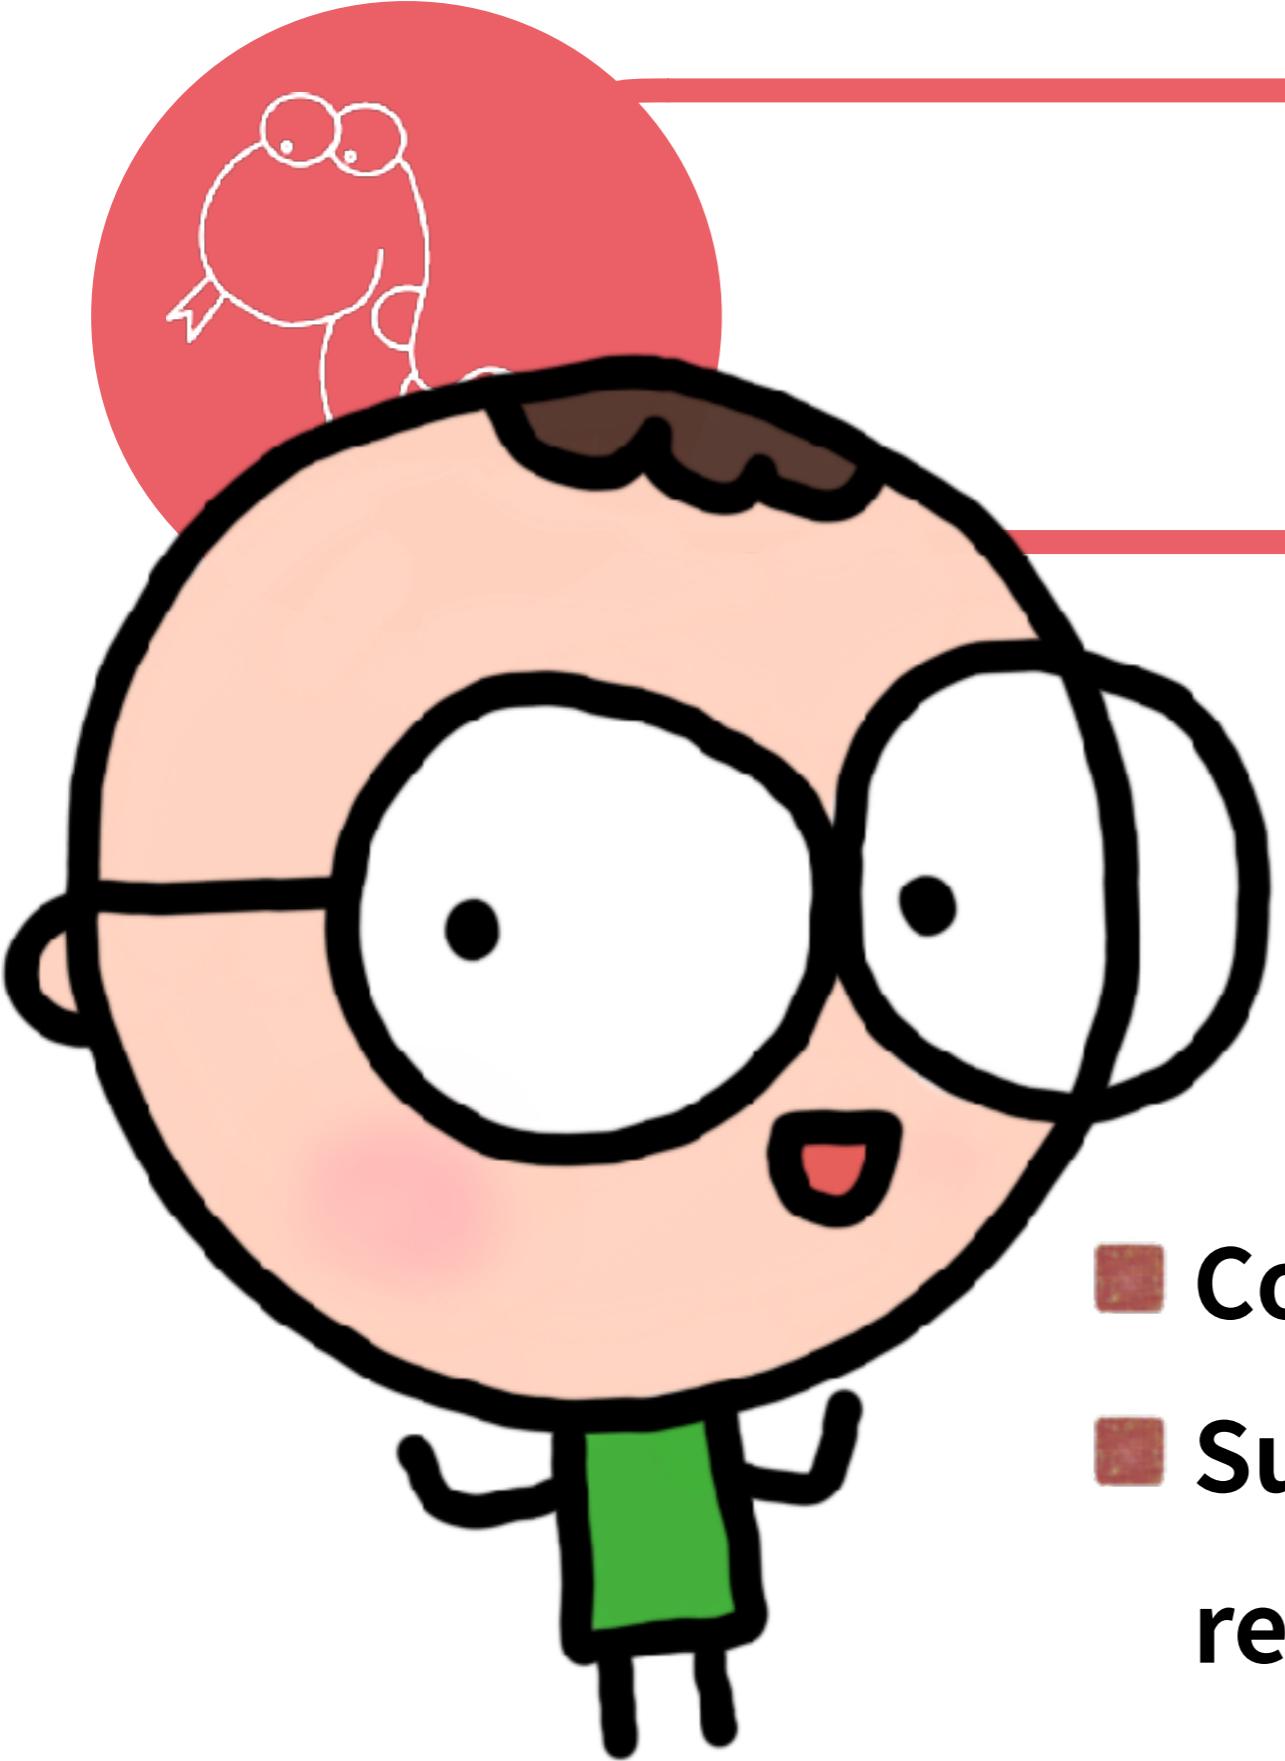


◀ 6

↑↓ 111

❤ 200

✉

A cartoon illustration of a character with large, round, white eyes and black outlines. The character wears black-rimmed glasses and a red baseball cap. A small white bird is perched on the rim of the cap. The character has a simple, rounded body and is wearing a green rectangular patch on their chest.

CNN

- Convolutional Neural Network
- Super star for image
recognition

image recognition

$f($  $) = \text{Formosan black bear}$

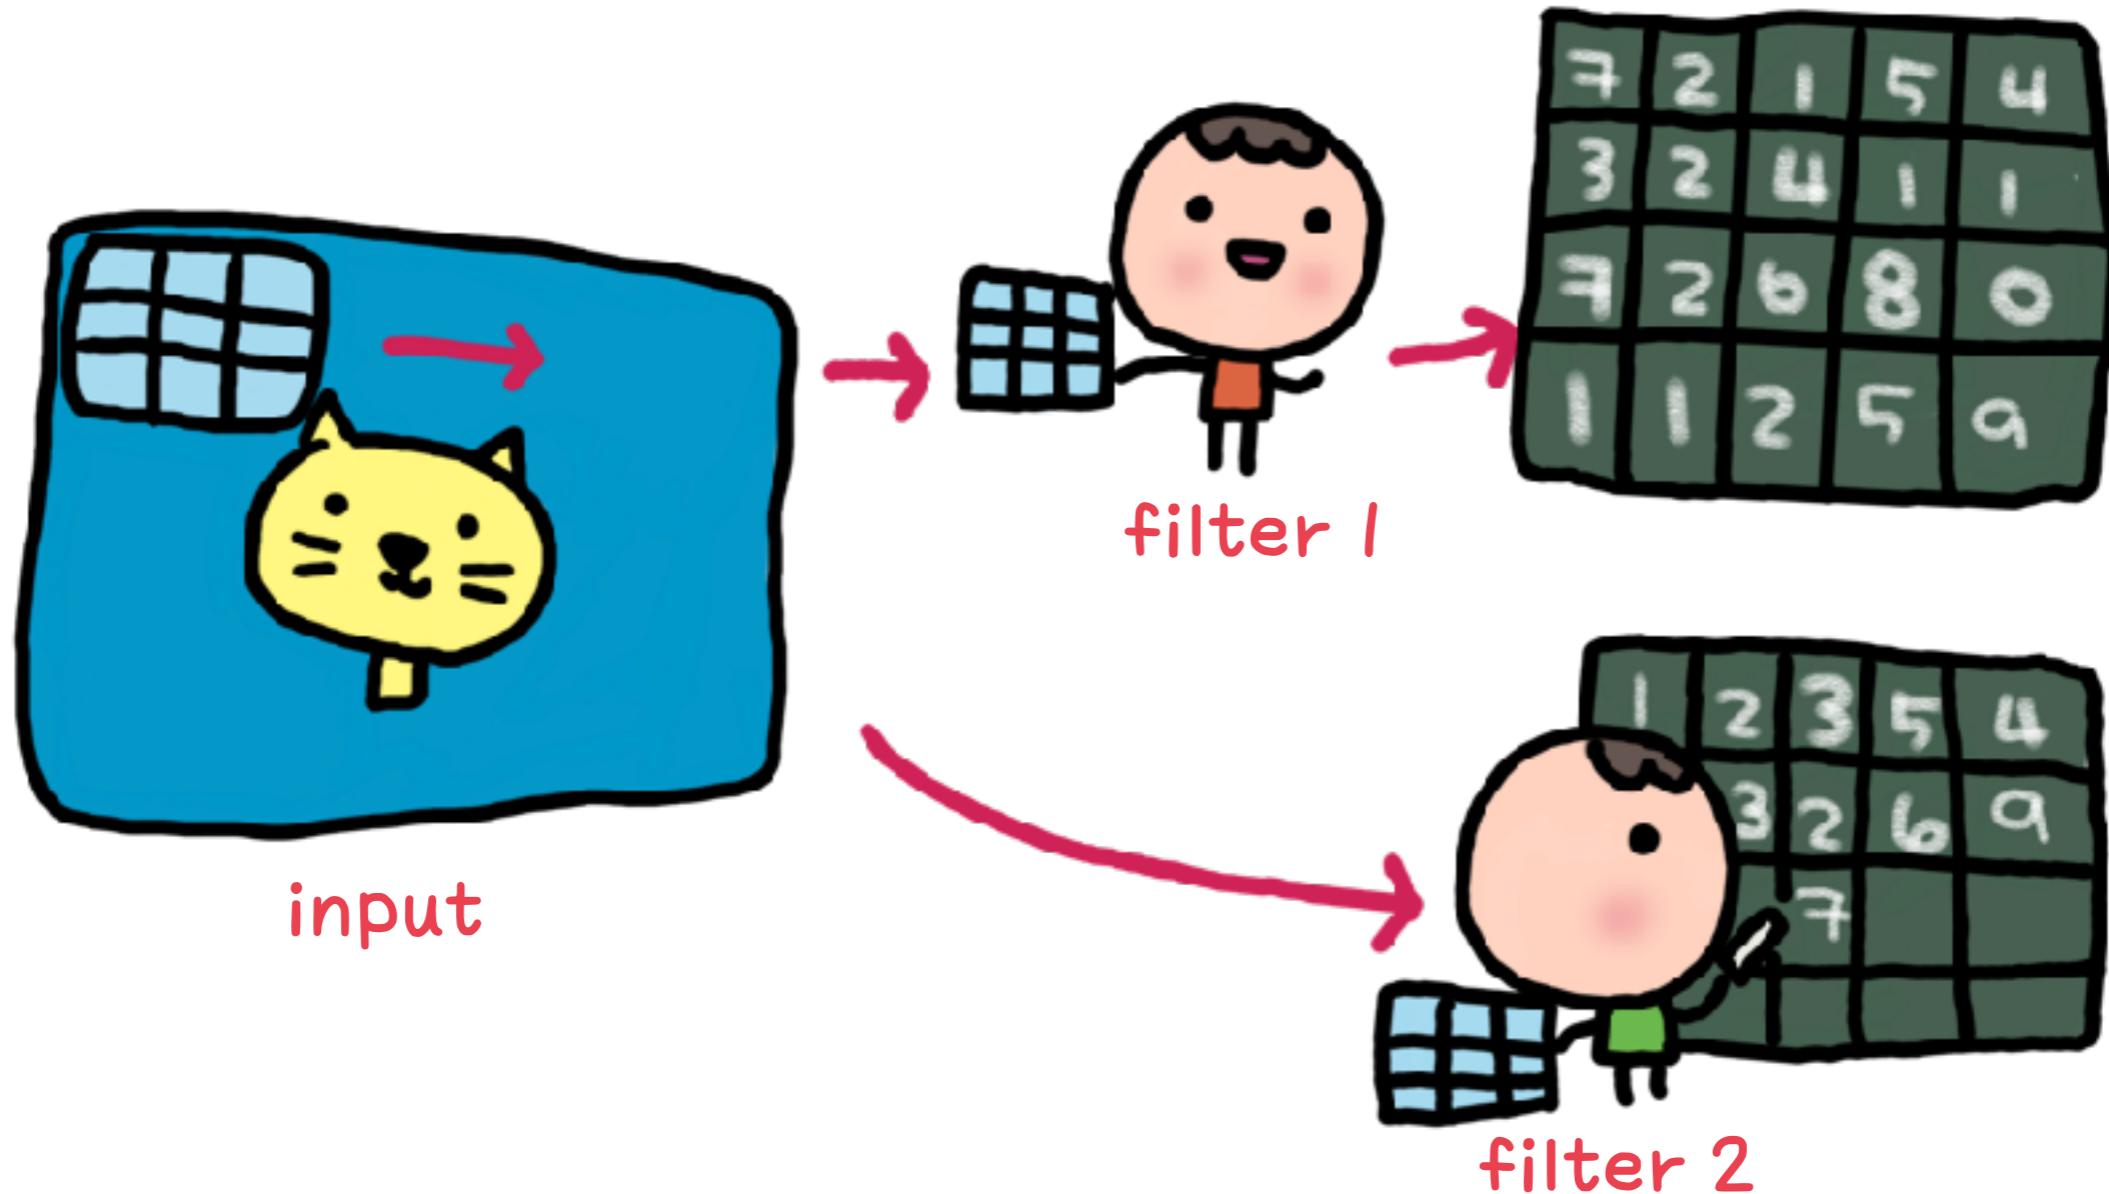
Playing Video Games

reinforcement learning



π (Current states) = “the best” action

Convolutional Neural Network (CNN)



1

Convolutional Layer

We need some “filters”

Say, 3x3 filters

Dot product

$$W = \begin{bmatrix} 1 & 0 & 1 \\ 2 & 1 & 3 \\ 1 & 1 & 2 \end{bmatrix}$$

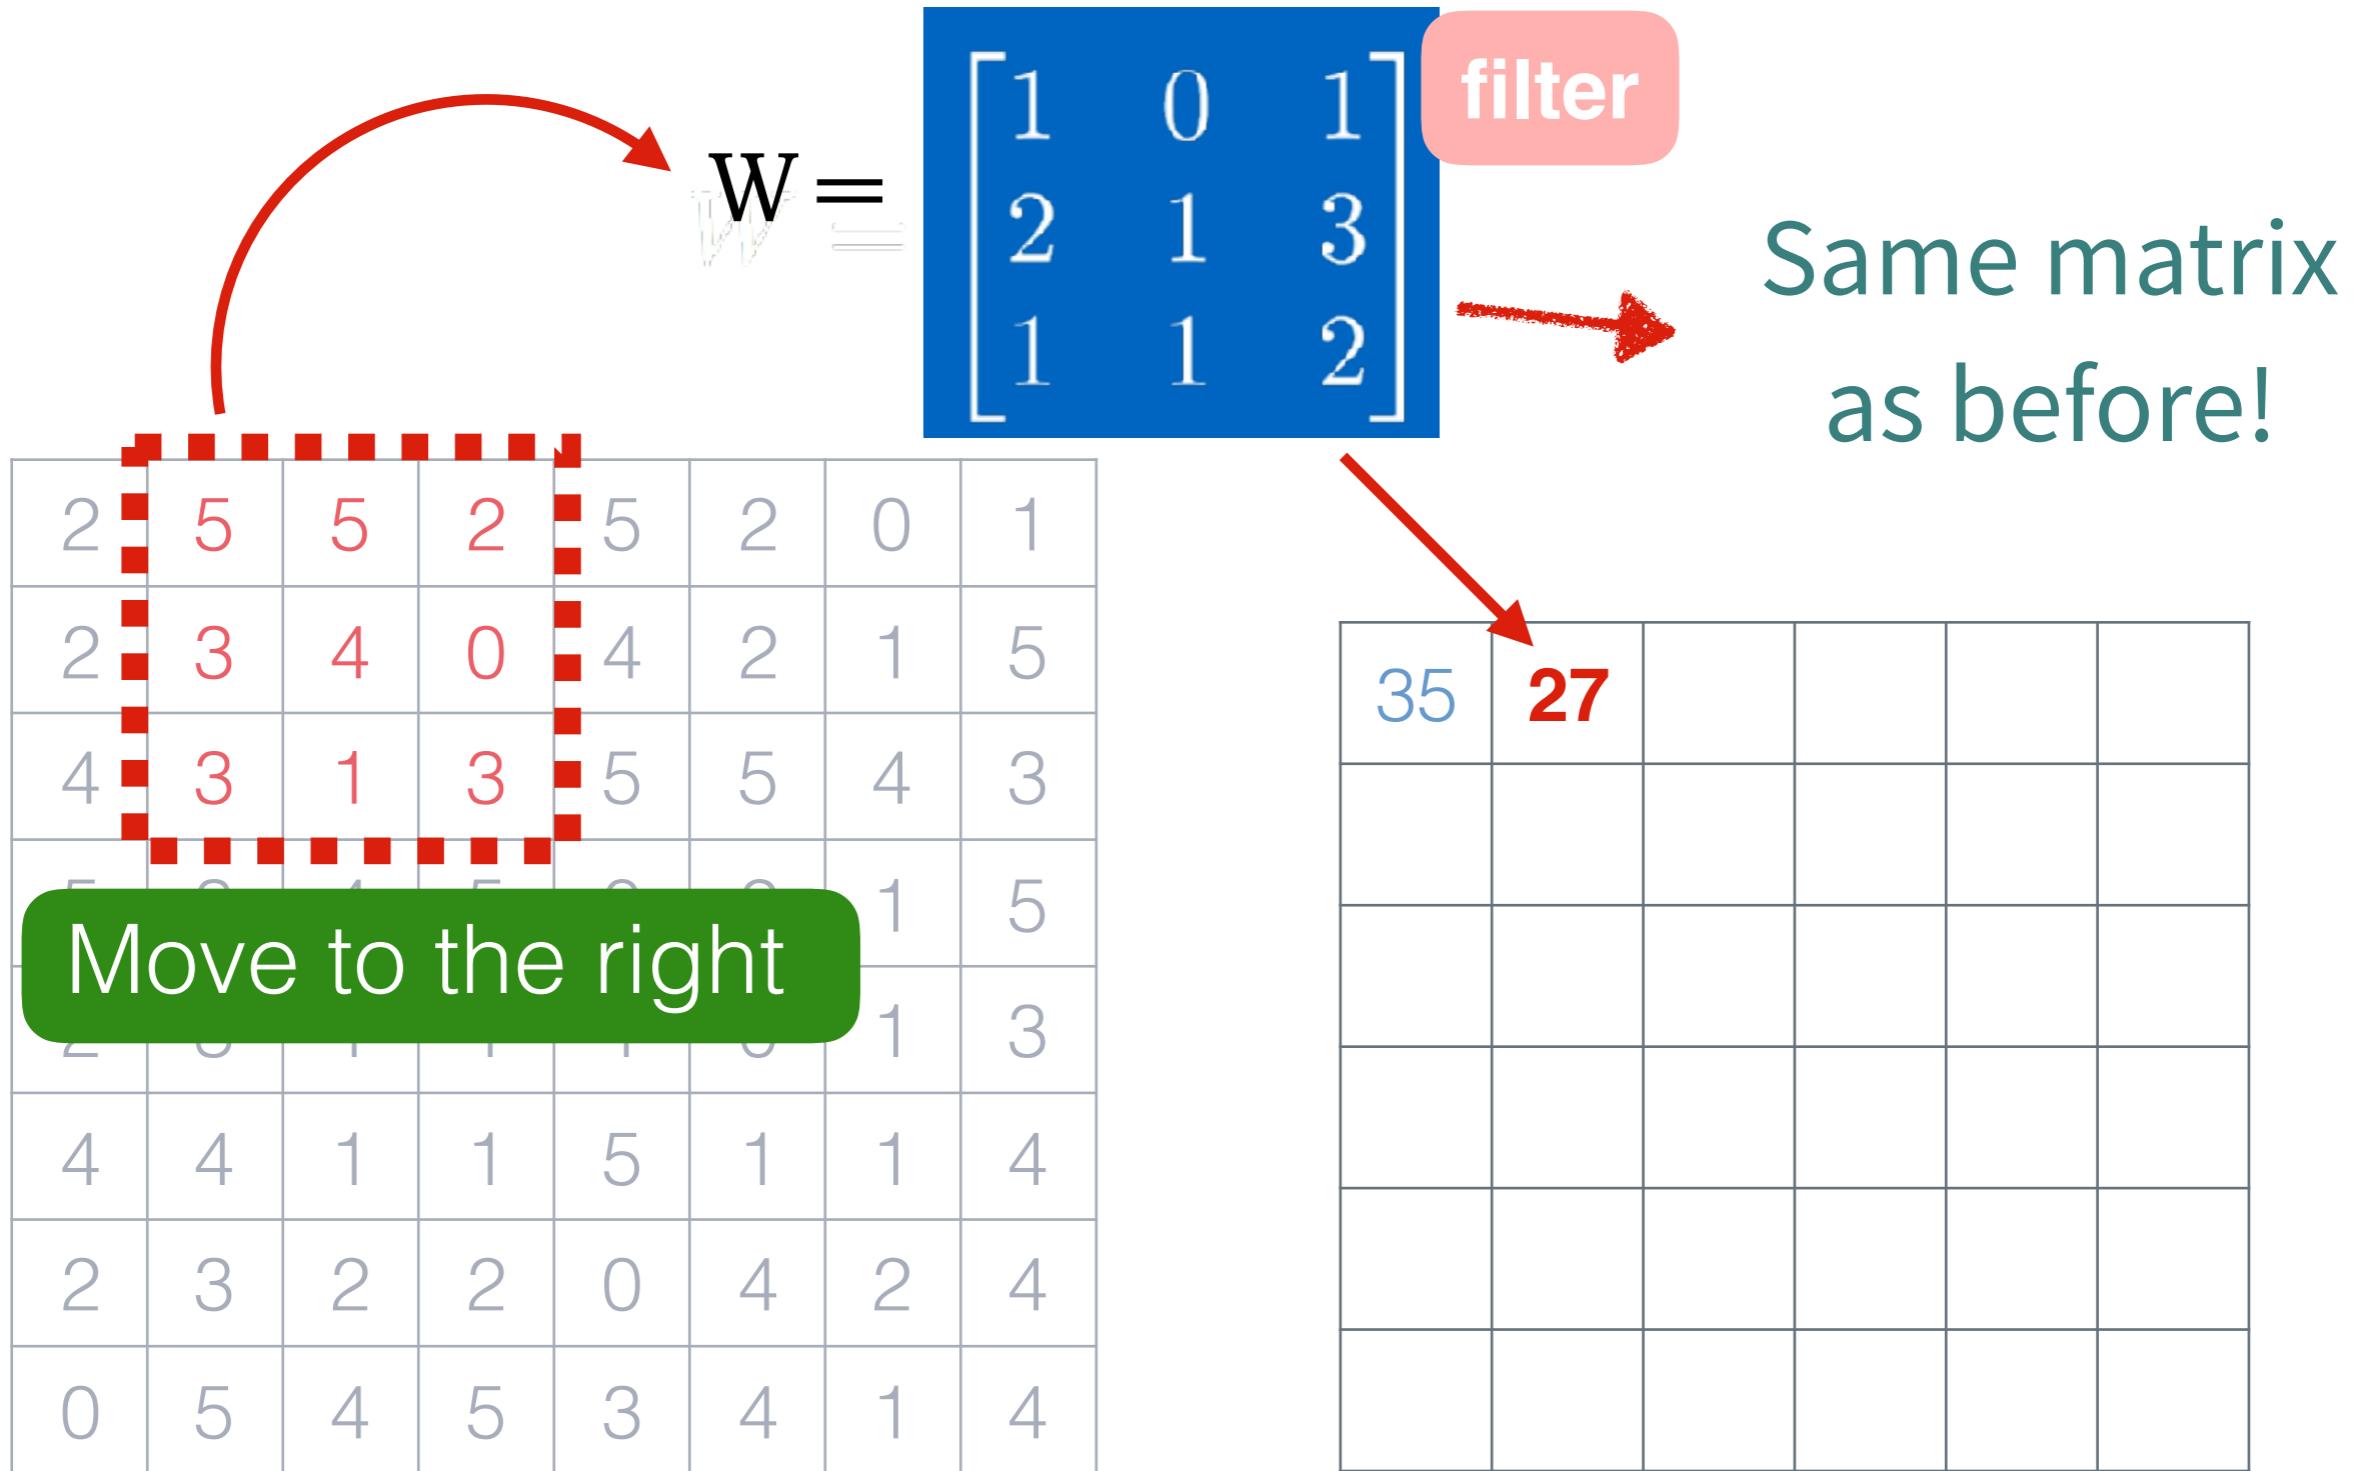
filter

This is learned.

2	5	5	2	5	2	0	1
2	3	4	0	4	2	1	5
4	3	1	3	5	5	4	3
5	3	4	5	0	2	1	5
2	3	1	1	1	0	1	3
4	4	1	1	5	1	1	4
2	3	2	2	0	4	2	4
0	5	4	5	3	4	1	4

35							

Think this is an image.



filter

$$W =$$

$$\begin{bmatrix} 1 & 0 & 1 \\ 2 & 1 & 3 \\ 1 & 1 & 2 \end{bmatrix}$$

2	5	5	2	5	2	0	1
2	3	4	0	4	2	1	5
4	3	1	3	5	5	4	3
5	3	4	5	0	2	1	5
2	3	1	1	1	0	1	3
4	4	1	1	5	1	1	4
2	3	2	2	0	4	2	4
0	5	4	5	3	4	1	4

35	27	44	32	36	38
36	36	37	36	36	43
37	37	23	26	17	35
29	25	22	18	14	27
27	25	24	21	24	32
31	38	27	34	25	40

All the way to the end

Neurons look like this

The points on the picture are one input layer neuron

$$W = \begin{bmatrix} 1 & 0 & 1 \\ 2 & 1 & 3 \\ 1 & 1 & 2 \end{bmatrix}$$

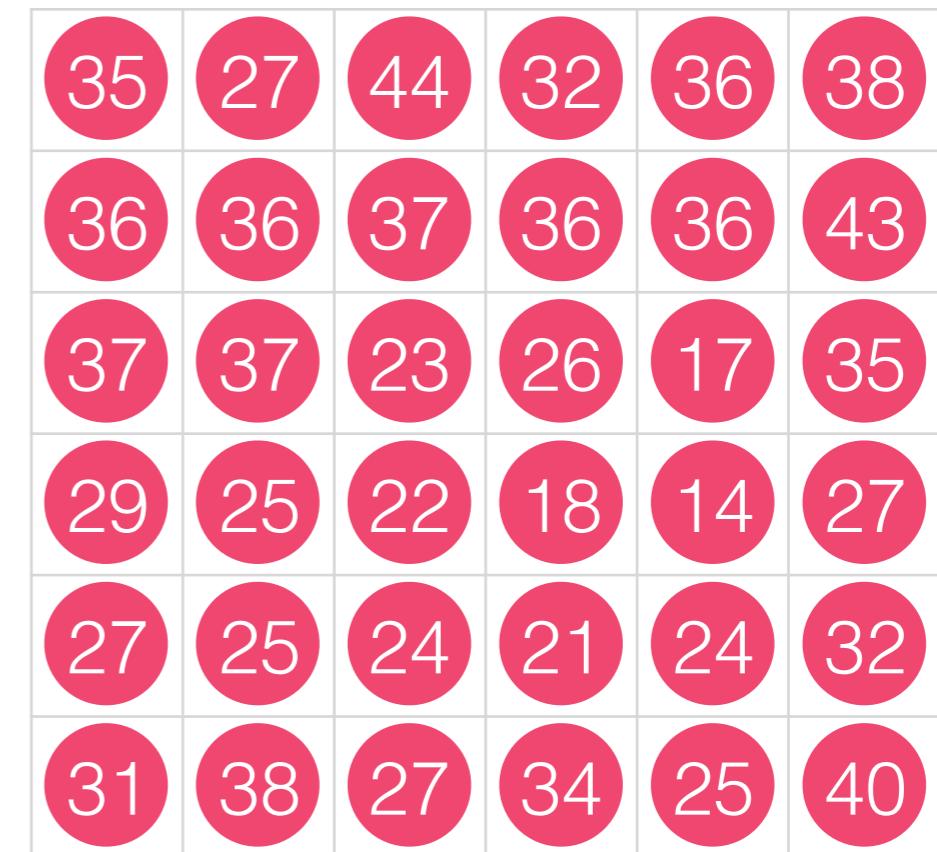
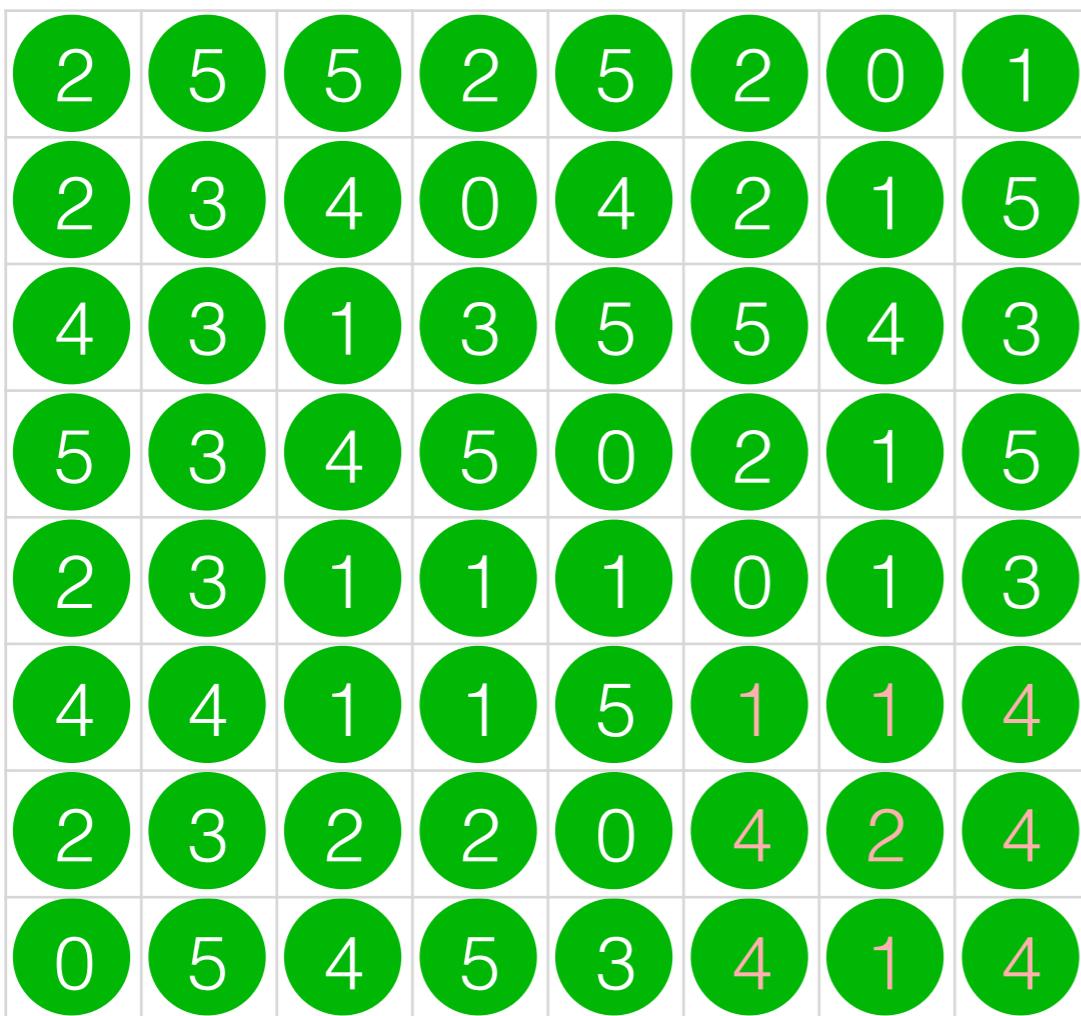
filter

2	5	5	2	5	2	0	1
2	3	4	0	4	2	1	5
4	3	1	3	5	5	4	3
5	3	4	5	0	2	1	5
2	3	1	1	1	0	1	3
4	4	1	1	5	1	1	4
2	3	2	2	0	4	2	4
0	5	4	5	3	4	1	4

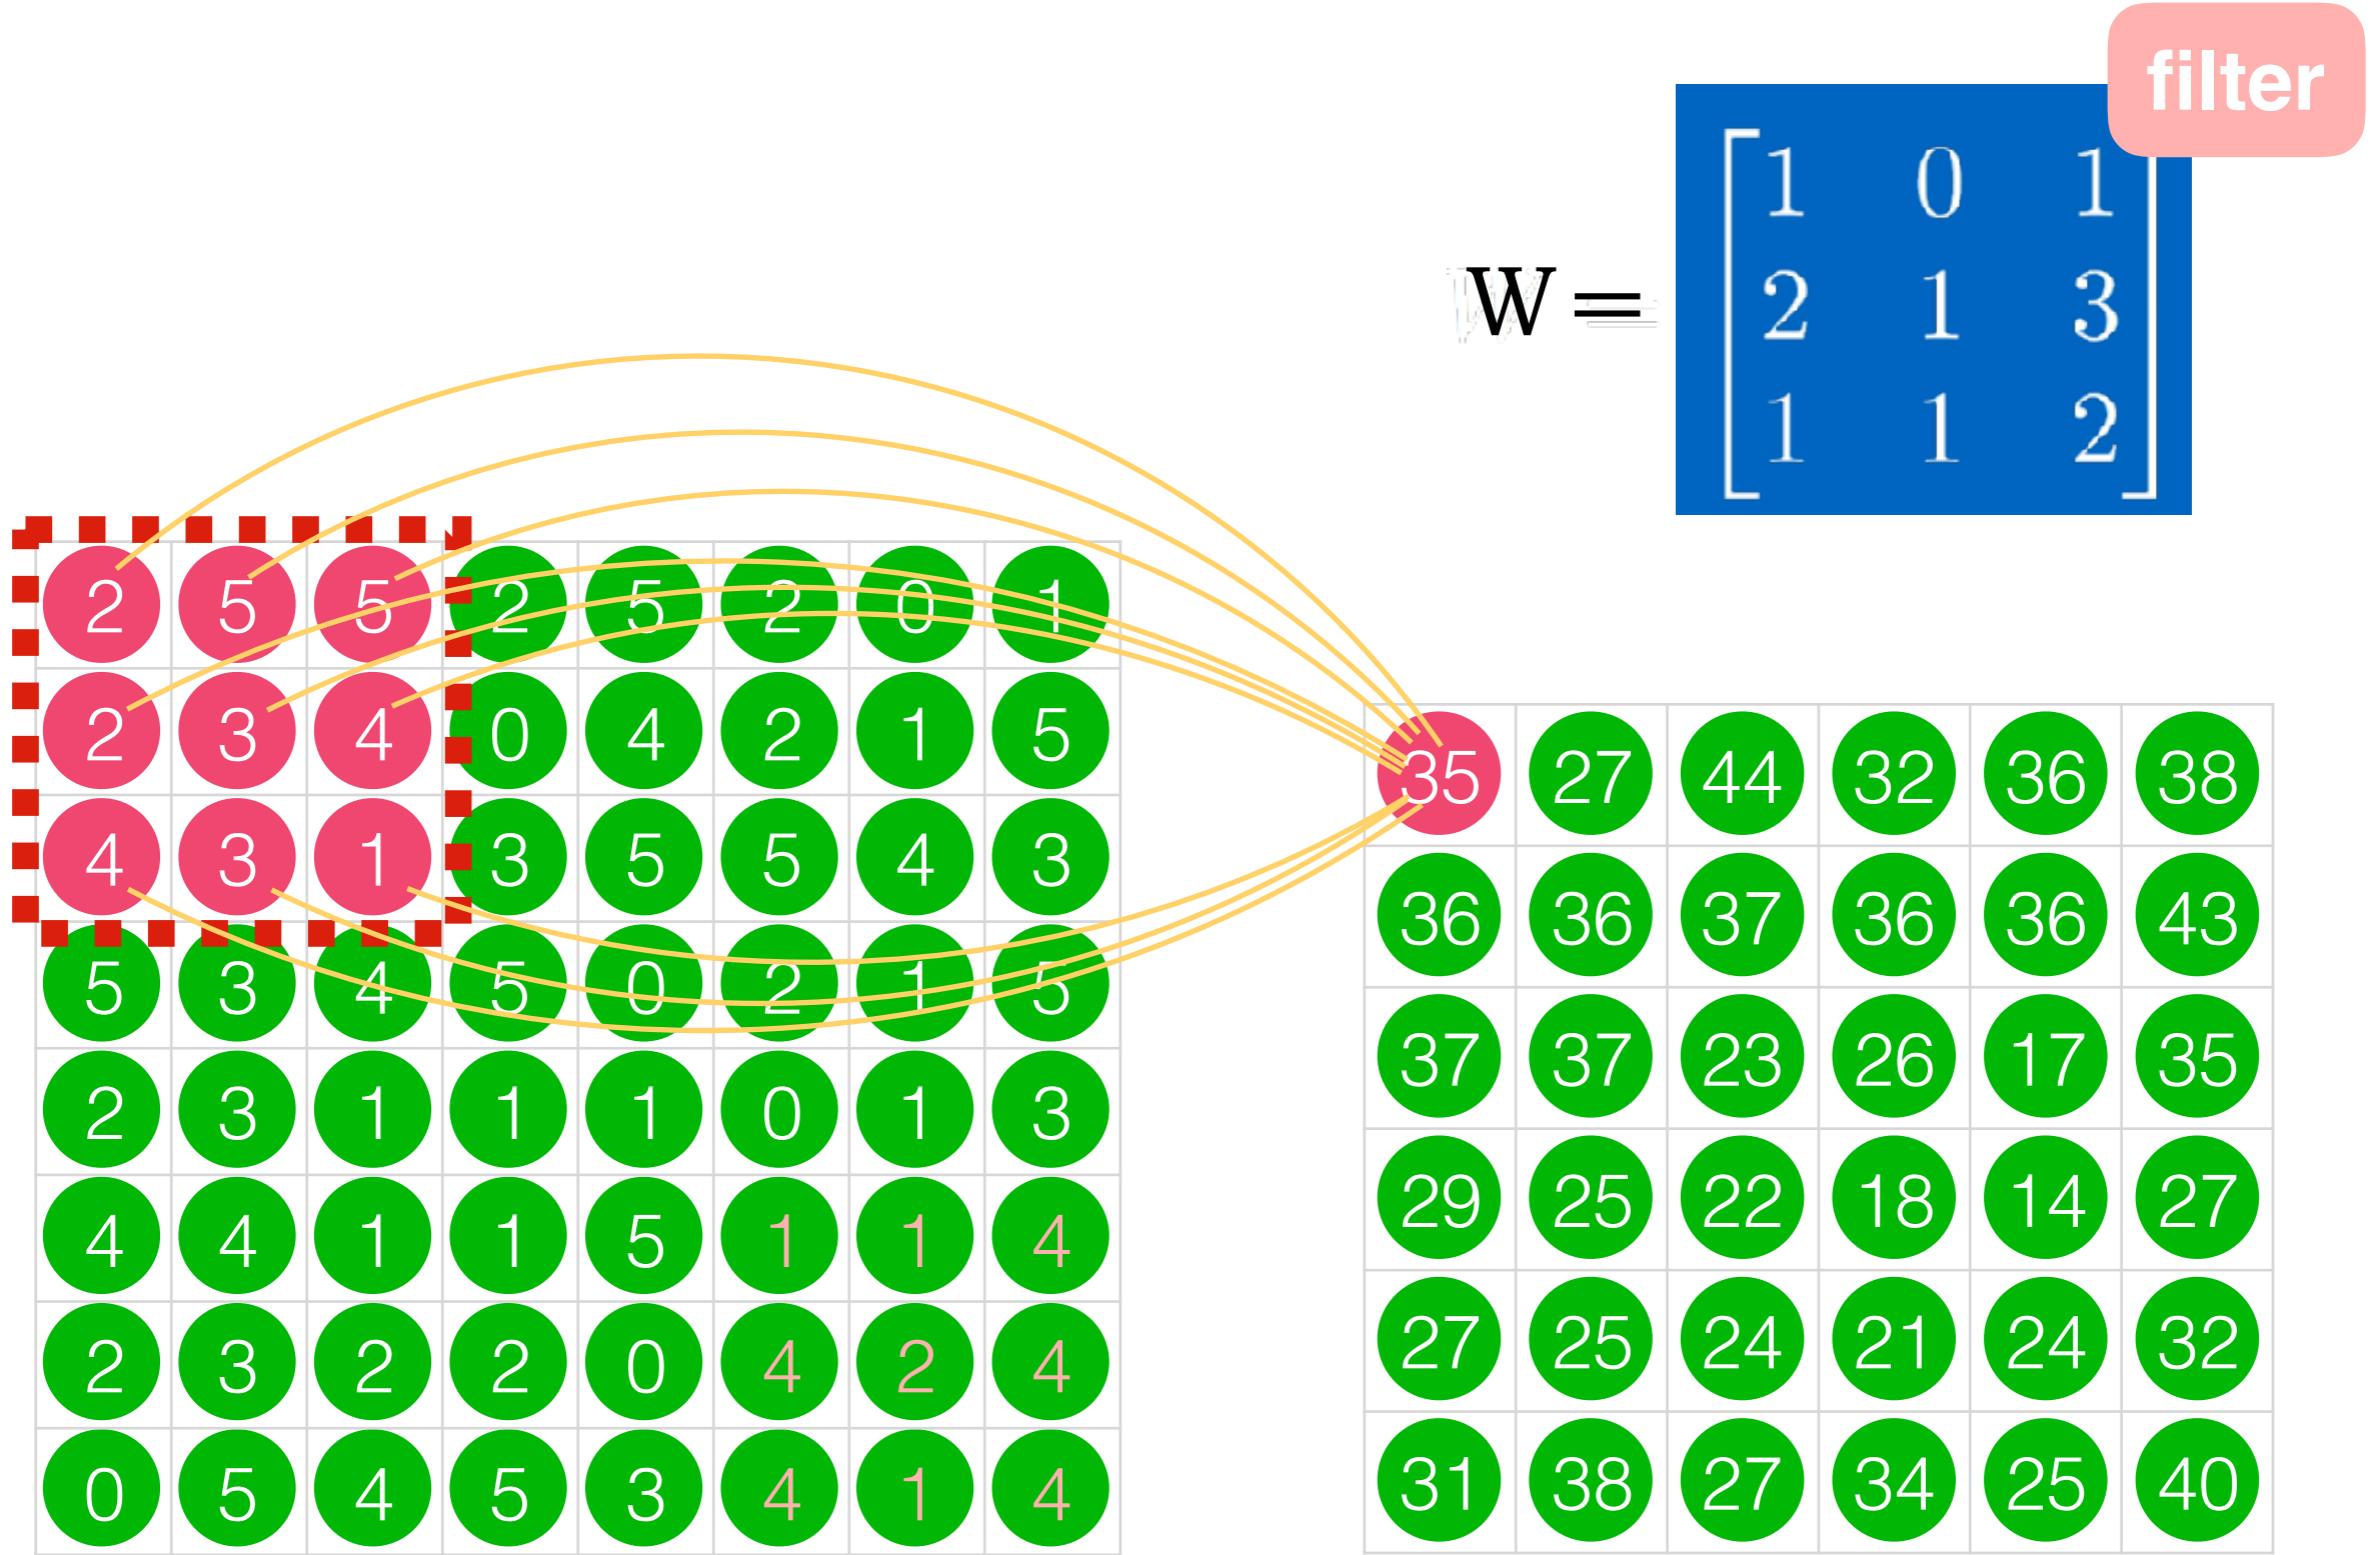
35	27	44	32	36	38
36	36	37	36	36	43
37	37	23	26	17	35
29	25	22	18	14	27
27	25	24	21	24	32
31	38	27	34	25	40

filter

$$W = \begin{bmatrix} 1 & 0 & 1 \\ 2 & 1 & 3 \\ 1 & 1 & 2 \end{bmatrix}$$



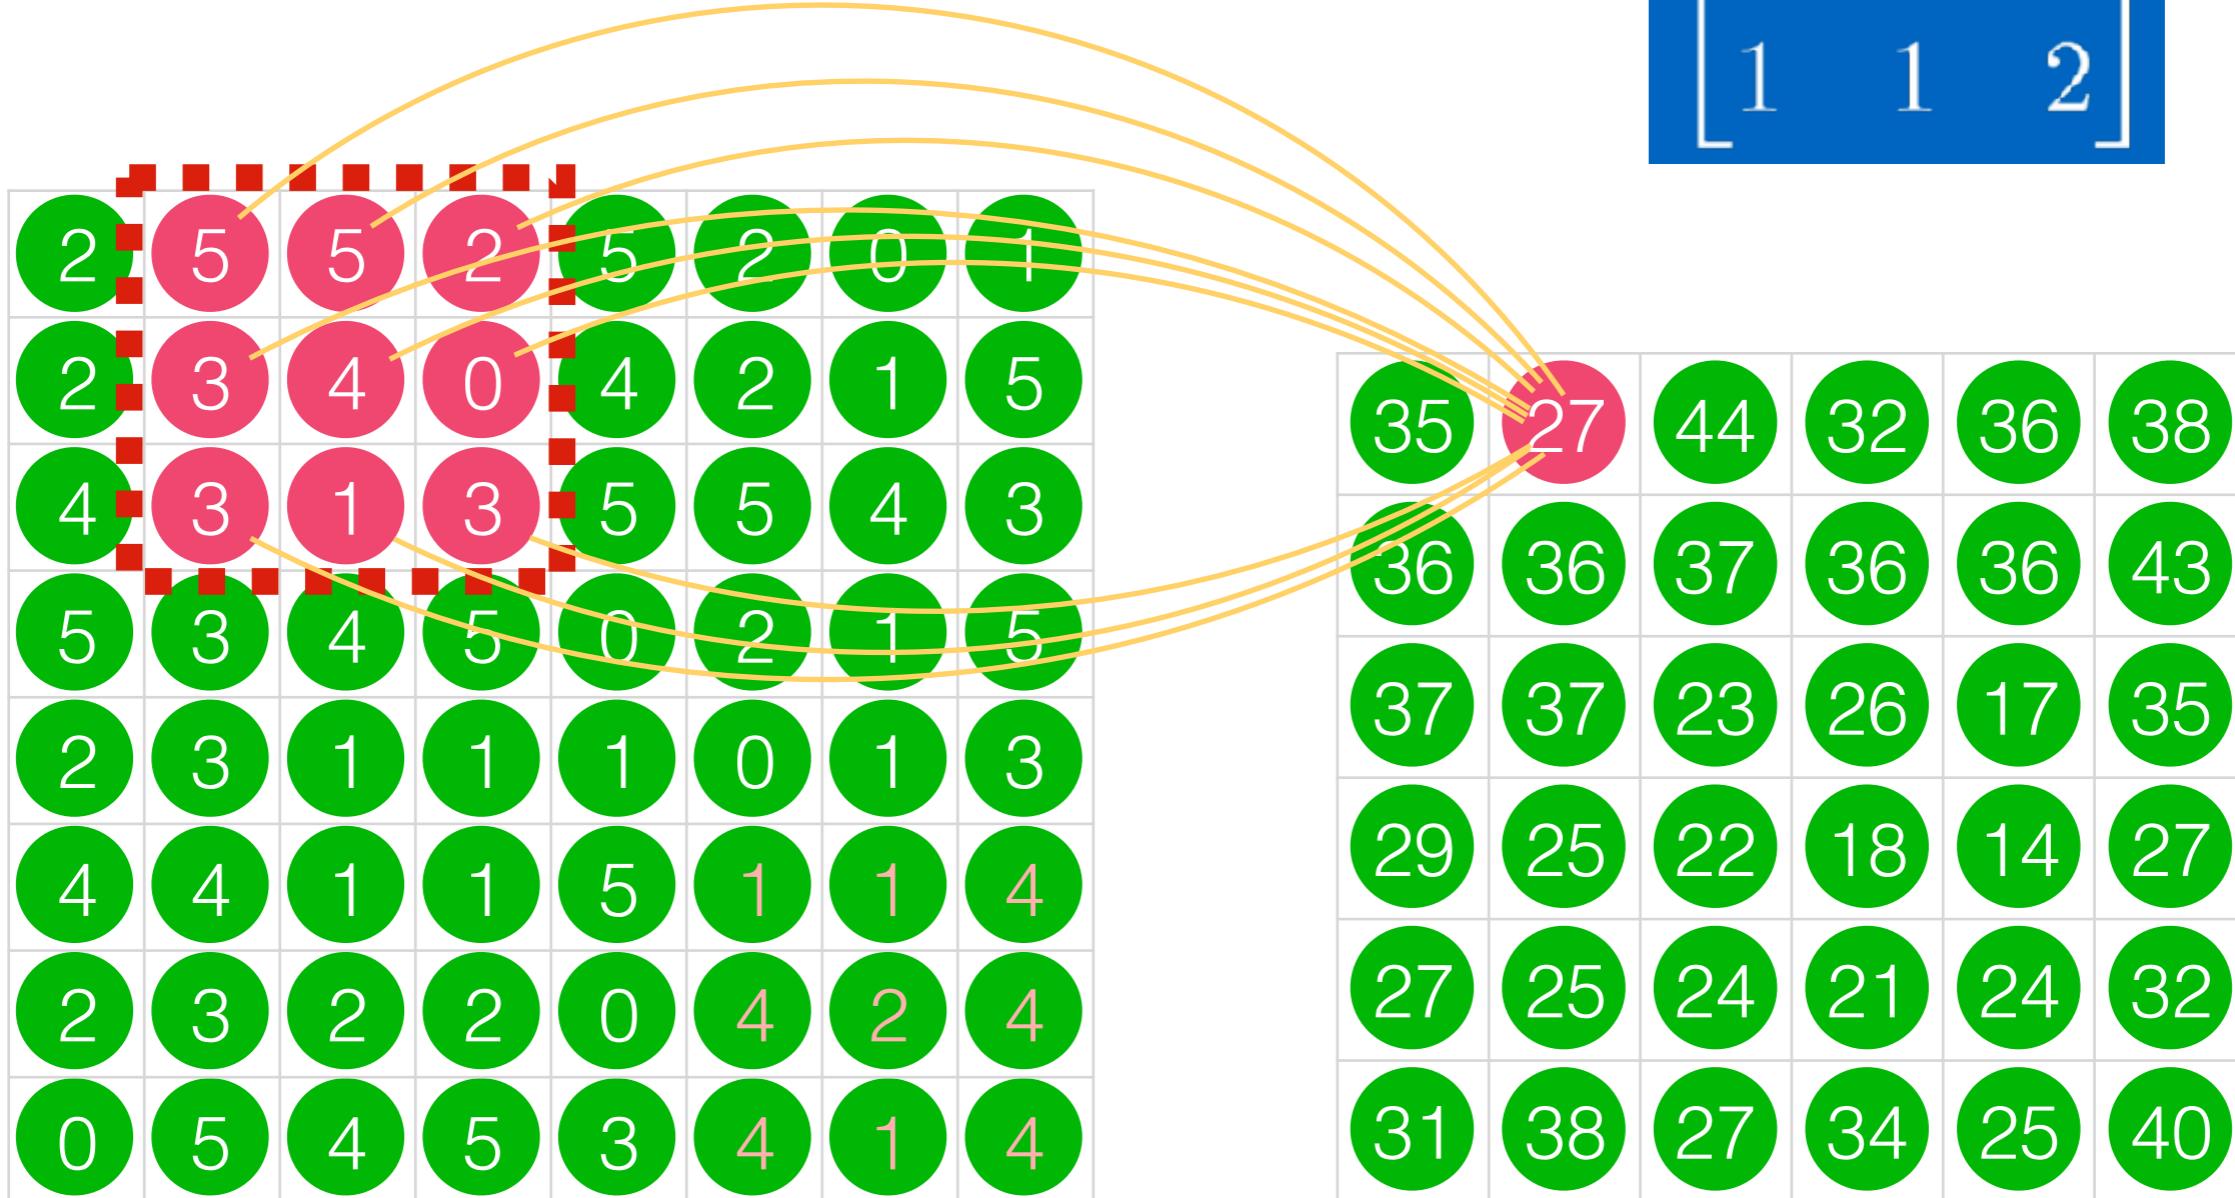
Conv layers also consist
lots of neural



The two layers are not completely connected

filter

$$W = \begin{bmatrix} 1 & 0 & 1 \\ 2 & 1 & 3 \\ 1 & 1 & 2 \end{bmatrix}$$



the weights are shared
(same as previous ones)

Finally, we get a 6x6 matrix,
and actually we usually
make it a 8x8 matrix.

And we have a lot of filters!

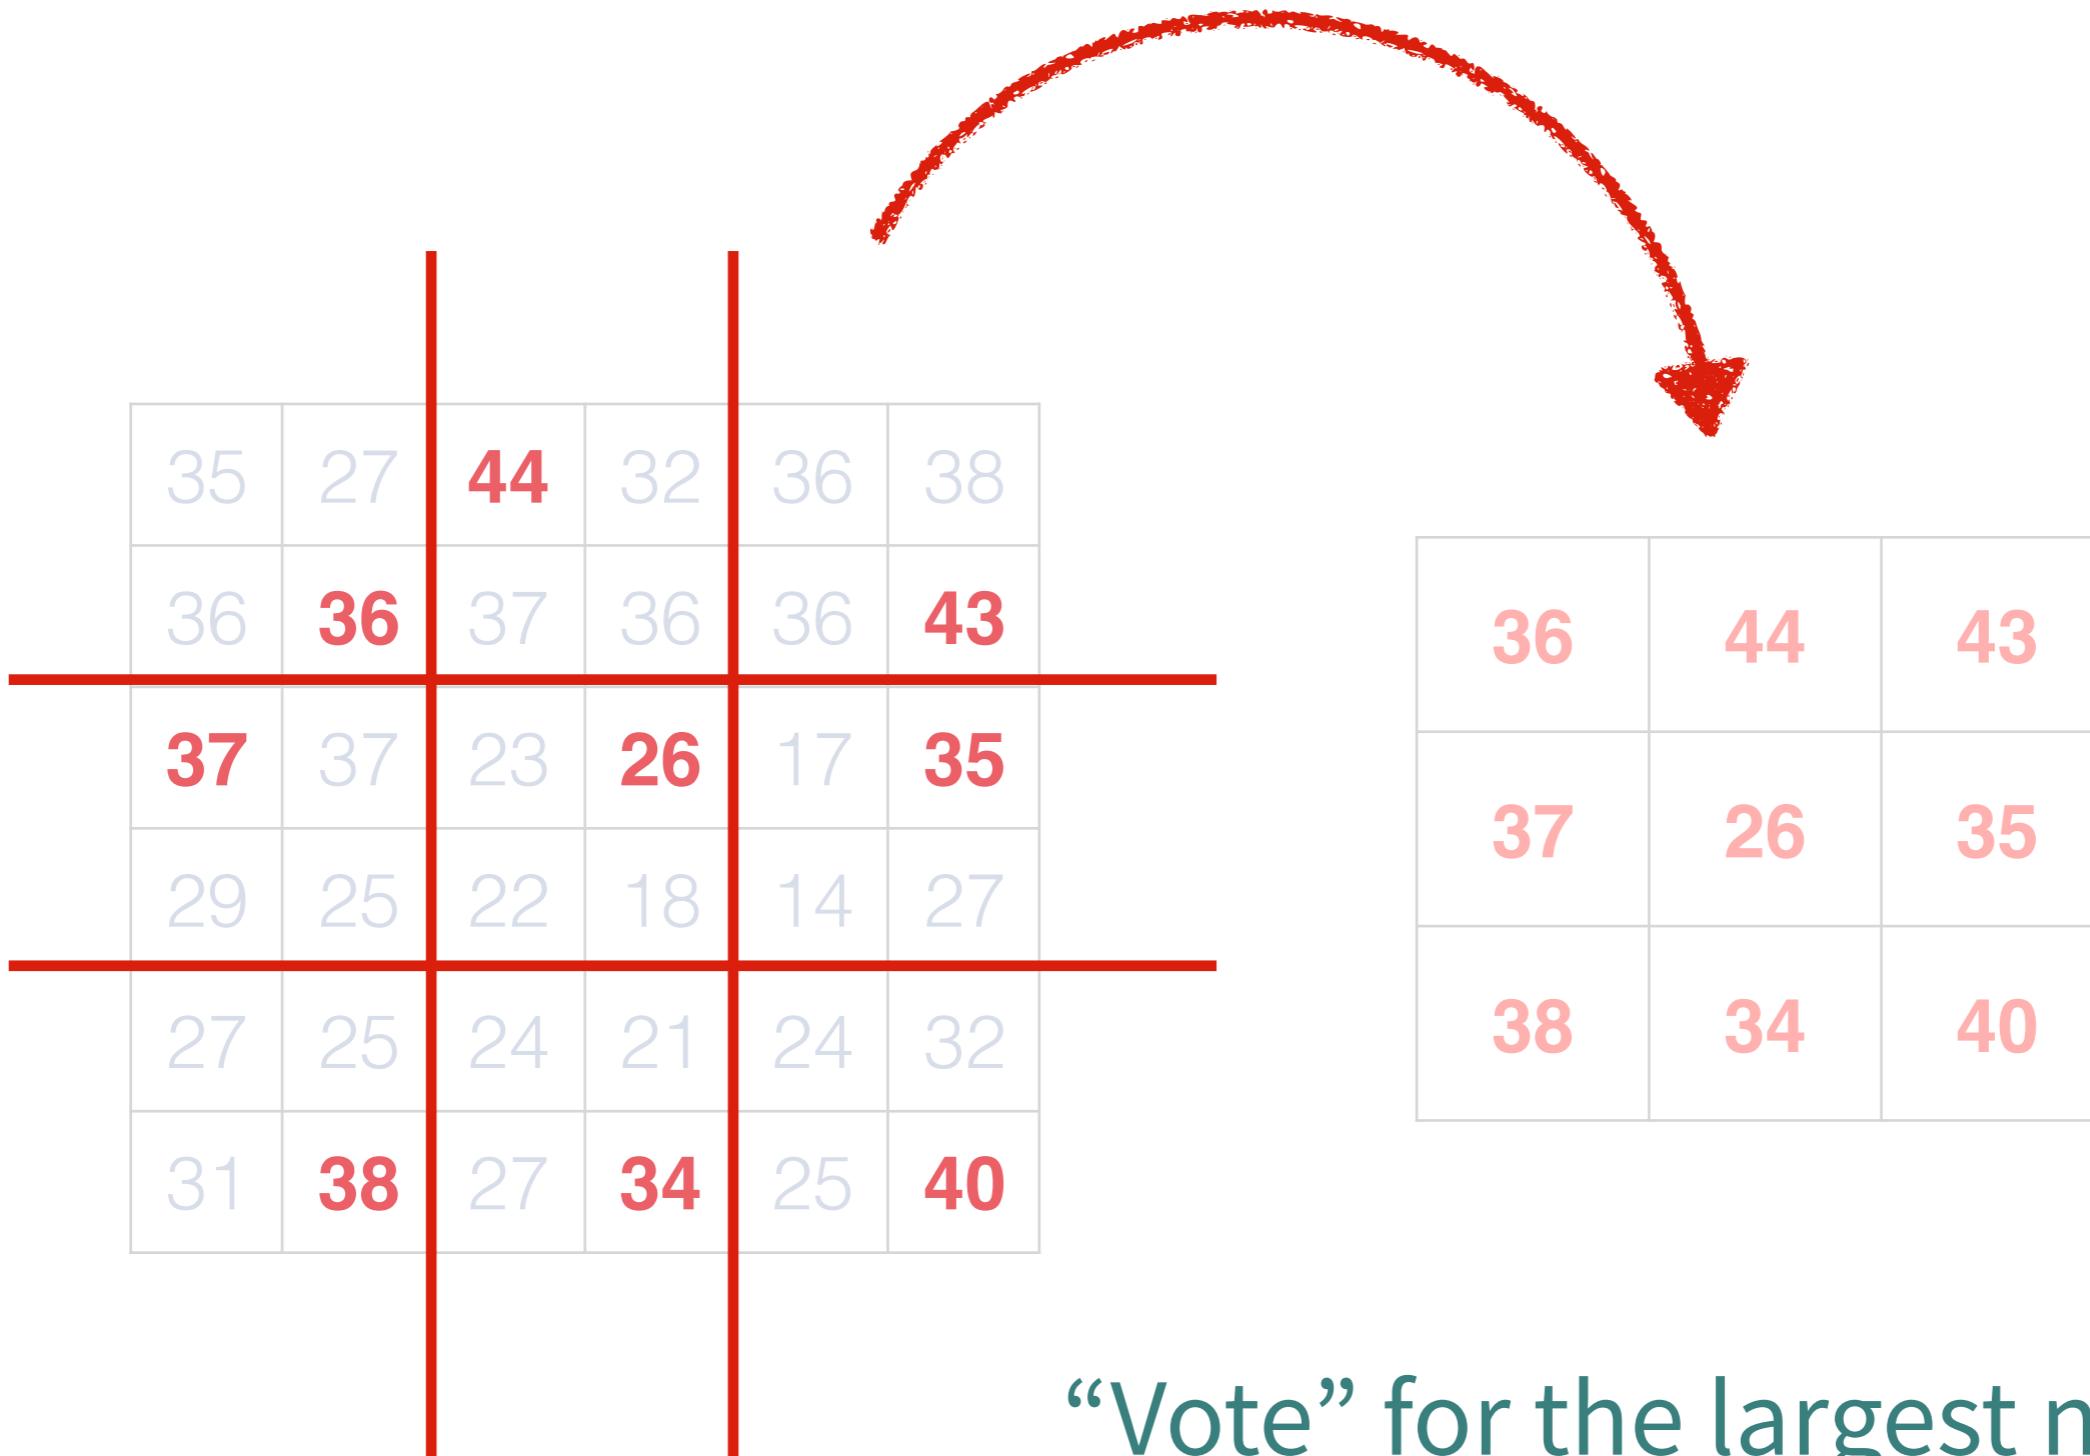
35	27	44	32	36	38
36	36	37	36	36	43
37	37	23	26	17	35
29	25	22	18	14	27
27	25	24	21	24	32
31	38	27	34	25	40

2

Max-Pooling Layer

Basically it is “voting”

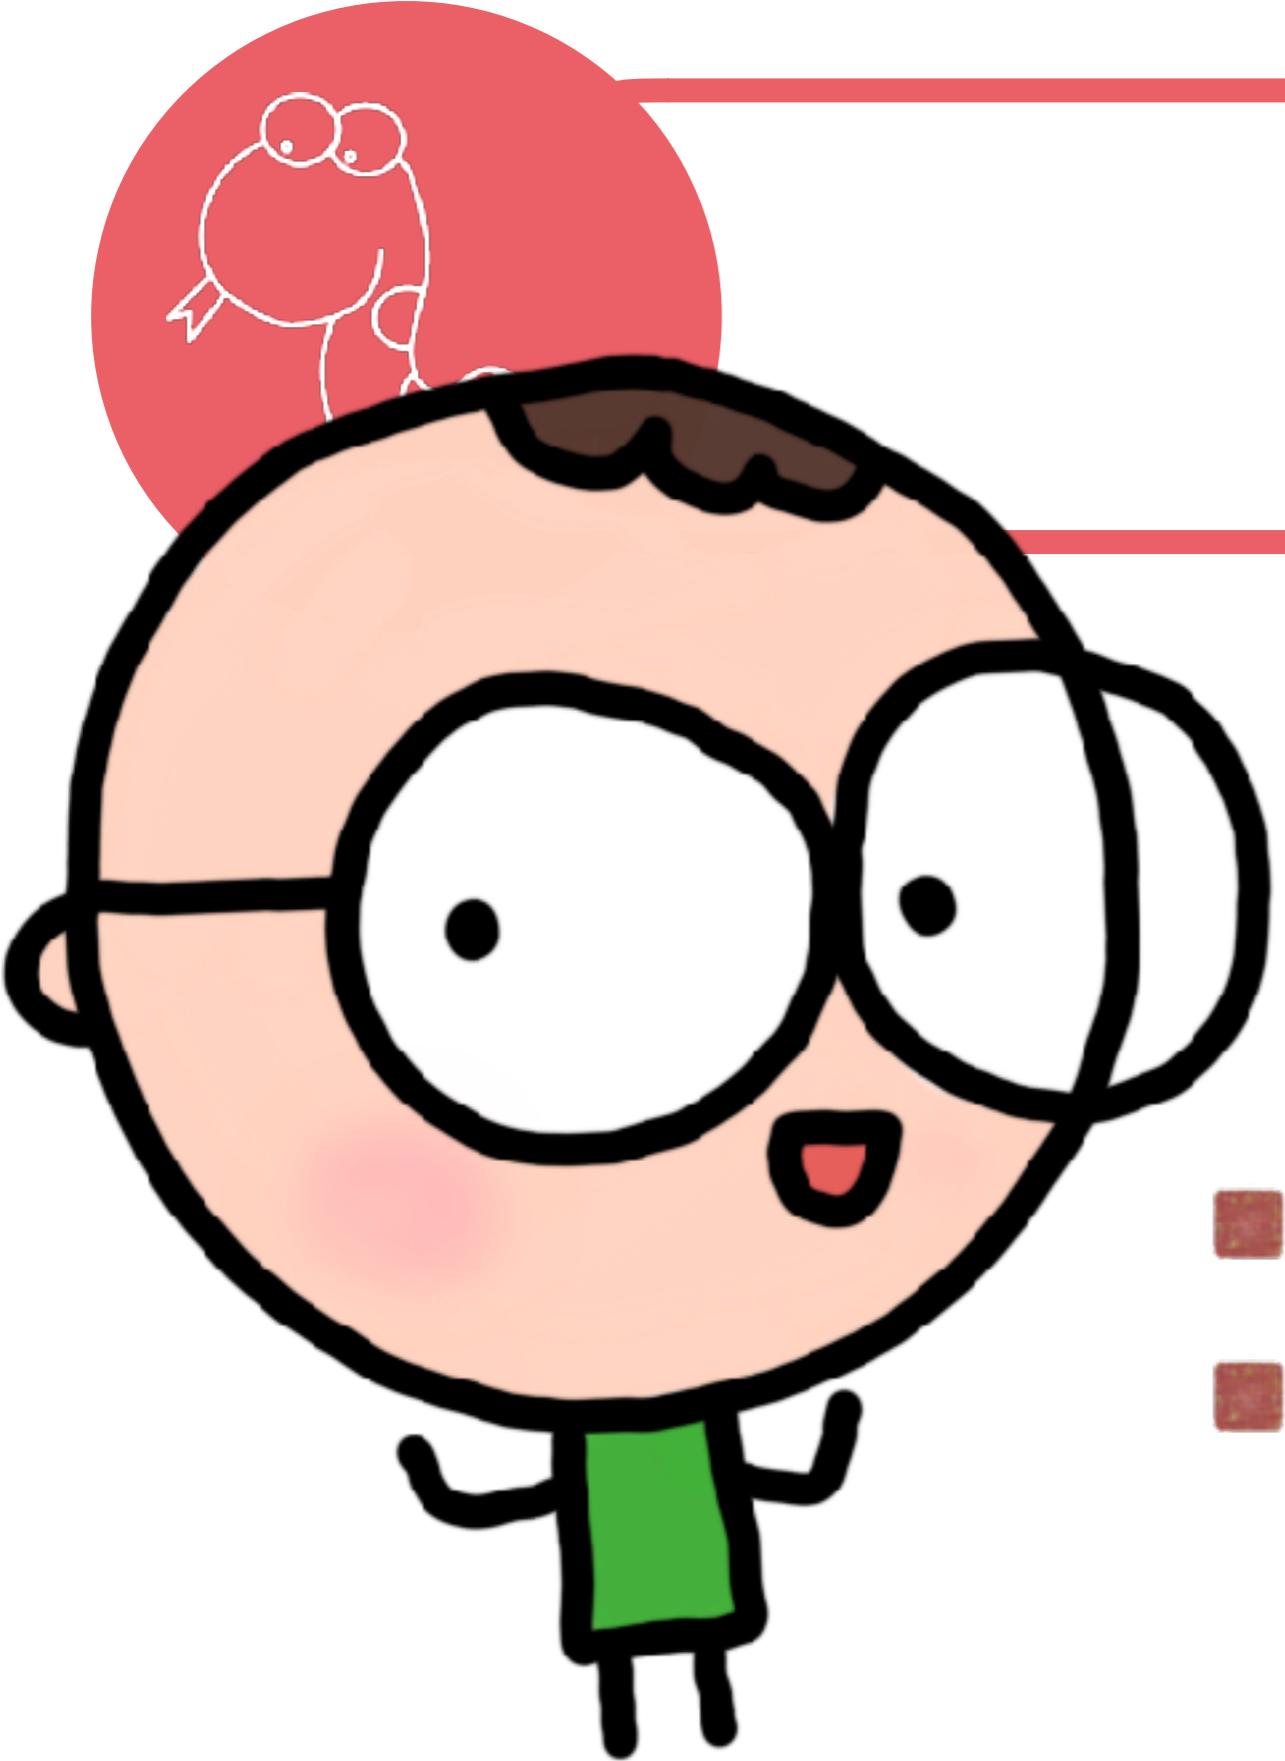
We have to decide the size of area to have one representative. Say, 2x2.



We can repeat and repeat...

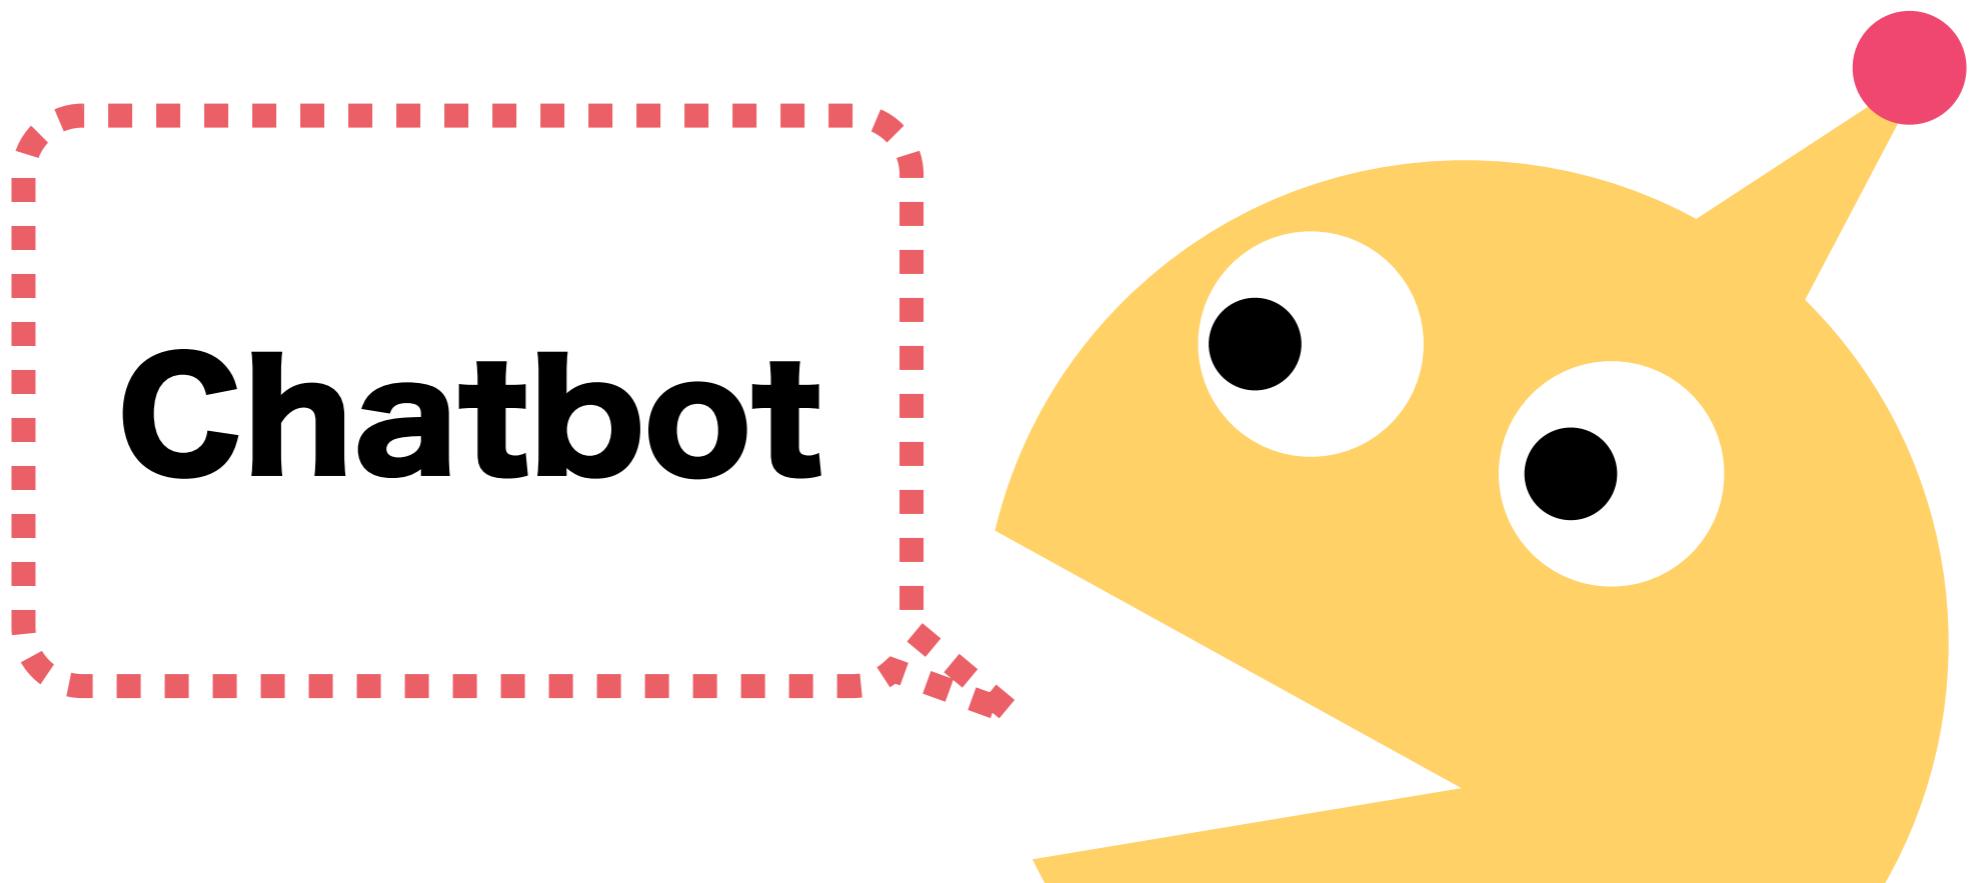
**convolution, max-pooling, convolution,
max-pooling...**

After we finished, we sent outputs to “normal” neural networks.

A cartoon illustration of a character with large, round, white eyes and black pupils. The character wears black-rimmed glasses and a red baseball cap. A small white bird with large eyes is perched on the rim of the cap.

RNN

- Recurrent Neural Network
- Neural Network with memory

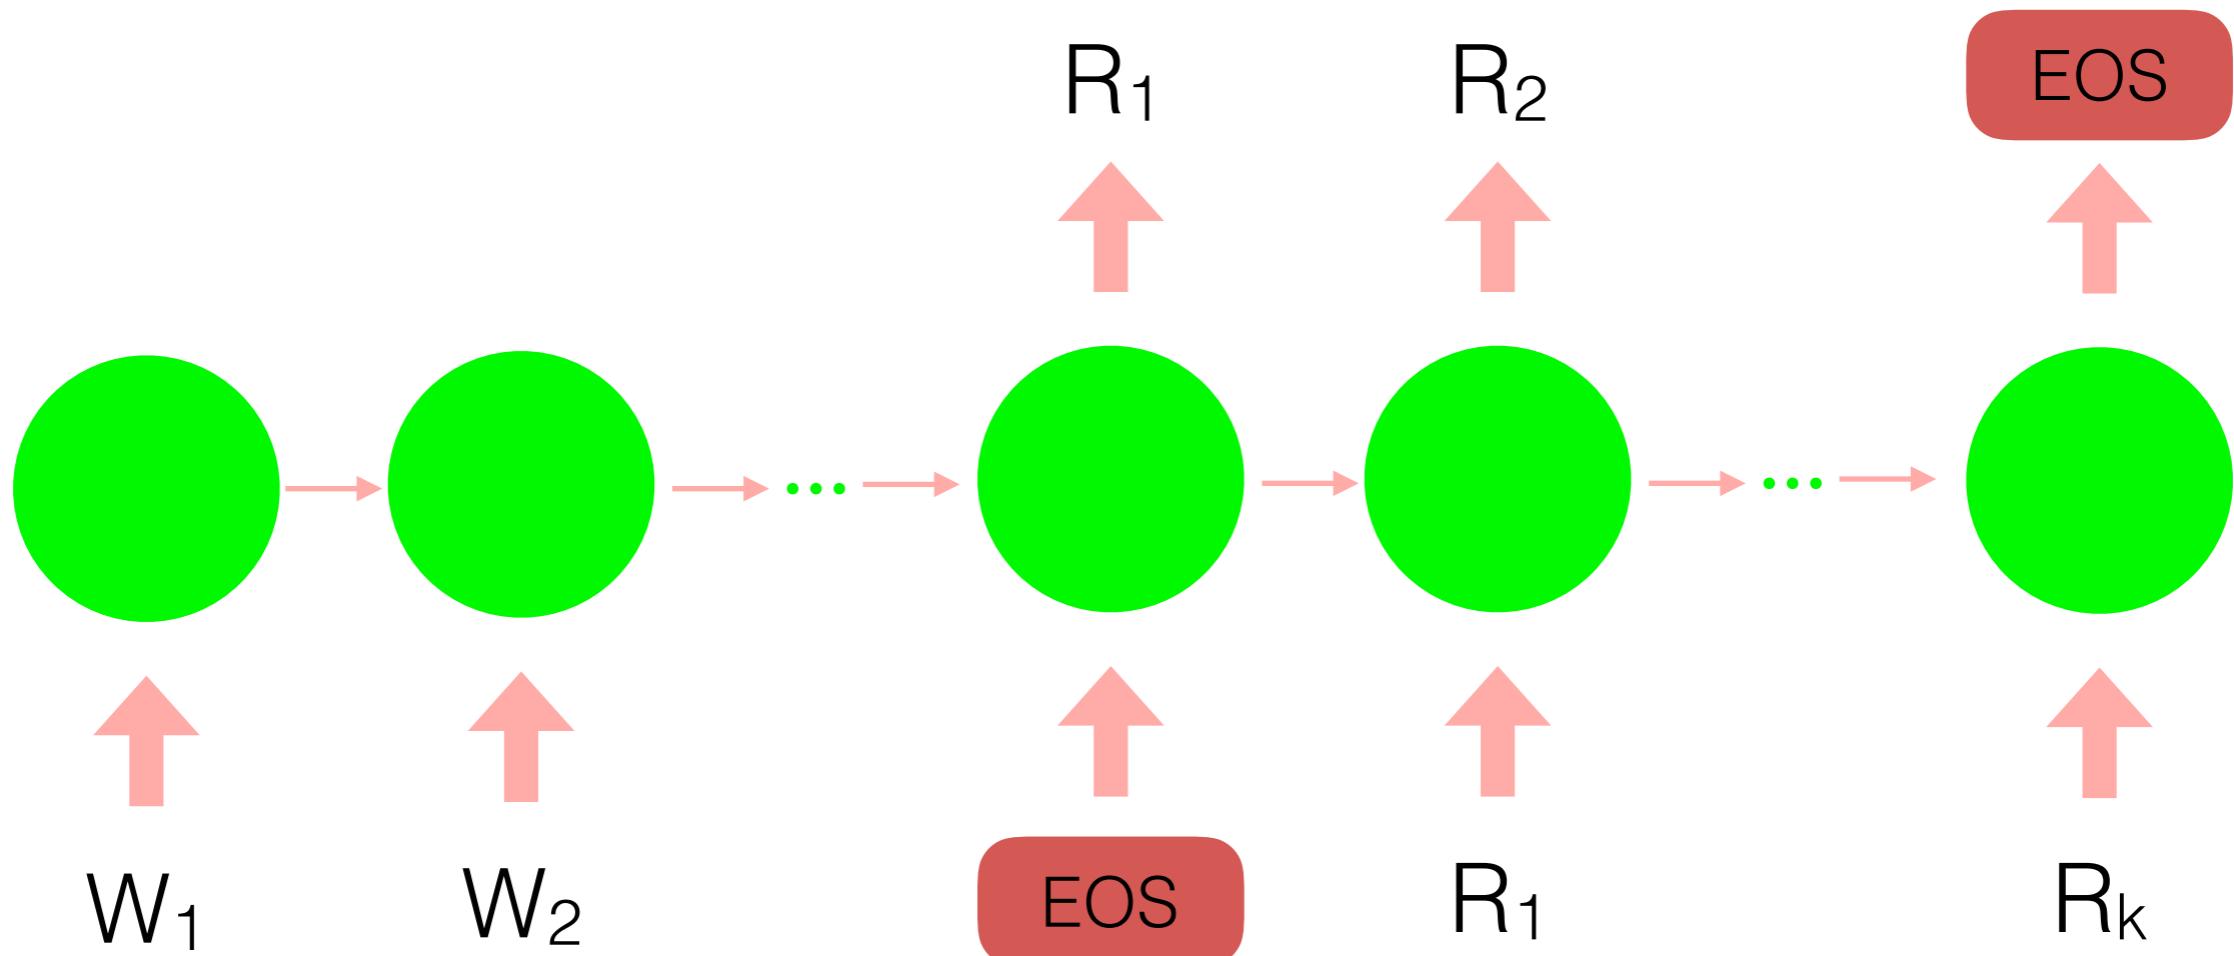


$f(\text{Current word}) = \text{Next word}$

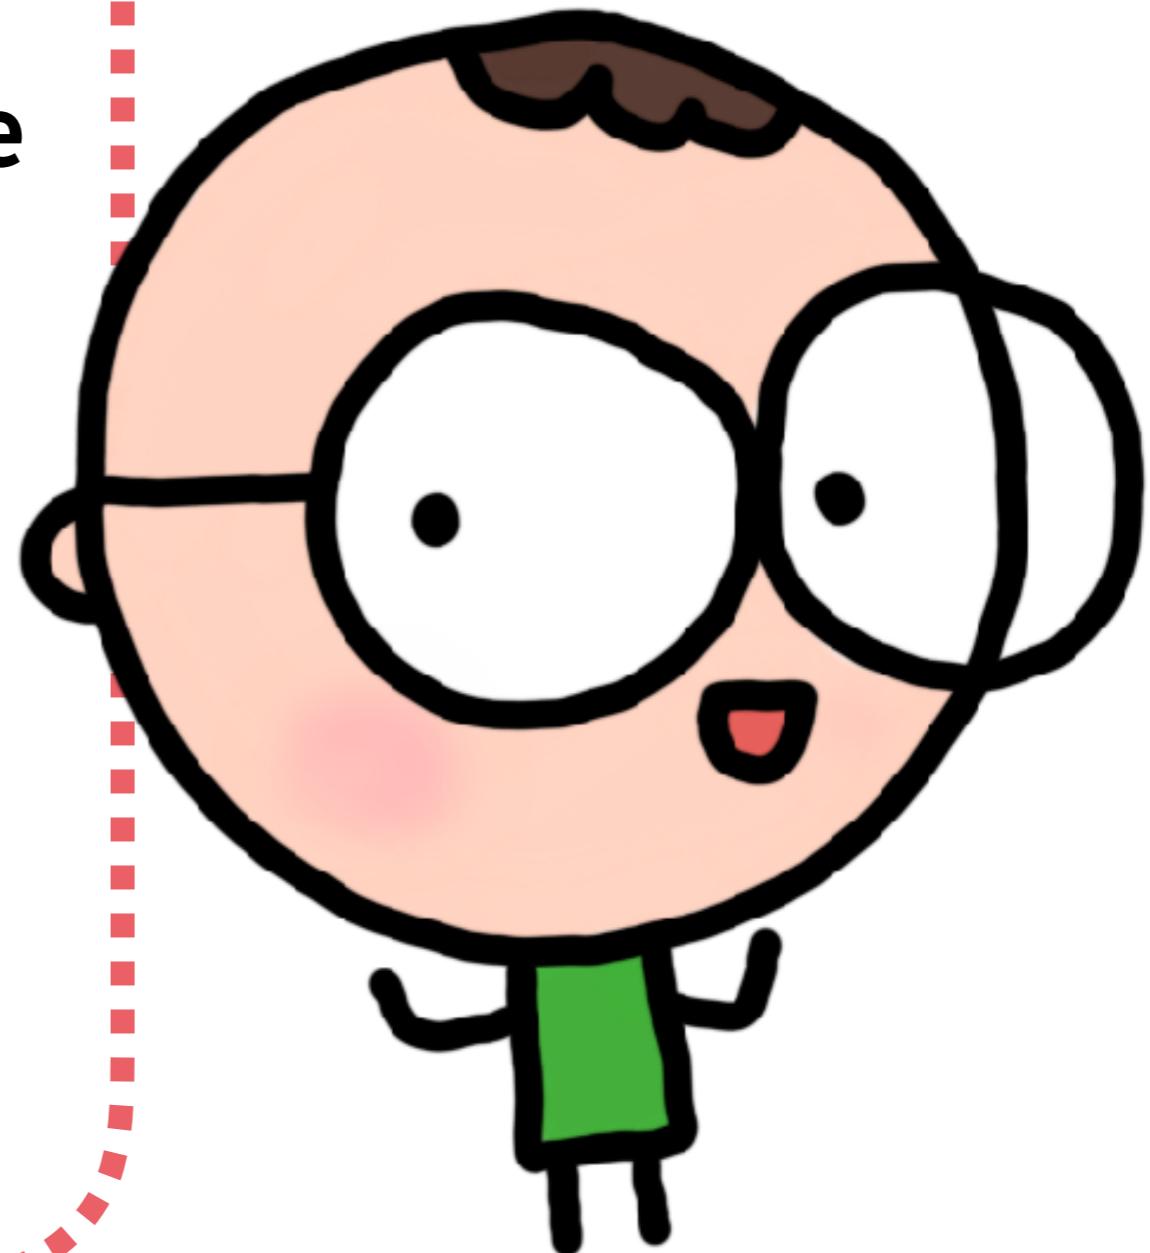
Current word →



→ **Next word**



In fact, the input does not have to be text, but it is also possible that the video (one by one) is output! The output can still be text, and the most common is probably to let the computer say what happened in the video.



Applications

- Translator
- Video Captioning Generator
- Context Generator
- Drawing



To prove study we see that $\mathcal{F}|_U$ is a covering of \mathcal{X}' , and \mathcal{T}_i is an object of $\mathcal{F}_{X/S}$ for $i > 0$ and \mathcal{F}_p exists and let \mathcal{F}_i be a presheaf of \mathcal{O}_X -modules on \mathcal{C} as a \mathcal{F} -module. In particular $\mathcal{F} = U/\mathcal{F}$ we have to show that

$$\widetilde{M}^\bullet = \mathcal{I}^\bullet \otimes_{\text{Spec}(k)} \mathcal{O}_{S,s} - i_X^{-1} \mathcal{F})$$

is a unique morphism of algebraic stacks. Note that

$$\text{Arrows} = (\text{Sch}/S)_{fppf}^{\text{opp}}, (\text{Sch}/S)_{fppf}$$

and

$$V = \Gamma(S, \mathcal{O}) \longmapsto (U, \text{Spec}(A))$$

is an open subset of X . Thus U is affine. This is a continuous map of X is the inverse, the groupoid scheme S .

Proof. See discussion of sheaves of sets. □

Andrej Karpathy use RNN to generate a book in “Stacks” (a deep topic in Algebraic Geometry)

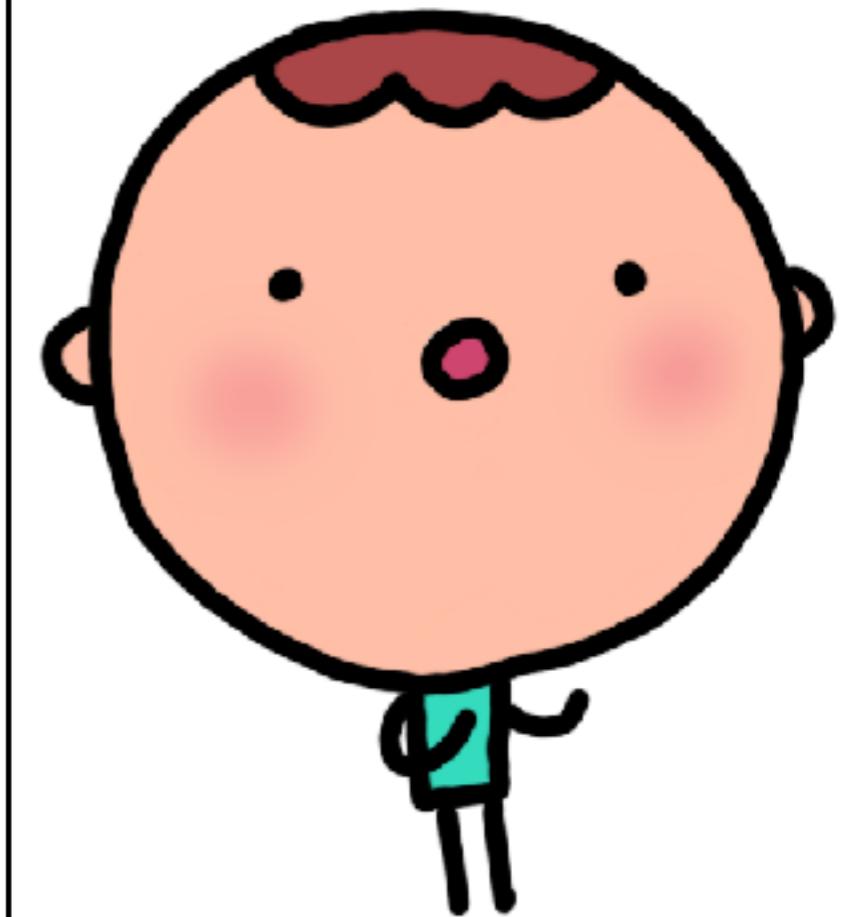
<http://karpathy.github.io/2015/05/21/rnn-effectiveness/>

PANDARUS:

Alas, I think he shall be come
approached and the day
When little strain would be attain'd into
being never fed,
And who is but a chain and subjects of
his death,
I should not sleep.

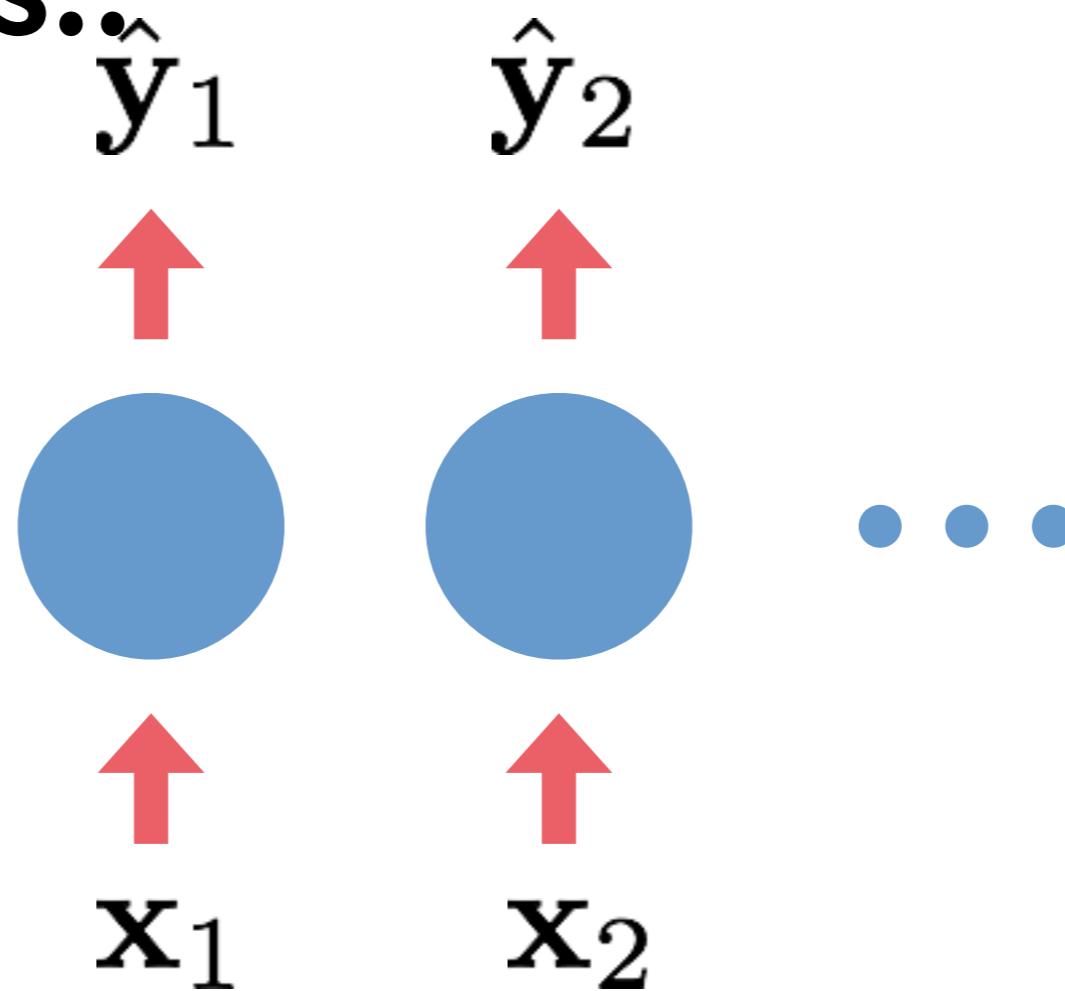
Second Senator:

They are away this miseries, produced
upon my soul,
Breaking and strongly should be
buried, when I perish
The earth and thoughts of many states.

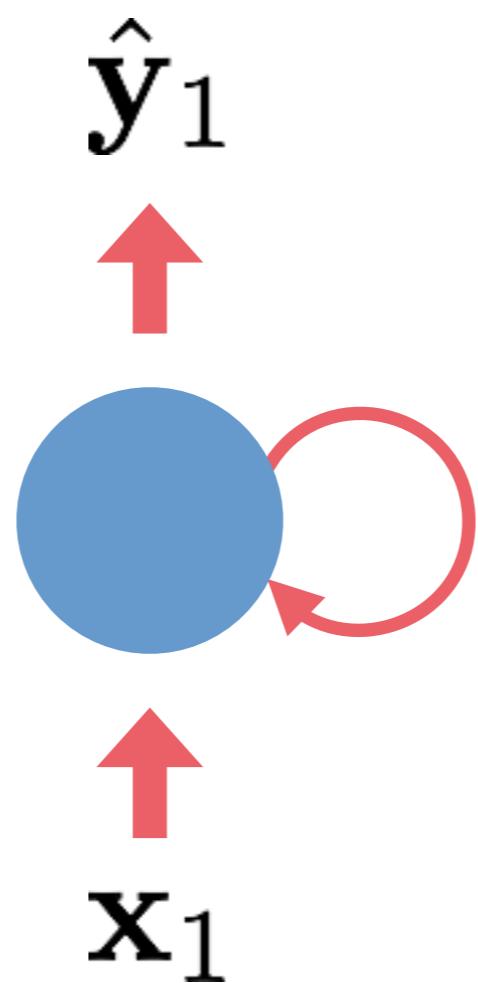


“Shakespeare”

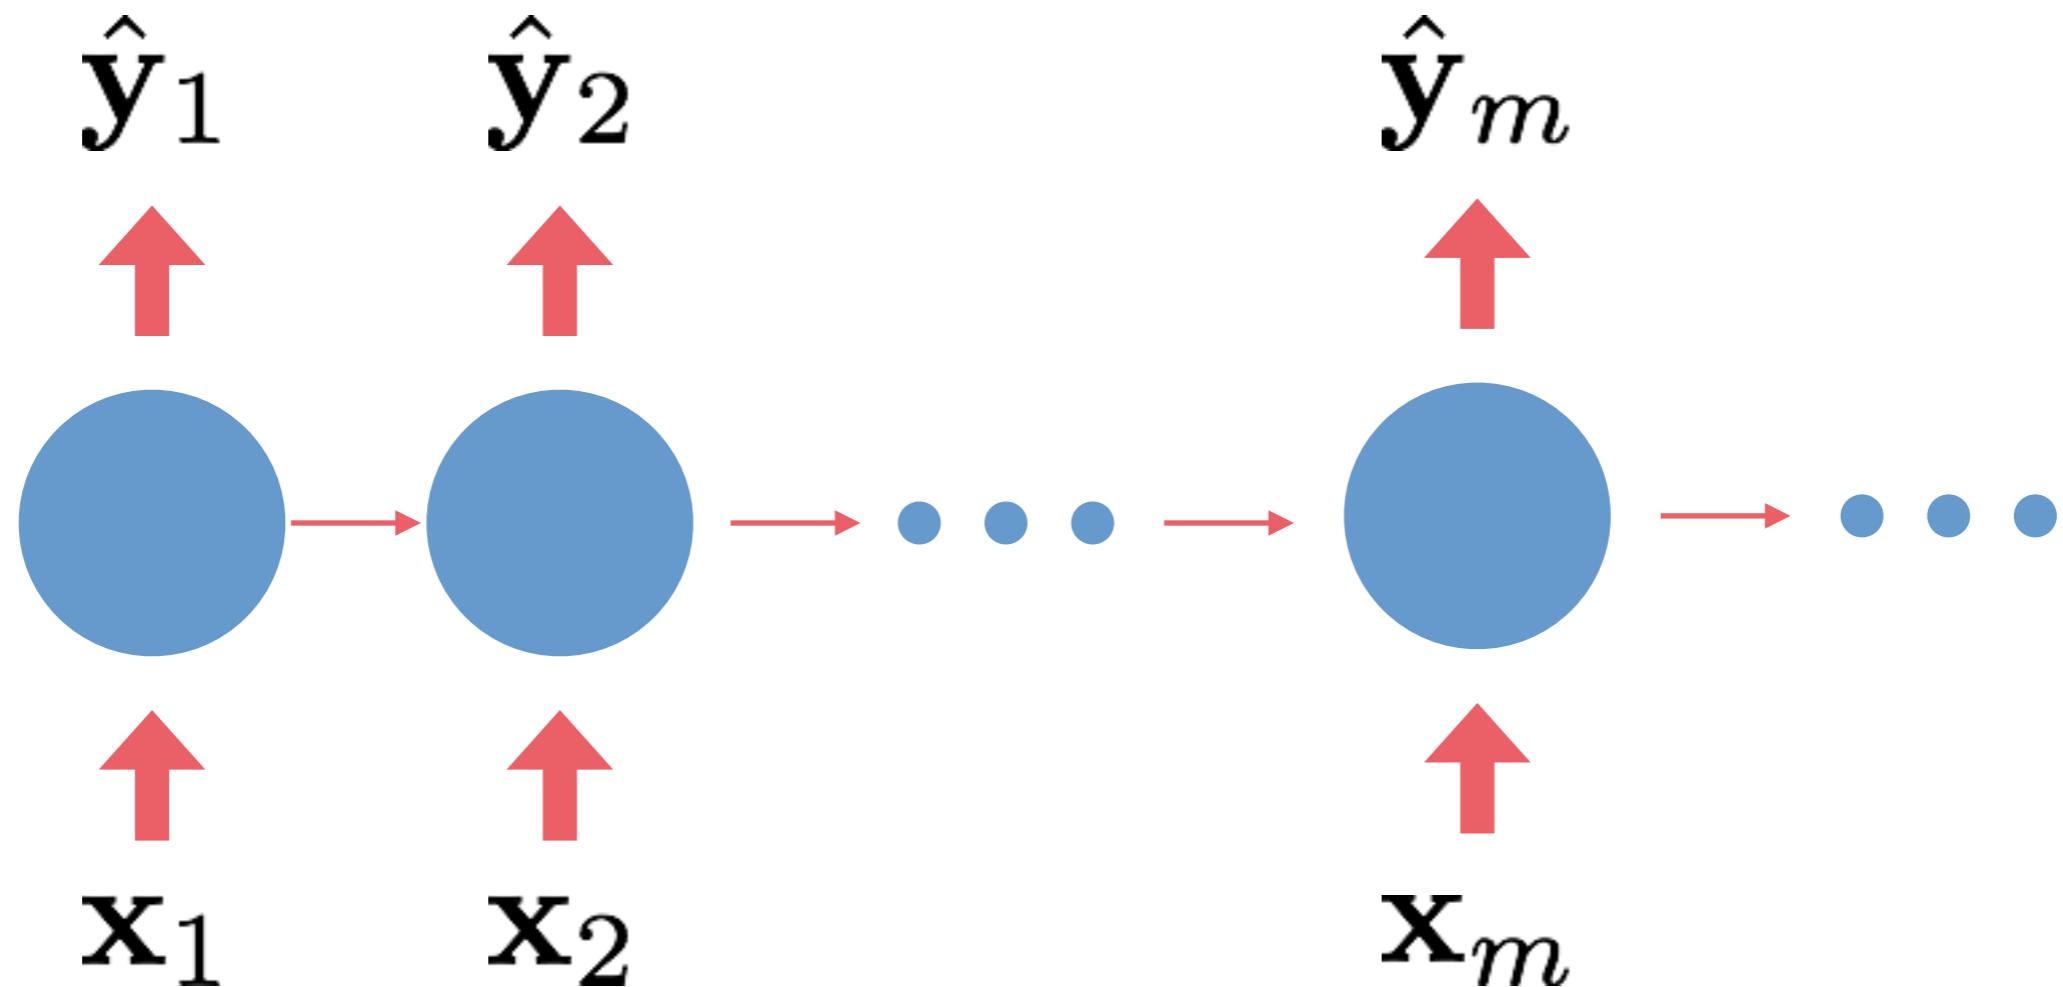
Usually, outputs of a Neural Network are not affected by the input orders..



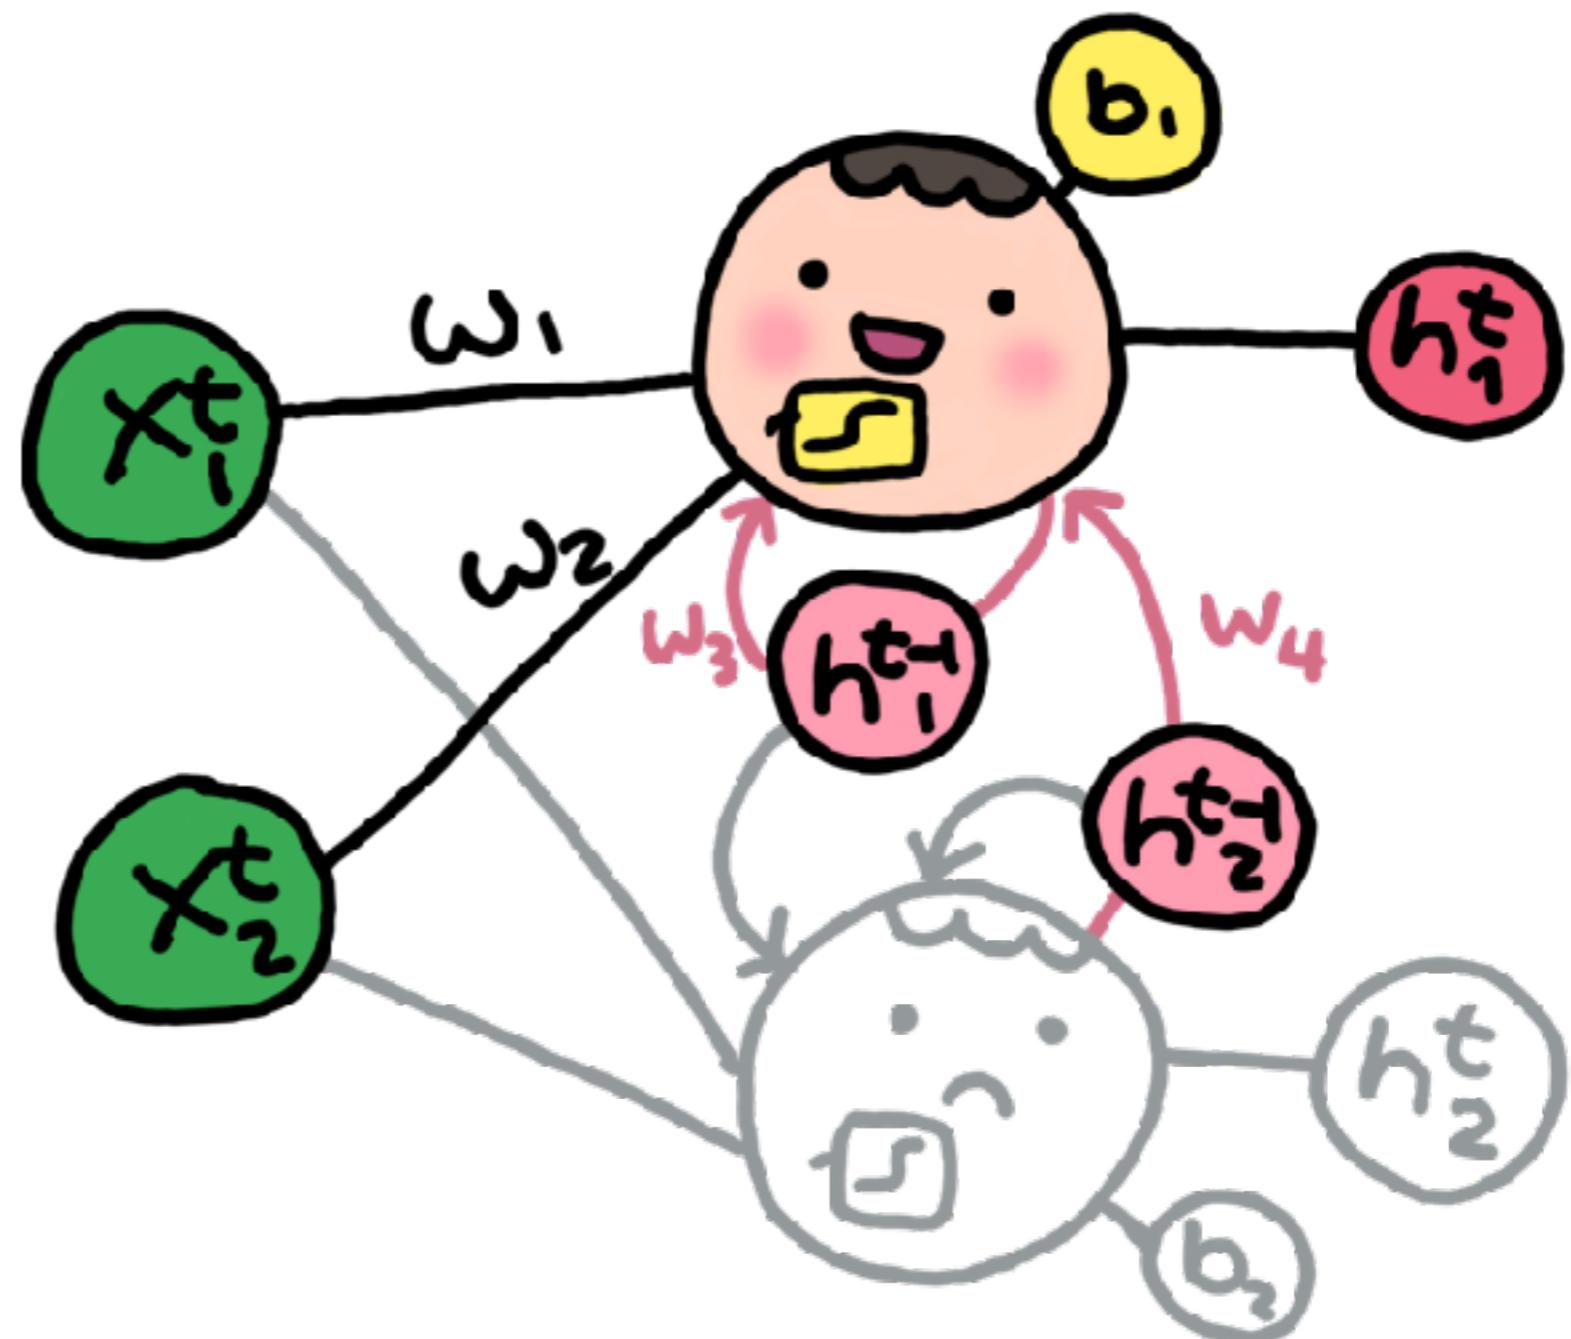
However, RNN cells will use previous outputs as part of inputs...



The “unfold” presentation.



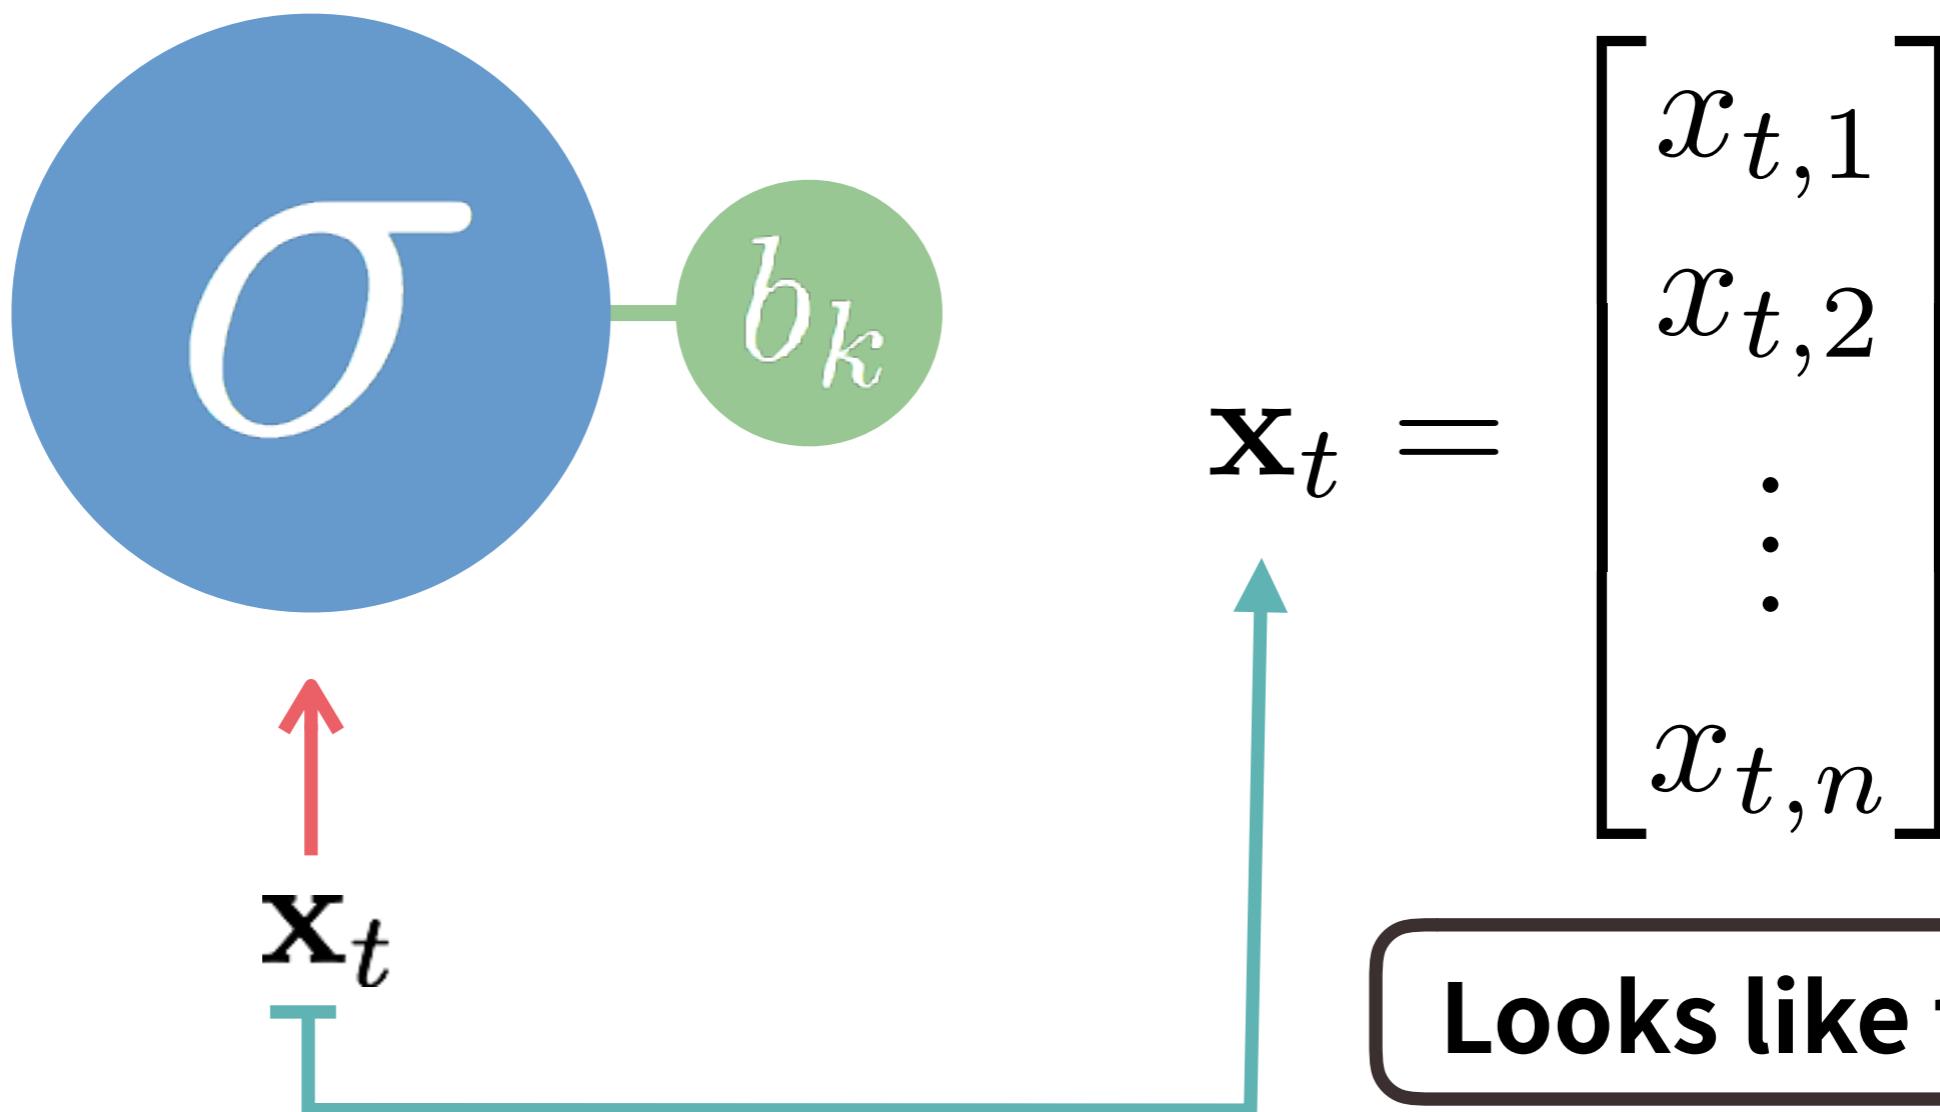
Recurrent Neural Network (RNN)



$$h_1^t = \sigma(w_1 x_1^t + w_2 x_2^t + w_3 h_1^{t-1} + w_4 h_2^{t-1} + b_1)$$

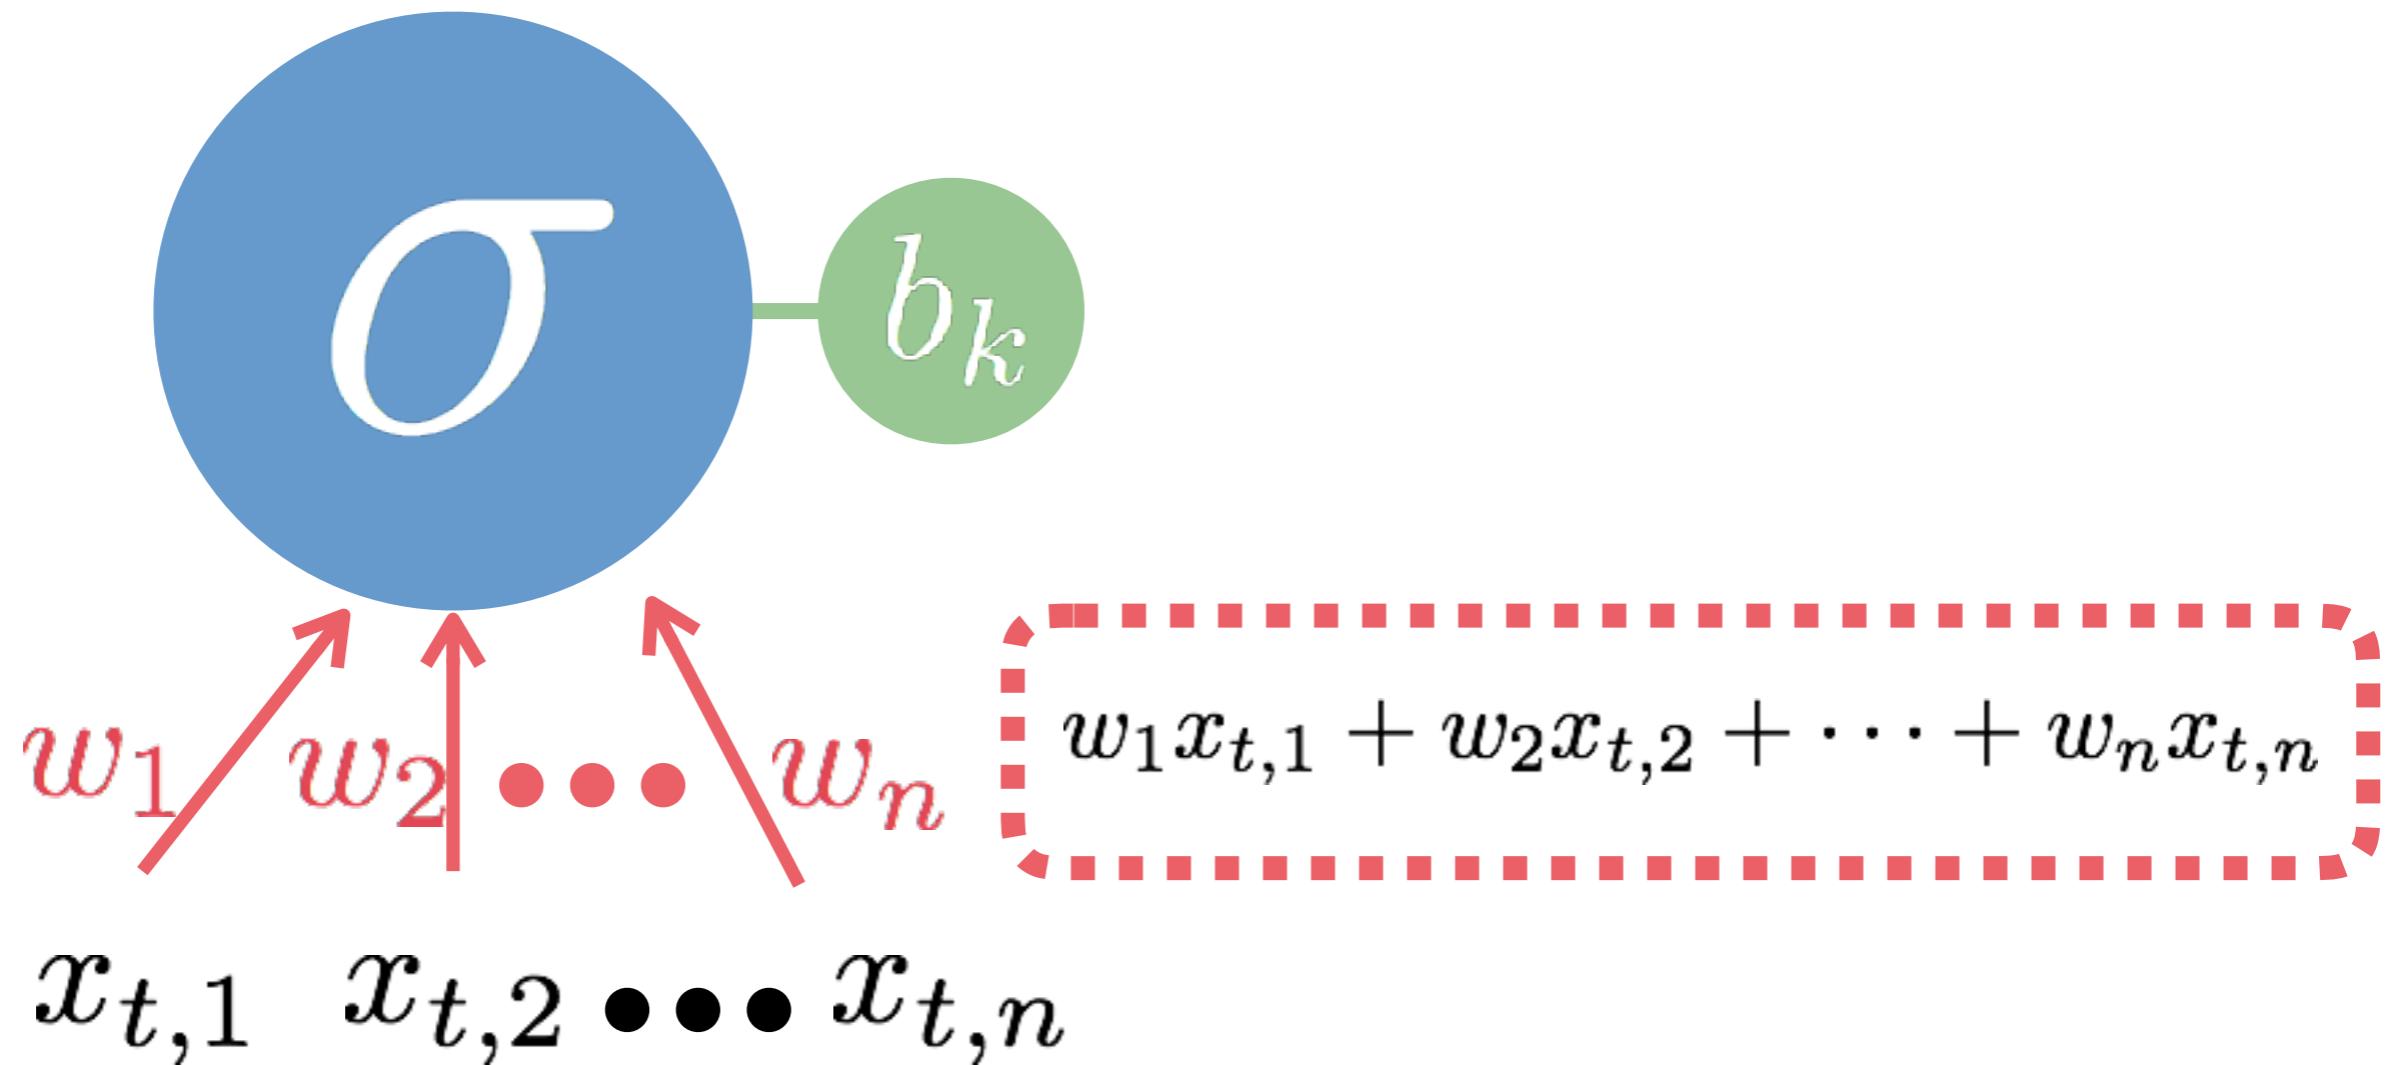
Remark

To make it easier for everyone to understand, we will use simpler representations. Please note that the inputs are vectors and will have respected weight; while the outputs are scalars.

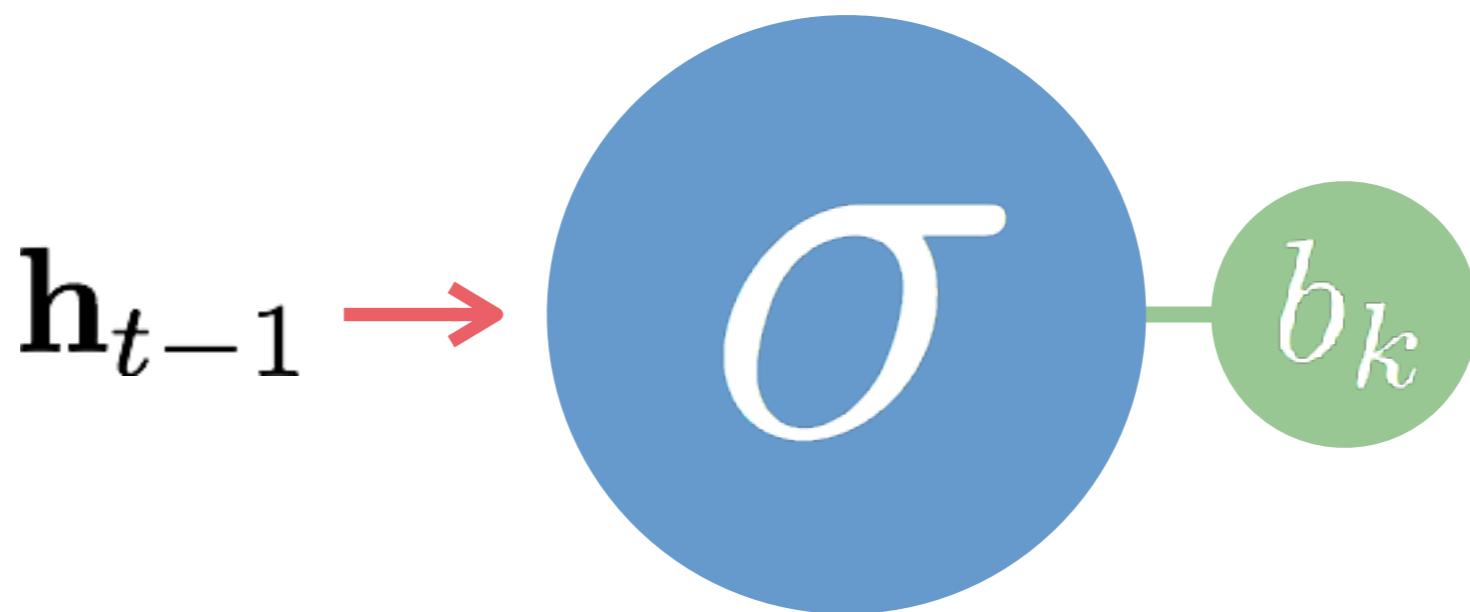




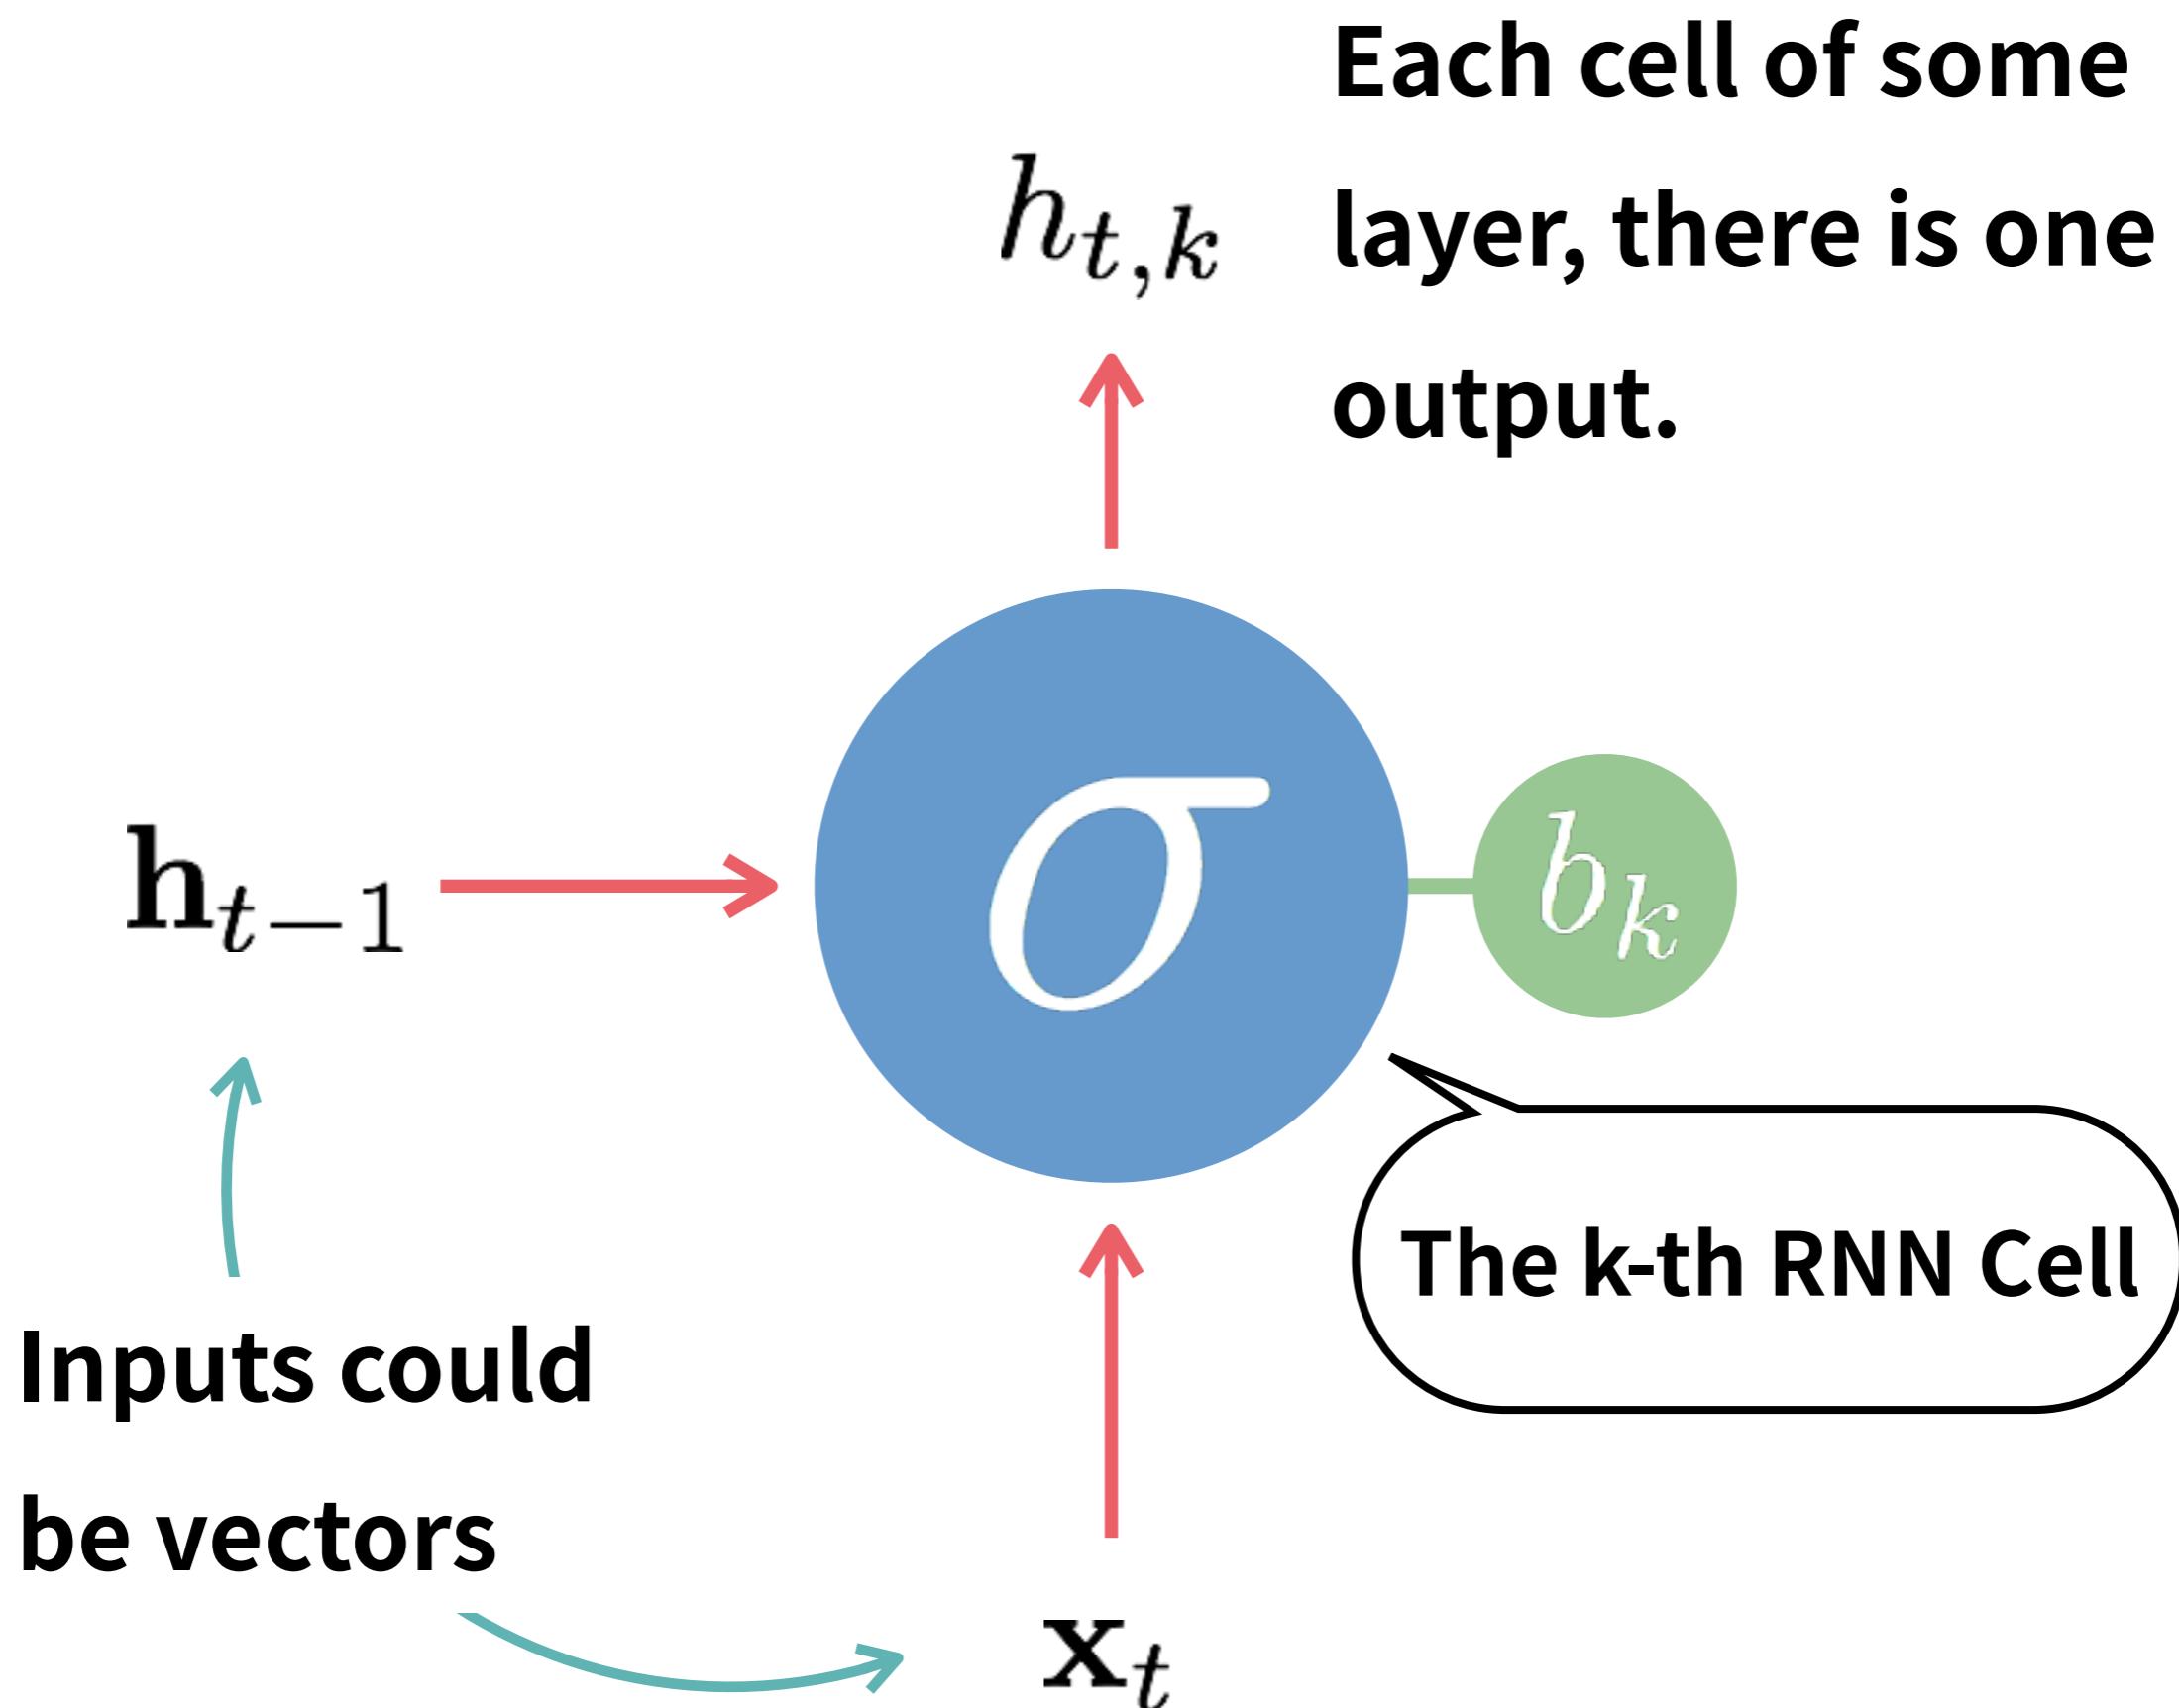
Actually look like this



Remark

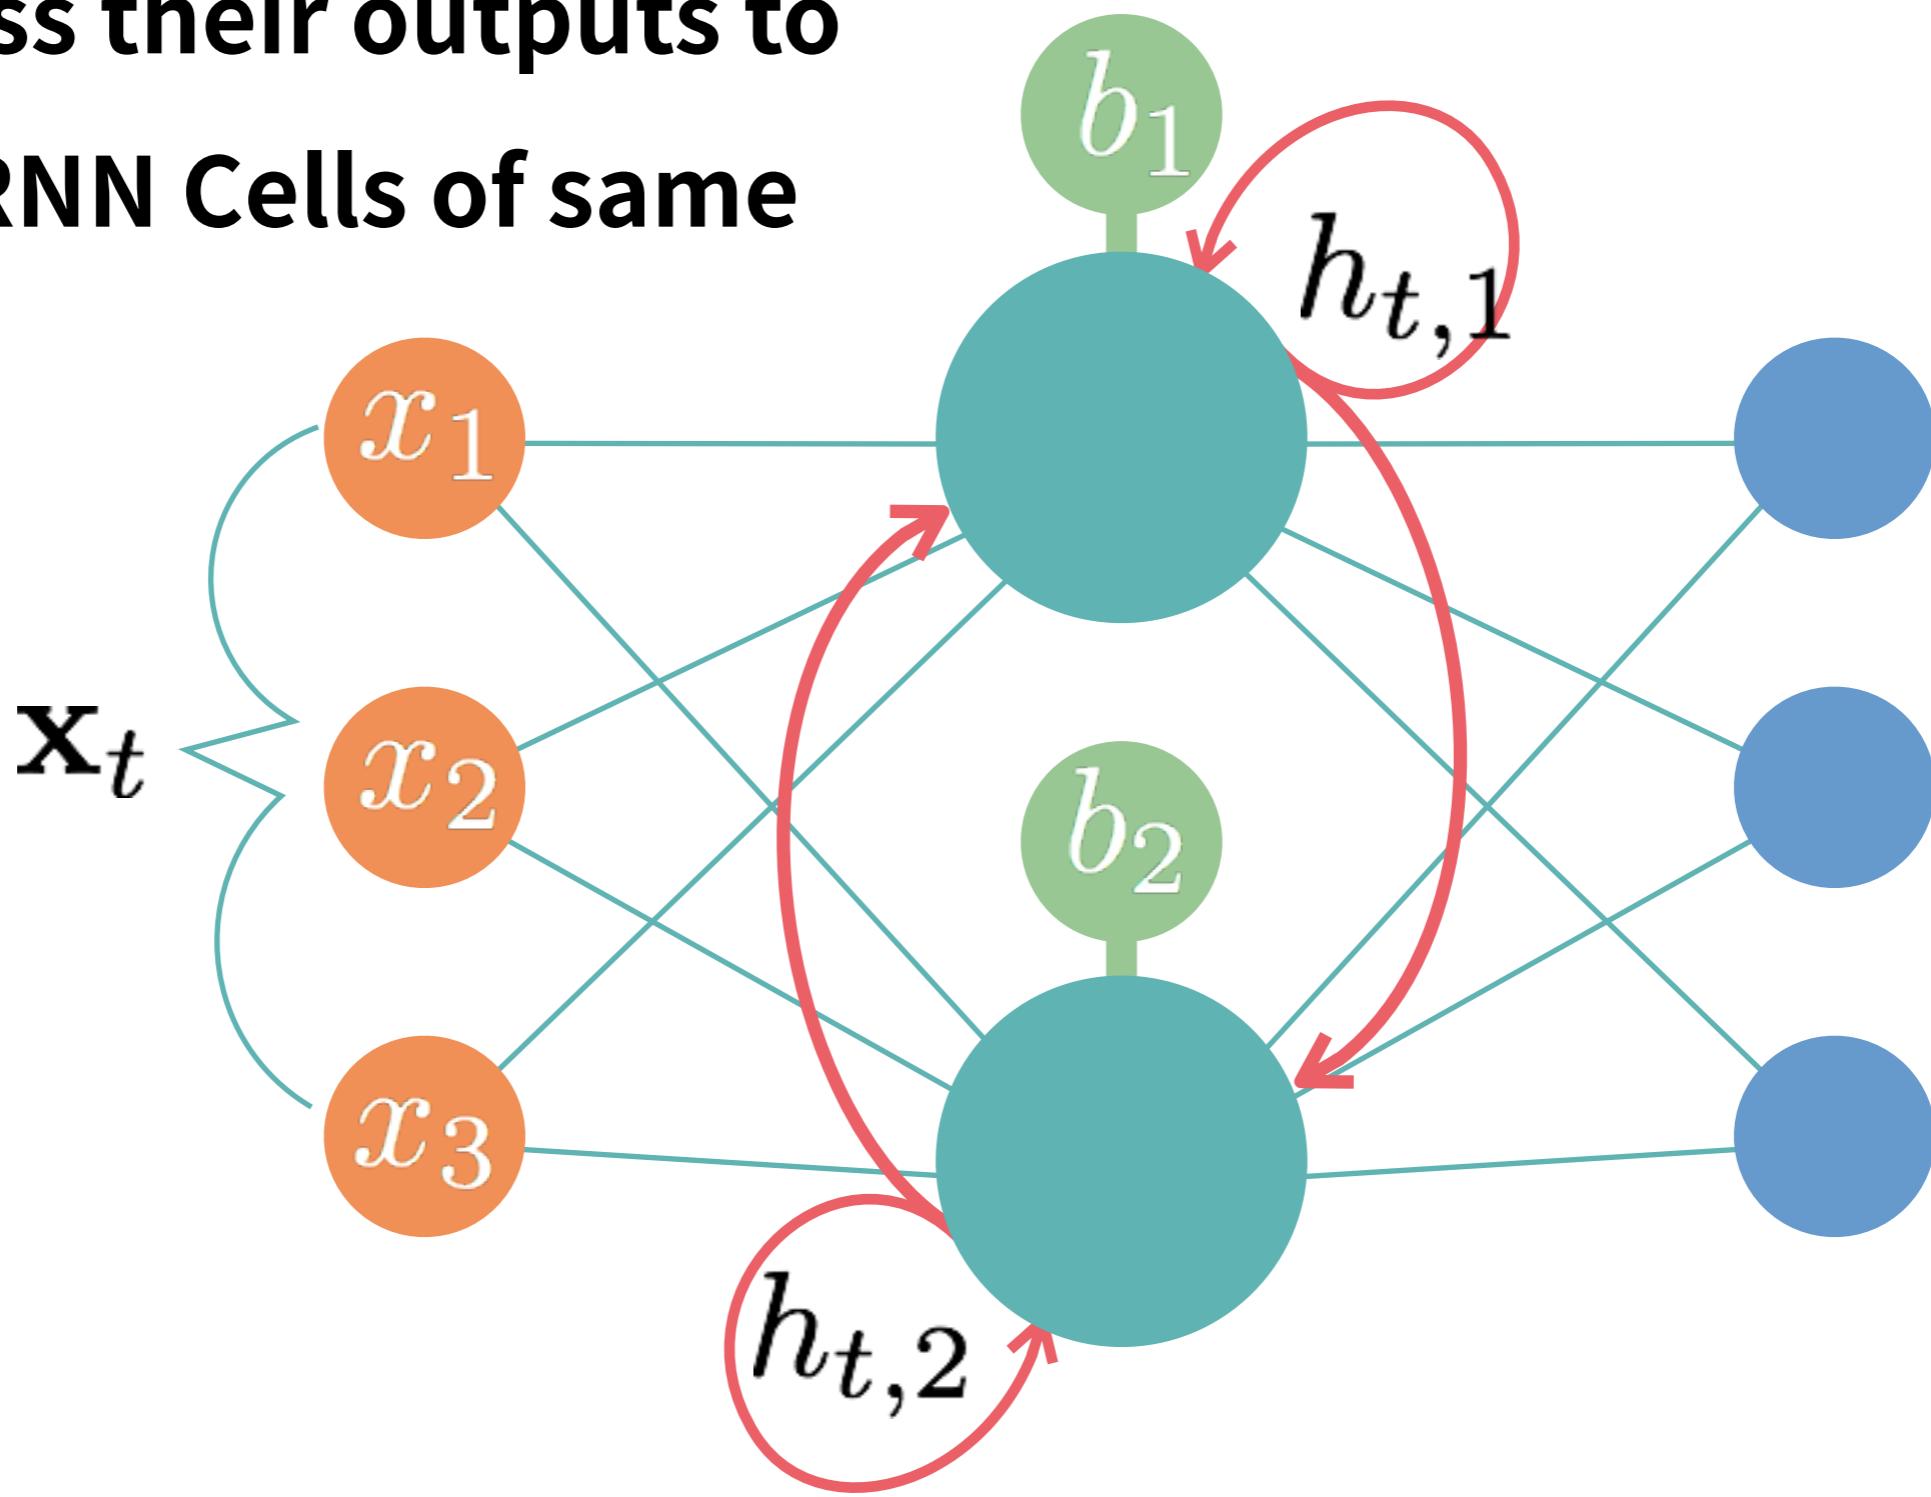


h's work similarly

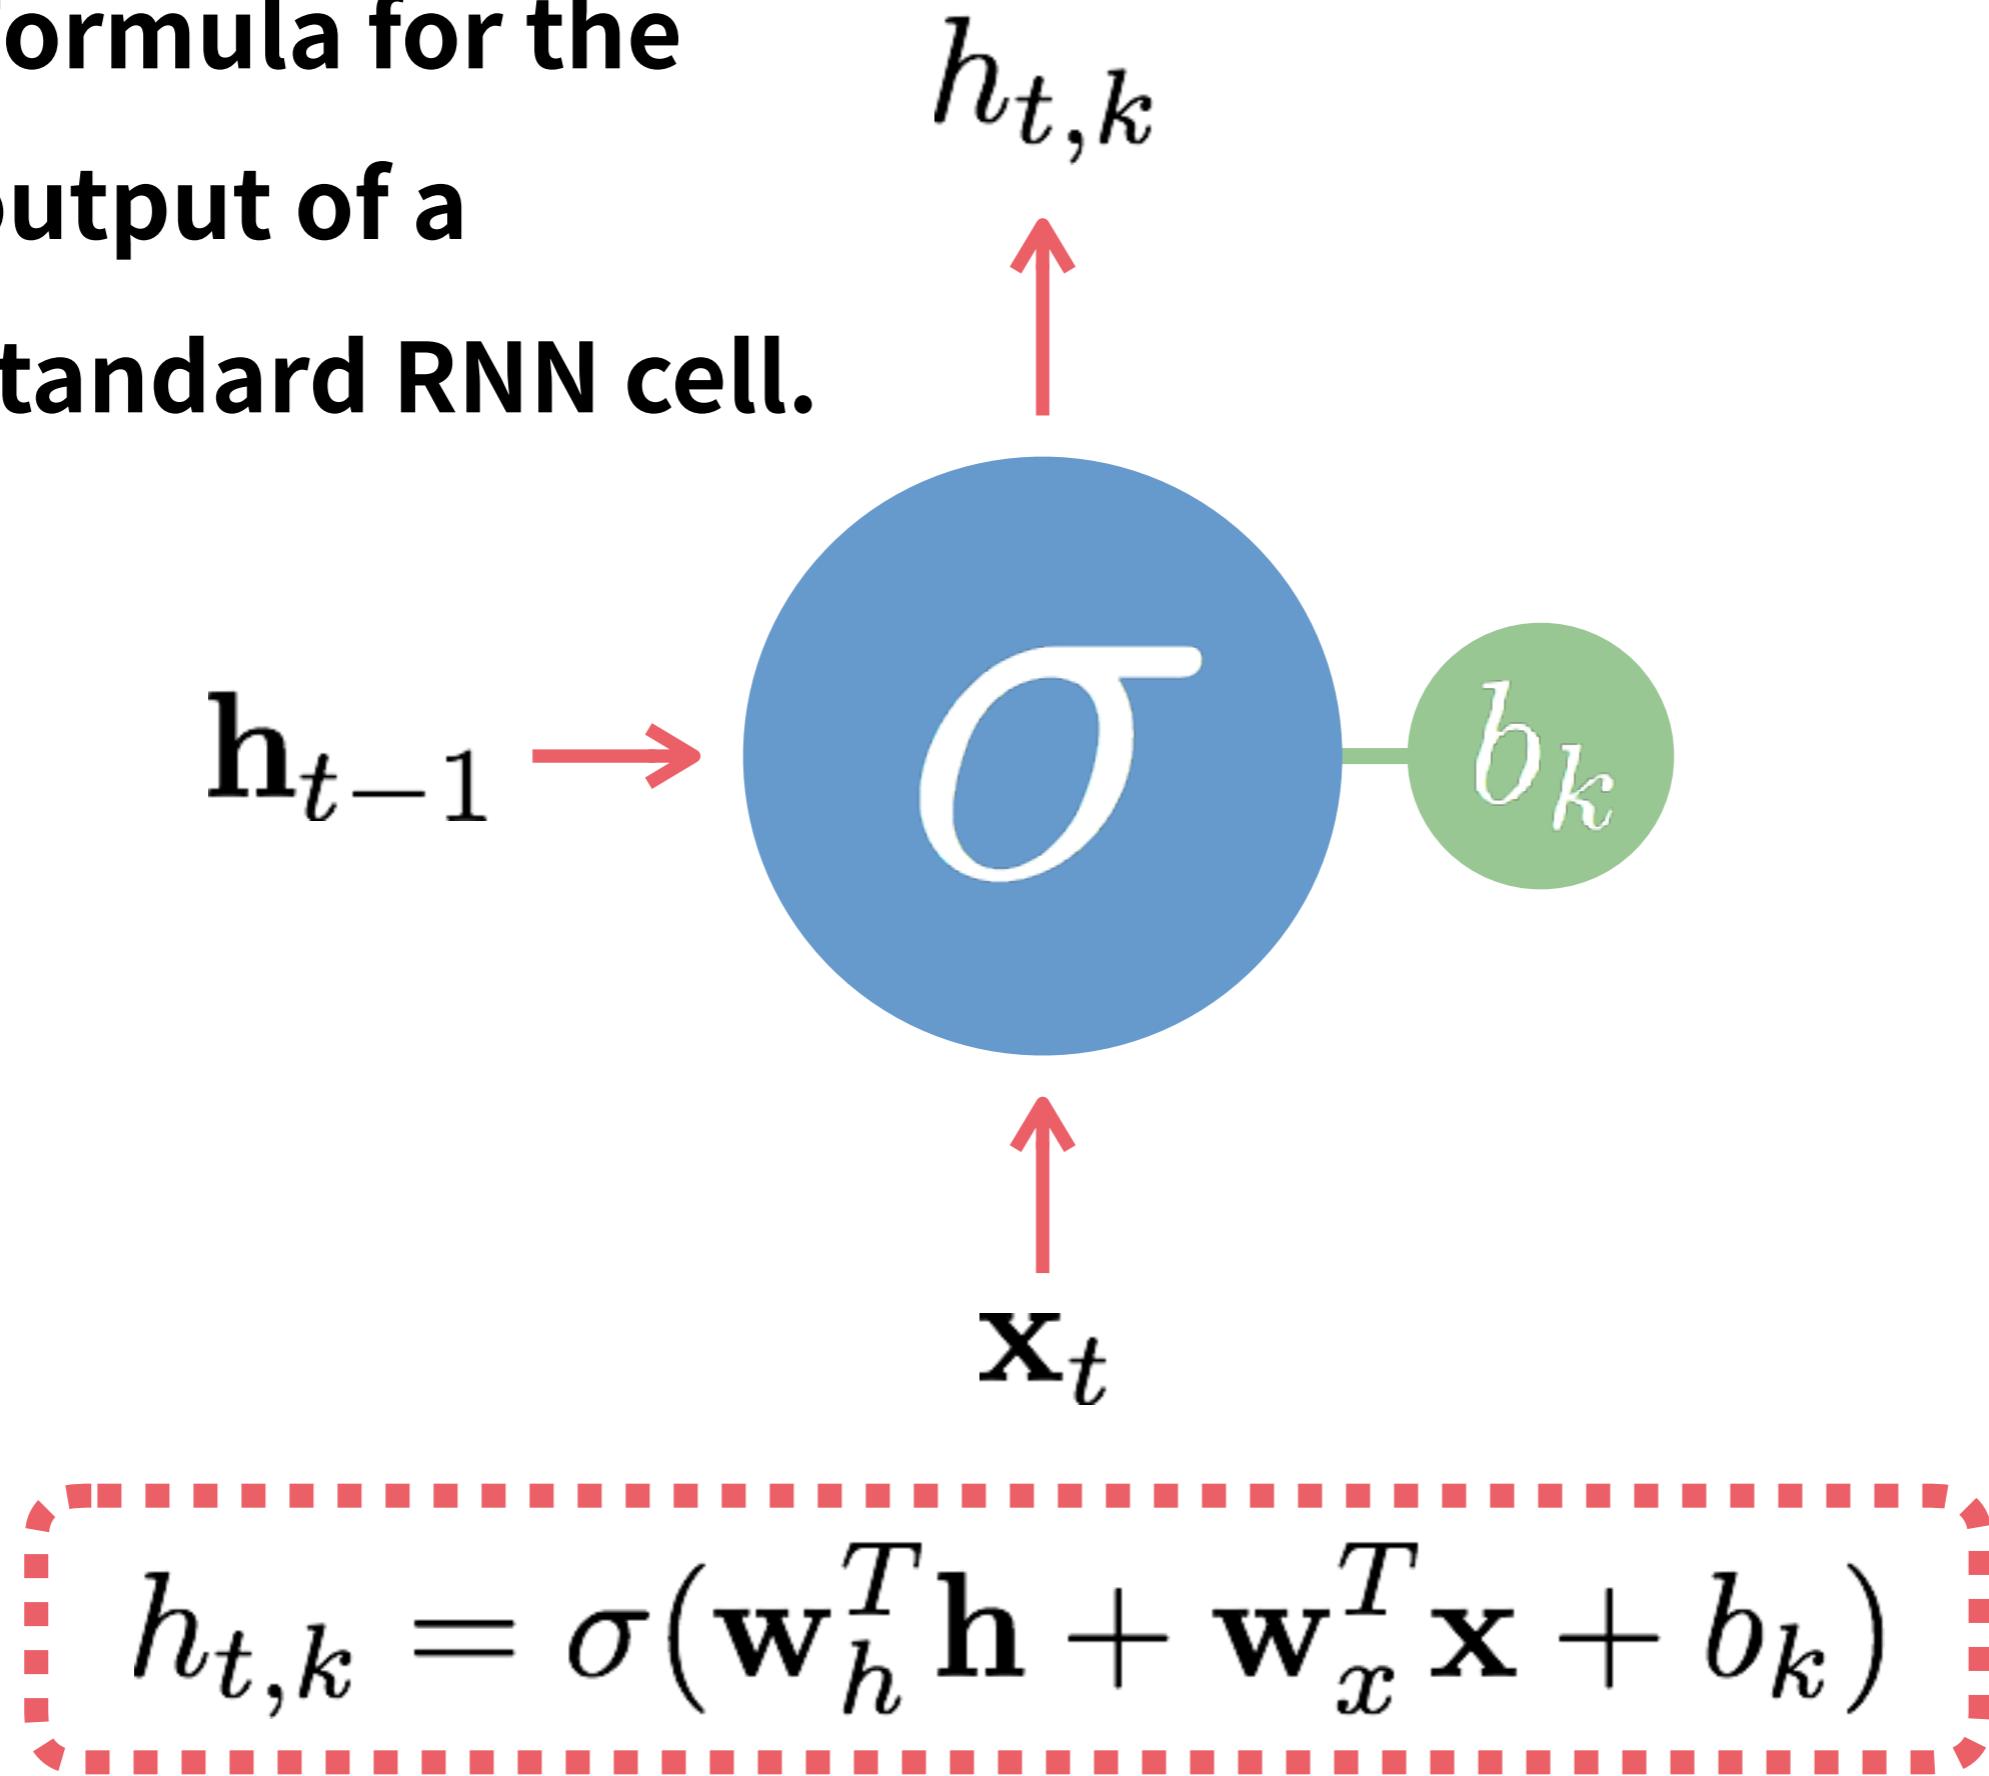


The connections of a RNN

layer. Note that RNN cells will pass their outputs to other RNN Cells of same layer.



**Formula for the
output of a
standard RNN cell.**



When we say “RNN,” most people
actually think of...



Long Short Term Memory

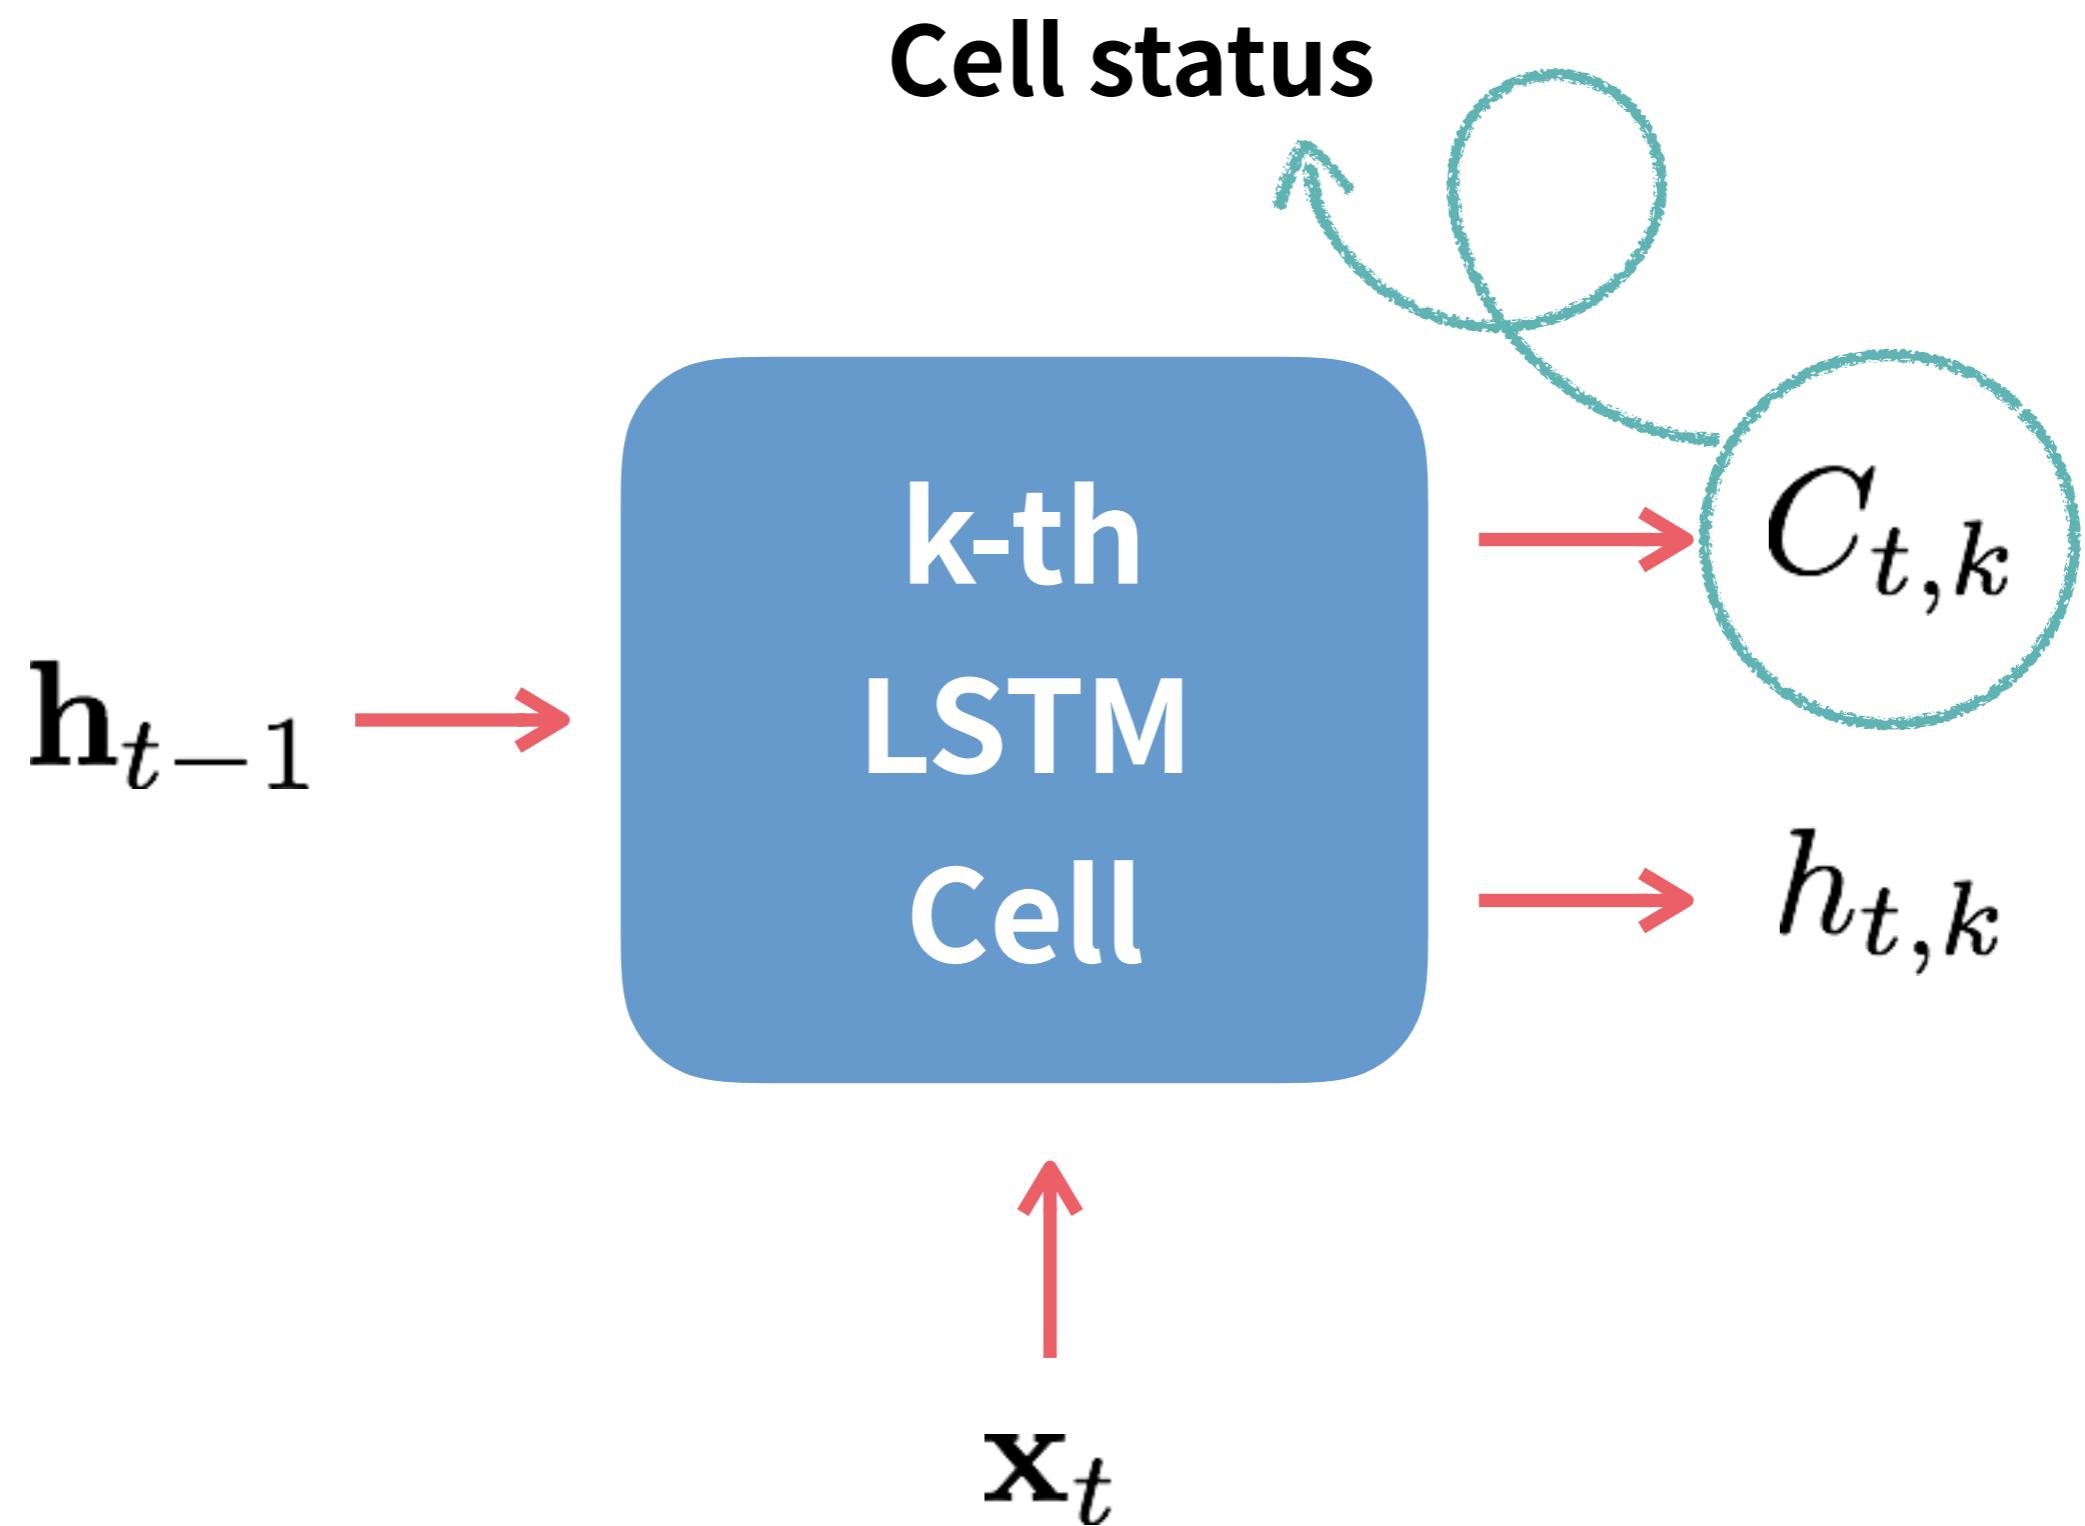
Gated Recurrent Unit





Long Short Term Memory

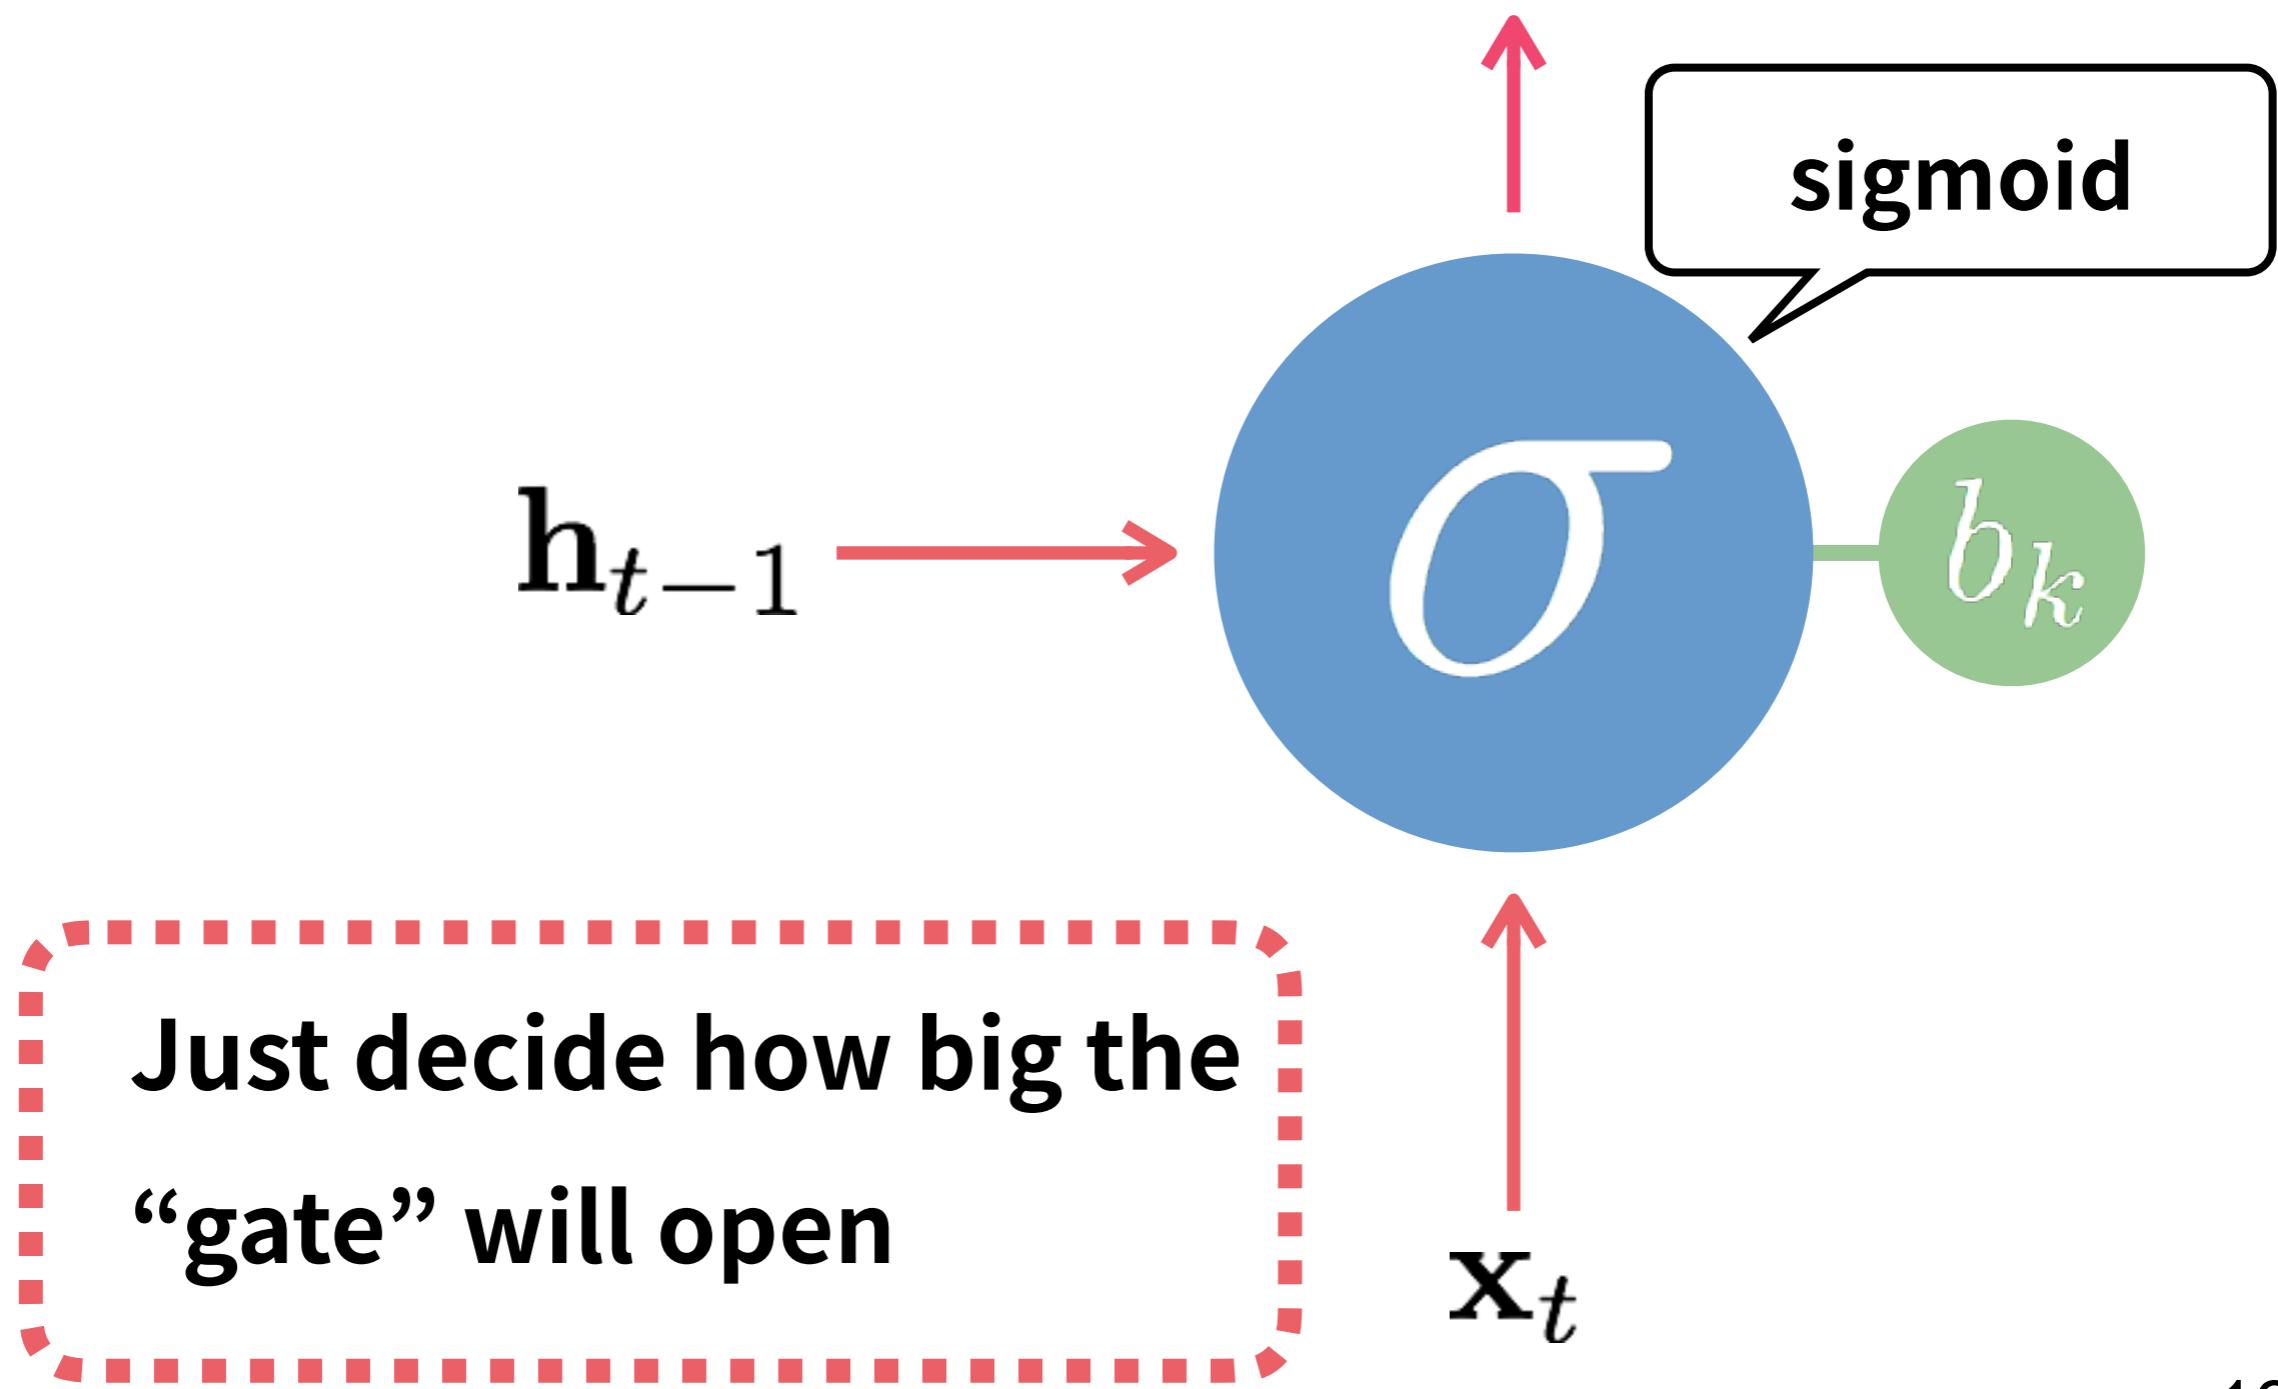
The Ace of RNNs



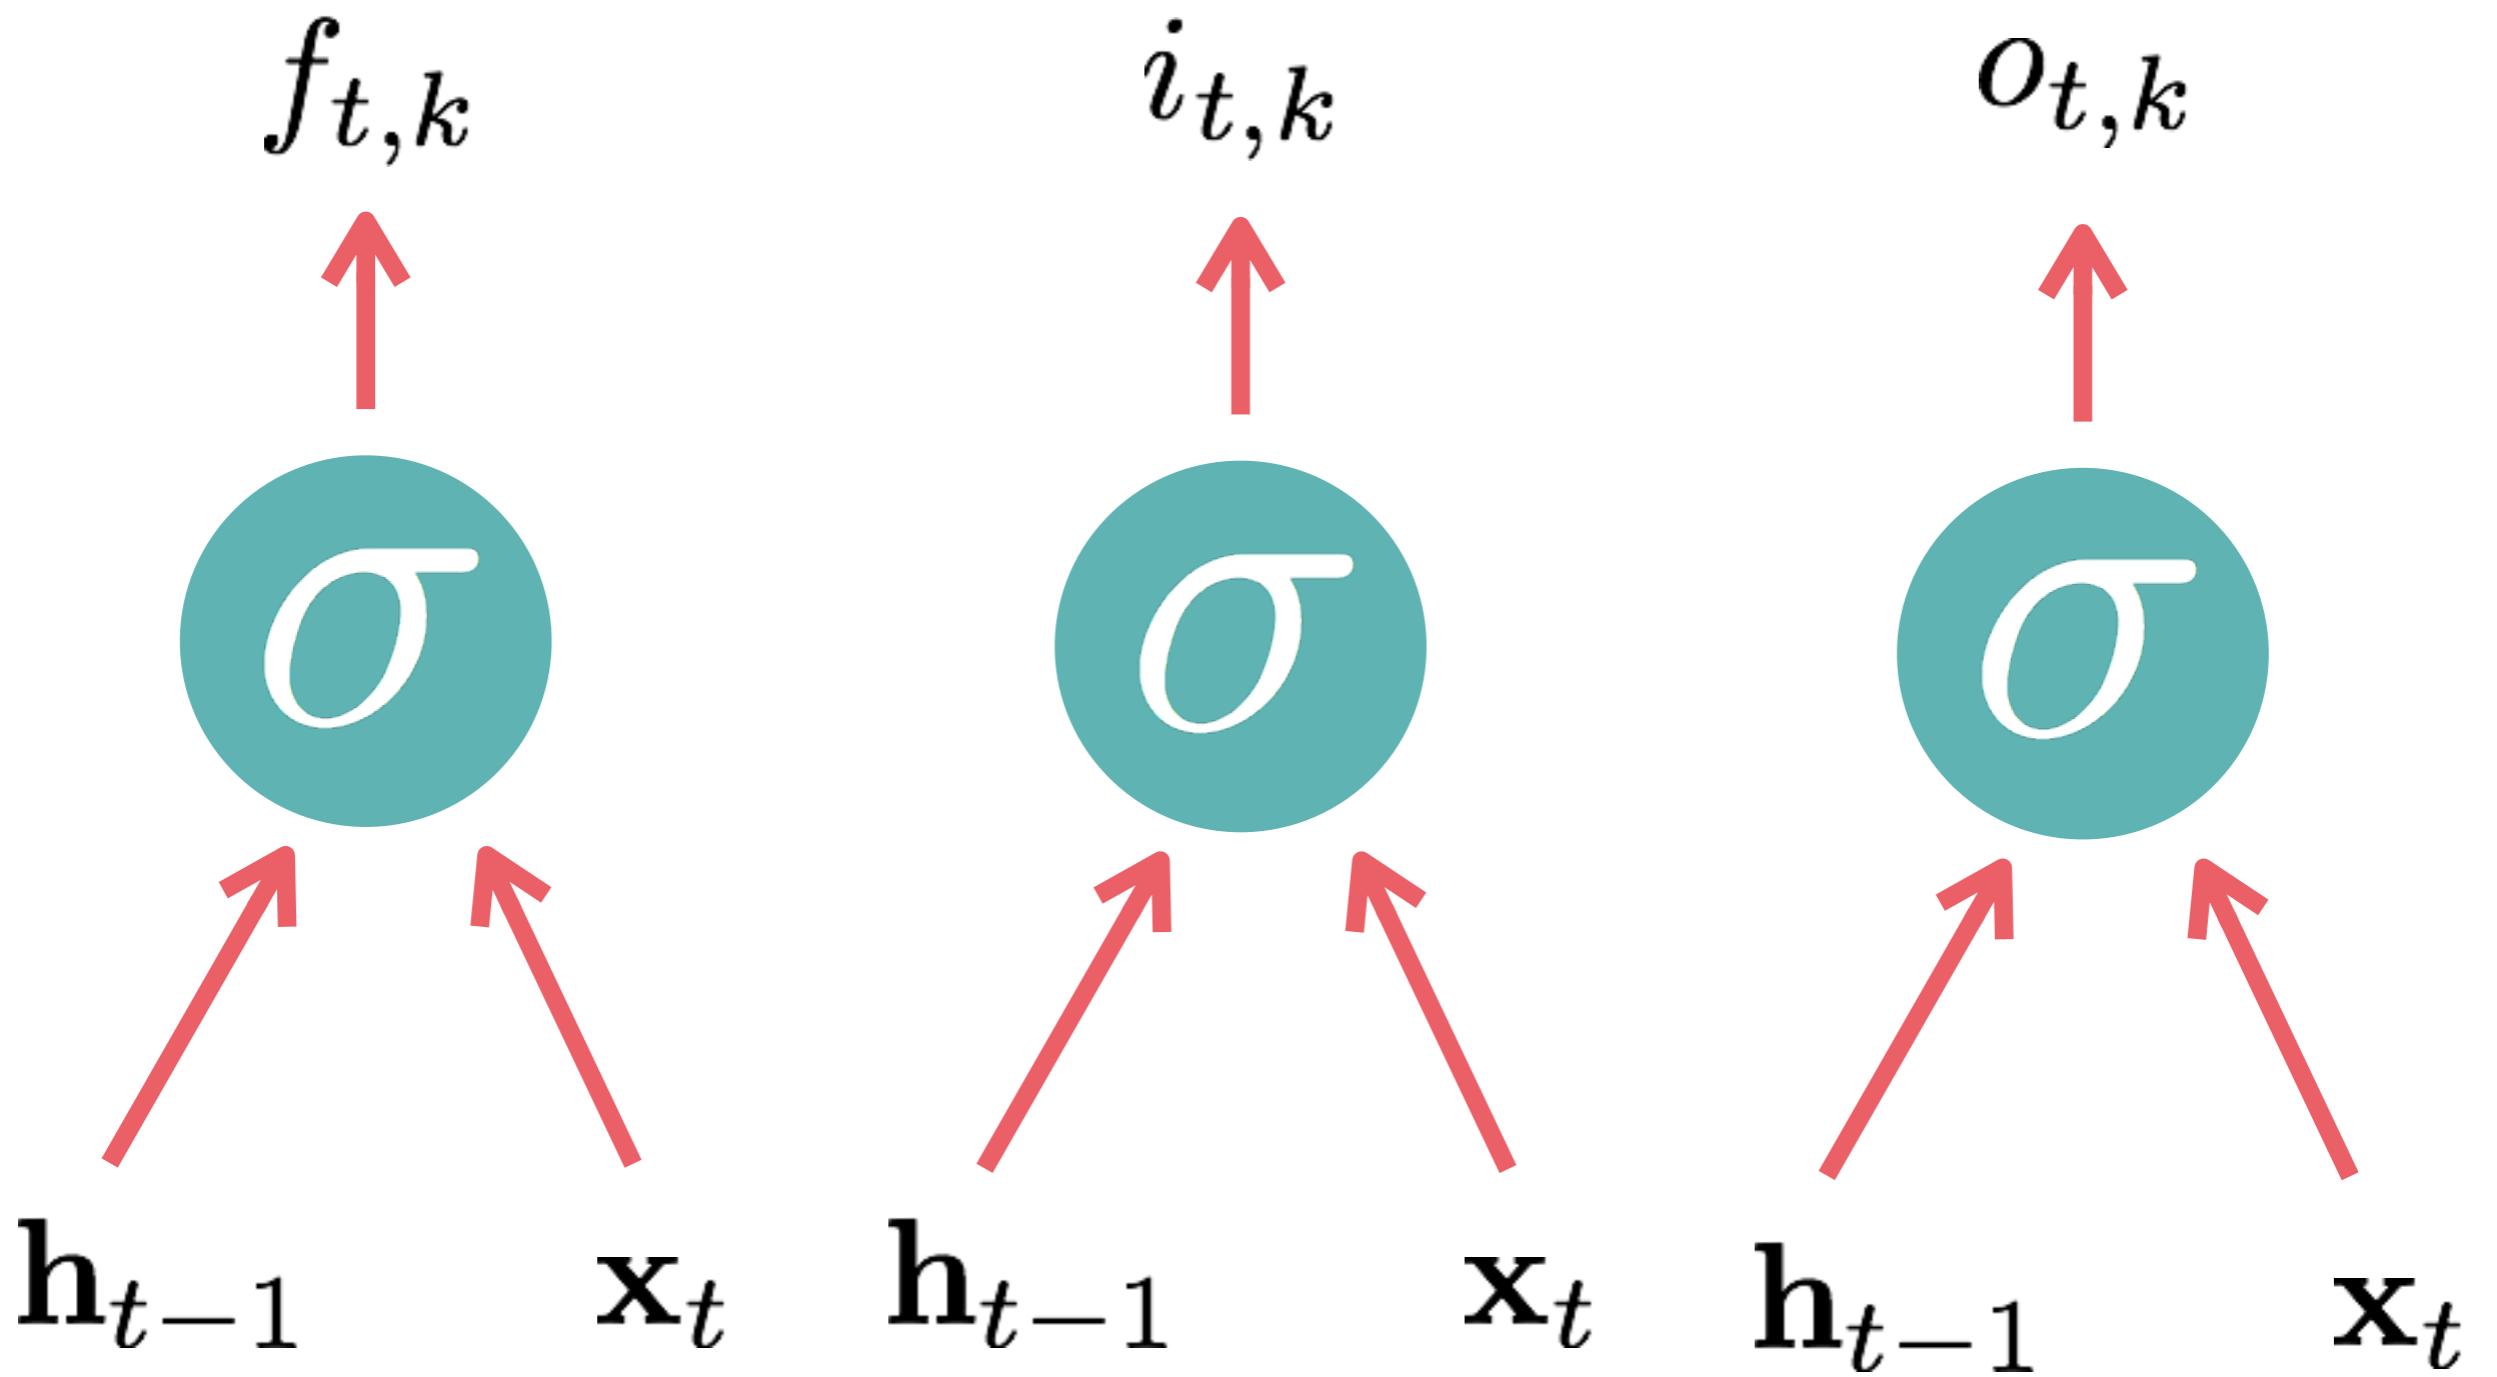
Gate

Important

The output is a number
between 0 and 1



**LSTMs has three
types of Gates**

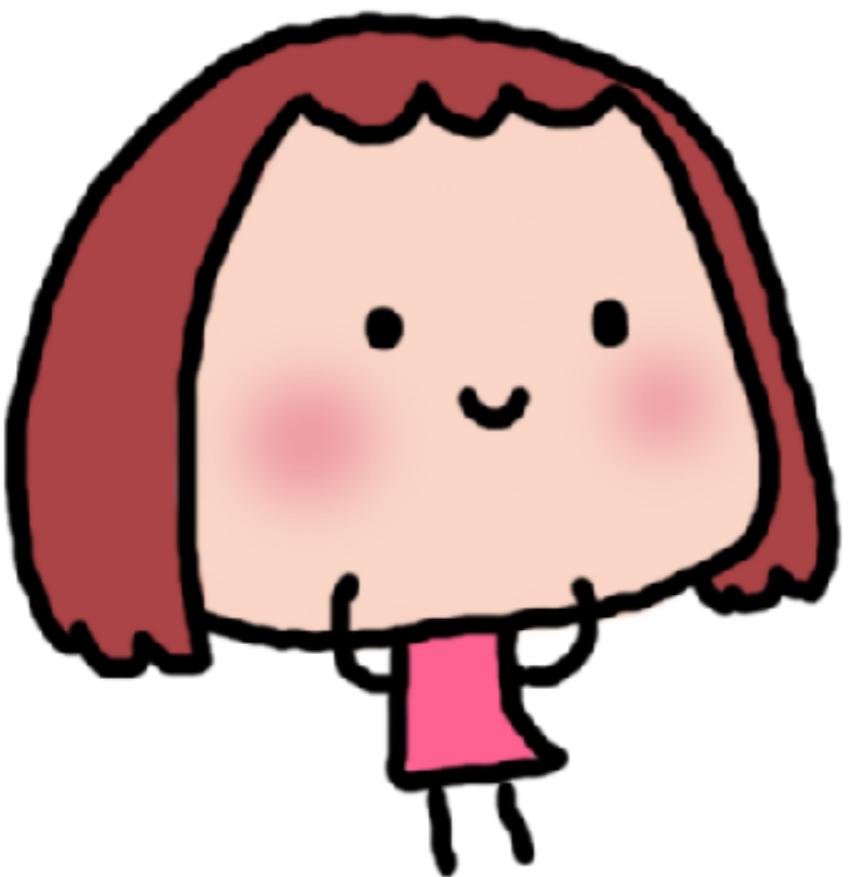


forget gate

input gate

output gate

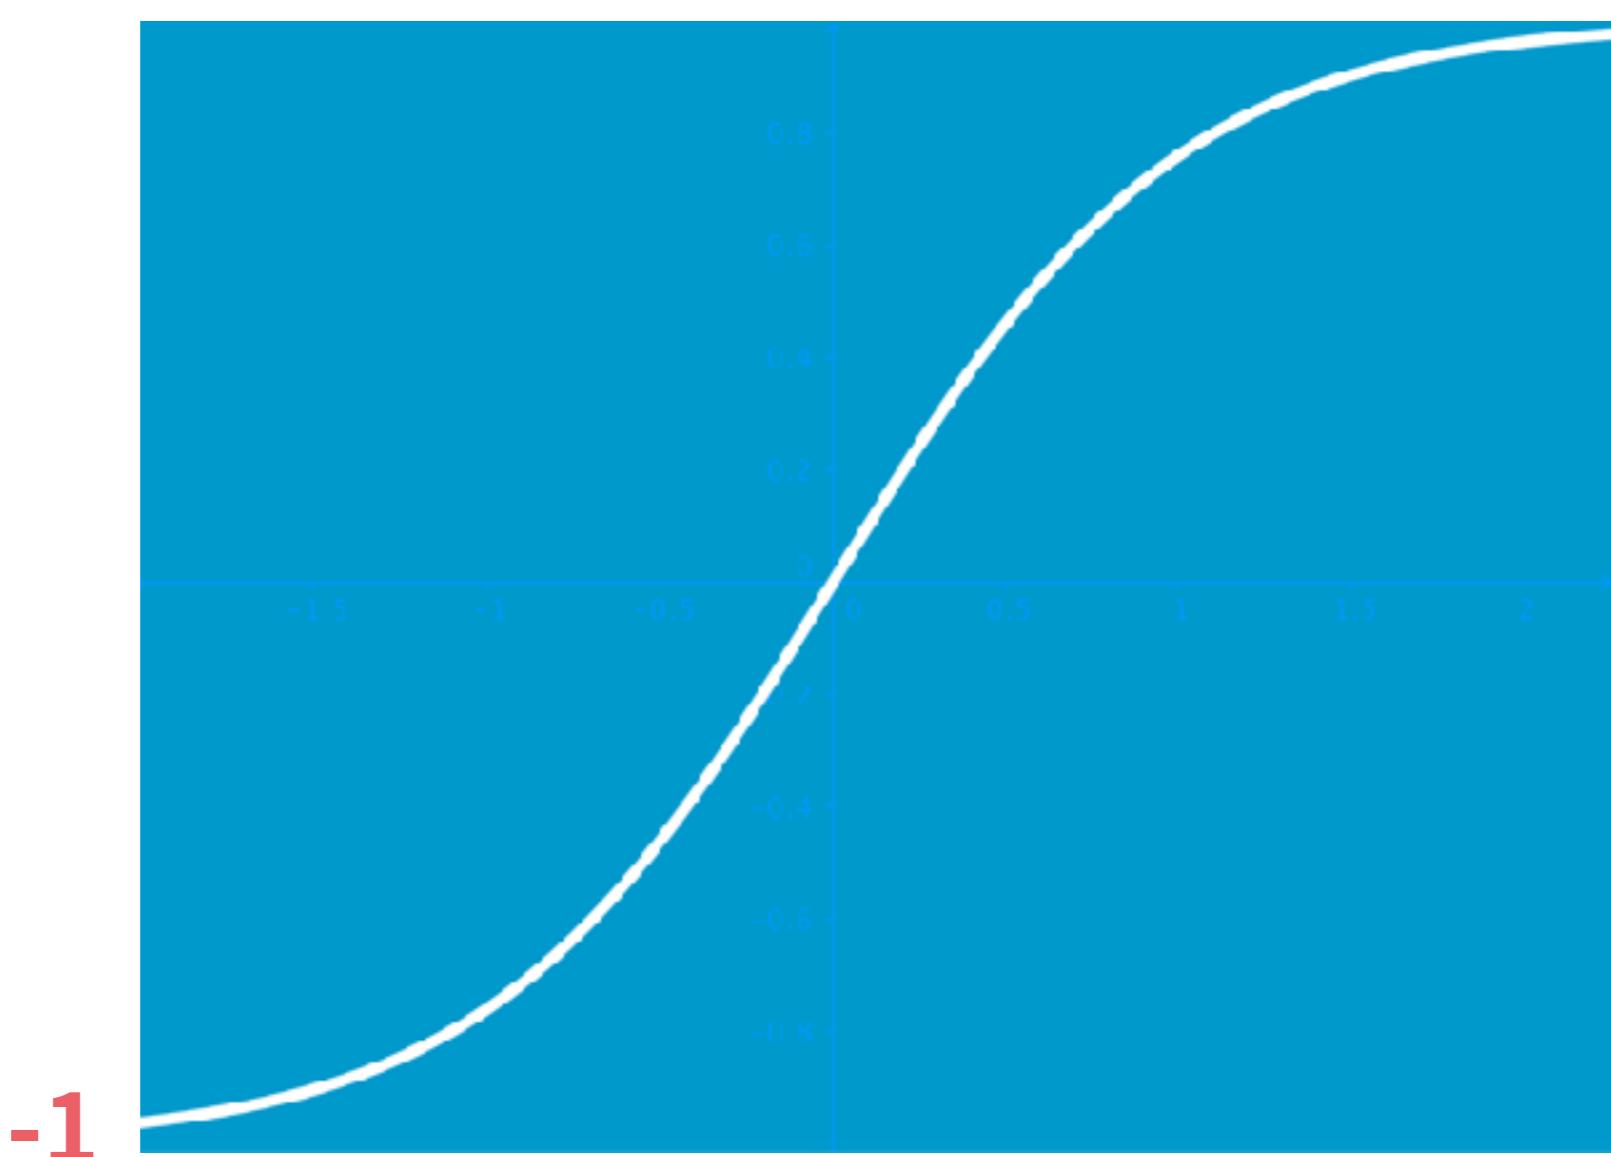
Recall



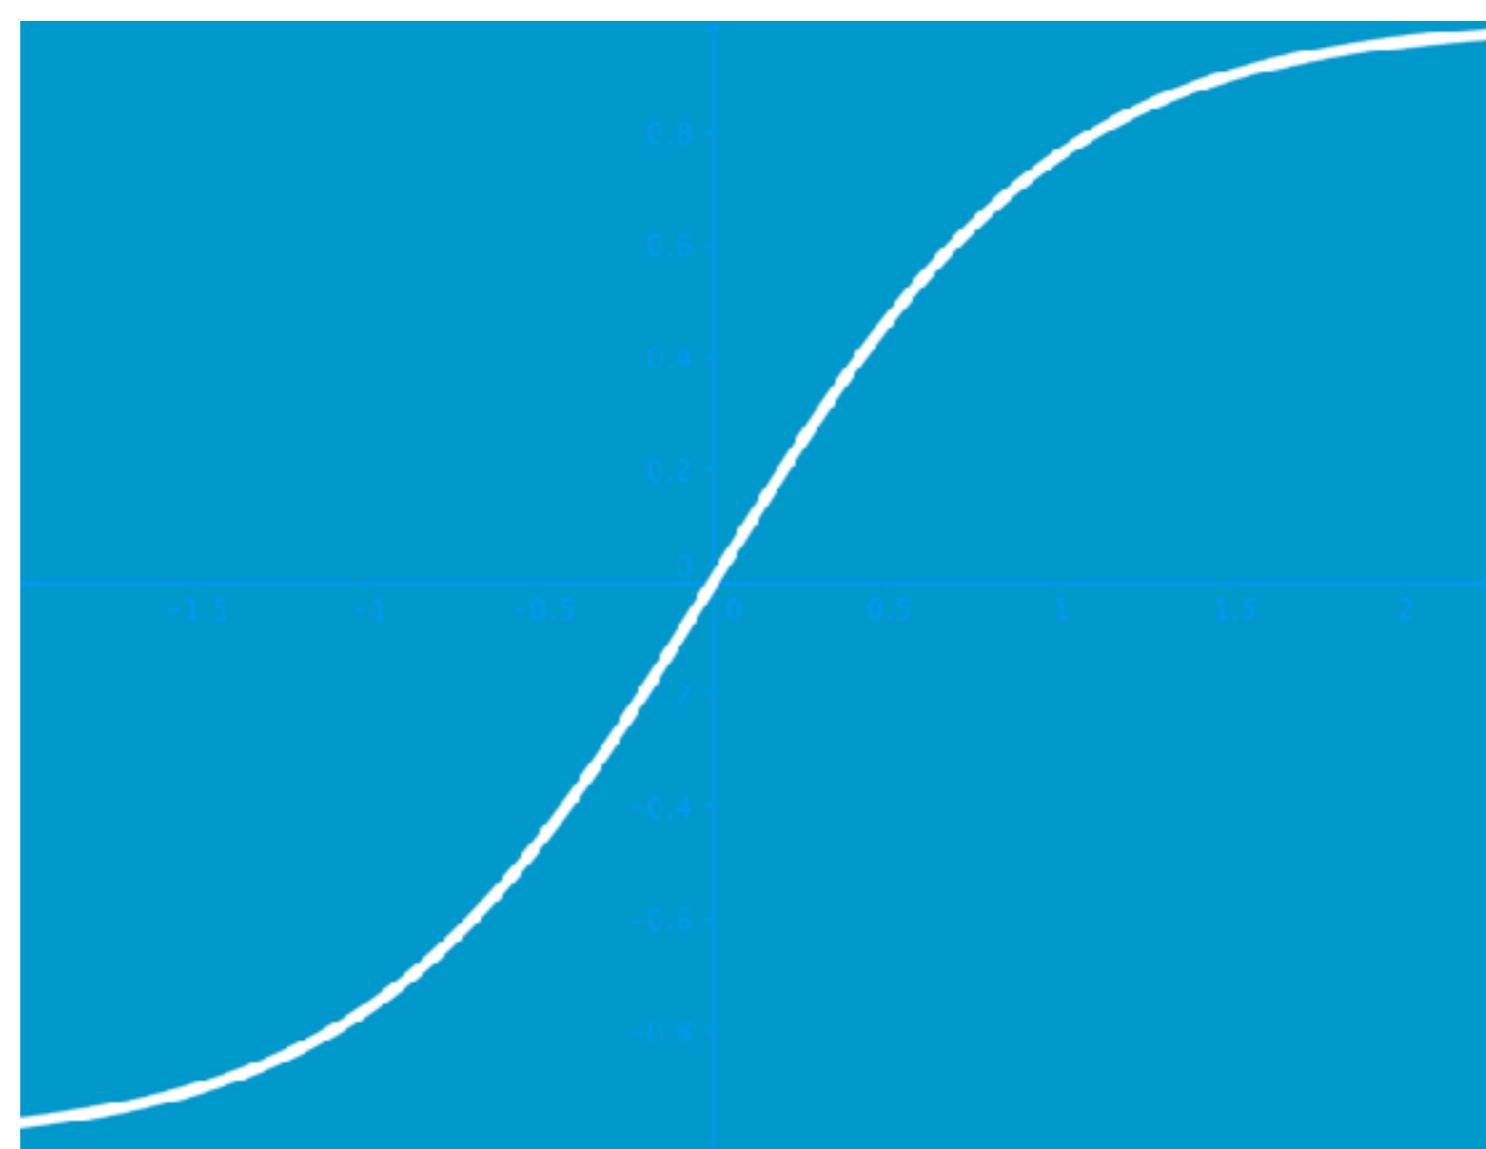
tanh

sigmoid

$$\tanh(x) = \frac{1 - e^{-x}}{1 + e^{-x}}$$

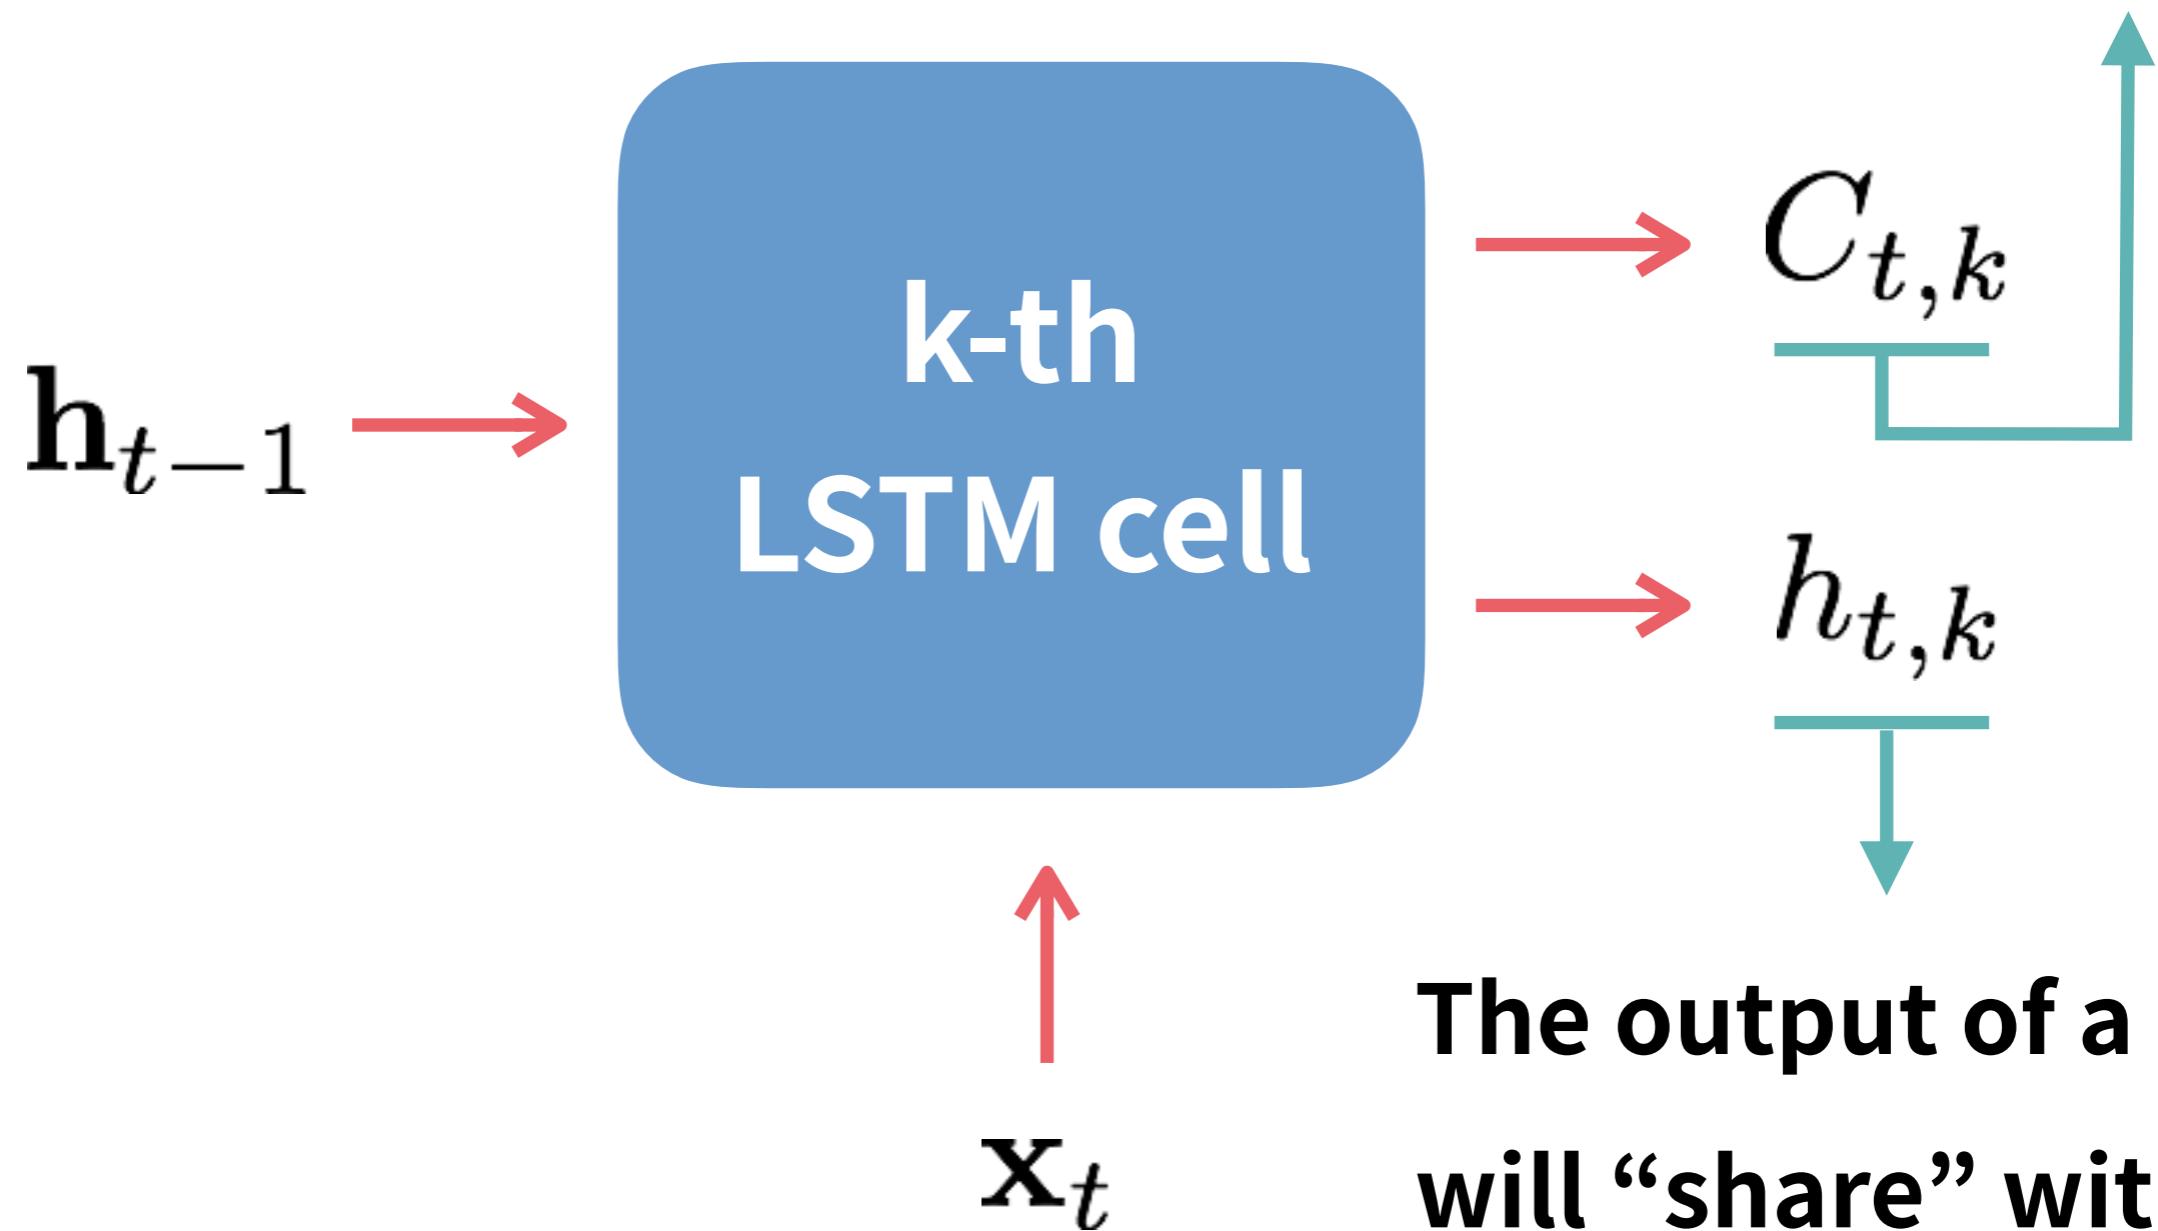


$$\sigma(x) = \frac{1}{1 + e^{-x}}$$

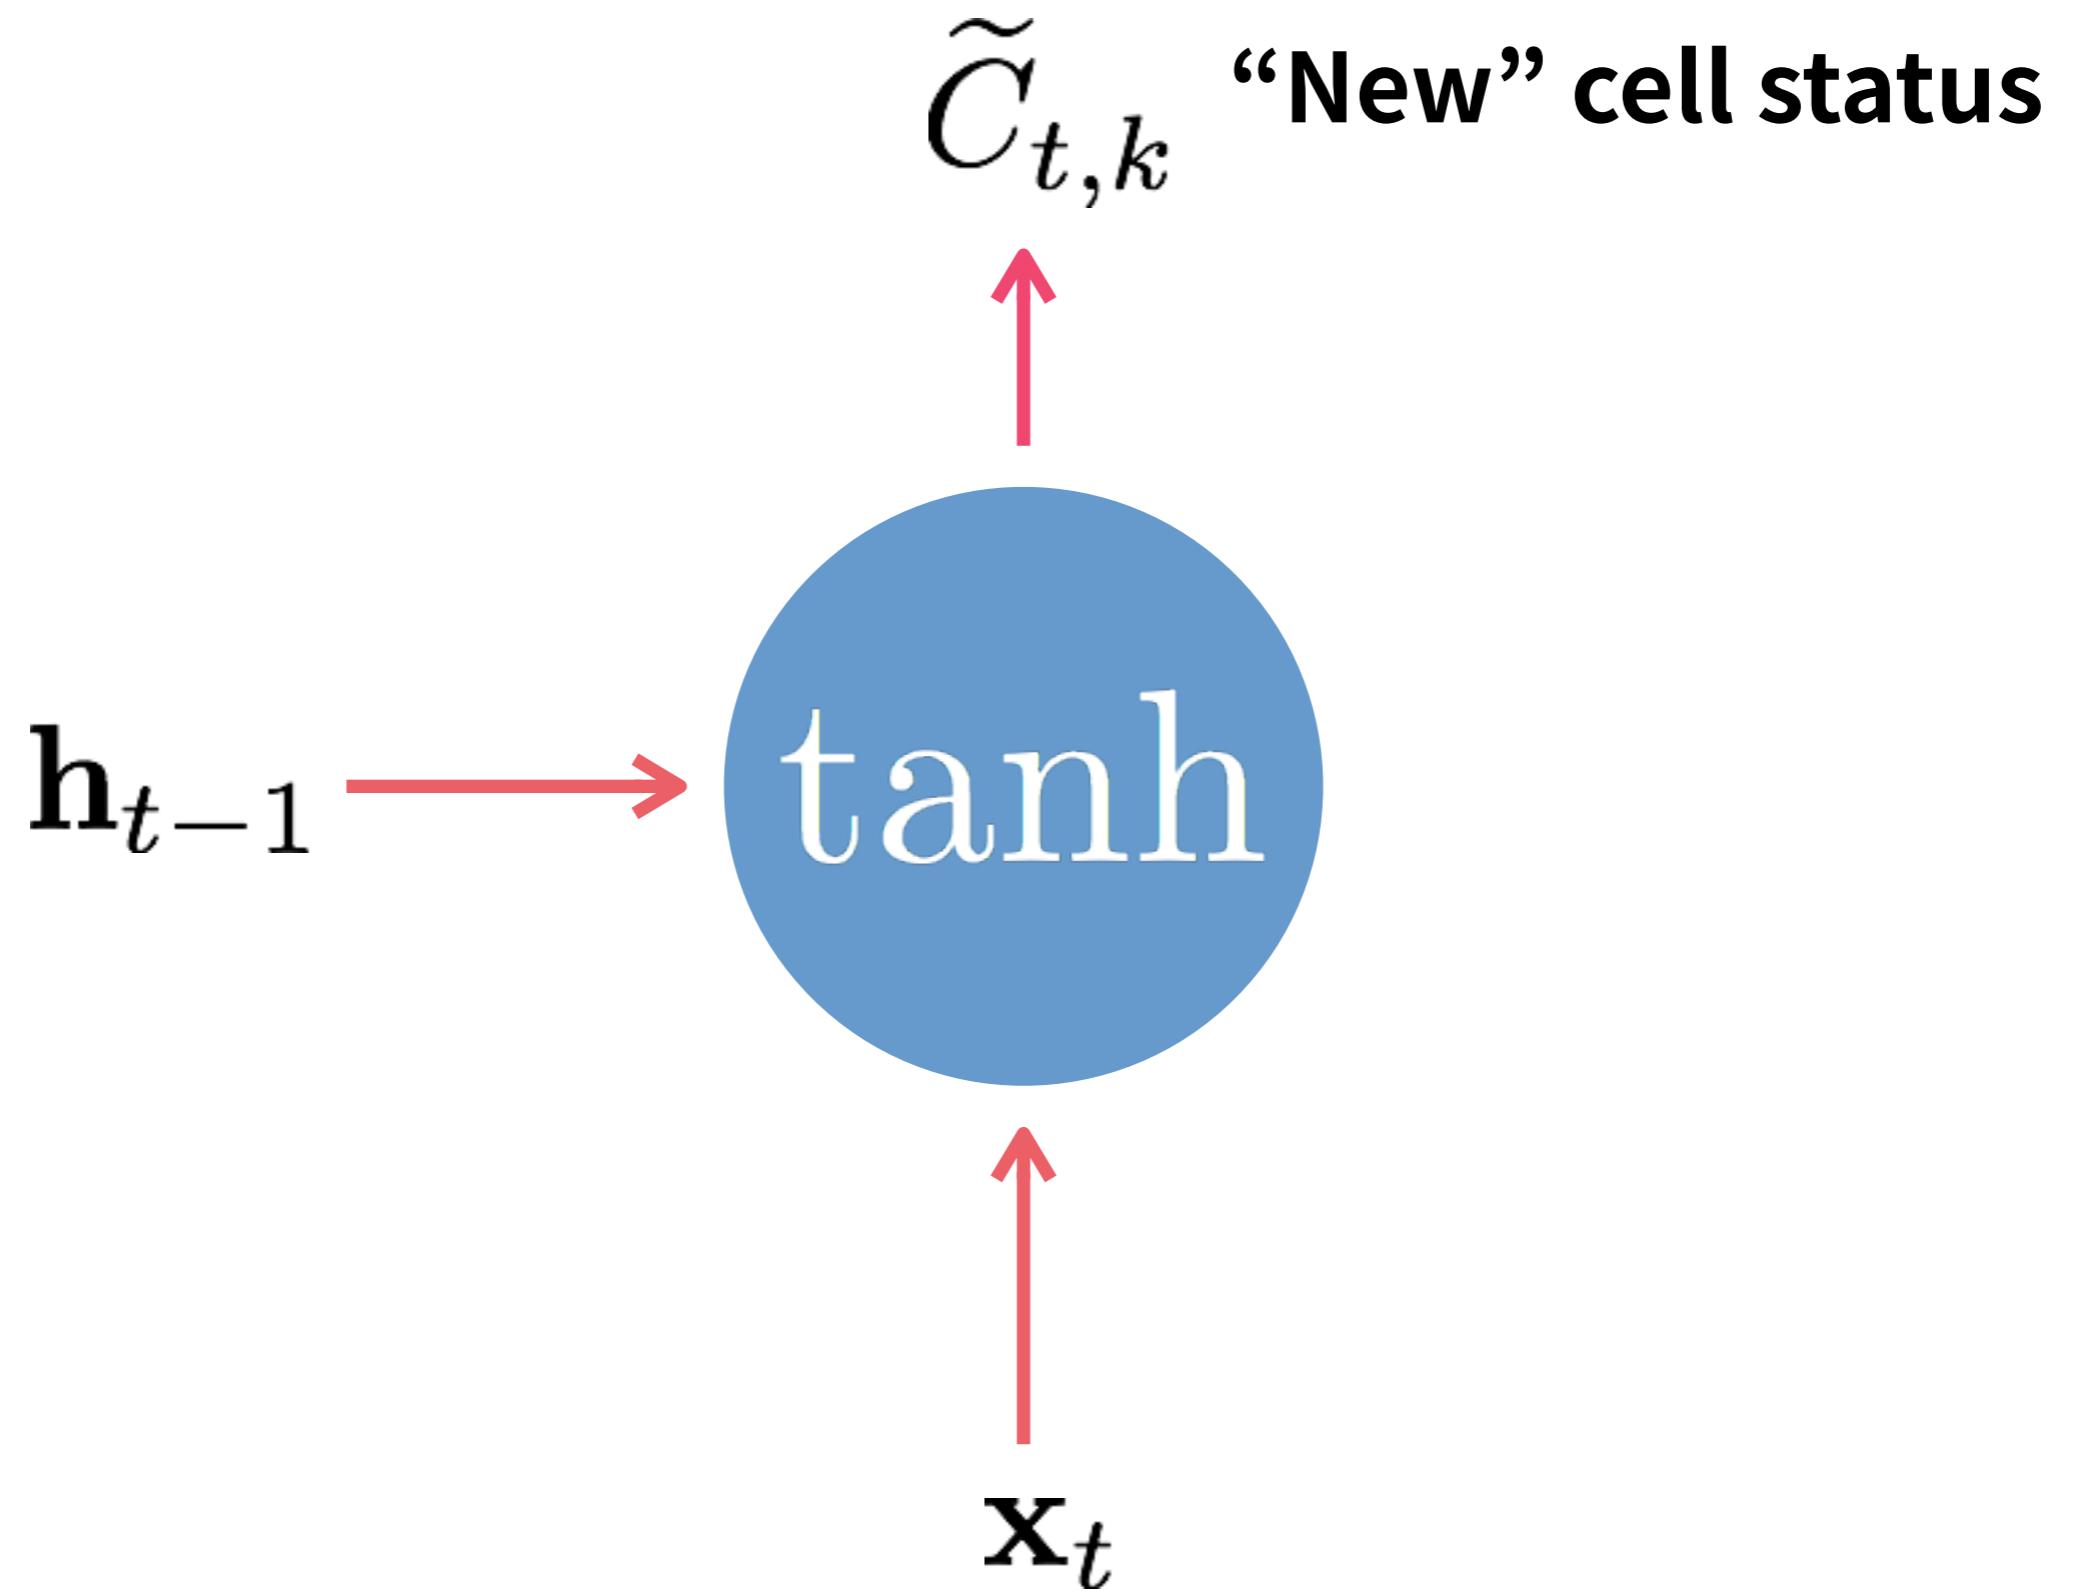


LSTM Again

The cell state is for
the cell only



The output of a cell
will “share” with
other cells



$$C_{t,k} = f_{t,k} C_{t,k-1} + i_{k,t} \tilde{C}_{t,k}$$

$$h_{t,k} = o_{t,k} \tanh(C_{t,k})$$

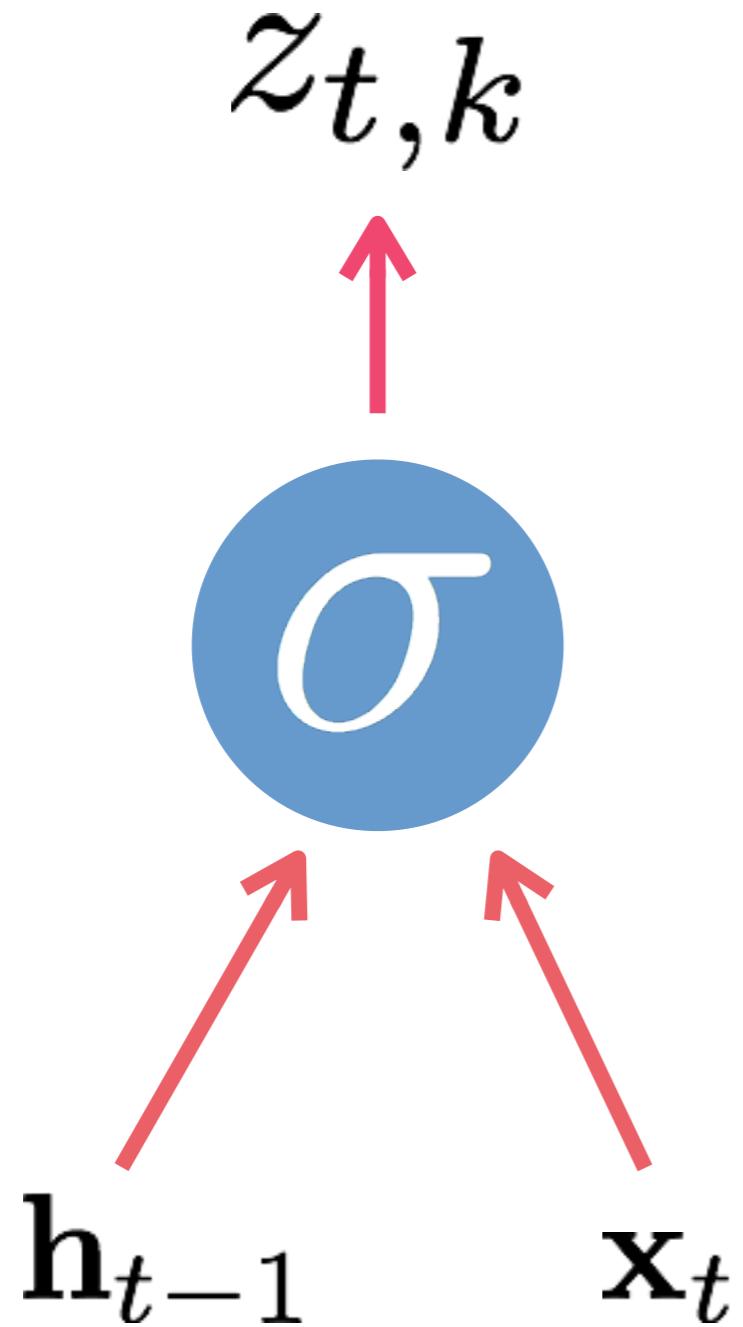
**Do we really need to make
things so complicated?**



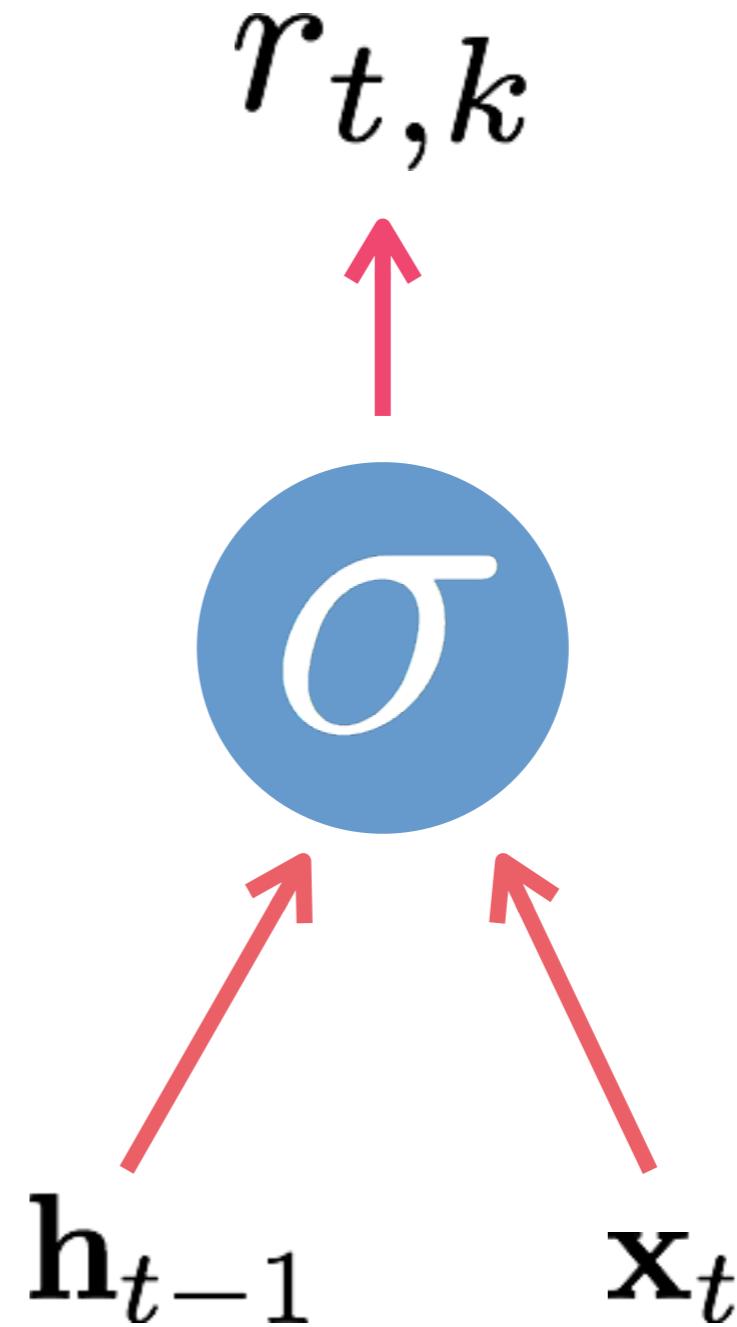
Gated Recurrent Unit
Simplified version of LSTM

Only 2 Gates

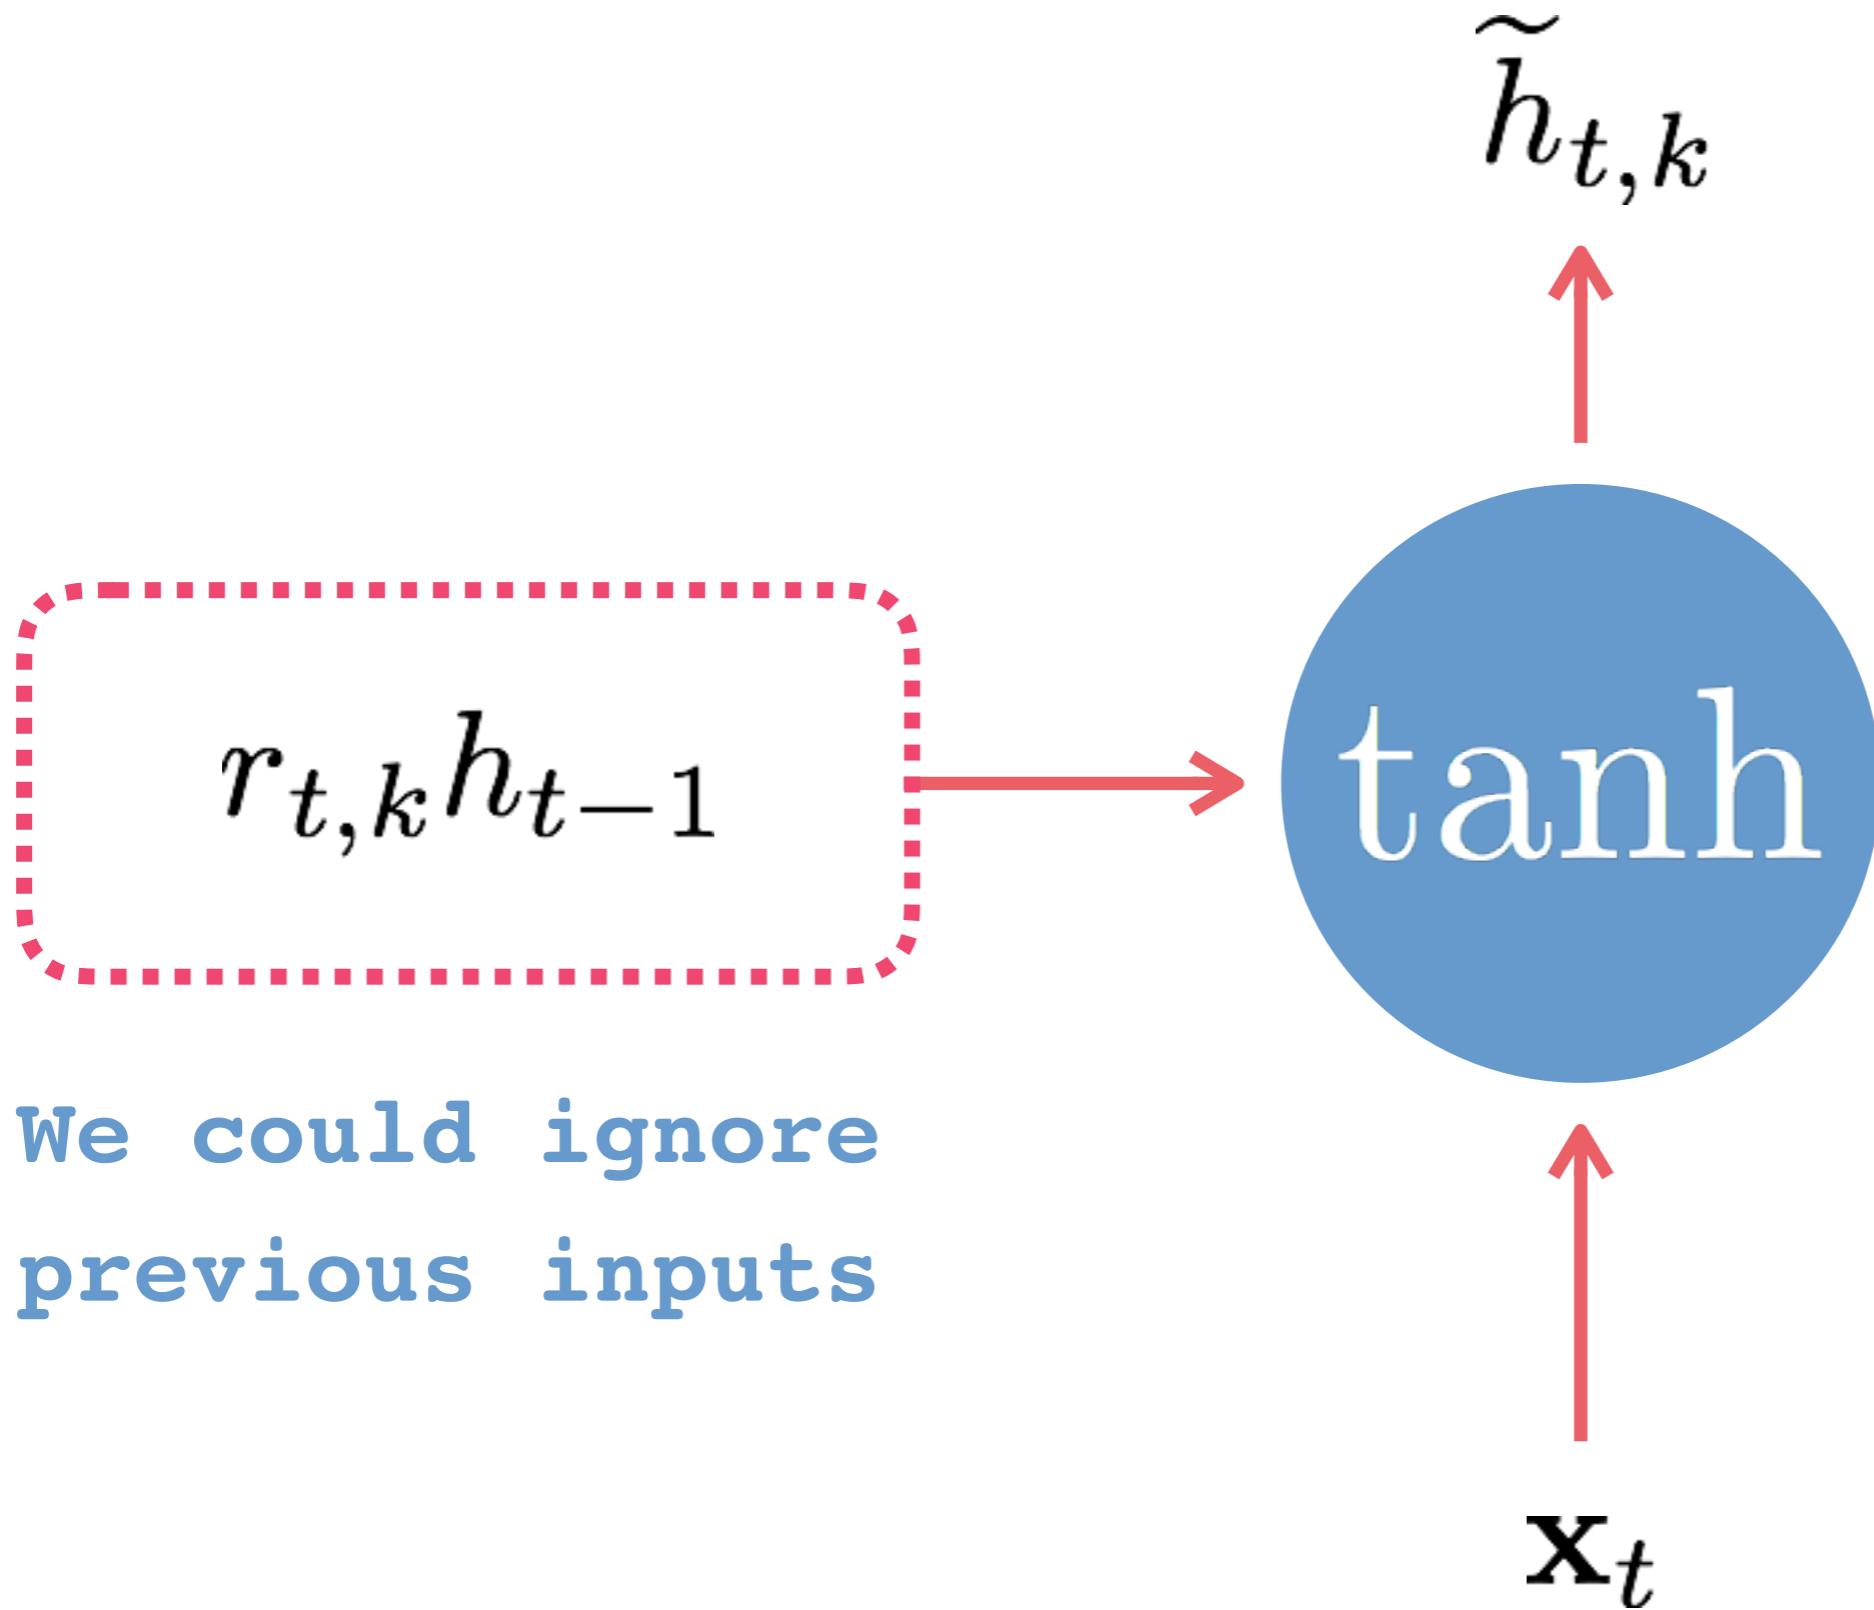
Although it has “gated” in the name



update gate



reset gate

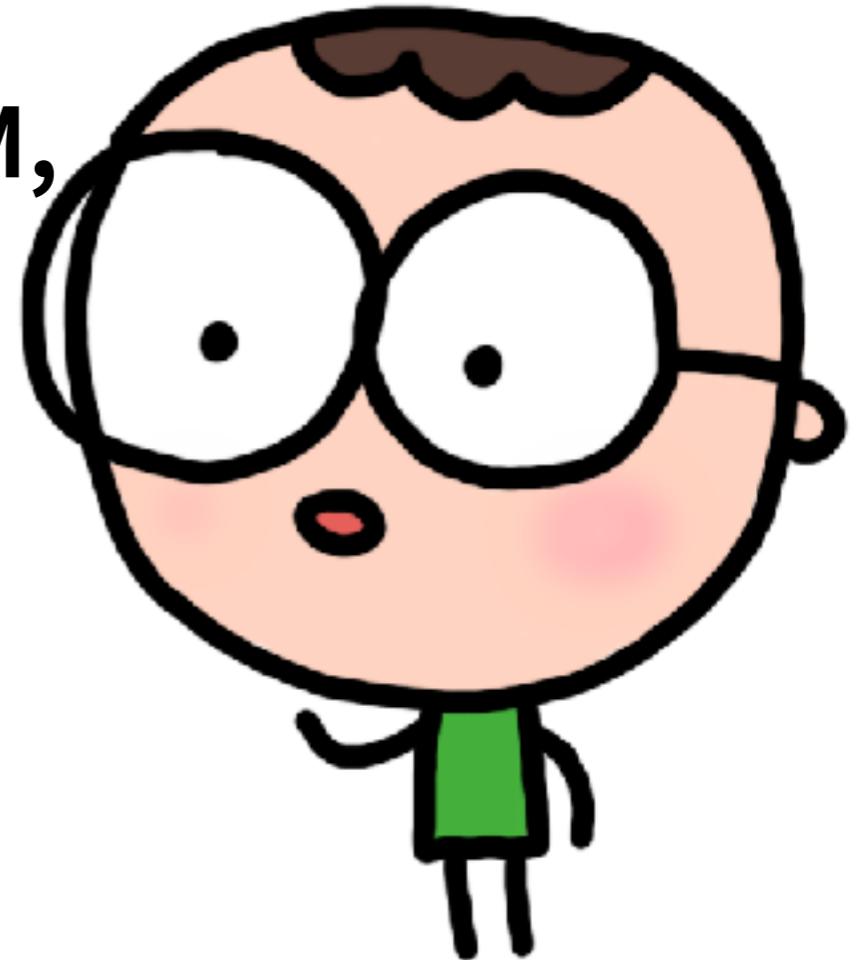


$$h_{t,k} = z_{t,k} h_{t-1,k} + (1 - z_{t,k}) \tilde{h}_{t,k}$$

Names of RNNs

Now talking about RNN, in fact, including the original RNN, LSTM, GRU and other variants.

In particular, the original RNN is called **Vanilla RNN**, and is **SimpleRNN** in Keras.



3

Design thinking,
creative problem
solving



Design Thinking Process (Stanford d.school)



Very close to the engineer's approach

“Our model is rarely successful
for the first time.”

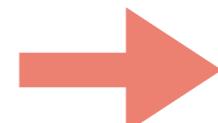
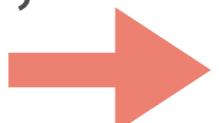
—Yi-Shin Chen, National Tsing Hua University

MLB Player Home Runs Prediction

Year t-1

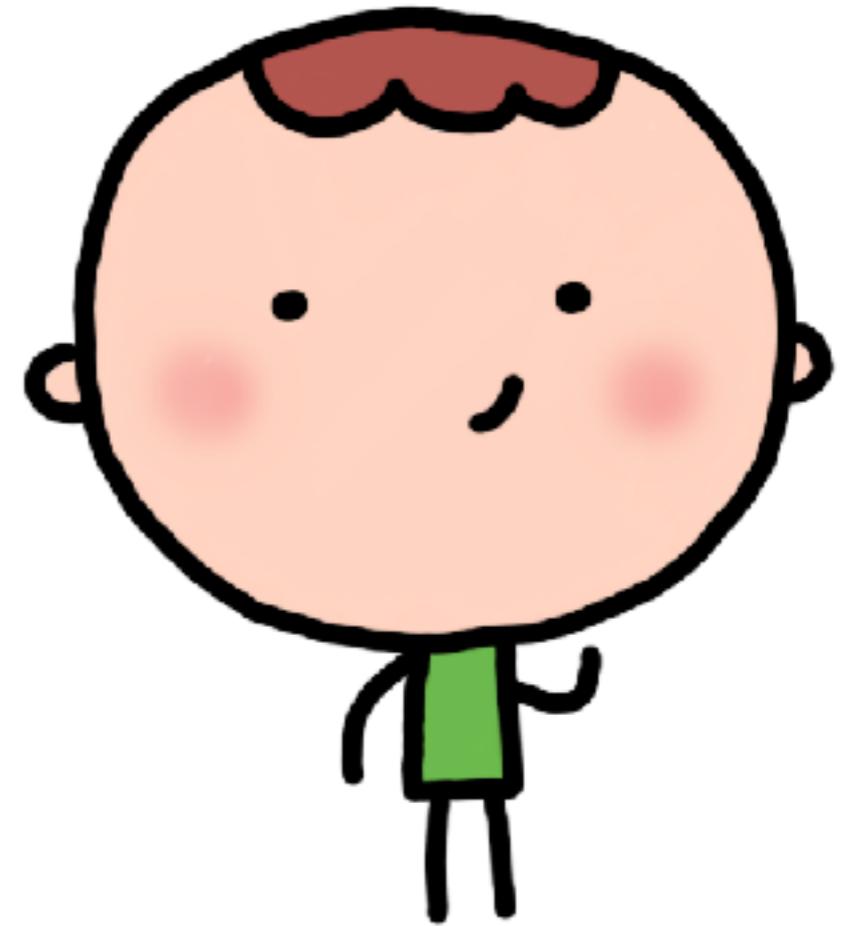
[Age, G, PA, AB, R, H, 2B, 3B, HR,
RBI, SB, BB, SO, OPS+, TB]

15 features



**Number of
homers in the
year of t**

- Use LSTM. Input the data in a period of 10 years and predict the home runs in the next year.
- Only one LSTM layer!



Don't guess the
exact number, guess
which interval!

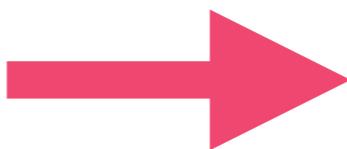


divided into 5 subintervals:

0-9, 10-19, 20-29, 30-39, 40+

One-Hot encoding

10-19



0	1	0-9
1	2	10-19
0	3	20-29
0	4	30-39
0	5	40+

2017 forecast result

Mike Trout (LAA)

Predicted 30-39

Actual 33

Kris Bryant (CHC)

Predicted 30-39

Actual 29

Mookie Betts (BOS)

Predicted 20-29

Actual 24

Daniel Murphy (WSH)

Predicted 20-29

Actual 23

Jose Altuve (HOU)

Predicted 20-29

Actual 24

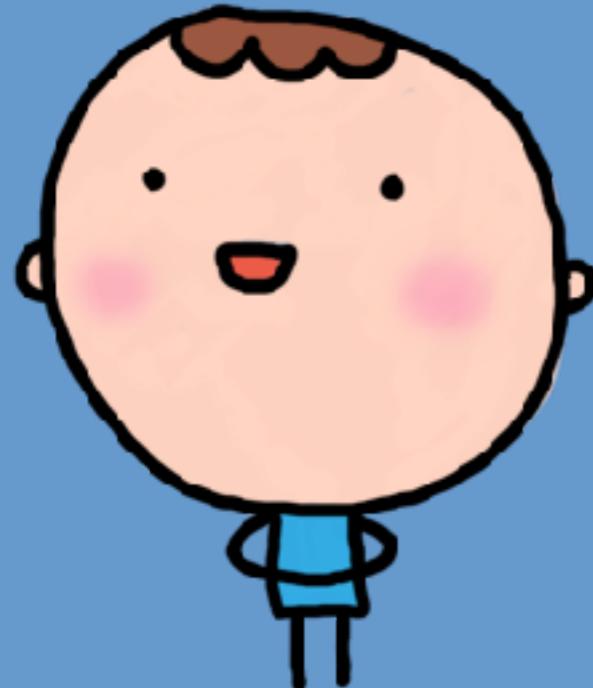
Corey Seager (LAD)

Predicted 20-29

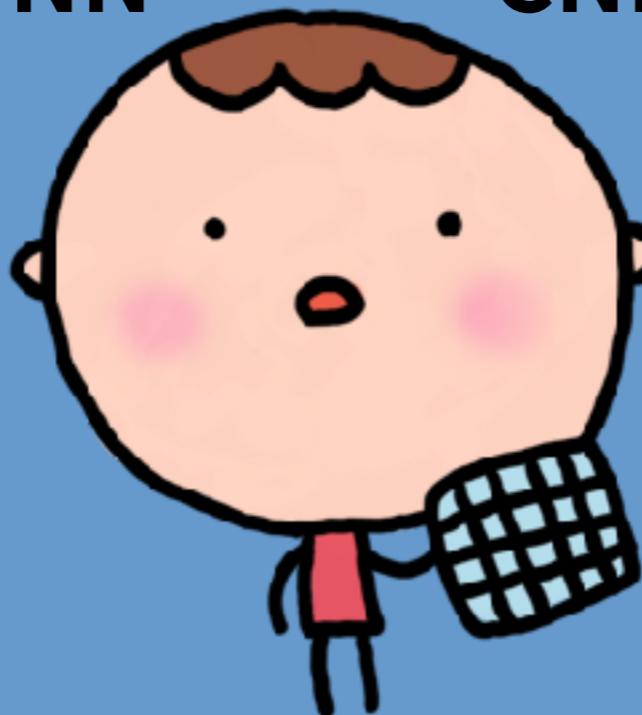
Actual 22

Fancy ways to ask questions!

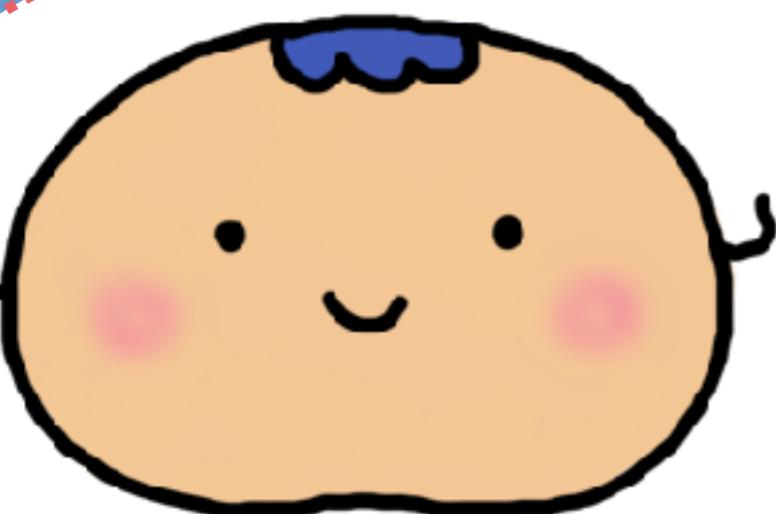
Standard NN



CNN



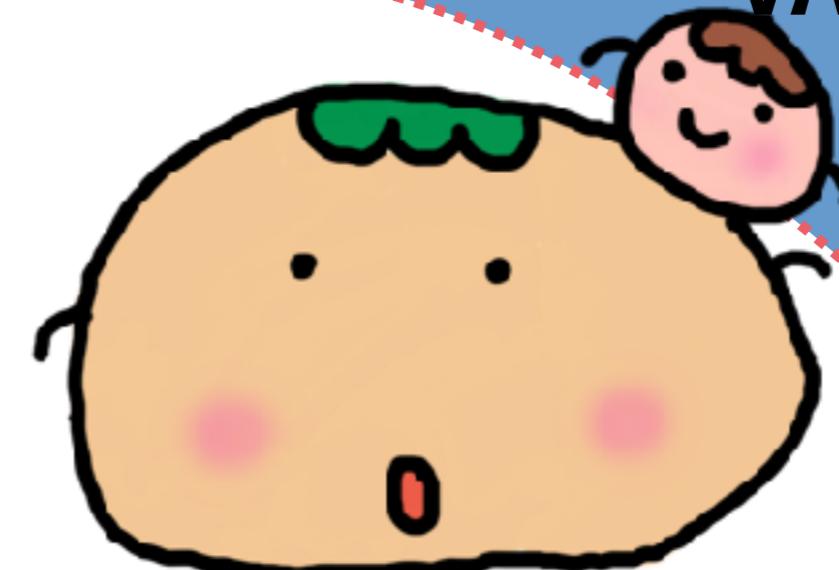
RNN



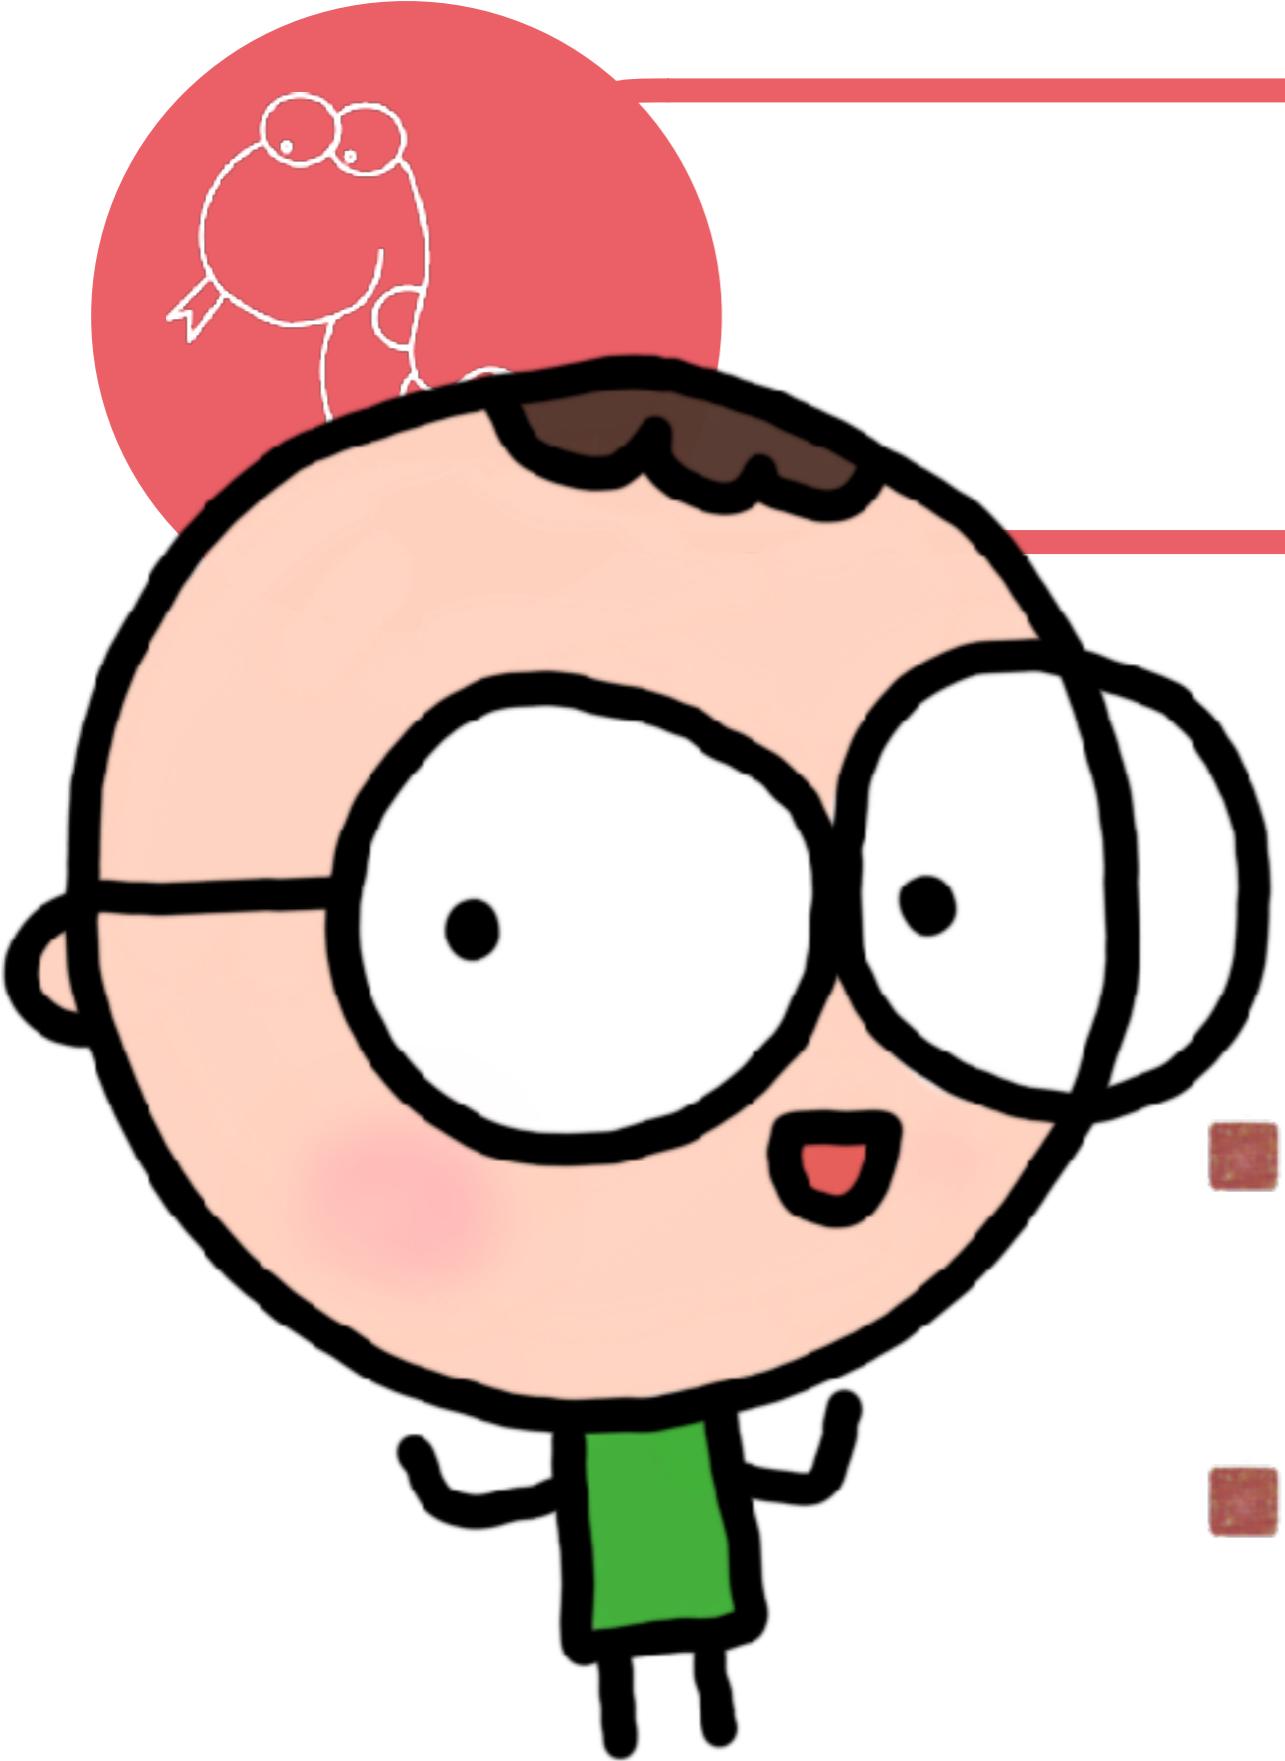
Reinforcement
Learning (RL)

Generative Adversarial
Network (GAN)

VAE



Capsule

A cartoon illustration of a character with large, round, white eyes and black pupils. The character wears black-rimmed glasses and a red baseball cap. A small white bird with large eyes is perched on the rim of the cap.

GAN

- Generative Adversarial Network
- Yann LeCun said GNN is the most promising model

“

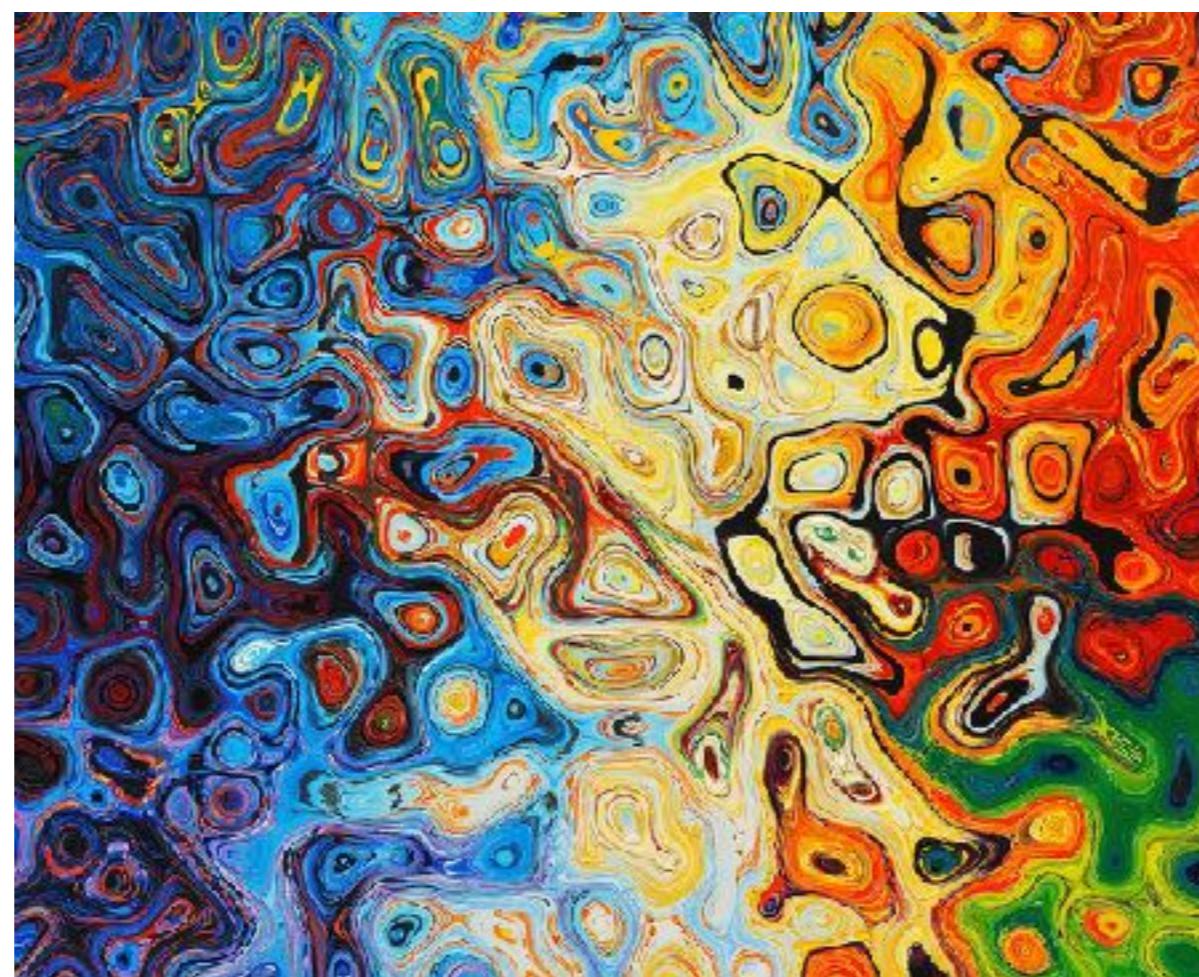
There are many interesting recent development in deep learning...

The most important one, in my opinion, is adversarial training (also called **GAN** for Generative Adversarial Networks).

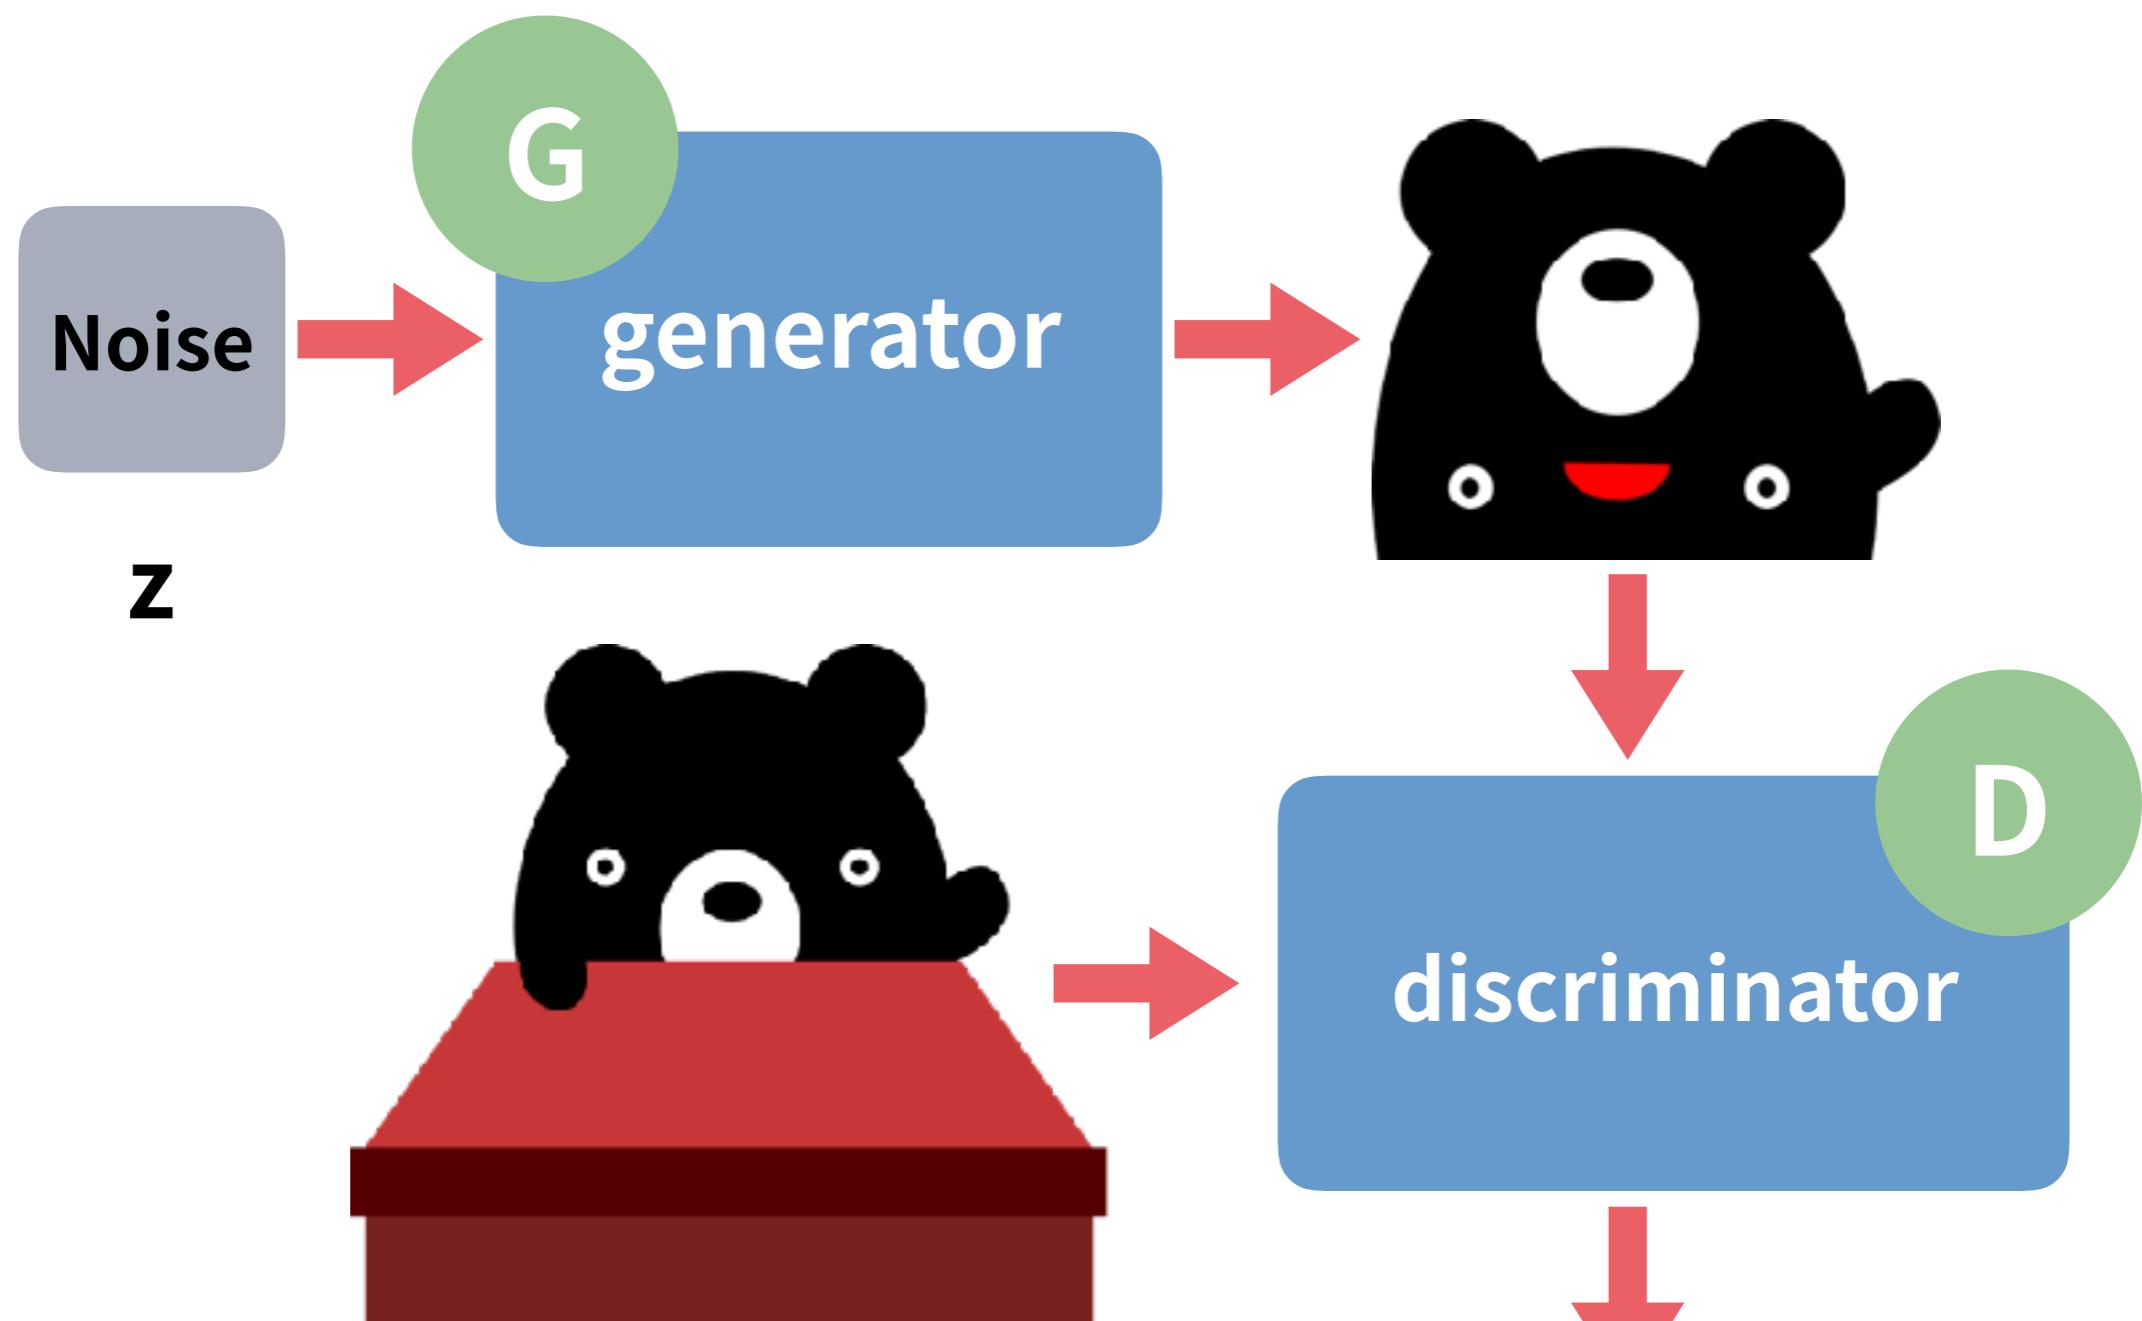
”

—Yan LeCun (楊立昆), 2016

The GAN Zoo

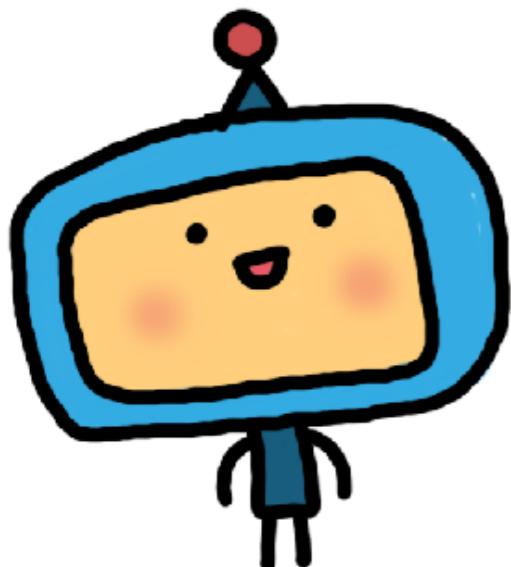


[https://github.com/hindupuravinash/
the-gan-zoo](https://github.com/hindupuravinash/the-gan-zoo)



- A GAN consists two neural networks,
- a Generator and a Discriminator.

D, G PK!

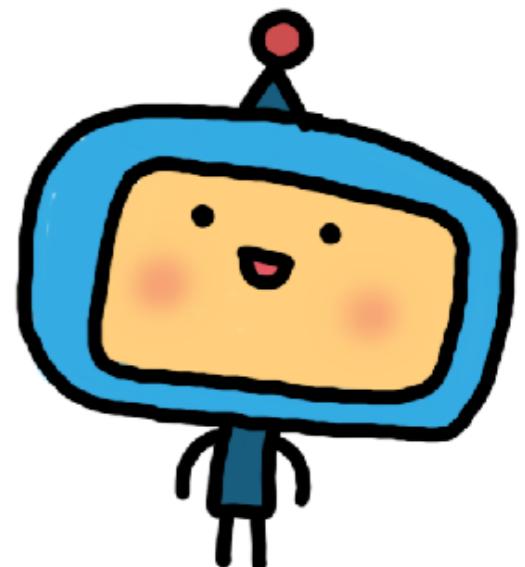


Generator
G

Want
 $D(G(\mathbf{z}))$
Close to 1



Want
 $D(G(\mathbf{z}))$
Close to 0



Discriminator
D

$D(\mathbf{x})$
Close to 1

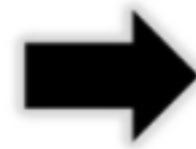
iGAN

Jun-Yan Zhu et al. (ECCV 2016)

<https://arxiv.org/abs/1609.03552>

“Generative Visual Manipulation on the Natural Image Manifold”

User edits



Generated images



— Color

— Sketch

Every one can draw!

<https://youtu.be/9c4z6YsBGQ0>

Progressive GAN

Karras et al. NVIDIA team, (ICLR 2018)
<https://arxiv.org/abs/1710.10196>

“Progressive Growing of GANs for Improved Quality, Stability, and Variation”

Karras-Aila-Laine-Lehtinen

Progressive Growing of GANs for Improved Quality, Stability, and Variation



- By a NVIDIA team
- Theano, Python 2, single GPU (Hey, it's NVIDIA)



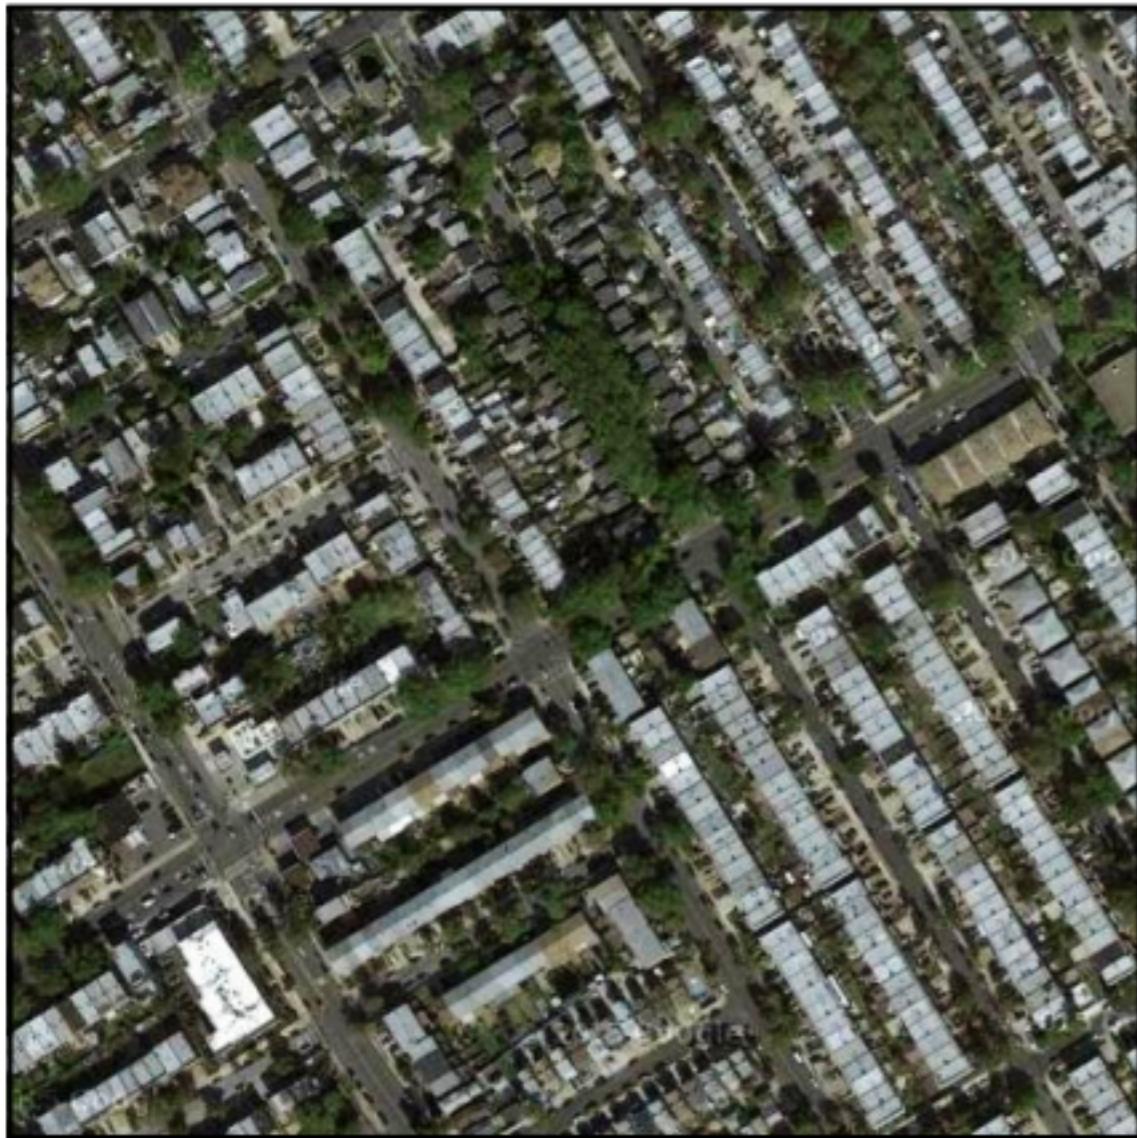
These are fake (1024x1024)

Pix2Pix

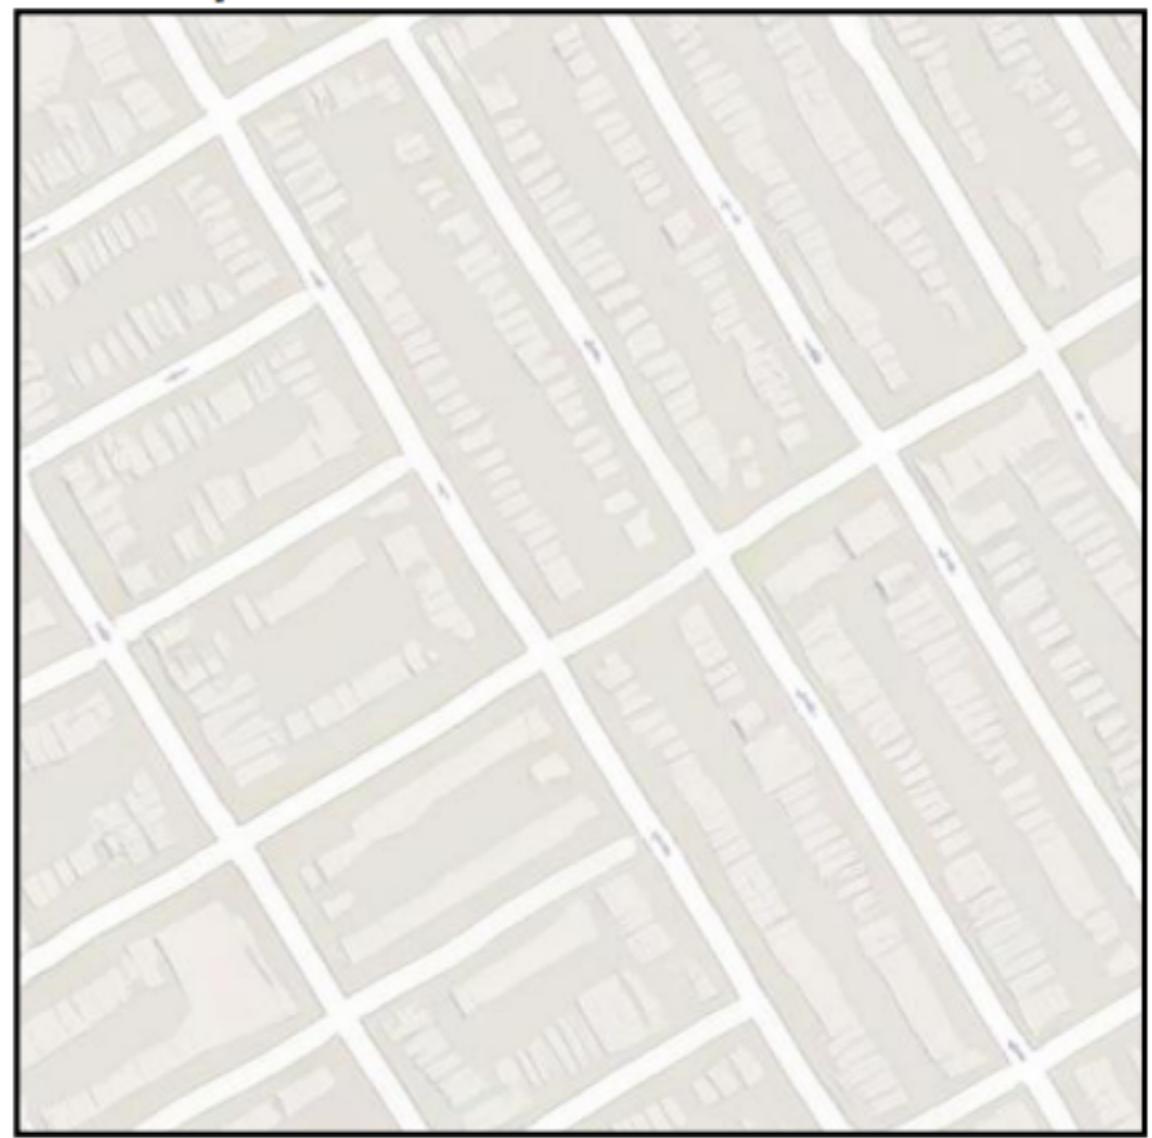
Isola, Zhu, et al., (CVPR 2017)

<https://arxiv.org/abs/1611.07004>

“Image-to-Image Translation with Conditional Adversarial Networks”



input



output

Pix2pix transfers satellite images to maps

* From the original paper of Isola, Zhu et al. (2017)

Labels to Street Scene



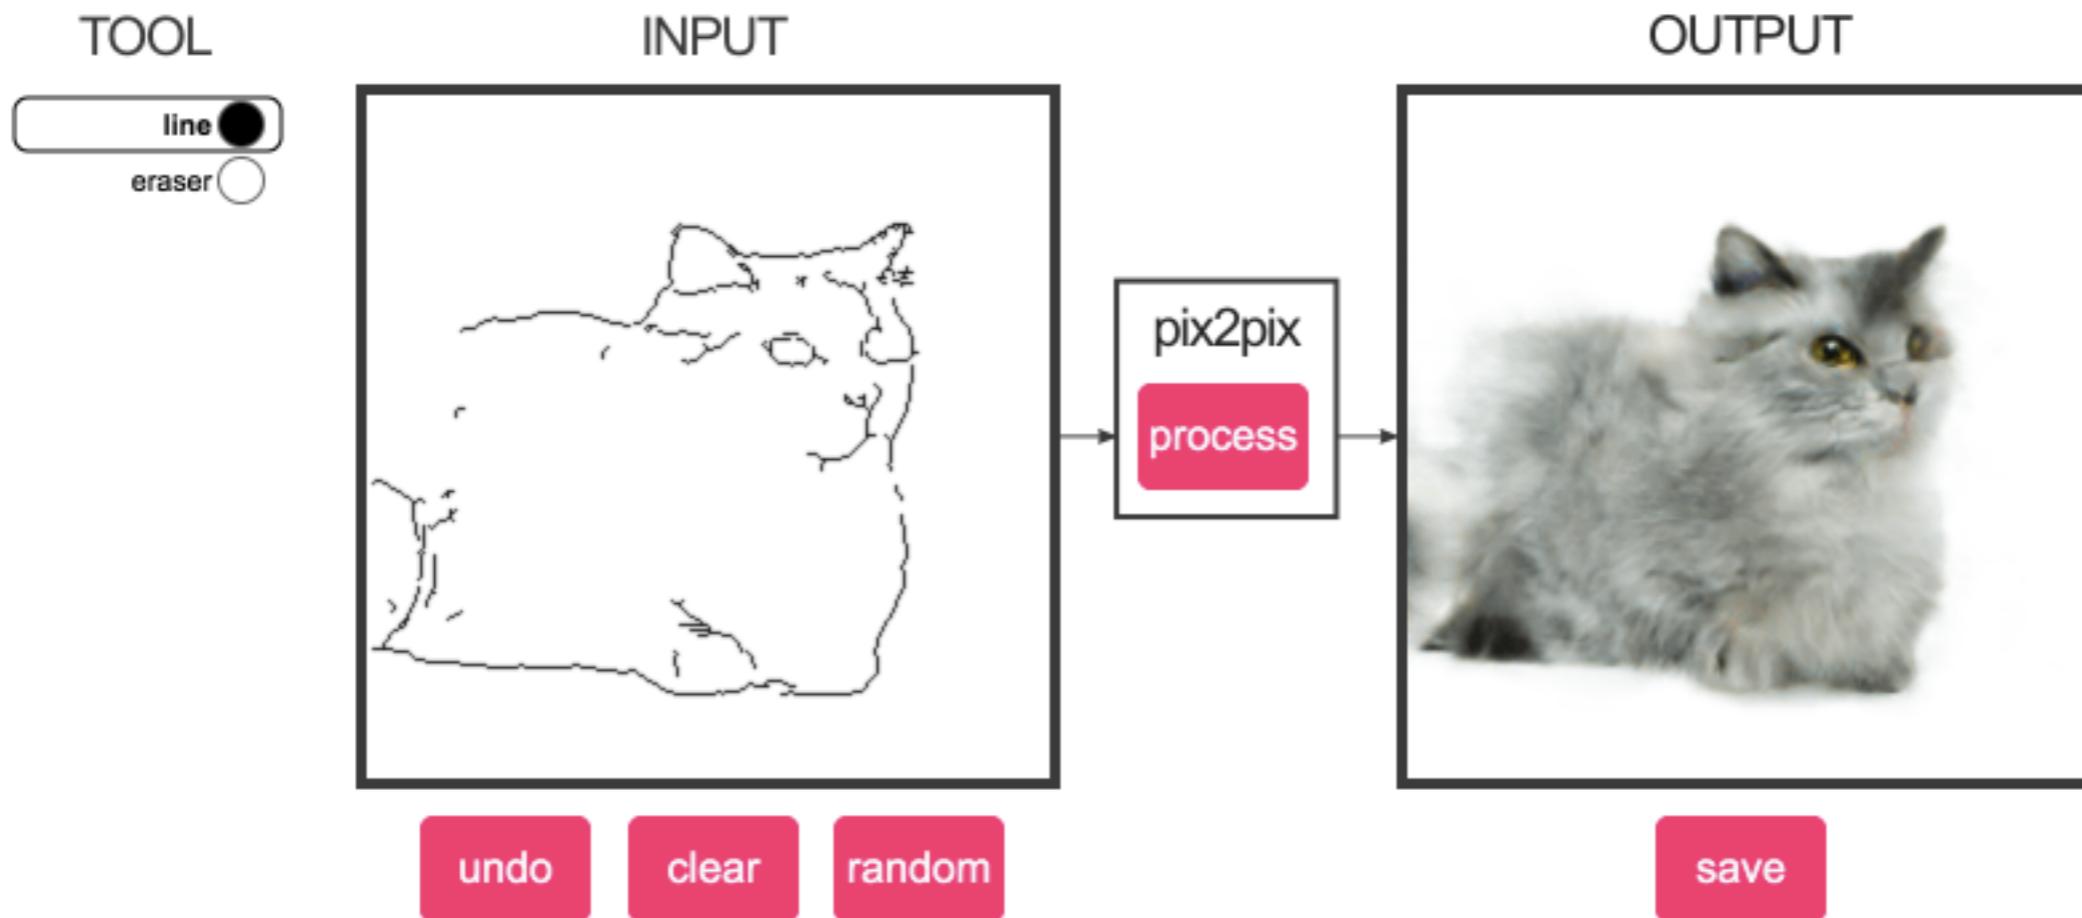
input

output

Pix2pix generates street views

* From the original paper of Isola, Zhu et al. (2017)

edges2cats



Pix2pix On-line Version

<https://affinelayer.com/pixsrv/>

* by Christopher Hesse

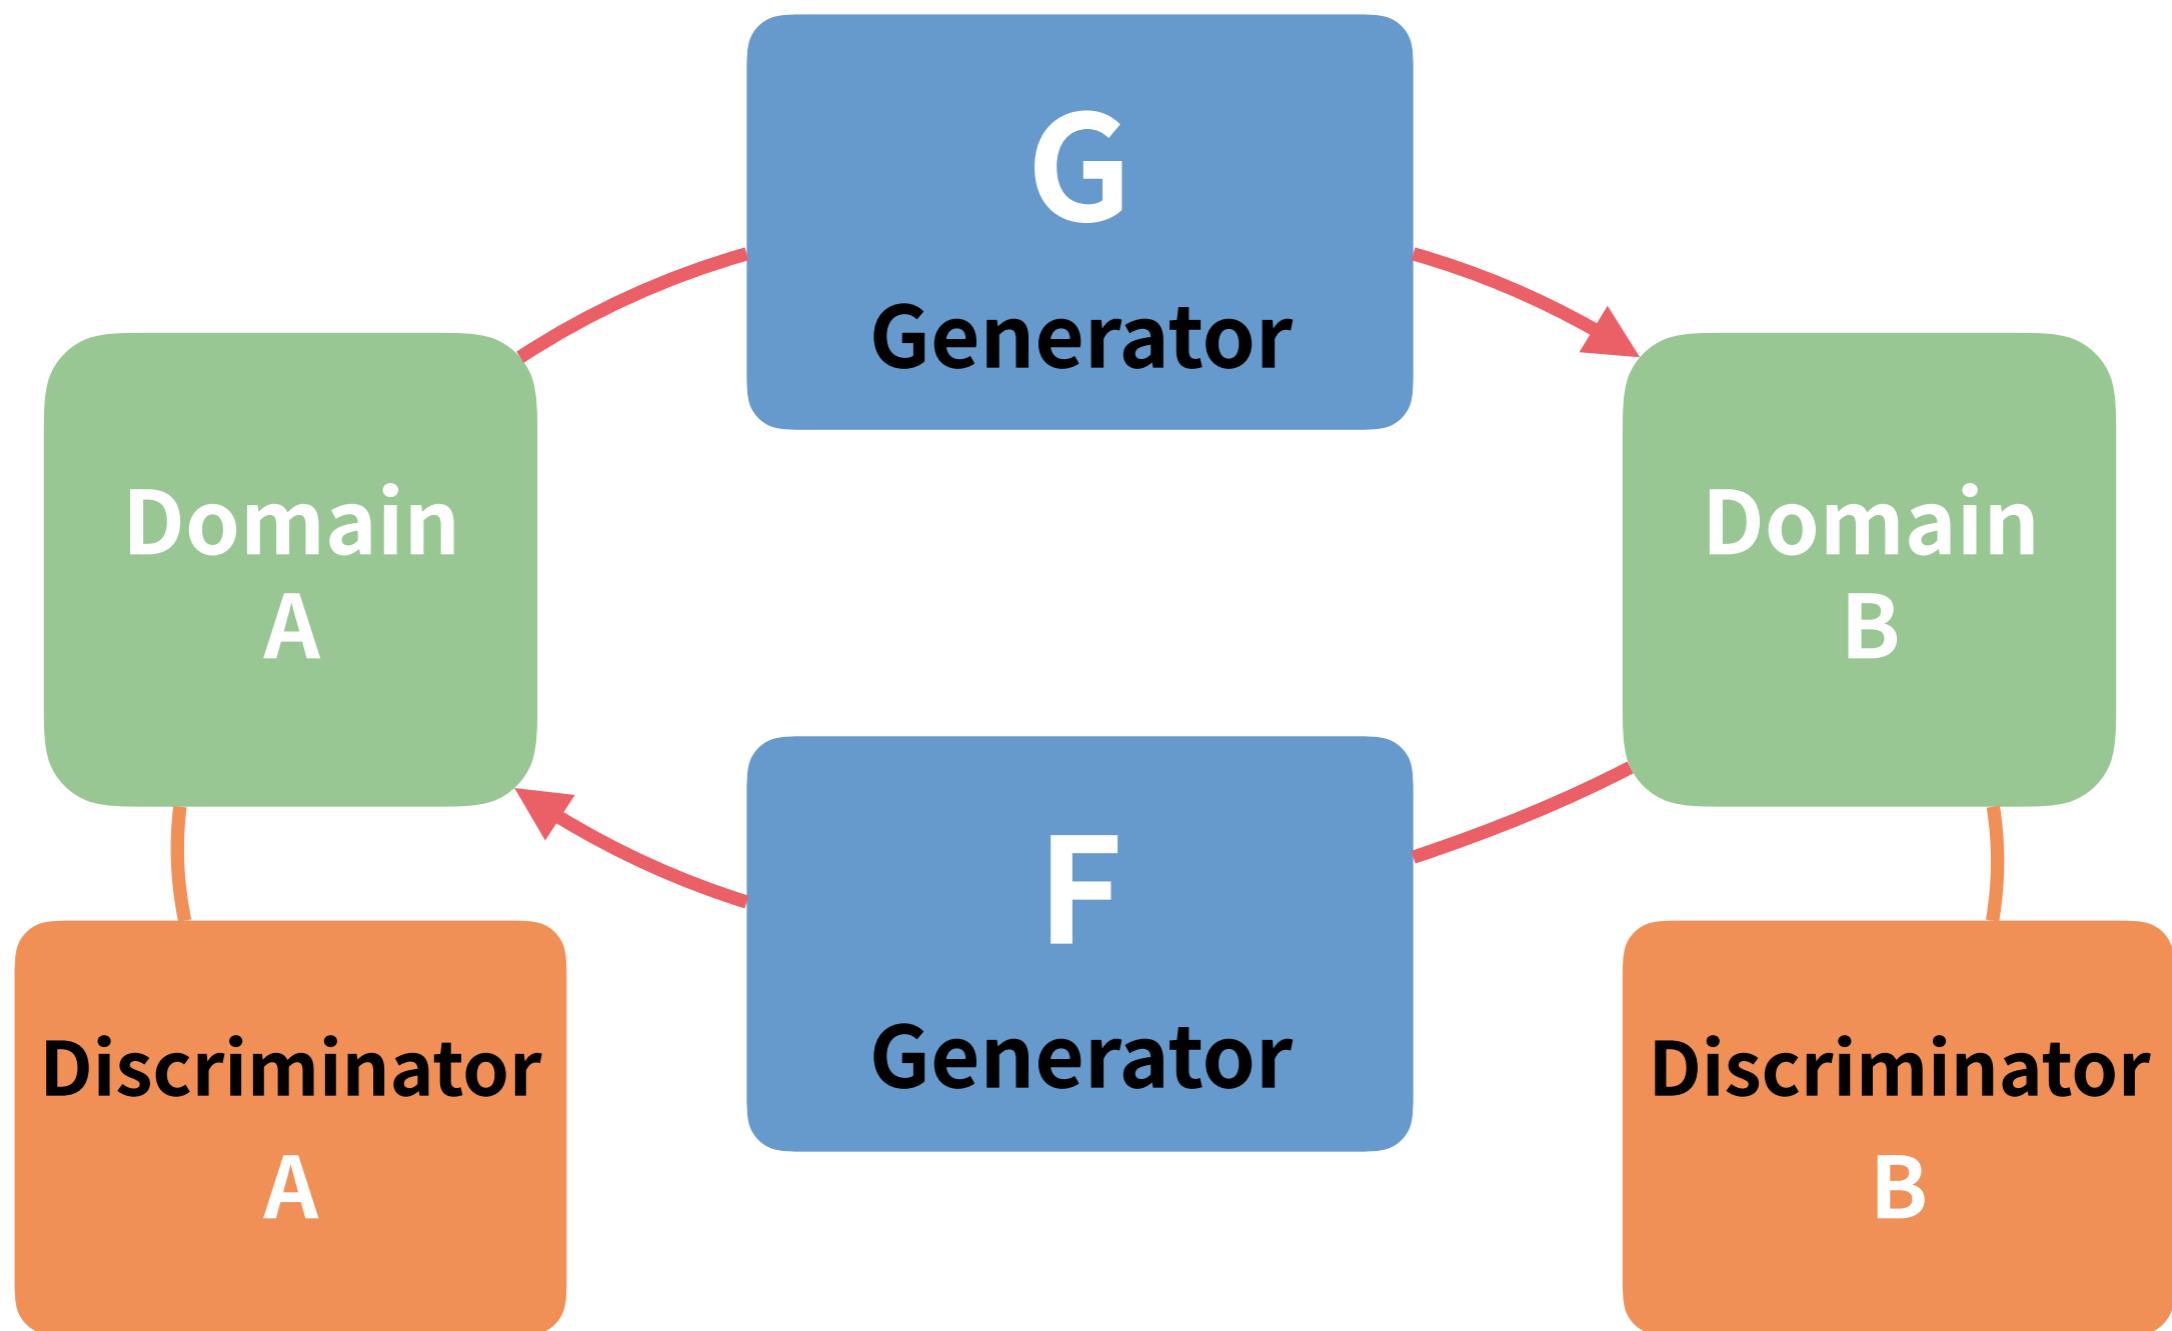
CycleGAN

Jun-Yan Zhu et al. (ICCV 2017)

<https://arxiv.org/abs/1703.10593>

“Unpaired Image-to-Image Translation using Cycle-Consistent Adversarial Networks”

CycleGAN



**Don't need to pair
our data sets!!**

And open infinitely many possibilities...



Horses to Zebras!!

<https://youtu.be/9reHvktowLY>



One example from authors of CycleGAN

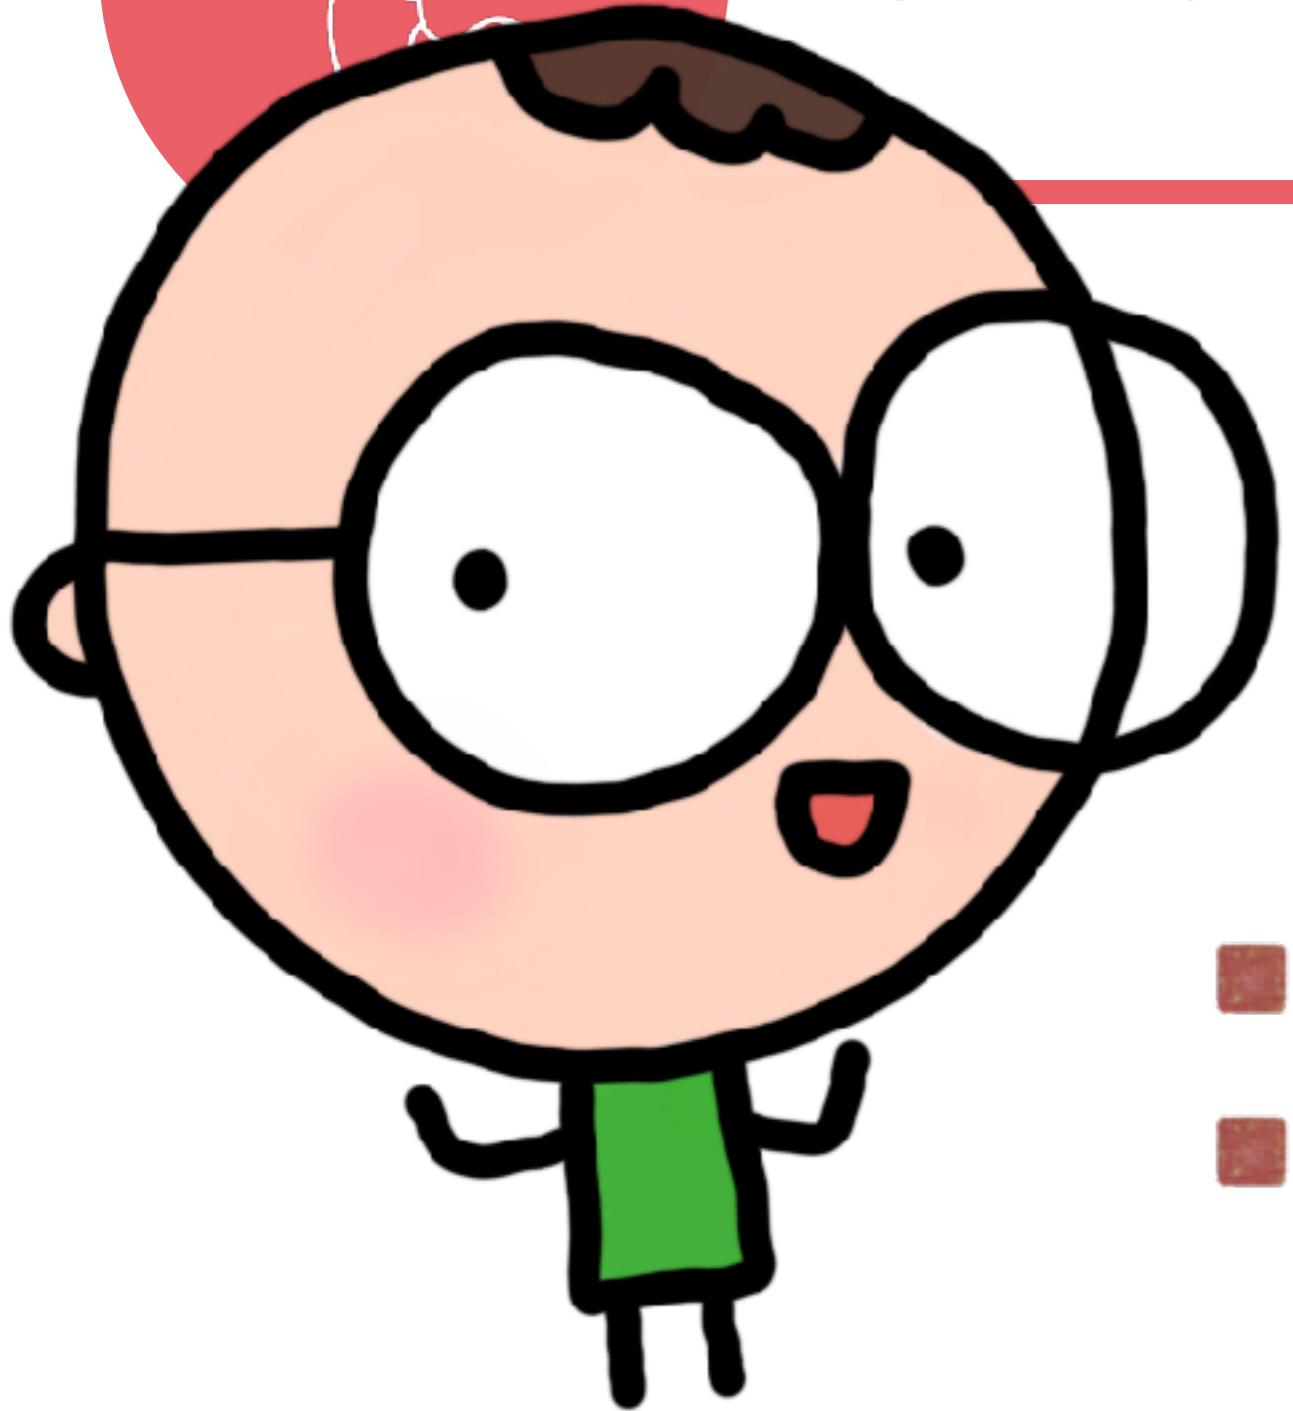


Face-off (Tzer-Jen Wei)

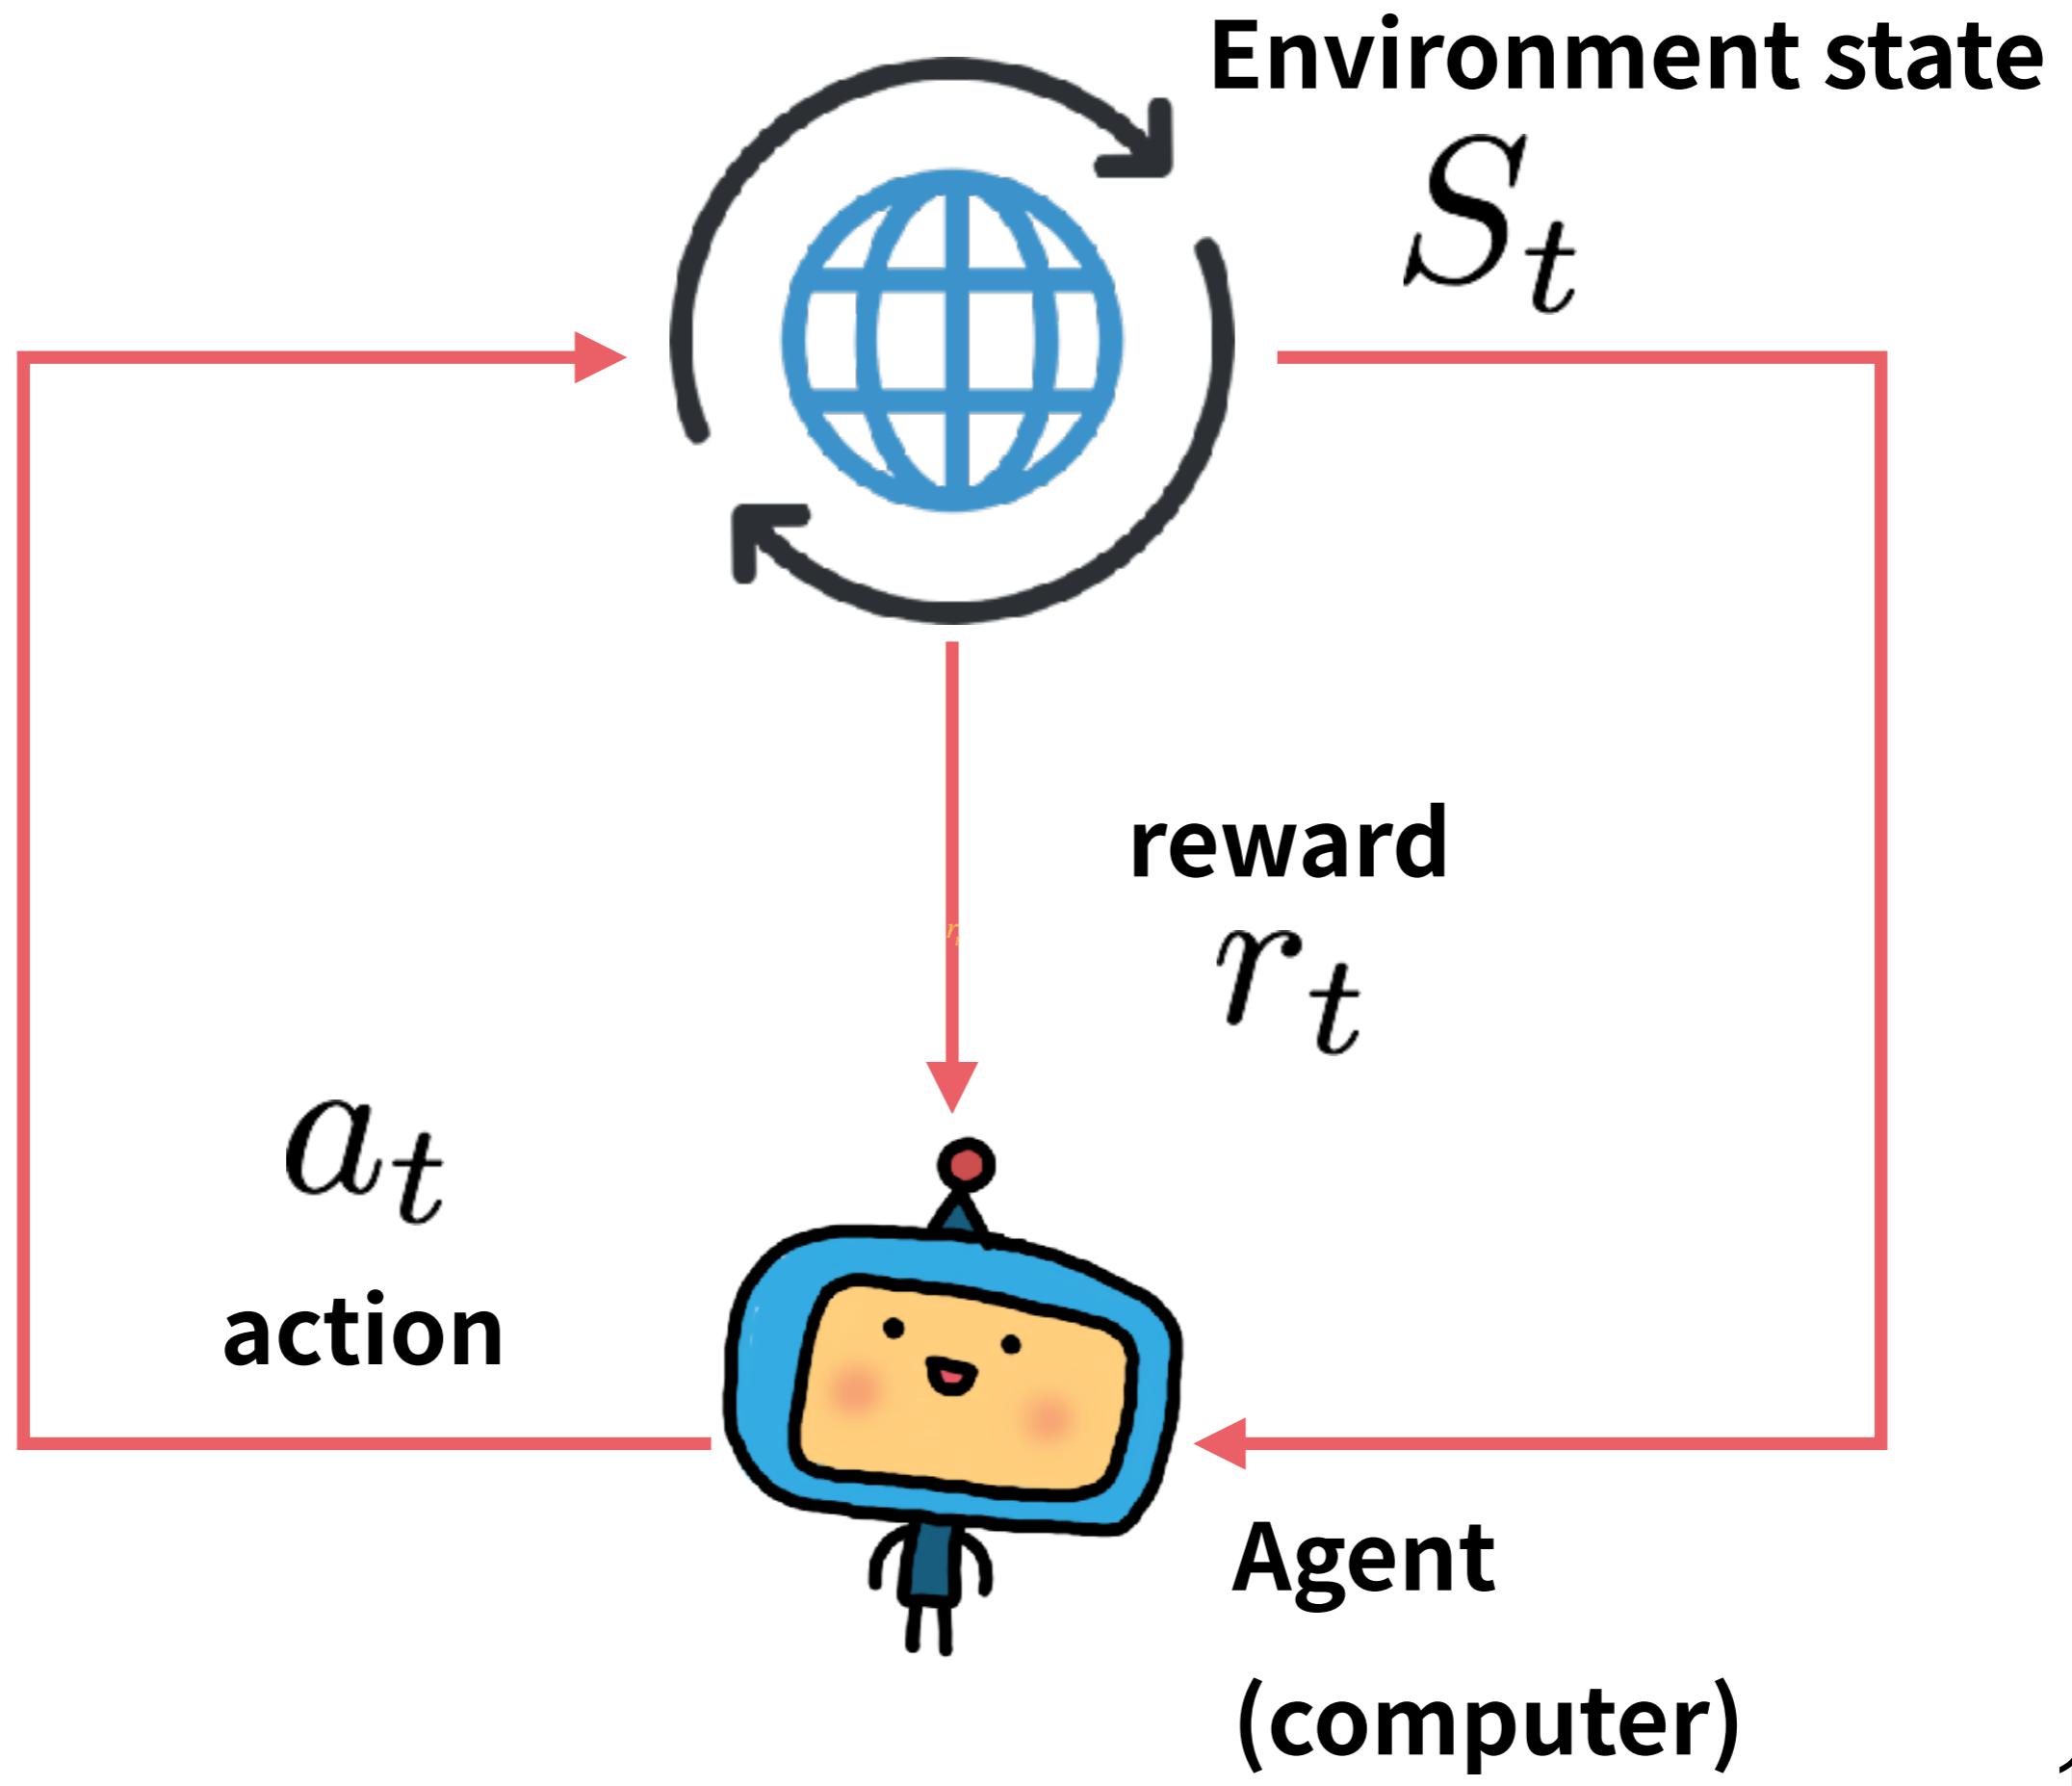
<https://youtu.be/Fea4kZq0oFQ>



Reinforcement Learning



- Reinforcement Learning
- AlphaGo key technique



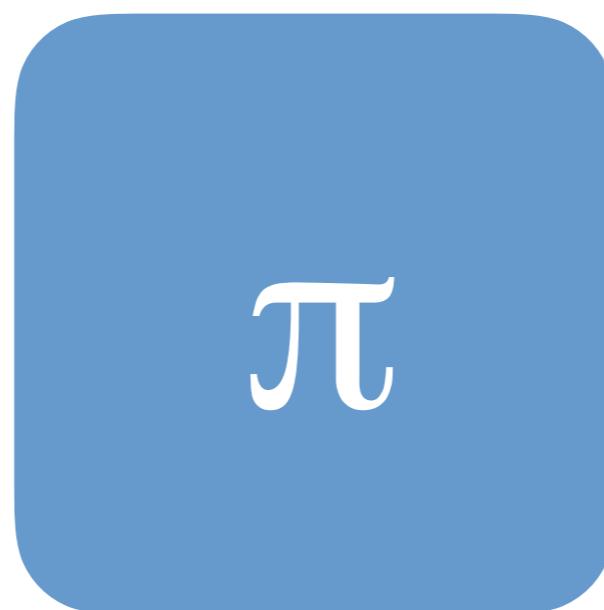


Let's play!

1

Policy Based

policy function



Left

or

Right

State S_t

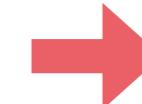
Action a_t

**Usually it is not easy to
learn directly...**

2

Value Based

Value function



Grade
(Usually
estimate
reward)

+

Action



Netflix AlphaGo Film (Highly Recommended)

Self-driving cars



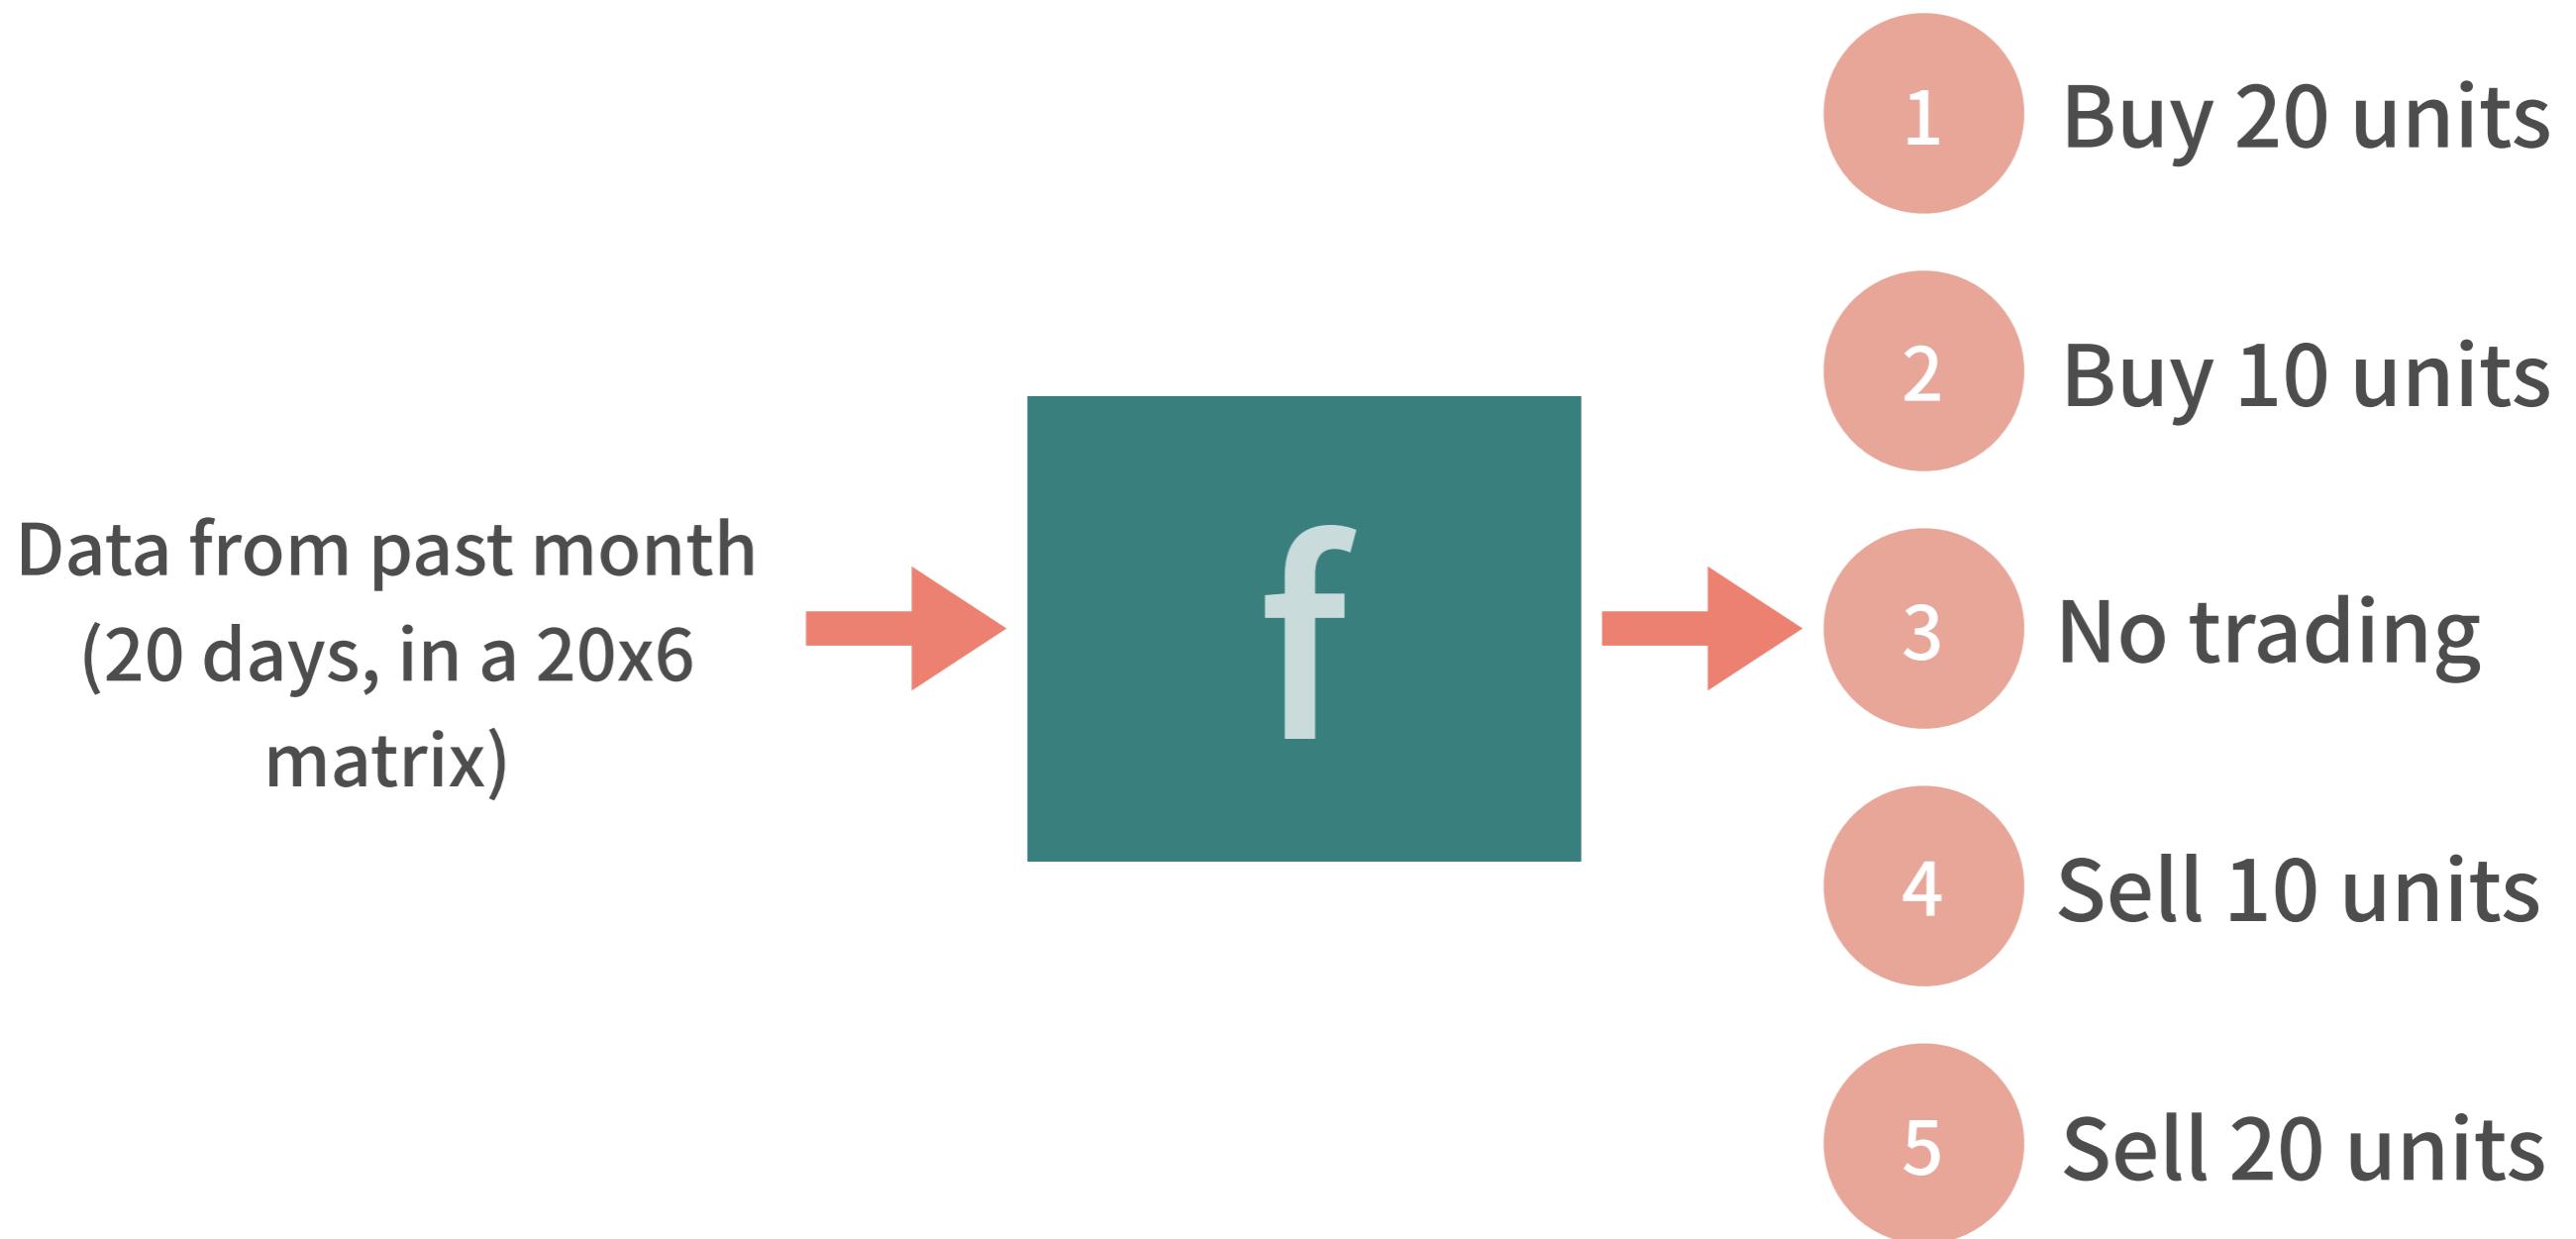
Automated trading system

- Fixed one ETF each time
- Start with \$20,000
- Work for one year (Sell all ETF at the end)
- Using reinforcement learning

* ETF data provided by the Global Intelligence



5 possible actions



Trading Results

	CDQN	Hold & buy
ETF1	17.71%	10.89%
ETF2	16.53%	12.6%
ETF3	16.3%	0.35%
ETF4	14.4%	13.25%
ETF5	14.3%	12.7%
ETF6	13.91%	13.37%
ETF7	13.17%	10.52%
ETF8	12.35%	17.07%
ETF9	11.68%	10.81%
ETF10	11.09%	8.14%

	CDQN	Hold & buy
ETF11	10.76%	5.26%
ETF12	10.19%	13.17%
ETF13	7.8%	1.42%
ETF14	6.23%	3.56%
ETF15	5.73%	4.61%
ETF16	3.78%	-12.76%
ETF17	2.85%	5.83%
ETF18	1.59%	-4.45%
ETF19	1.07%	-18.09%
ETF20	-0.59%	-0.75%